

MH4921  
Supervised Independent Study II

**Linux Virtual Private Network (VPN)  
Lab**

**Brandon Goh Wen Heng**

Academic Year 2018/19

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Overview</b>	<b>1</b>
<b>3</b>	<b>Exploration</b>	<b>1</b>
3.1	Virtual Machine (VM) Setup . . . . .	1
3.1.1	VM Configuration . . . . .	1
3.2	Task 1: Creating Host-to-Host Tunnel with TUN/TAP . . . . .	2
3.2.1	Setting Up Tunnel Point A . . . . .	3
3.2.2	Setting Up Tunnel Point B . . . . .	4
3.2.3	Establishing Routing Path . . . . .	5
3.2.4	Using The Tunnel . . . . .	5
3.3	Task 2: Creating Host-to-Gateway Tunnel . . . . .	8
3.3.1	Setting Up IP Forwarding . . . . .	8
3.3.2	Getting Around The Limitation of NAT . . . . .	8
3.3.3	Setting Up Firewall . . . . .	9
3.3.4	Bypassing Firewall . . . . .	10
<b>4</b>	<b>Appendix A</b>	<b>14</b>

# 1 Introduction

Organisation, Internet Service Providers (ISPs) and countries often block users from accessing certain external sites through what is known as egress filtering. This is used within organisations to reduce distractions, countries may make use of it to censor foreign web sites. However, these firewalls can easily be bypassed and there are multiple services/applications that aid in the circumvention.

The most common technology that is used to bypass these egress filtering are Virtual Private Networks (VPNs). VPNs are also commonly available on smartphone devices to help bypass these egress firewalls.

## 2 Overview

This lab will focus on how the VPN works and how it can be used to bypass egress firewalls. A VPN usually depends on two components, IP tunnelling and encryption. Tunnelling is crucial in helping to bypass these firewalls while the encryption helps to protect the content that is being transmitted through the VPN tunnel. For simplicity, encryption is not in the scope of this lab.

## 3 Exploration

Implementing a simple VPN for Linux is not trivial and is broken into multiple segments, with detailed observations for each step in the process. By setting up the VPN, a tunnel is established between two systems. When system A is behind a firewall and wants to access restricted resources, it will be blocked by egress filtering. Therefore, the packets are routed through to system B where it will send the packets out to the Internet. Similarly, the response from the Internet will be routed through to system B before sending it back to system A via the tunnel. This is how the VPN helps system A to bypass the firewall.

### 3.1 Virtual Machine (VM) Setup

#### 3.1.1 VM Configuration

Two VMs are required for this lab. VM1 is behind a firewall while VM2 is outside the firewall. The objective is to help VM1 bypass the firewall so it is able to access blocked sites on the Internet.

VM1 and VM2 are connected to a separate NAT adapter, so both can access the Internet using the corresponding NAT servers. To emulate the Internet, where both VM1 and VM2 can communicate to each other, both VMs are connected to a host-only network adapter. This network emulates the facilitation of communication between the two VMs. The network setup is depicted in figure 1 for easier reference.

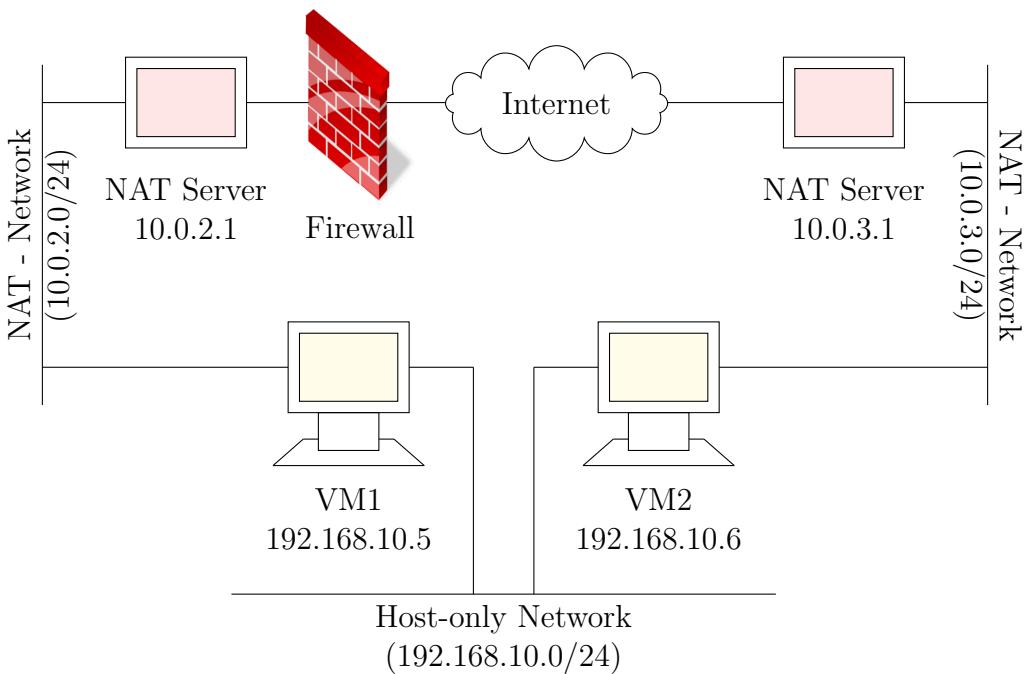


Figure 1: Network Topology

### 3.2 Task 1: Creating Host-to-Host Tunnel with TUN/TAP

TUN/TAP are essential and is widely implemented in modern computing systems. TUN and TAP are virtual network kernel drivers and are implemented entirely by software. TAP (as in network tap) simulates an Ethernet device and operates with layer-2 packets such as Ethernet frames. TUN (as in network TUNnel) simulates a network layer device and operates with layer-3 packets such as IP packets. With TUN/TAP, virtual network interfaces can be created.

A user-space program usually attaches to the TUN/TAP virtual network interface. Packets sent by an operating system via a TUN/TAP network interface are delivered to the user-space program. When packets are sent by the program via a TUN/TAP network interface, the packets are injected into the operating system network stack. On the operating system, the packets appear to originate from an external source through the virtual network interface.

When a program attaches to a TUN/TAP interface, the IP packets that the computer sends to this interface will be piped into the program. IP packets that are sent by the program instead will be piped into the computer, as if they came from the Internet through this virtual network interface. Standard `read()` and `write()` function calls can be made to send or receive packets to and from the virtual interface accordingly.

The source code to create a simple TUN/TAP has been provided by Davide Brini<sup>1</sup>

---

<sup>1</sup><http://backreference.org/2010/03/26/tuntap-interface-tutorial>

and has been adapted by SEED Labs for this lab. The tutorial for the source code also shows how two computers can be connected using the TUN tunnelling technique. The adapted TUN/TAP code (**simpletun**) has been attached to Appendix A for reference.

The **simpletun** program can run as both a client and a server. To run as a client, the **-c** is used. To run it as a server instead, the **-s** flag is used. The following list details the requirements to create a tunnel between two computers using the **simpletun** program.

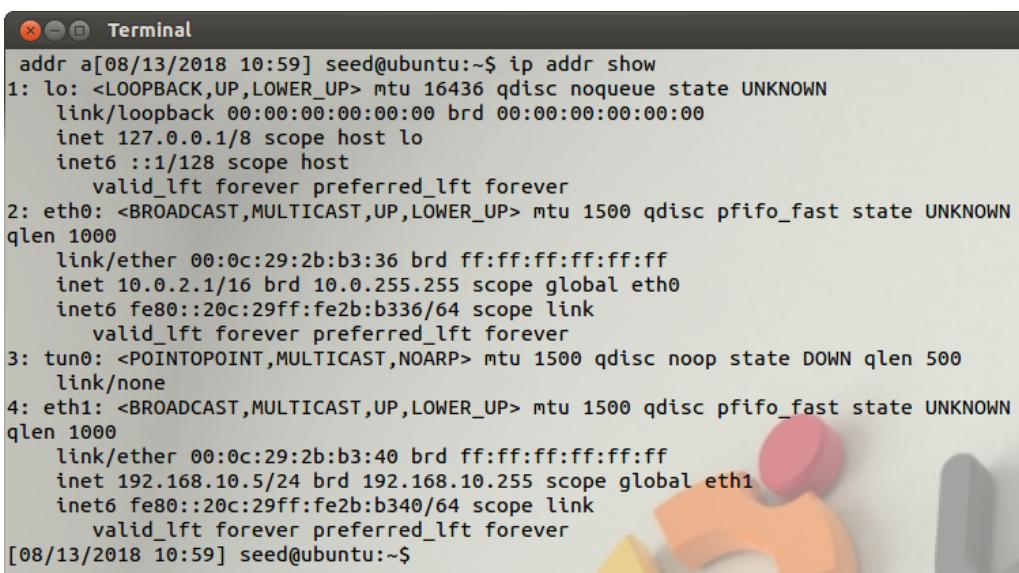
### 3.2.1 Setting Up Tunnel Point A

Tunnel point A is set to be the server side of the tunnel, which is on VM1. Once the tunnel between the two systems have been established, there is no difference between client and the server as the concept is only meaningful during the establishment of connection between two systems.

The following command is executed in terminal (with root privileges).

```
# ./simpletun -i tun0 -s -d
```

At this point, the virtual network device **tun0** has been created and the system now has multiple network interfaces. It will not show up if **ifconfig** is used as the interface is currently not active. Instead, the command **ip addr show** will list all physical and virtual network interfaces attached to the system regardless of status, as shown in Figure 2.



```
addr a[08/13/2018 10:59] seed@ubuntu:~$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 16436 qdisc noqueue state UNKNOWN
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            inet6 ::1/128 scope host
                valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UNKNOWN
    qlen 1000
    link/ether 00:0c:29:2b:b3:36 brd ff:ff:ff:ff:ff:ff
        inet 10.0.2.1/16 brd 10.0.255.255 scope global eth0
            inet6 fe80::20c:29ff:fe2b:b336/64 scope link
                valid_lft forever preferred_lft forever
3: tun0: <POINTOPOINT,MULTICAST,NOARP> mtu 1500 qdisc noop state DOWN qlen 500
    link/none
4: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UNKNOWN
    qlen 1000
    link/ether 00:0c:29:2b:b3:40 brd ff:ff:ff:ff:ff:ff
        inet 192.168.10.5/24 brd 192.168.10.255 scope global eth1
            inet6 fe80::20c:29ff:fe2b:b340/64 scope link
                valid_lft forever preferred_lft forever
[08/13/2018 10:59] seed@ubuntu:~$
```

Figure 2: Virtual Network Details

The new virtual network device is not fully configured, as evident by the absence of an IP address. The IP address that we will be assigning will be from the reserved IP address space 10.0.0.0/8. It is also important to note that the **Terminal** window that

was used to create the virtual network interface is currently waiting for connections and cannot be used currently. Another window needs to be opened to assign an IP address to the virtual network interface tun0, using the following commands.

```
# ip addr add 10.0.3.15/24 dev tun0
# ifconfig tun0 up
```

If `ifconfig` is used now, it appears as a valid network interface with its configuration displayed.

```
[08/13/2018 11:01] root@ubuntu:/home/seed# ifconfig
eth0      Link encap:Ethernet HWaddr 00:0c:29:b3:36
          inet addr:10.0.2.1 Bcast:10.0.255.255 Mask:255.255.0.0
          inet6 addr: fe80::20c:29ff:fe2b:b336/64 Scope:Link
            UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
            RX packets:4863 errors:0 dropped:0 overruns:0 frame:0
            TX packets:2126 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:1000
            RX bytes:5670417 (5.6 MB) TX bytes:216115 (216.1 KB)
            Interrupt:19 Base address:0x2000

eth1      Link encap:Ethernet HWaddr 00:0c:29:b3:40
          inet addr:192.168.10.5 Bcast:192.168.10.255 Mask:255.255.255.0
          inet6 addr: fe80::20c:29ff:fe2b:b340/64 Scope:Link
            UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
            RX packets:26 errors:0 dropped:0 overruns:0 frame:0
            TX packets:126 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:1000
            RX bytes:4900 (4.9 KB) TX bytes:21463 (21.4 KB)
            Interrupt:19 Base address:0x2400

lo        Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
            UP LOOPBACK RUNNING MTU:16436 Metric:1
            RX packets:272 errors:0 dropped:0 overruns:0 frame:0
            TX packets:272 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:0
            RX bytes:28567 (28.5 KB) TX bytes:28567 (28.5 KB)

tun0     Link encap:UNSPEC HWaddr 00-00-00-00-00-00-00-00-00-00-00-00-00-00-00-00
          inet addr:10.0.3.15 P-t-P:10.0.3.15 Mask:255.255.255.0
            UP POINTOPOINT RUNNING NOARP MULTICAST MTU:1500 Metric:1
            RX packets:0 errors:0 dropped:0 overruns:0 frame:0
            TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:500
            RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

[08/13/2018 11:02] root@ubuntu:/home/seed#
```

Figure 3: Valid Network Interfaces

### 3.2.2 Setting Up Tunnel Point B

Tunnel point B is used as the client side of the tunnel, which is to be on VM2. The setup is similar to VM1. The first command used will connect our client to the server (VM1), which is tunnel point A.

```
# ./simpletun -i tun0 -c 192.168.10.5 -d
```

Executing the line above will immediately trigger the Terminal window on both VM1 and VM2 to state that the connection between each other has been established. This shows that the configuration is correct.

```

[08/12/2018 06:36] root@ubuntu:/home/seed/Desktop# simpletun -i tun0 -s -d
Successfully connected to interface tun0
SERVER: Client connected from 192.168.10.6

[08/12/2018 07:53] root@ubuntu:/home/seed# simpletun -i tun0 -c 192.168.10.5 -d
Successfully connected to interface tun0
CLIENT: Connected to server 192.168.10.5

```

(a) Server Response  
(b) Client Response

Figure 4: Both Systems Connected

Again, the previous command will prevent any further input in the current window, hence the following commands to initialise the virtual network interface must be done in a separate window.

```
# ip addr add 10.0.3.16/24 dev tun0
# ifconfig tun0 up
```

Once it has been executed, the virtual network interface will have all the configurations required to be a functional network adapter.

### 3.2.3 Establishing Routing Path

By the previous steps, a tunnel has been established between the two systems. However before it can be used to transmit data, the routing path needs to be set up so that both machines will direct the intended outgoing traffic through the tunnel. The following routing table directs all packs for the 10.0.3.0/24 network through the interface `tun0`, where the packets will travel through the tunnel. The following must be executed on **both** VMs:

```
# route add -net 10.0.3.0 netmask 255.255.255.0 dev tun0
```

### 3.2.4 Using The Tunnel

Using VM1, we are now able to access VM2 through the `tun0` adapter. Likewise, VM2 can access VM1 through the `tun0` adapter as well. The tunnel connection can be tested using the `ssh` or `ping` programs.

To use `ssh`, the `ssh` server must be started first.

```
$ sudo service ssh start
```

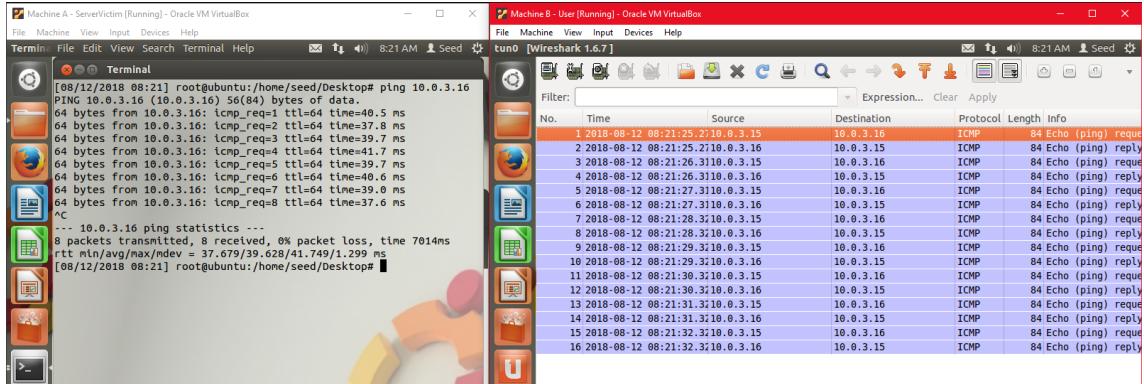
To test the tunnel from VM1:

```
$ ping 10.0.3.16
$ ssh 10.0.3.16
```

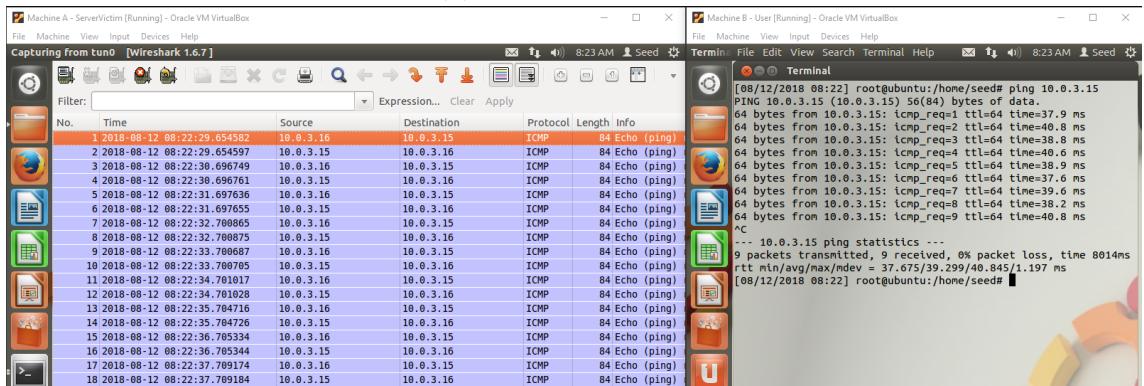
To test the tunnel from VM2:

```
$ ping 10.0.3.15
$ ssh 10.0.3.15
```

If the tunnel is successful, both `ssh` and `ping` will display valid responses from the other system. The figure below shows the packet capture from Wireshark for `ssh` in both directions is proof that the tunnel works as expected.

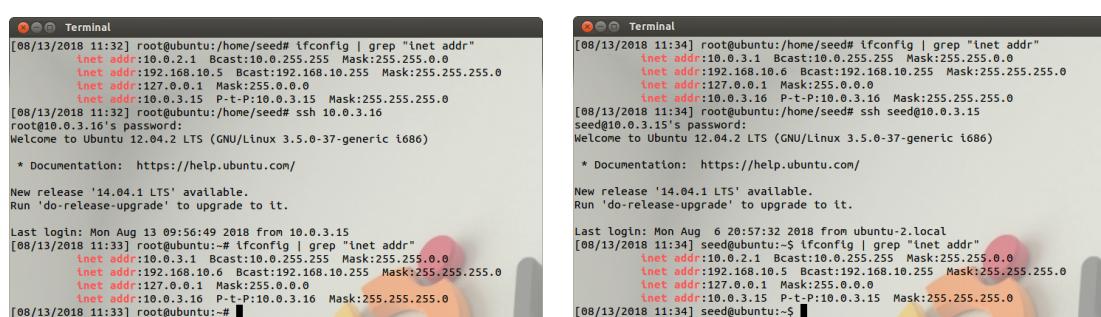


(a) Ping from A to B



(b) Ping from B to A

Figure 5: Both Systems Contactable



(a) SSH from A to B

(b) SSH from B to A

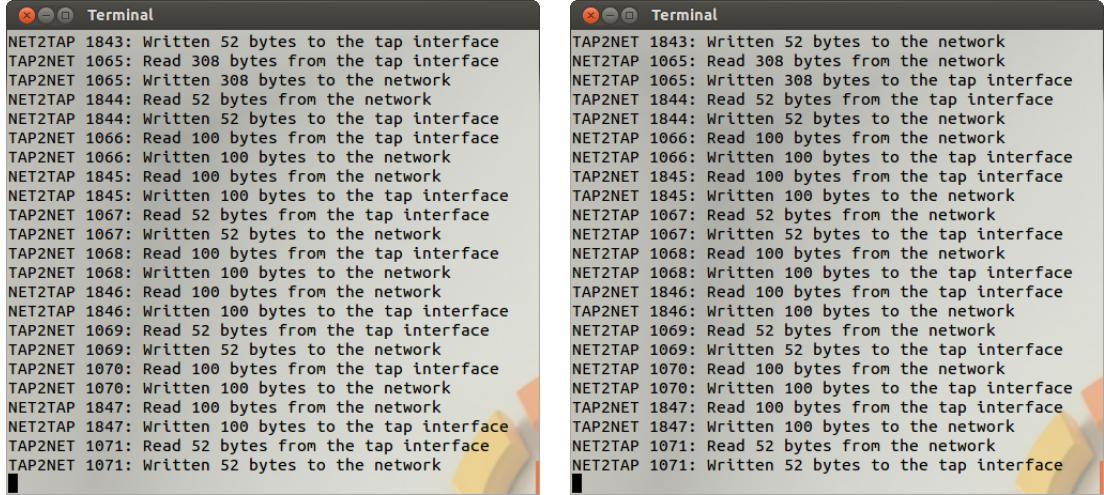
Figure 6: Both Systems Connected

Also of interest while the two systems are connected is that data that goes through the tunnel are displayed in the first Terminal window that was opened in VM1 and VM2 (the window that was used to start the `simpletun` program). Both display the data that was sent to/from the tap interface and the network. It is also noteworthy

to notice that the logs on both systems corroborate with the details mentioned in the introduction, which has been summarised into the table below with the corresponding screenshots from the VMs.

VM1	VM2
TAP2NET <id>: Read x bytes from the tap interface	NET2TAP <id>: Read x bytes from the network
TAP2NET <id>: Written x bytes to the network	NET2TAP <id>: Written x bytes to the tap interface
NET2TAP <id>: Read x bytes from the network	TAP2NET <id>: Read x bytes to the tap interface
NET2TAP <id>: Written x bytes to the tap interface	TAP2NET <id>: Written x bytes to the network

Table 1: Differences between logs of both VMs



```

Terminal
NET2TAP 1843: Written 52 bytes to the tap interface
TAP2NET 1065: Read 308 bytes from the tap interface
TAP2NET 1065: Written 308 bytes to the network
NET2TAP 1844: Read 52 bytes from the network
NET2TAP 1844: Written 52 bytes to the tap interface
TAP2NET 1066: Read 100 bytes from the tap interface
TAP2NET 1066: Written 100 bytes to the network
NET2TAP 1845: Read 100 bytes from the network
NET2TAP 1845: Written 100 bytes to the tap interface
TAP2NET 1067: Read 52 bytes from the tap interface
TAP2NET 1067: Written 52 bytes to the network
TAP2NET 1068: Read 100 bytes from the tap interface
TAP2NET 1068: Written 100 bytes to the network
NET2TAP 1846: Read 100 bytes from the network
NET2TAP 1846: Written 100 bytes to the tap interface
TAP2NET 1069: Read 52 bytes from the tap interface
TAP2NET 1069: Written 52 bytes to the network
TAP2NET 1070: Read 100 bytes from the tap interface
TAP2NET 1070: Written 100 bytes to the network
NET2TAP 1847: Read 100 bytes from the network
NET2TAP 1847: Written 100 bytes to the tap interface
TAP2NET 1071: Read 52 bytes from the tap interface
TAP2NET 1071: Written 52 bytes to the network

```

(a) Logs from VM1

```

Terminal
TAP2NET 1843: Written 52 bytes to the network
NET2TAP 1065: Read 308 bytes from the network
NET2TAP 1065: Written 308 bytes to the tap interface
TAP2NET 1844: Read 52 bytes from the tap interface
TAP2NET 1844: Written 52 bytes to the network
NET2TAP 1066: Read 100 bytes from the network
NET2TAP 1066: Written 100 bytes to the tap interface
TAP2NET 1845: Read 100 bytes from the tap interface
TAP2NET 1845: Written 100 bytes to the network
NET2TAP 1067: Read 52 bytes from the network
NET2TAP 1067: Written 52 bytes to the tap interface
NET2TAP 1068: Read 100 bytes from the network
NET2TAP 1068: Written 100 bytes to the tap interface
TAP2NET 1846: Read 100 bytes from the tap interface
TAP2NET 1846: Written 100 bytes to the network
NET2TAP 1069: Read 52 bytes from the network
NET2TAP 1069: Written 52 bytes to the tap interface
NET2TAP 1070: Read 100 bytes from the network
NET2TAP 1070: Written 100 bytes to the tap interface
TAP2NET 1847: Read 100 bytes from the tap interface
TAP2NET 1847: Written 100 bytes to the network
NET2TAP 1071: Read 52 bytes from the network
NET2TAP 1071: Written 52 bytes to the tap interface

```

(b) Logs from VM2

Figure 7: Logs in Sync

To this point, the connection of both VMs using a tunnel has been successful. Figure 8 is a diagrammatic representation on the current state of the network between the two VMs.

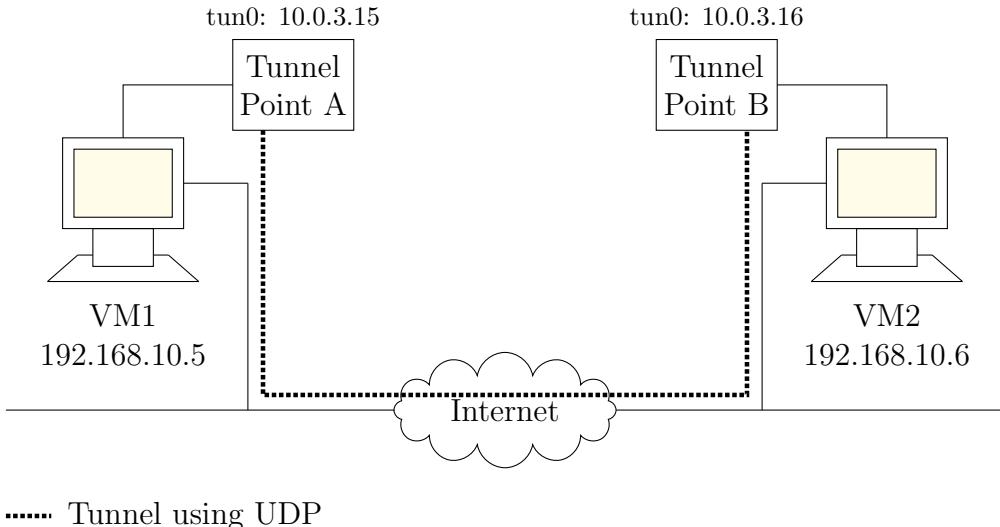


Figure 8: Tunnel Network over Emulated Internet

### 3.3 Task 2: Creating Host-to-Gateway Tunnel

Completing the previous task is insufficient as accessing VM2 through the tunnel is a secondary objective. Instead, the packet sent through the VPN tunnel are to be routed out to the Internet. In other words, VM2 must function as a gateway. This will make our tunnel a host-to-gateway tunnel. As a continuation of the previous task, the following additional tasks must be completed.

#### 3.3.1 Setting Up IP Forwarding

Unless the system is explicitly configured, a computer will only act as a host and not a gateway. To do so, the command below enables IP forwarding, allowing the computer to behave like a gateway.

```
$ sudo sysctl net.ipv4.ip_forward=1
```

#### 3.3.2 Getting Around The Limitation of NAT

When the destination sends the reply packets back to the machine on the private network, it will reach the NAT first (as the source IP of all outgoing packets are changed to the NAT's external IP address). The NAT will usually replace the destination IP address with the IP address of the original packet and send it to the system to whoever owns the IP address.

Before the NAT sends out the packet, it needs to know the MAC address of the machine that owns the IP address 10.0.3.15. To solve the limitation, an extra NAT is created on VM2 so that all packets sent out of VM2 will have VM2's IP address as its source IP. To reach the Internet, the packets will need to go through the NAT adapter that has already been provisioned when the VMs were set-up. The following sets of commands will enable the NAT on VM2.

The first step is to clean all of `iptables` rules.

```
$ sudo iptables -F  
$ sudo iptables -t nat -F
```

Next, we need to add a rule on the postROUTING position to the NAT adapter, in this instance with `eth0` (this network adapter may change depending on the interface identified within the VM). The network adapter selected for postROUTING must be outward-facing.

```
$ sudo iptables -t nat -A POSTROUTING -j MASQUERADE -o eth0
```

### 3.3.3 Setting Up Firewall

The firewall is set up on VM1 to block the access of a target website. Setting up the firewall requires superuser privileges, as with the setup of the VPN tunnel. With reference to the previously completed Firewall lab, popular social media sites Facebook and Twitter are blocked to prevent distractions within organisations.

Setting up the firewall on VM1 invokes an issue where the firewall should not be able to block packets over the virtual network interface or the VPN will not be able to function at all. Therefore, the firewall rule cannot be set before the routing, nor can it be set on the virtual interface.

The firewall rule only needs to be set for the physical network interface so it will not affect the virtual network interface. The following `iptables` command can help to overcome this problem by preventing the packets with the affected IP addresses from going through the physical network adapter `eth0`.

```
# iptables -t mangle -A POSTROUTING -d 155.69.7.173/24 -o eth0 -j DROP
```

In the above code, the NTU site (including NTULearn, iNTU etc.) are all blocked as one IP address serves the entire domain. Since NTU is the only site is blocked by the firewall rule, accessing other sites is not an issue.

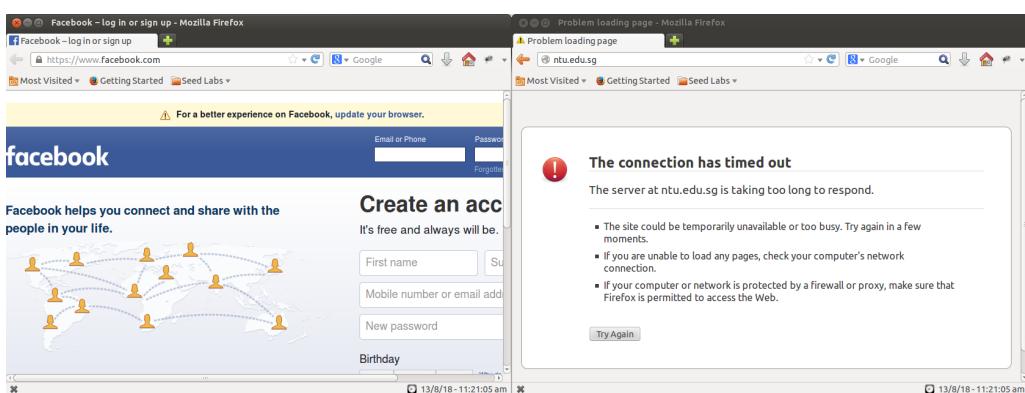


Figure 9: NTU Site Blocked

To manually check whether the IP address was properly stored in `iptables`, the following line of code can be used.

```
# iptables -L -t mangle --line-numbers
```

```
[08/13/2018 12:06] root@ubuntu:/home/seed# iptables -L -t mangle --line-numbers
Chain PREROUTING (policy ACCEPT)
num  target     prot opt source          destination
Chain INPUT (policy ACCEPT)
num  target     prot opt source          destination
Chain FORWARD (policy ACCEPT)
num  target     prot opt source          destination
Chain OUTPUT (policy ACCEPT)
num  target     prot opt source          destination
Chain POSTROUTING (policy ACCEPT)
num  target     prot opt source          destination
1   DROP       all  --  anywhere        155.69.7.0/24
[08/13/2018 12:06] root@ubuntu:/home/seed#
```

Figure 10: List of `mangle` entries

### 3.3.4 Bypassing Firewall

To bypass the egress filtering by the firewall, a SSH tunnel is established between the two VMs using the existing VPN tunnel. For web filtering, the dynamic port forwarding (`-D`) switch can be used as packets will automatically be forwarded to the destinations.

```
ssh -D 3000 seed@10.0.3.16
```

After SSH has been established, the browser must be configured to make use of the dynamic port forwarding mechanism. To do so, we must navigate to the following: `Edit >Preferences >Advanced >Network >Settings`. The “Manual Proxy Configuration” option must be selected. As we are using SSH, the SOCKS proxy should be used. For the host, it should be `localhost` with port 3000. SOCKS v5 needs to be selected as well. The rest of the options will remain as the default. Figure 11 is a screenshot of the proxy settings in Firefox.

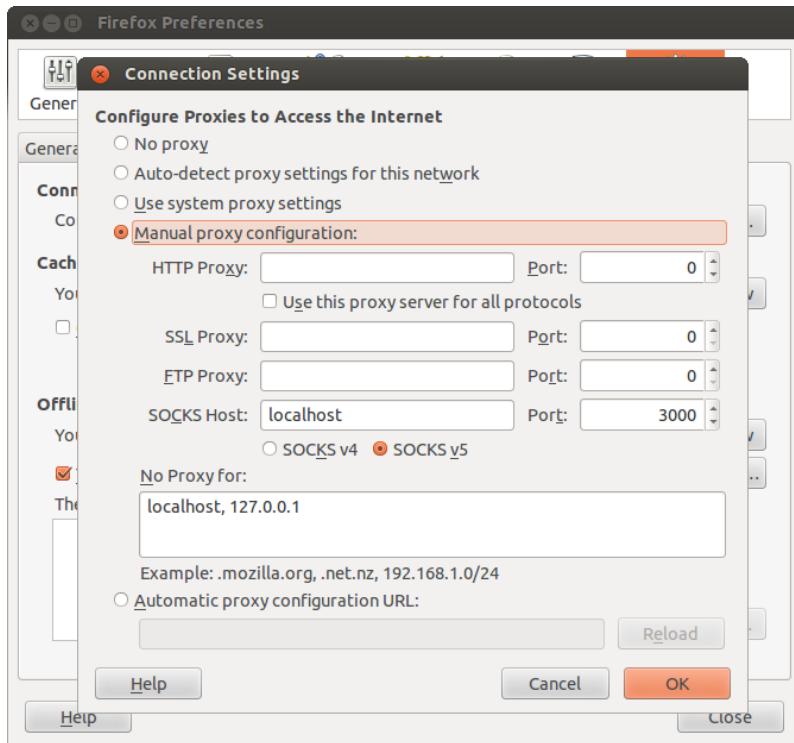


Figure 11: Proxy Configuration

Once completed, the *Connection Settings & Preferences* windows can be closed and the new settings will take effect immediately. If the proxy has been properly configured, then refreshing the site will allow the page to load up successfully, bypassing the firewall rule.

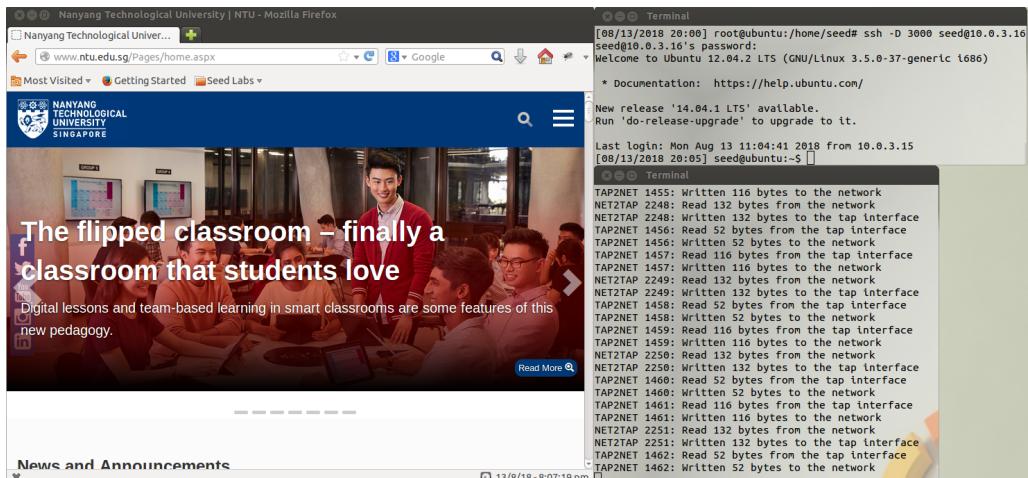


Figure 12: Firewall Bypassed For Websites

To demonstrate how applications can bypass the firewall, **Telnet** is used as a classic example. To block the **Telnet** protocol, the packet filtering firewall **ufw** can be used and it complements **iptables**. The following lines are executed in superuser privileges.

```
# ufw deny out to any port 23
# ufw deny in from any port 23
# ufw disable
# ufw enable
```

If there is any attempt to initiate Telnet from VM1, then the Telnet requests will timeout and no connection will be established since all packets are dropped by ufw.

The screenshot shows a terminal window titled "Terminal". The command entered was "telnet 10.0.3.1". The output shows the connection attempt failing with the message "telnet: Unable to connect to remote host: Connection timed out".

Figure 13: Telnet Timeout

To use Telnet, the SSH tunnel is made use of again. However, this time static port forwarding is used.

```
$ ssh -L 3000:10.0.3.16:23 seed@10.0.3.16
```

When the SSH tunnel has been established, the Telnet program has to make to use of the static port forwarding rule. To do so, Telnet connects to the source port and `localhost` to initiate the connection.

```
$ telnet localhost 3000
```

Successful connection through the SSL tunnel will prompt the user login and password of the remote system, in this case VM2. When `ifconfig | grep "inet addr"` is executed, the IP address that should show up will be for VM2.

The screenshot shows two terminal windows. The left window shows the command "telnet localhost 3000" being run, followed by the Telnet login process where the user "seed" logs in. The right window shows the command "ssh -L 3000:10.0.3.16:23 seed@10.0.3.16" being run, followed by the SSH session where the user "seed" logs in. Below these, the command "ifconfig | grep 'inet addr'" is run, showing the IP configuration for the interface.

Figure 14: Proof of Successful Telnet via ifconfig Checking

Using Wireshark to analyse the packets between the two VMs, it can be noticed that packets that are from VM1 are sent to `tun0` with the source and the destination being the host-only network IP address (`eth1`). This packet is “repackaged” and sent through the tunnel by SSH. This is reflected as the source and destination IP of the packet is the IP address of the TUN/TAP that was set in task 1. Again when the data is sent back, the packet headers are modified by the TUN/TAP adapter before it travels through the tunnel.

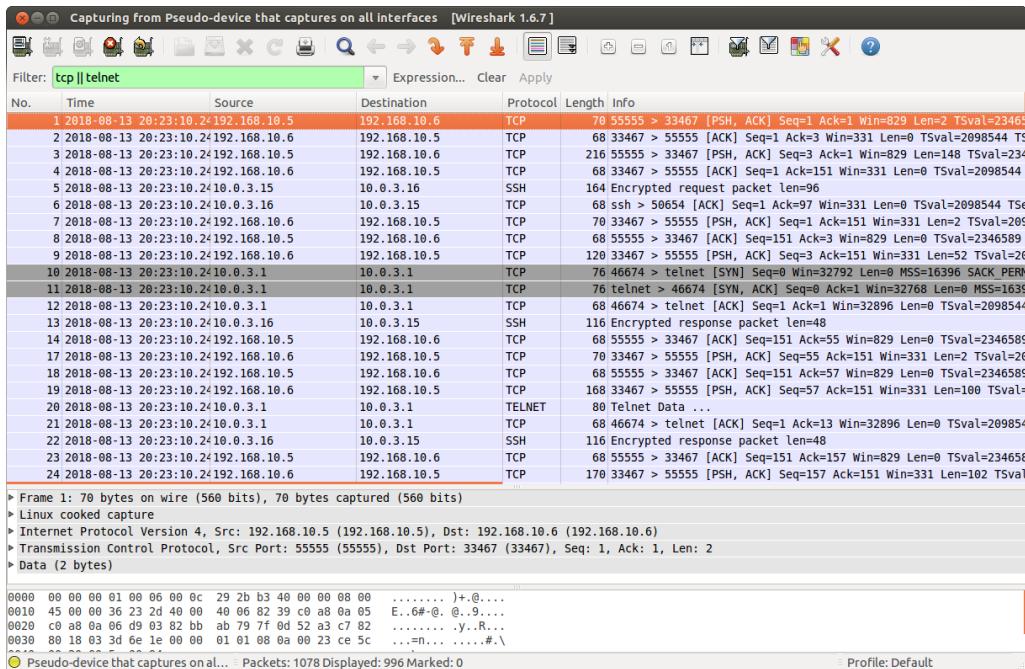


Figure 15: Packet Flow in Wireshark

## 4 Appendix A

```
1  ****
2  * simpletun.c
3  *
4  * A simplistic, simple-minded, naive tunnelling program using tun/tap
5  * interfaces and TCP. Handles (badly) IPv4 for tun, ARP and IPv4 for
6  * tap. DO NOT USE THIS PROGRAM FOR SERIOUS PURPOSES.
7  *
8  * You have been warned.
9  *
10 * (C) 2009 Davide Brini.
11 *
12 * DISCLAIMER AND WARNING: this is all work in progress. The code is
13 * ugly, the algorithms are naive, error checking and input validation
14 * are very basic, and of course there can be bugs. If that's not enough,
15 * the program has not been thoroughly tested, so it might even fail at
16 * the few simple things it should be supposed to do right.
17 * Needless to say, I take no responsibility whatsoever for what the
18 * program might do. The program has been written mostly for learning
19 * purposes, and can be used in the hope that is useful, but everything
20 * is to be taken "as is" and without any kind of warranty, implicit or
21 * explicit. See the file LICENSE for further details.
22 ****
23
24 #include <stdio.h>
25 #include <stdlib.h>
26 #include <string.h>
27 #include <unistd.h>
28 #include <sys/socket.h>
29 #include <linux/if.h>
30 #include <linux/if_tun.h>
31 #include <sys/types.h>
32 #include <sys/ioctl.h>
33 #include <sys/stat.h>
34 #include <fcntl.h>
35 #include <arpa/inet.h>
36 #include <sys/select.h>
37 #include <sys/time.h>
38 #include <errno.h>
39 #include <stdarg.h>
40
41 /* buffer for reading from tun/tap interface, must be >= 1500 */
42 #define BUFSIZE 2000
43 #define CLIENT 0
44 #define SERVER 1
45 #define PORT 55555
46
47 /* some common lengths */
48 #define IP_HDR_LEN 20
49 #define ETH_HDR_LEN 14
50 #define ARP_PKT_LEN 28
```

```

51
52 int debug;
53 char *progname;
54
55 /*****
56 * tun_alloc: allocates or reconnects to a tun/tap device. The caller      *
57 *           needs to reserve enough space in *dev.                      *
58 *****/
59 int tun_alloc(char *dev, int flags) {
60
61     struct ifreq ifr;
62     int fd, err;
63
64     if( (fd = open("/dev/net/tun", O_RDWR)) < 0 ) {
65         perror("Opening /dev/net/tun");
66         return fd;
67     }
68
69     memset(&ifr, 0, sizeof(ifr));
70
71     ifr.ifr_flags = flags;
72
73     if (*dev) {
74         strncpy(ifr.ifr_name, dev, IFNAMSIZ);
75     }
76
77     if( (err = ioctl(fd, TUNSETIFF, (void *)&ifr)) < 0 ) {
78         perror("ioctl(TUNSETIFF)");
79         close(fd);
80         return err;
81     }
82
83     strcpy(dev, ifr.ifr_name);
84
85     return fd;
86 }
87
88 /*****
89 * cread: read routine that checks for errors and exits if an error is   *
90 *           returned.                                              *
91 *****/
92 int cread(int fd, char *buf, int n){
93
94     int nread;
95
96     if((nread=read(fd, buf, n))<0){
97         perror("Reading data");
98         exit(1);
99     }
100    return nread;
101 }
102

```

```

103 /*****
104 * cwrite: write routine that checks for errors and exits if an error is *
105 *          returned.                                              */
106 *****/
107 int cwrite(int fd, char *buf, int n){
108
109     int nwrite;
110
111     if((nwrite=write(fd, buf, n))<0){
112         perror("Writing data");
113         exit(1);
114     }
115     return nwrite;
116 }
117
118 /*****
119 * read_n: ensures we read exactly n bytes, and puts those into "buf".      *
120 *          (unless EOF, of course)                                         */
121 *****/
122 int read_n(int fd, char *buf, int n) {
123
124     int nread, left = n;
125
126     while(left > 0) {
127         if ((nread = cread(fd, buf, left))==0){
128             return 0 ;
129         }else {
130             left -= nread;
131             buf += nread;
132         }
133     }
134     return n;
135 }
136
137 /*****
138 * do_debug: prints debugging stuff (doh!)                                     */
139 *****/
140 void do_debug(char *msg, ...){
141
142     va_list argp;
143
144     if(debug){
145         va_start(argp, msg);
146         vfprintf(stderr, msg, argp);
147         va_end(argp);
148     }
149 }
150
151 /*****
152 * my_err: prints custom error messages on stderr.                           */
153 *****/
154 void my_err(char *msg, ... ) {

```

```

155
156     va_list argp;
157
158     va_start(argp, msg);
159     vfprintf(stderr, msg, argp);
160     va_end(argp);
161 }
162
163 /*****
164 * usage: prints usage and exits.
165 *****/
166 void usage(void) {
167     fprintf(stderr, "Usage:\n");
168     fprintf(stderr, "%s -i <ifacename> [-s|-c <serverIP>] [-p <port>] [-u|-a]
169     ↪ [-d]\n", progname);
170     fprintf(stderr, "%s -h\n", progname);
171     fprintf(stderr, "\n");
172     fprintf(stderr, "-i <ifacename>: Name of interface to use (mandatory)\n");
173     fprintf(stderr, "-s|-c <serverIP>: run in server mode (-s), or specify server
174     ↪ address (-c <serverIP>) (mandatory)\n");
175     fprintf(stderr, "-p <port>: port to listen on (if run in server mode) or to
176     ↪ connect to (in client mode), default 55555\n");
177     fprintf(stderr, "-u|-a: use TUN (-u, default) or TAP (-a)\n");
178     fprintf(stderr, "-d: outputs debug information while running\n");
179     fprintf(stderr, "-h: prints this help text\n");
180     exit(1);
181 }
182
183 int main(int argc, char *argv[]) {
184     int tap_fd, option;
185     int flags = IFF_TUN;
186     char if_name[IFNAMSIZ] = "";
187     int header_len = IP_HDR_LEN;
188     int maxfd;
189     uint16_t nread, nwrite, plength;
190 //     uint16_t total_len, ethertype;
191     char buffer[BUFSIZE];
192     struct sockaddr_in local, remote;
193     char remote_ip[16] = "";
194     unsigned short int port = PORT;
195     int sock_fd, net_fd, optval = 1;
196     socklen_t remotelen;
197     int cliserv = -1; /* must be specified on cmd line */
198     unsigned long int tap2net = 0, net2tap = 0;
199
200     /* Check command line options */
201     while((option = getopt(argc, argv, "i:sc:p:uahd")) > 0){
202         switch(option) {
203             case 'd':

```

```

204     debug = 1;
205     break;
206     case 'h':
207         usage();
208         break;
209     case 'i':
210         strncpy(if_name,optarg,IFNAMSIZ-1);
211         break;
212     case 's':
213         cliserv = SERVER;
214         break;
215     case 'c':
216         cliserv = CLIENT;
217         strncpy(remote_ip,optarg,15);
218         break;
219     case 'p':
220         port = atoi(optarg);
221         break;
222     case 'u':
223         flags = IFF_TUN;
224         break;
225     case 'a':
226         flags = IFF_TAP;
227         header_len = ETH_HDR_LEN;
228         break;
229     default:
230         my_err("Unknown option %c\n", option);
231         usage();
232     }
233 }

234 argv += optind;
235 argc -= optind;
236
237 if(argc > 0){
238     my_err("Too many options!\n");
239     usage();
240 }
241

242 if(*if_name == '\0'){
243     my_err("Must specify interface name!\n");
244     usage();
245 }
246 else if(cliserv < 0){
247     my_err("Must specify client or server mode!\n");
248     usage();
249 }
250 else if((cliserv == CLIENT)&&(*remote_ip == '\0')){
251     my_err("Must specify server address!\n");
252     usage();
253 }
254 /* initialize tun/tap interface */
255 if ( (tap_fd = tun_alloc(if_name, flags | IFF_NO_PI)) < 0 ) {

```

```

256     my_err("Error connecting to tun/tap interface %s!\n", if_name);
257     exit(1);
258 }
259
260 do_debug("Successfully connected to interface %s\n", if_name);
261
262 if ( (sock_fd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
263     perror("socket()");
264     exit(1);
265 }
266
267 if(cliserv==CLIENT){
268     /* Client, try to connect to server */
269
270     /* assign the destination address */
271     memset(&remote, 0, sizeof(remote));
272     remote.sin_family = AF_INET;
273     remote.sin_addr.s_addr = inet_addr(remote_ip);
274     remote.sin_port = htons(port);
275
276     /* connection request */
277     if (connect(sock_fd, (struct sockaddr*) &remote, sizeof(remote)) < 0){
278         perror("connect()");
279         exit(1);
280     }
281
282     net_fd = sock_fd;
283     do_debug("CLIENT: Connected to server %s\n", inet_ntoa(remote.sin_addr));
284
285 } else {
286     /* Server, wait for connections */
287
288     /* avoid EADDRINUSE error on bind() */
289     if(setsockopt(sock_fd, SOL_SOCKET, SO_REUSEADDR, (char *)&optval,
290                   sizeof(optval)) < 0){
291         perror("setsockopt()");
292         exit(1);
293     }
294
295     memset(&local, 0, sizeof(local));
296     local.sin_family = AF_INET;
297     local.sin_addr.s_addr = htonl(INADDR_ANY);
298     local.sin_port = htons(port);
299     if (bind(sock_fd, (struct sockaddr*) &local, sizeof(local)) < 0){
300         perror("bind()");
301         exit(1);
302     }
303
304     if (listen(sock_fd, 5) < 0){
305         perror("listen()");
306         exit(1);
307     }

```

```

307
308 /* wait for connection request */
309 remotelen = sizeof(remote);
310 memset(&remote, 0, remotelen);
311 if ((net_fd = accept(sock_fd, (struct sockaddr*)&remote, &remotelen)) < 0){
312     perror("accept()");
313     exit(1);
314 }
315
316 do_debug("SERVER: Client connected from %s\n", inet_ntoa(remote.sin_addr));
317 }
318
319 /* use select() to handle two descriptors at once */
320 maxfd = (tap_fd > net_fd)?tap_fd:net_fd;
321
322 while(1) {
323     int ret;
324     fd_set rd_set;
325
326     FD_ZERO(&rd_set);
327     FD_SET(tap_fd, &rd_set); FD_SET(net_fd, &rd_set);
328
329     ret = select(maxfd + 1, &rd_set, NULL, NULL, NULL);
330
331     if (ret < 0 && errno == EINTR){
332         continue;
333     }
334
335     if (ret < 0) {
336         perror("select()");
337         exit(1);
338     }
339
340     if(FD_ISSET(tap_fd, &rd_set)){
341         /* data from tun/tap: just read it and write it to the network */
342
343         nread = cread(tap_fd, buffer, BUFSIZE);
344
345         tap2net++;
346         do_debug("TAP2NET %lu: Read %d bytes from the tap interface\n", tap2net,
347             → nread);
348
349         /* write length + packet */
350         plength = htons(nread);
351         nwrite = cwrite(net_fd, (char *)&plength, sizeof(plength));
352         nwrite = cwrite(net_fd, buffer, nread);
353
354         do_debug("TAP2NET %lu: Written %d bytes to the network\n", tap2net, nwrite);
355     }
356
357     if(FD_ISSET(net_fd, &rd_set)){
358         /* data from the network: read it, and write it to the tun/tap interface.

```

```

358     * We need to read the length first, and then the packet */
359
360     /* Read length */
361     nread = read_n(net_fd, (char *)&plenlength, sizeof(plenlength));
362     if(nread == 0) {
363         /* ctrl-c at the other end */
364         break;
365     }
366
367     net2tap++;
368
369     /* read packet */
370     nread = read_n(net_fd, buffer, ntohs(plenlength));
371     do_debug("NET2TAP %lu: Read %d bytes from the network\n", net2tap, nread);
372
373     /* now buffer[] contains a full packet or frame, write it into the tun/tap
374      ↵ interface */
375     nwrite = cwrite(tap_fd, buffer, nread);
376     do_debug("NET2TAP %lu: Written %d bytes to the tap interface\n", net2tap,
377             ↵ nwrite);
378
379     return(0);
380 }
```